

*Microcontroller-Based  
Digital System Design  
featuring the Motorola 68HC12*

***PRELIMINARY Edition***  
*of*  
*Chapters 2 & 3*

*David G. Meyer*

*Copyright 2001 by D. G. Meyer*

### Copyright Notice

All rights reserved. No part of this Lecture-Workbook or Text may be reproduced, in any form or by any means, without permission in writing from the author.

## *Preface*

The purpose of this book is to teach students how to design and implement a microcontroller-based digital system. As such, it contains material that might typically be covered in a sequence of two courses: (1) a junior-level “microprocessor” course covering the basics of how a microprocessor works, how to program it to perform basic functions, and how to interface it to various external devices using integrated peripherals; and (2) a senior-level “digital system design project” course covering more advanced topics on microprocessor programming and interfacing, along with a series of practical system design considerations. Note that a background in basic digital system design is a necessary prerequisite, ideally obtained during the student’s sophomore year. While there are a number of reasonably good texts currently available that provide such an introduction, one of the best (and my long-time personal favorite) is John F. Wakerly’s *Digital Design Principles and Practices (Third Edition)*, Prentice Hall, 2000.

A unique feature of *Microcontroller-Based Digital System Design* (sub-titled *Bigger Bytes of Digital Wisdom*, or *Bigger Bytes* for short) is the availability of what I refer to as a “Lecture Workbook”, i.e., a set of lecture slides (provided in PowerPoint™ format) with *carefully chosen* portions to be annotated or completed in class. The Lecture Workbook concept is based on the premise that notes taken during a classroom lecture serve more than mere archival of information – an encoding process occurs in the student’s brain as he/she writes. By focusing this encoding process on *key words* or *selected aspects* of hardware/software design, the time and effort spent in class can be optimized. A special set of PowerPoint™ slides, which include an *animated, successive annotation* of the Lecture Workbook slides (including completed exercises), is available for instructor use. (The “skeleton” slides can also be made into overhead transparencies and annotated “manually”, for those instructors who prefer that mode of presentation.)

Another student- and instructor-friendly feature is the availability of an “Exercise Workbook” that contains a set of (full-size) printable homework problems in PDF format along with solutions to selected exercises. Also included are a number of source files that are to be completed as part of these problems. Individual students can print out selected problems and complete them in a structured, “easy-to-grade” fashion.

The availability of a complete ‘Lab Workbook’ – based on a low-cost evaluation board (EVB) available directly from Motorola University Support – is another feature of this text. The Motorola EVBs have a small prototyping area that makes them ideal not only for introductory courses on microcontrollers, but also for use in senior design projects.

## ***Table of Contents***

2	DESIGN OF A SIMPLE COMPUTER	
2.1	Computer Design Basics	3
2.2	Simple Computer Big Picture	5
2.3	Simple Computer Floor Plan	7
2.4	Simple Computer Programming Example	9
2.5	Simple Computer Block Diagram	15
2.6	Instruction Execution Tracing	18
3.7	Bottom-Up Implementation of Simple Computer	24
3.7.1	<i>Memory</i>	24
3.7.2	<i>Program Counter</i>	28
3.7.3	<i>Instruction Register</i>	30
3.7.4	<i>Arithmetic Logic Unit</i>	31
3.7.5	<i>Instruction Decoder and Micro-sequencer</i>	35
3.8	System Timing Analysis	40
3.9	Simple Computer Extensions	42
3.9.1	<i>Input/Output Instructions</i>	42
3.9.2	<i>Transfer-of-Control Instructions</i>	47
3.9.3	<i>Multiple Execute Cycle Instructions</i>	50
3.9.4	<i>Stack Manipulation Instructions</i>	53
3.9.5	<i>Subroutine Linkage Instructions</i>	58
3.9.6	<i>Other Possibilities</i>	63
2.10	Summary and References	64
	Problems	65
3	INTRODUCTION TO MICROCONTROLLER ARCHITECTURE AND PROGRAMMING MODEL	
3.1	Differing World Views	2
3.2	Characteristics That Distinguish Microprocessors	4
3.3	Taxonomy of Microprocessors	6
3.4	Choosing an Education-Appropriate Microprocessor	9
3.5	Tools of the Trade	12
3.6	Motorola 68HC12 Architecture and Programming Model	26
3.10	Addressing Modes	30
3.7.1	<i>Non-Indexed Modes</i>	31
3.7.2	<i>Indexed Modes</i>	33
3.7.3	<i>Addressing Mode Summary</i>	38
3.8	Motorola 68HC12 Instruction Set Overview	40
3.8.1	<i>Data Transfer Group Instructions</i>	40
3.8.2	<i>Arithmetic Group Instructions</i>	46
3.8.3	<i>Logical Group Instructions</i>	57
3.8.4	<i>Transfer-of-Control Group Instructions</i>	64
3.8.5	<i>Machine Control Group Instructions</i>	76
3.8.6	<i>Special Group Instructions</i>	79
3.9	Summary and References	82
	Problems	83

## CHAPTER 2

### DESIGN OF A SIMPLE COMPUTER

Before we launch into the details associated with a relatively complex, contemporary microcontroller, it will be helpful for us to examine the design and implementation of a *simple* computer. In particular, the overall approach – based on a *top-down* specification of functionality, followed by a *bottom-up* implementation of the various functional blocks – will prove useful to our basic understanding of how a “real” microcontroller works.

*top-down,  
bottom-up*

In Chapter 1, we reviewed a number of digital system building blocks. This included combinational elements such as decoders, priority encoders, and multiplexers as well as sequential elements such as latches and flip-flops. We then reviewed how these combinational and sequential elements can be combined to build digital systems. We also reviewed how digital systems could be specified using a hardware description language and subsequently implemented using *programmable logic devices* (PLDs).

*programmable  
logic devices*

Our purpose here is to apply this background to the design of a simple computer. Before we go any further, though, some basic definitions are in order. First, what is a *computer*? What distinguishes computers from random combinations of logic or from simple “light flashing” state machines? Simply stated, a computer is a device that *sequentially executes a stored program*. The program executed is typically called *software* if it is a user-programmable (“general purpose”) computer system; or called *firmware* if it is a single-purpose, non-user-programmable system (also referred to as a “turn-key” system). A given program consists of a *series of instructions* that the machine understands. Instructions are simply bit patterns that tell the computer what operation to perform on specified data. That a program is *stored* implies the existence of *memory*. To perform the series of instructions stored in memory, two basic operations need to be performed. First, an instruction must be *fetched* (read) from memory. Second, that instruction must be *executed*, e.g., two numbers are added together to produce a result. The memory that is used to store a program can take many different forms – ranging from removable media devices such as CD-ROMs to patterns in the metal layer of an integrated circuit. While the physical implementation of the memory in which the program is

*computer*

*stored program*

*software*

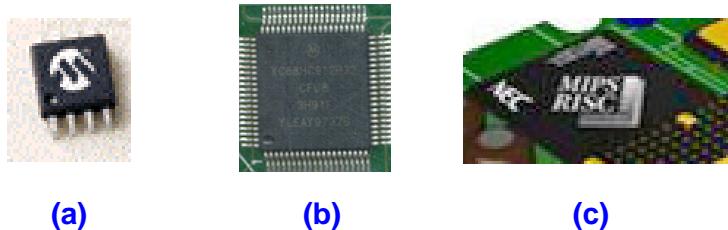
*firmware*

*memory*

stored may vary, the information stored in memory is interpreted (i.e., fetched and executed) the same way.

Given the basic definition of a computer, above, what is a *microprocessor*? Classically, it is a single-chip embodiment of the major functional blocks of a computer. Today, though, the term “microprocessor” is often applied to a wide range of single- and multi-chip computational devices, ranging from “mainframes on a chip” (used in personal computers and workstations) to small dedicated controllers (used in a wide variety of “intelligent” devices). They can range in physical size from packages with several hundred pins to packages with only a few pins; some examples are illustrated in Figure 2-1. They can range in cost from less than one dollar to hundreds of dollars. The simple computer we will be designing here can be implemented using a modest-size PLD; we could therefore rightfully call this single-chip embodiment of our simple computer a “microprocessor.”

**microprocessor**



(a)

(b)

(c)

**Figure 2-1** Contrasting contemporary microprocessors: (a) an 8-bit PIC microcontroller; (b) a 16-bit Motorola 68HC12 microcontroller; and (c) a 64-bit MIPS microprocessor.

Finally, what is a *microcontroller*, and how does it differ from a microprocessor? Typically a microcontroller integrates, in addition to a microprocessor, a number of *peripheral devices* that are commonly used in control-type applications onto a single integrated circuit (and are thus often referred to as “single-chip microcontrollers”). Peripheral devices get their name from the fact that they provide interfaces with devices that are external (i.e., “peripheral”) to the computer. For example, a common series of operations often performed in control applications is: (1) input analog signals from sensors, (2) process them according to some algorithm, (3) and output analog control voltages to actuators. A device that digitizes an analog input voltage is called an analog-to-digital (A-to-D) converter. Conversely, a device that produces an analog output voltage based on a digital code is called a digital-to-analog (D-to-A) converter. A-to-D and D-to-A converters are examples of peripherals one might find integrated onto a microcontroller chip.

**microcontroller**

**peripheral devices**

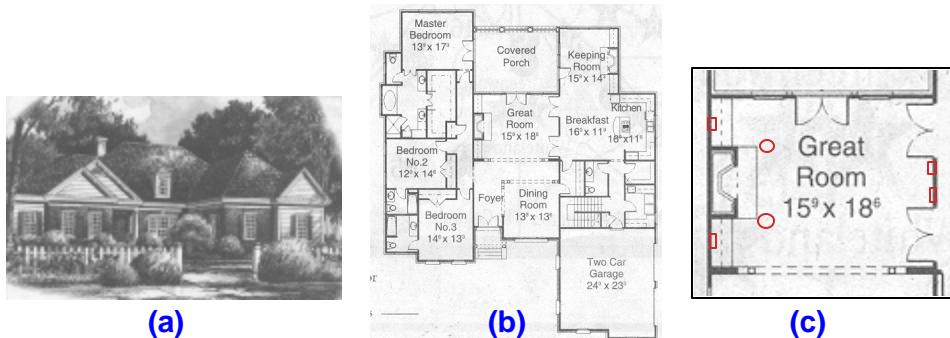
Other common peripherals include communication controllers, timer modules, and pulse-width modulation (PWM) generators. Later, we will see a variety of applications for all of these integrated peripherals.

## 2.1 Computer Design Basics

How can we apply what we have learned thus far about basic digital system building blocks toward building a simple computer? Basically, what we need is some way to structure and break down this design problem, because now it is a somewhat bigger than drawing a single state transition diagram or filling out a truth table. We will need a structured approach that enables us to take a written description of the functions performed by our simple computer and create a high-level block diagram. Based on this diagram, we can proceed to define what each block does, and ultimately design the circuitry required to implement each block.

Before starting this process, though, we need to define what we mean by the *structure* of a computer. “Architecture” is a word commonly used to depict the arrangement and interconnection of a computer’s functional blocks. While some might argue that this definition of computer architecture is a bit simplistic, it will serve our purposes for the discussion that follows.

Before starting to design our simple computer, let us first consider a “real world” analogy: building a house. Where is the logical place to start? Probably with a “big picture” – i.e., an *exterior elevation* or *plan view* of the entire project. Of course, the floor plan and exterior elevation are greatly influenced by the size, shape, and grade of the lot chosen for the house. Once we know the physical constraints dictated by our choice of lot, we can then begin to develop a floor plan. At this stage we can define the overall “functionality” of the house, i.e., the purpose of each room. Once we have defined the functionality of each room, the next step is to determine their arrangement and interconnection. Once we have a working floor plan, we can begin to embellish it with a number of details – for example, the location and size of windows, the location of light fixtures and their associated wall switches, the location of power outlets, the routing of plumbing, etc. The important thing to note from this analogy is that we have described a top-down design process: starting with a “big picture”, and progressively embellishing it with layers of details. Figure 2-2 depicts such a progression.



**Figure 2-2** Top-down design of a house: (a) the “big picture”, (b) the floor plan, (c) details of a particular room.

Once all the design specifications have been formulated, how would we proceed to build our house? From the ground up – assuming we have adequate financing, of course. We have to dig a hole first (perhaps analogous to going into debt), then pour a foundation, “stick-build” the basic structure, put a roof on it, complete the exterior walls, and finally embellish each room with its finishing details. Note that the *order* in which this “bottom up” implementation proceeds is quite important – certainly one would not wish to start hanging drywall before the roof is in place, or run plumbing lines before the floor joists are in place. Clearly, there is a structured, ordered way in which the entire process must take place – an approach strikingly similar to the one we will follow in designing our simple computer.

What would be a good name for the overall process described above? Ignoring the financial aspects for a moment, we could aptly call it the top-down specification of functionality followed by bottom-up implementation of each basic step (or “block”). More succinctly, we could call it *top-down specification and bottom-up implementation*. This is the process we will apply to the design and implementation of our simple computer.

First, a disclaimer. The initial machine we design will be very, very simple. It will be an 8-bit machine with just a few instructions. Further, there will be a single *instruction format* (layout of bit patterns) as well as a single *addressing mode* (way that the processor accesses operands in memory). By the time we finish this “first phase” design, however, we will find out that even this rather simple machine is fairly complex in terms of implementation details.

Once we have mastered our simple computer, we will then add “modern conveniences” such as input and output (or “I/O”), transfer of control instructions, stack manipulation instructions, and subroutine

**top-down specification**

**bottom-up implementation**

**instruction format**

**addressing mode**

linkage instructions. We will have the makings of a “socially redeeming” computer once we get done, plus have a firm footing upon which to understand the architecture and instruction set of a “real” computer.

*socially  
redeeming*

## 2.2 Simple Computer Big Picture

Just as one might begin the design of a house by sketching an exterior elevation view, we will begin the design of our simple computer with a “big picture” of its control console. In the “old days” (which was actually not so long ago), computers had lots of lights and switches on their front panels. The Digital Equipment Corporation PDP-8 (the first commercial “minicomputer”), illustrated in Figure 2-3, was a good example of such a computer. The Intellect 8 microcomputer system (one of the first commercially-available microprocessor development systems) from Intel, based on the 8008 microprocessor, was another example. Frankly, these ground-breaking computer systems were a lot more interesting (and fun) to watch “crunch numbers” than today’s computers...and a lot less irritating than the “this application has performed an illegal function and will be shut down” message we’ve all become accustomed to today.

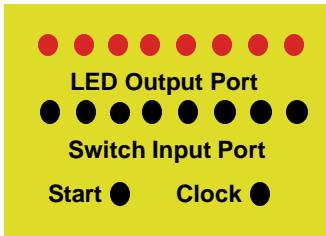
*old days*

*minicomputer*

*crunch numbers*



**Figure 2-3** World's first “desktop” minicomputer, the PDP-8.



**Figure 2-4** Our simple computer console.

Our computer’s console, then, will have some lights that indicate the result of the most recent computation along with some switches that will be used to input data. A “START” pushbutton will be included to get the machine into a known initial state (in preparation for “running” a program), and a “CLOCK” pushbutton will be included to facilitate debugging (as we manually clock the machine from state-to-state). An “artist’s conception” of our simple computer’s console is shown in Figure 2-4.

Returning to the “house analogy” for a moment, the floor plan of a computer is basically its *instruction set* and *programming model*. The instruction set is simply the list of operations that the computer performs. There are five fundamental groups (or categories) of machine instructions: data transfer, arithmetic, logical (or “Boolean”), transfer of control, and machine control. (Some computers include a sixth group dedicated to specific applications, e.g., multimedia extensions or graphics support.) The *addressing modes* that instructions can use to access operands in memory are also a key aspect of a computer’s instruction set.

The *programming model* of a computer is the software writer’s view of the machine. Basically, it tells what resources are available for the programmer’s use, in particular, the machine’s registers. A register is simply a “memory location” within the processor that can be used to store intermediate results and/or as an operand (or as a *pointer* to an *operand*) used in a computation.

As alluded to above, the programming model and instruction set of our computer will be relatively simple. Initially there will only be one register, called the accumulator (or “A” register), so-named because it is the register in which the result of computations accumulate. Our computer will also include several condition code bits: a zero flag (ZF), negative flag (NF), overflow flag (VF), and carry/borrow flag (CF). Before we complete this chapter, we will add a stack pointer register and discuss the role of index registers.

The instructions executed by our simple computer will be of the fixed-length variety (i.e., all 8-bits in size, hence its designation as an “8-bit” computer) that consist of two fixed-length fields. The upper 3-bits of each instruction will indicate the operation to be performed, and is therefore called the *operation code field* (or “opcode” field). The lower 5-bits will indicate the memory address in which the operand is located (or, a result is to be stored). The 5-bit memory address dictates a maximum memory size of  $2^5 = 32$  locations. For those who have become jaded by multi-megabyte programs that appear to do trivial things, this may not seem like much memory! Fortunately, though, it will be enough to illustrate basic principles of instruction execution, despite being too small to contain a “practical” (i.e., useful and socially redeeming) program.

In addition to fixed-field decoding, another simplification in our initial design will be a single *addressing mode*. An addressing mode is the mechanism (or “function”) used to generate what is often called the

*instruction set*

*programming model*

*addressing modes*

*pointer*

*condition code bits*

*ZF*

*NF*

*VF*

*CF*

*opcode field*

*addressing mode*

effective address of an operand, i.e., the actual address in memory where an operand is stored. The addressing mode our machine will support might aptly be called “absolute” addressing, based on the fact that this 5-bit field directly indicates the effective address in memory where the operand is stored. It is important to note at this point that not all manufacturers of microprocessors agree on the names ascribed to certain addressing modes. What we have just referred to as an “absolute” addressing mode is typically called “extended” (by Motorola) or “direct” (by Intel).

*effective  
address*

*absolute  
addressing  
mode*

One other bit of terminology worth mentioning before delving into the instruction set concerns the number of addresses a given instruction (or more generally, a machine) can accommodate. Our simple computer here could be described as a “two address” machine, which means that two different locations (at two different addresses) are used in a given operation, e.g., ADD. In our computer, one location will be the “A” register (the accumulator), and the other will be contained in memory. Note that a “side-effect” of such an arrangement is that the result of the computation will overwrite one of the operands, here the value in the “A” register (the operand in memory will be unaffected). As one might guess, there are a lot of variations in instruction format and addressing capability, ranging from single-address instructions to three-address (or more) instructions.

*two-address  
machine*

## 2.3 Simple Computer Floor Plan

We are now ready to introduce the “floor plan” (instruction set) of our simple computer. Note that we will initially define six of the eight possible instructions afforded by our 3-bit opcode field. We will save the last two opcode bit patterns to define some extensions to our instruction set later in this chapter. Our simple computer’s instruction set is given in Table 2-1.

**Table 2-1** Simple computer instruction set.

<b>Opcode</b>	<b>Mnemonic</b>	<b>Function Performed</b>
0 0 0	LDA <i>addr</i>	Load A with contents of location <i>addr</i>
0 0 1	STA <i>addr</i>	Store contents of A at location <i>addr</i>
0 1 0	ADD <i>addr</i>	Add contents of <i>addr</i> to contents of A
0 1 1	SUB <i>addr</i>	Subtract contents of <i>addr</i> from contents of A
1 0 0	AND <i>addr</i>	AND contents of <i>addr</i> with contents of A
1 0 1	HLT	Halt – Stop, discontinue execution

The first two instructions, “LDA” and “STA”, are examples of data transfer group instructions. As their assembly mnemonics imply, these instructions transfer data between the “A” register (accumulator) and memory. For the “load A” (LDA) instruction, the source of the data is memory location *addr*, and the destination is the “A” register. For the “store A” (STA) instruction, it is just the opposite: here, *addr* indicates the location in memory where the value in A (also referred to as the *contents of A*) is to be stored. As it turns out, “load” and “store” instructions are the “most popular” instructions in any machine’s instruction set, often comprising as much as 30% of the compiled code for typical applications.

A “shorthand” notation we will use throughout the remainder of this text is the use of parenthesis to indicate “the contents of” a particular register or memory location. This allows us to describe what an LDA instruction does as simply “ $(A) \leftarrow (addr)$ ” and what an STA does as “ $(addr) \leftarrow (A)$ ”. An important point to note in both cases is that the *source* of the data transfer – i.e., *(addr)* for LDA and *(A)* for STA – *does not change* (or, *is unaffected*) as a result of the instruction execution.

Continuing down the list of available instructions, we next find two arithmetic group instructions: ADD and SUB. The ADD instruction performs the operation  $(A) \leftarrow (A) + (addr)$  using radix (or two’s complement) arithmetic, and sets the condition code bits based on the result obtained. (Details on radix arithmetic and condition codes can be found in the review material presented in Chapter 1.) The SUB instruction performs the operation  $(A) \leftarrow (A) - (addr)$  and sets the condition code bits accordingly. Recall that there is an important difference regarding how the carry flag (CF) is affected in an addition versus a subtraction. Following an ADD, the carry flag is the carry out of the most significant (or *sign*) position; whereas following a SUB, the carry flag is the complement of the carry out of the sign position (based on its interpretation as a *borrow*). Because of this difference between ADD and SUB, the CF bit is sometimes referred to as the “carry/borrow” flag – which is the way we will formally refer to it. If what we just described seems a bit “fuzzy”, now would be a good time to review the material in Chapter 1.

Moving down the chart, we find that our next instruction, AND, is from the logical (or “Boolean”) group. Because logical group instructions perform bit-wise operations, they are sometimes referred to as *bit manipulation* instructions. At minimum, most microprocessors worth their silicon generally have at least three Boolean instructions: AND,

*data transfer  
group instructions  
assembly  
mnemonics*

*LDA  
STA*

*arithmetic group  
instructions*

*ADD  
SUB*

*two's complement  
arithmetic*

*carry/borrow  
flag*

*logical group  
instructions  
bit manipulation  
instructions*

OR, and NOT (many also include XOR). Our simple computer, however, will just implement the first of these operations, which can be described using the notation  $(A) \leftarrow (A) \cap (\text{addr})$ , where the “ $\cap$ ” symbol is used to denote the bit-wise logical AND of the two operands to produce the corresponding result bits.

*AND  
OR  
NOT  
XOR*

No instruction set would be complete without a way to stop the machine. Our sixth (and final, for now) instruction, HLT (for “halt”) serves this purpose. The HLT instruction is an example of a *machine control group* instruction. Execution of the HLT instruction will “freeze” the machine at its current point in the program being executed, and prevent the machine from fetching or executing any additional instructions until it is restarted (by pressing the START pushbutton described previously).

*HLT  
machine control  
group instructions*

## 2.4 Simple Computer Programming Example

To better understand how our simple computer operates, we will “walk through” the execution of a short program. This program will exercise each instruction in our simple computer’s repertoire. An important point to consider before proceeding is that it would be rather difficult to design a “simple” computer that directly interprets the instruction mnemonics (i.e., LDA, STA, etc.) we have defined. Rather, it is much easier to design a machine that directly interprets bit patterns (0’s and 1’s) that represent these instructions. This means that, before we can place our program in memory, we must translate the instruction mnemonics into bit patterns (“code”) the machine understands, called *machine code*. This translation process is called *assembly*, since machine code is created directly (“assembled”) based on instruction mnemonics. As one might guess, instruction mnemonics are typically referred to as *assembly level mnemonics*, or simply *assembly language*. A software program that translates assembly level mnemonics into machine code is called an *assembler*. If one is unfortunate enough to perform the translation by hand, the process is called *hand assembly*.

*machine code  
assembly  
language  
hand assembly*

Fortunately, most computer programming is done at a higher level of abstraction, using *high-level languages* such as “C”. Here, a *compiler* program is used to translate code written in high-level language into assembly code. An assembler program is then used to translate the compiler’s output into machine code for the target processor. We will find, though, that a firm grasp of assembly language programming techniques is essential for effectively utilizing the resources integrated

*high-level  
language  
compiler*

into a modern microcontroller. Once we master assembly-level programming, we'll consider how to program a microcontroller using "C". But to get there, we need to start at the "basic bit" level – so let's return to the illustrative simple computer program in Table 2-2.

**Table 2-2** Programming example.

Addr	Instruction	Comments
00000	LDA 01011	Load A with contents of location 01011
00001	ADD 01100	Add contents of location 01100 to A
00010	STA 01101	Store contents of A at location 01101
00011	LDA 01011	Load A with contents of location 01011
00100	AND 01100	AND contents of 01100 with contents of A
00101	STA 01110	Store contents of A at location 01110
00110	LDA 01011	Load A with contents of location 01011
00111	SUB 01100	Subtract contents of location 01100 from A
01000	STA 01111	Store contents of A at location 01111
01001	HLT	Stop – discontinue execution

One of the first things we need to know is *where* in memory our program needs to be located. The logical thing to do is place our program at the *beginning* of memory, i.e., starting at location  $00000_2$ . We can then design the circuitry that, after the START pushbutton is pressed, begins fetching instructions from memory at location  $00000_2$ . Recalling that instructions are of fixed length (8 bits) and that memory locations are 8-bits wide, we realize that consecutive instructions will occupy consecutive memory locations. We can then imagine a "pointer" that tells us which instruction is to be executed, and that gets incremented after each instruction is fetched. Such a pointer is typically referred to as either an *instruction pointer* or a *program counter*.

*instruction pointer*  
*program counter*

A "snapshot" of what our short program looks like in memory prior to execution is provided in Figure 2-5 (just the "first half" of memory, from locations  $00000_2$  to  $01111_2$  is shown). The lightly shaded part corresponds to the assembled machine code. Referring back to Table 2-2, note that the first instruction (at address  $00000_2$ ) is *load accumulator* (LDA) with the contents of memory location  $01011_2$ . Since the 3-bit opcode for LDA is "000", this instruction is encoded as the bit pattern "000 01011" in memory. Stated another way, the instruction "LDA 01011" has been *assembled* into the machine code "000 01011". We could go through a similar "hand assembly" process for the rest of the instructions that comprise the program, up to and

including the HLT instruction at location  $01001_2$  (note that the address field of this instruction is not used, and is shown here to be “00000”).

Location	Contents
00000	00001011
00001	01001100
00010	00101101
00011	000001011
00100	10001100
00101	00101110
00110	000001011
00111	01101100
01000	00101111
01001	10100000
01010	
01011	10101010
01100	01010101
01101	
01110	
01111	

### Beam in the Bits, Scotty!

One important detail we will ignore for the moment is how these bit patterns get loaded into memory. In a later chapter, we'll discuss how to write what's called a "loader" program, which – as its name implies – does just that. For now, assume Scotty (of *Star Trek* fame, for those of you much younger than the author) has used a molecular beam transporter to "beam the bits" into memory.

**Figure 2-5** Memory snapshot prior to program execution.

The operands used by each arithmetic (ADD, SUB) or logical (AND) operation will be stored at locations  $01011_2$  and  $01100_2$  (in the darker shaded area of Figure 2-5); note that we have initialized these two locations to *arbitrarily chosen values*. The results of each operation (ADD, AND, SUB) will be stored in three consecutive locations, starting at location  $01101_2$ . Note that our computer's memory will contain a mix of instructions and data (operands and results).

### No Stopping It Now

What happens if the HLT instruction is omitted? Perhaps even worse than “not stopping”, the computer will start *executing data*, which, as one might imagine, is not a pretty sight (or, stated less formally, causes “bits to fly all over the place”) and, at best, leads to *very strange* program behavior. Any “honest” programmer (not to be confused with an honest politician), however, will confess that he/she has inadvertently done this “at least once...”

*executing data*

*honest  
programmer*

Given that our computer only understands 0's and 1's rather than the more human-friendly assembly mnemonics, the question that begs is: “How is our computer able to distinguish between instructions and data?” The hopefully obvious answer is: “It can't!” Rather, it has to be

told which locations contain instructions and which contain data. The convention we will use to make this distinction is that our programs will always start at location  $00000_2$  and continue until they reach a “halt” (HLT) instruction; any locations following the HLT instruction may be used for data (operands or results).

Location	Contents
00000	00001011
00001	01001100
00010	00101101
00011	00001011
00100	10001100
00101	00101110
00110	00001011
00111	01101100
01000	00101111
01001	10100000
01010	
01011	10101010
01100	01010101
01101	11111111
01110	
01111	

**Add:**

$$\begin{array}{r}
 10101010 \\
 +01010101 \\
 \hline
 11111111
 \end{array}$$

CF = 0

NF = 1

VF = 0

ZF = 0

← Add

**Figure 2-6** Result after executing the first three instructions.

We are now ready to step through the execution of this program. Referring back to Table 2-2, we see that the purpose of the first three instructions is to add the two operands (at locations  $01011_2$  and  $01100_2$ , respectively) and store the result at location  $01101_2$ . As illustrated in Figure 2-6, the result obtained will be  $11111111_2$  (recall that this is the 8-bit representation for “-1” in two’s complement notation). Also, the negative flag (NF) will be set to “1”, the carry flag (CF) will be cleared to “0”, the overflow flag (VF) will be cleared to “0”, and the zero flag (ZF) will be cleared to “0”.

### Self-Perpetrating Programs

It is entirely possible to contrive a program that writes data into locations that contain instructions yet to be executed. The name “self-modifying code” has been used to describe such a creation. A self-modifying program, as one might guess, could prove to be excruciatingly difficult to debug. In a word, don’t try this at home! (And, don’t try to convince your boss that you’ve invented a new way to write “interesting” programs!).

*self-modifying  
code*

Again referring back to Table 2-2, we see that the purpose of the next three instructions is to logically AND the two operands and store the result at location  $01110_2$ . Note that, for the AND operation, the carry flag (CF) and overflow flag (VF) are meaningless, and therefore should be *unaffected* by the execution of the AND instruction. The result obtained, however, may be negative (in a two's complement sense) or zero, so the negative flag (NF) and zero flag (ZF) should be affected. A snapshot of memory following execution of the three AND-related instructions is provided in Figure 2-7. Note that, since the result obtained is  $00000000_2$ , the zero flag is set to "1".

Location	Contents
00000	00001011
00001	01001100
00010	00101101
00011	00001011
00100	10001100
00101	00101110
00110	00001011
00111	01101100
01000	00101111
01001	10100000
01010	
01011	10101010
01100	01010101
01101	11111111
01110	00000000
01111	

AND:

$$\begin{array}{r}
 10101010 \\
 \underline{\text{C}01010101} \\
 00000000
 \end{array}$$

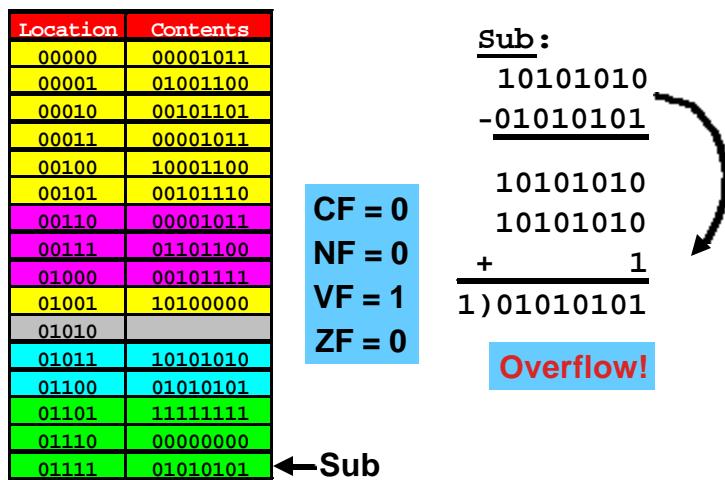
**CF = <unaffected>**  
**NF = 0**  
**VF = <unaffected>**  
**ZF = 1**

← AND

Figure 2-7 Result after executing the “middle” three instructions.

The purpose of the next group of three instructions is to take the difference of the two operands at locations  $01011_2$  and  $01100_2$ . Specifically, we are going to subtract (SUB) the operand at location  $01100_2$  from the operand at location  $01011_2$ , and place the result at location  $01111_2$ . Recall from Chapter 1 that a radix subtraction is realized by forming the two's complement of the subtrahend (here, the operand at location  $01100_2$ ) and adding it to the minuend (the operand at location  $01011_2$ ). Further, the easiest way to generate the radix complement of a signed number is to add one to its *diminished radix complement* (or *ones' complement*). Figure 2-8 shows what happens. Note that, while the result  $01010101_2$  will be stored at location  $01111_2$ , it will be invalid because overflow has occurred (denoted by VF set to "1"). Note also that CF (the carry/borrow flag) is cleared to "0" due to its interpretation here as a *borrow flag* – recall that, following a subtract operation, CF is set to the *complement* of the carry out of the sign position (which in this case was "1"). A borrow flag of "0" following a

subtract operation essentially means that “no borrow is propagated forward.”



**Figure 2-8** Result after executing the last group of three instructions.

### Bumbling Borrows

Perhaps the single-most issue that causes students consternation is that of the carry/borrow flag. The interpretation of a “carry propagated forward” following an addition is no problem; but when it gets to subtraction, all “bits are off” (pardon the very bad pun). Here, the proper interpretation is as a “borrow propagated forward” to the next-most significant group of digits in an extended precision subtraction. The borrow flag (still called CF), when set, is basically telling that next group of digits to “reduce its result by one” because the previous stage “has borrowed from it.” The best real-world analogy that comes to mind is that of a statement from your friendly, local banking institution listing the service charge they have extracted from your account for the privilege of serving you. The point is: since they have already taken the money, you need to adjust your idea of how much money you have left!

Before we leave this last block of code, yet another question that comes to mind is: “How should error conditions like overflow be handled?” As one might guess, we will need some “new” instructions that allow us to test the state of the various condition codes (here, VF) and transfer control to a different part of the program (typically called an “exception handler”) if an error has occurred. Before we finish this chapter, we will learn how to implement such “conditional transfer of control” instructions.

The final instruction in our short program, HLT, simply tells our computer to “stop executing”. Once the program has stopped, we could presumably look at the contents of each location to determine the results of the program execution. What we should find is the memory image depicted in Figure 2-8 (note that memory location  $01010_2$  was unused by our example program and may contain a “random” value).

## 2.5 Simple Computer Block Diagram

Now that we know *how* our simple computer works, we are ready to consider the functional blocks necessary to *make* it work. Basically we want to build what appears to be a “big state machine” that performs the calculations just done by hand. At a fundamental level, there are two basic steps associated with the processing of each instruction. The first step is to read the instruction from memory, called an *instruction fetch cycle*. The second step is to extract the opcode and address fields from the instruction just fetched and perform the operation specified by the opcode on the data located at the specified address; this step is referred to as an *instruction execute cycle*.

*instruction  
fetch cycle*

*instruction  
execute cycle*

*memory unit*

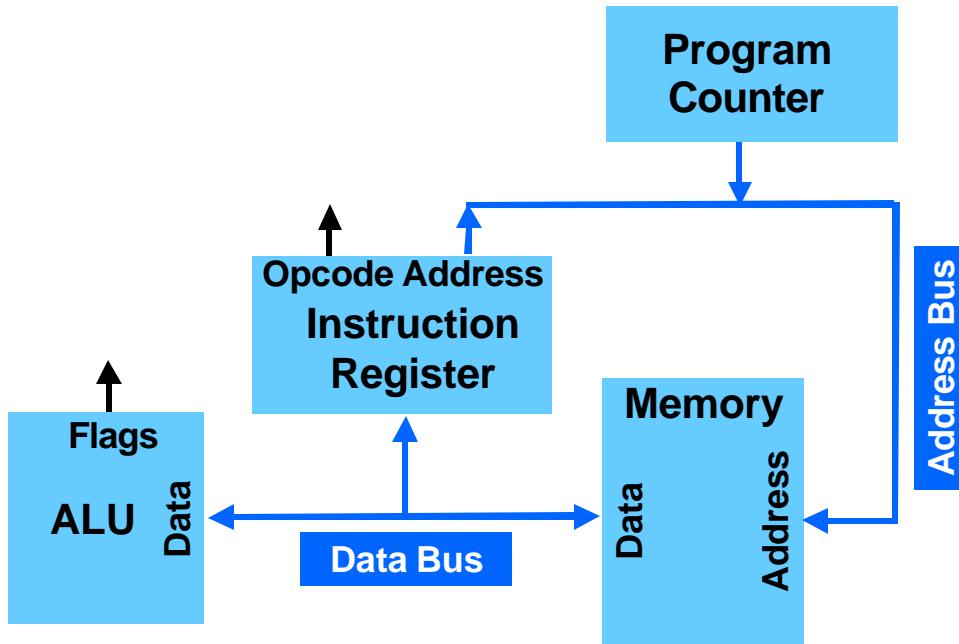
What are the basic functional blocks, then, that are necessary to implement the simple computer described here? Clearly, a memory unit – for storing instructions and data – is one of the major functional blocks necessary. This memory unit needs to be capable of reading the contents of a specified location (indicated on its address lines) as well as writing a new value to a specified location.

Another major functional block needed is one that will keep track of which instruction is next in line to be executed. In our simple computer, the instructions are stored in consecutive memory locations, starting at location  $00000_2$ . What is needed is a pointer that keeps track of which instruction is next. Because this block is nothing more than a binary counter, we will call it the *program counter* (PC).

*program counter  
PC*

Once it is fetched from memory, a place is needed to temporarily “stage” an instruction while the opcode field is decoded and the address field is extracted. We can think of this block as a place to hold the instruction just fetched while it is being “digested”. While more creative, biologically inspired names for it are certainly possible, we will simply call this functional block the *instruction register* (IR).

*instruction register  
IR*



**Figure 2-9** Simple computer core block diagram.

Next we realize the need for a functional block that performs the arithmetic and logical operations we have defined in the simple computer's instruction set. Not surprisingly, this block is usually called an *arithmetic logic unit*, or simply ALU. Note that the accumulator ("A" register) and condition code bits (CF, NF, VF, ZF) are part of the ALU.

*arithmetic logic unit  
ALU*

Finally, we realize that our simple computer needs a "manager" – a functional block that orchestrates the activities of all the other functional blocks delineated above. This "manager" is responsible for indicating whether a fetch or an execute cycle is to be performed and, once an instruction is fetched, for decoding the opcode field of that instruction and telling the other blocks in the system what to do in order to execute it. Because our simple computer's "manager" controls the sequencing of events that, taken together, constitute the completion of a machine instruction, we often refer to the state machine part of the manager's personality as a *micro-sequencer* (similar to, perhaps, but not to be confused with a "micro-manager"). And because decoding the opcode field of the instruction is an essential part of the sequencing process, we award our simple computer's manager the grand and glorious name: *instruction decoder and micro-sequencer* (IDMS). This more extravagant sounding name helps prevent images of "kicking bits around" that might be associated with a "manager" (think baseball).

*manager*

*micro-sequencer*

*IDMS*

Returning to the “house” analogy for a moment, what we have just done is “define the rooms” of the “structure” (or system) we wish to build. What we have not yet done, however, is interconnect the functional blocks into a working “floor plan”. In order to do this, we need an understanding of the “traffic patterns” (here, of address, data, and control information) that need to flow among the various functional blocks.

Starting with the memory unit, we note that a series of address lines tell which location is being accessed; the collection of address lines is referred to as the *address bus*. (Recall that a *bus* is a set of signal lines that have a *common purpose*.) At the location in memory accessed, data can be *read* (output) or *written* (input); the memory’s data lines (and the associated data bus) must therefore be *bi-directional*. Further, control signals need to be supplied to the memory unit that tell whether or not it is *enabled* to respond (or *selected*), and, if enabled to respond, whether it should perform a read operation or a write operation.

*address bus*

*bi-directional*

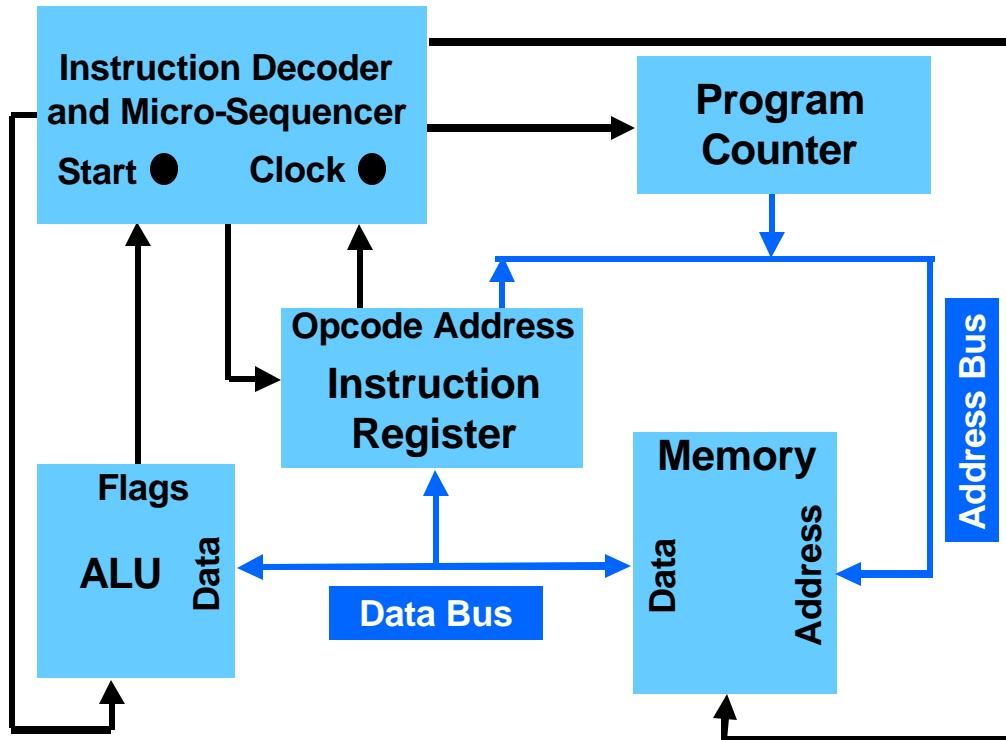
Next, we realize that the program counter (PC) will supply the instruction address to memory during a fetch cycle, and that the instruction register (IR) will be used to temporarily stage the instruction after it has been read from memory. Further, on an execute cycle, the IR will supply the operand address to memory, and the destination (or source) of the data in this transaction is the “A” register of the ALU. Thus, there are two potential sources of address information – the PC and the IR – on the address bus. Since only one device can “talk” on the bus at a given instant in time, we will need to provide each of these functional blocks with *three-state output capability* – and it will be our “manager’s” job to keep them from talking at the same time!

*three-state output  
capability*

Further, there are two potential destinations of data read from memory. On a fetch cycle, an instruction destined for the IR is read from memory. On an execute cycle, an operand destined for the ALU is read from memory (alternately, data in the ALU is destined for memory if an STA instruction is being executed). Again, we note the need for three-state buffers in all the functional blocks involved with driving the data bus.

Putting this all together, the “core” of our simple computer is depicted in Figure 2-9. Left on their own, however, these functional blocks are incapable of doing anything “intelligent”, let alone successfully executing instructions. Hence the need for a “manager” – the instruction decoder and micro-sequencer – to tell each block what to

do when. As such, the IDMS can aptly be thought of as the “heart” of the machine. The simple computer augmented with an IDMS is shown in Figure 2-10.



**Figure 2-10** Complete simple computer block diagram.

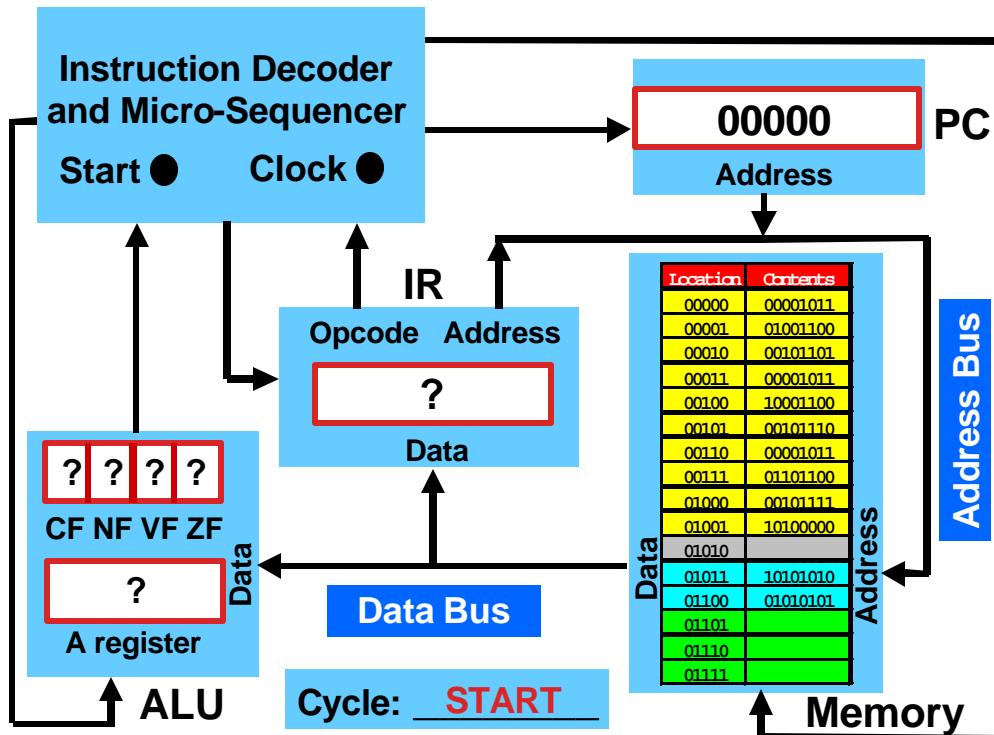
We now have a complete “floor plan” for our “house”, that we have specified in a top-down fashion. Before actually building it, though, let’s make sure we understand how the “rooms” work together.

## 2.6 Instruction Execution Tracing

To get a better idea of how the various functional blocks of our simple computer work in concert to process instructions, we will return to our short program of Table 2-2 and use a technique called *instruction tracing* to help us visualize the flow of information. On a cycle-by-cycle basis, we will examine the address and data paths as well as the bit patterns in each register for the first three instructions of this short program. Recall that we used the term “micro-sequencer” because there is a sequence of events associated with processing an instruction: here, a fetch cycle followed by an execute cycle.

*instruction  
tracing*

The instruction trace worksheet in Figure 2-11 sets the stage for this exercise, which shows the initial state of the machine after START is pressed. Note that there are several things we will keep track of as our machine executes the program. In particular, we will be monitoring what happens to the PC, IR, and “A” register as well as the contents of memory. We will also practice naming each cycle as it occurs.



**Figure 2-11** Instruction trace worksheet for machine state after START is pressed, prior to first fetch cycle.

Recall that pressing the START pushbutton places the machine in a known initial state: the PC is reset to “00000” and the state counter (in the IDMS) is set to “fetch”. Note that the initial state of the IR and ALU may be “random” and that memory is initialized to the values indicated (although at this point we “don’t care” what is in the unused location 01010<sub>2</sub> or the locations where the results will be stored, 01101<sub>2</sub>–01111<sub>2</sub>).

During the first fetch cycle, shown in Figure 2-12, the instruction at memory location 00000<sub>2</sub> is read and placed in the IR. As the IR is being loaded with the instruction, the PC is incremented by one (i.e., once the fetch of the current cycle is complete, the PC is pointing to the *next* instruction to execute). Note that the values in each register are those obtained *after* the “fetch LDA” cycle is *complete*.

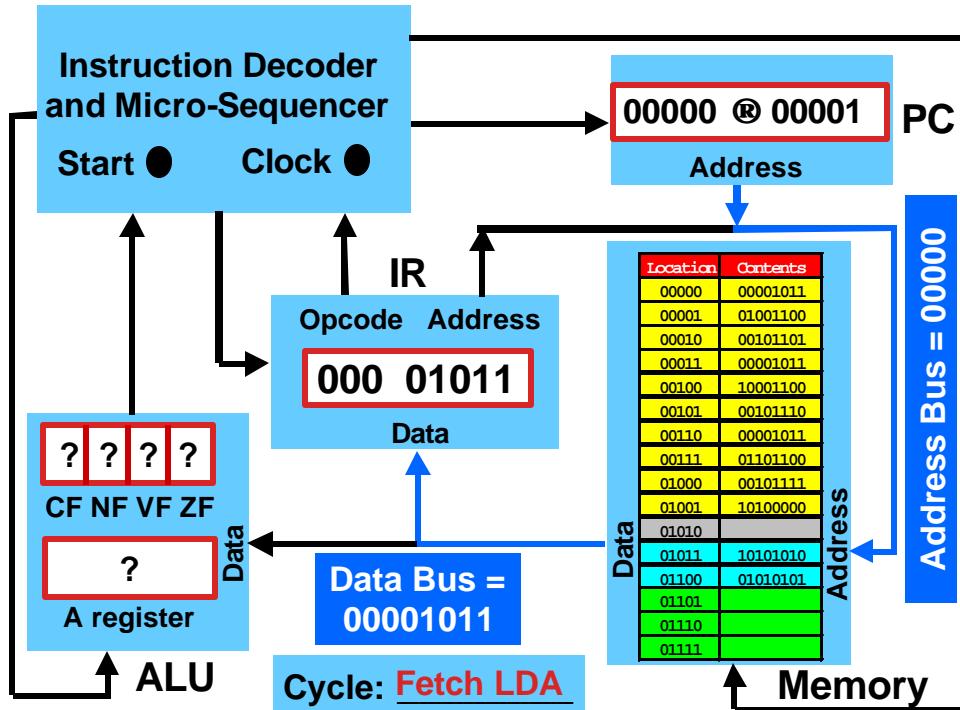


Figure 2-12 Instruction trace worksheet for first fetch cycle.

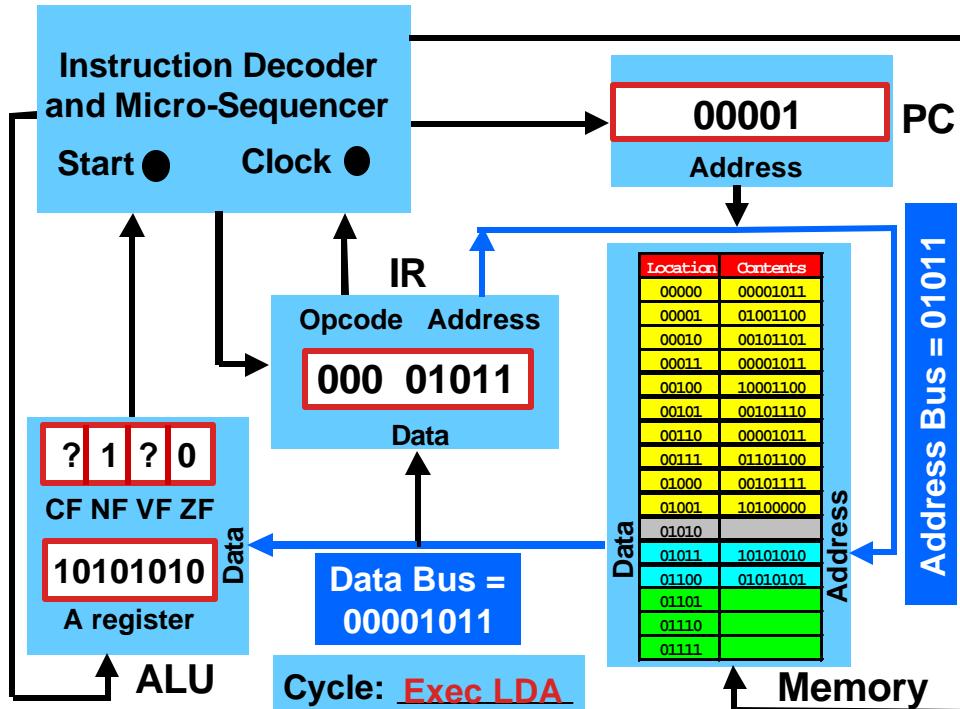


Figure 2-13 Instruction trace worksheet for first execute cycle.

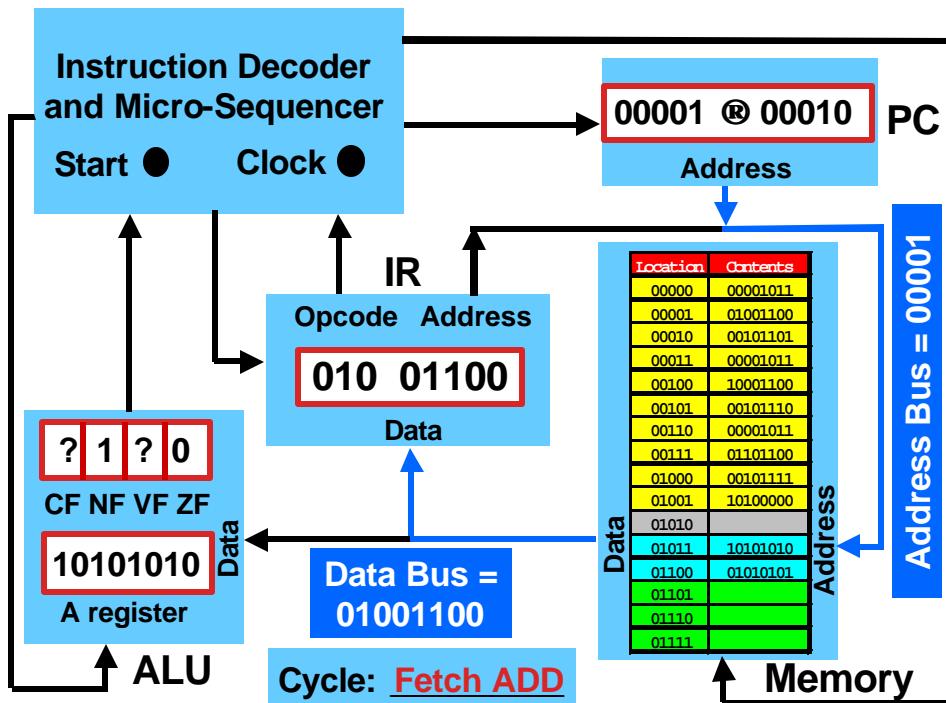


Figure 2-14 Instruction trace worksheet for second fetch cycle.

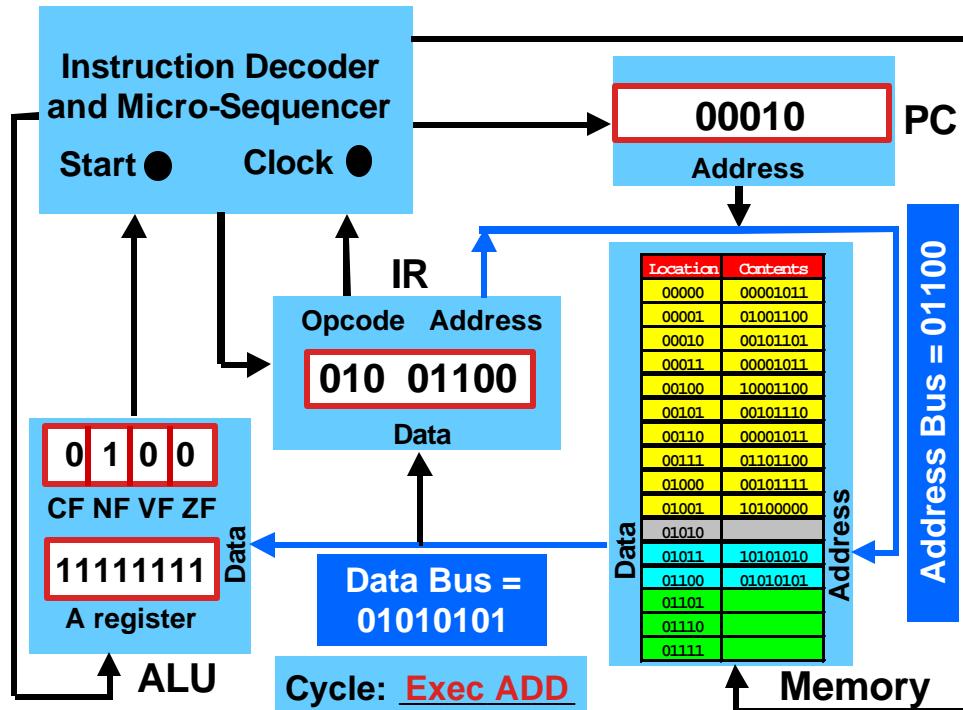


Figure 2-15 Instruction trace worksheet for second execute cycle.

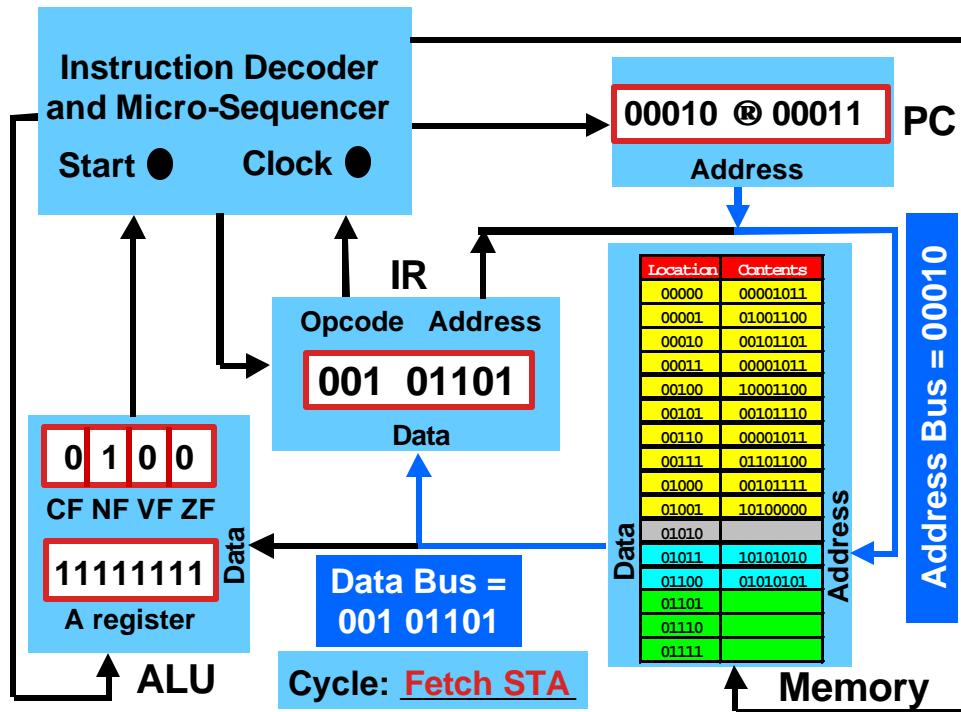


Figure 2-16 Instruction trace worksheet for third fetch cycle.

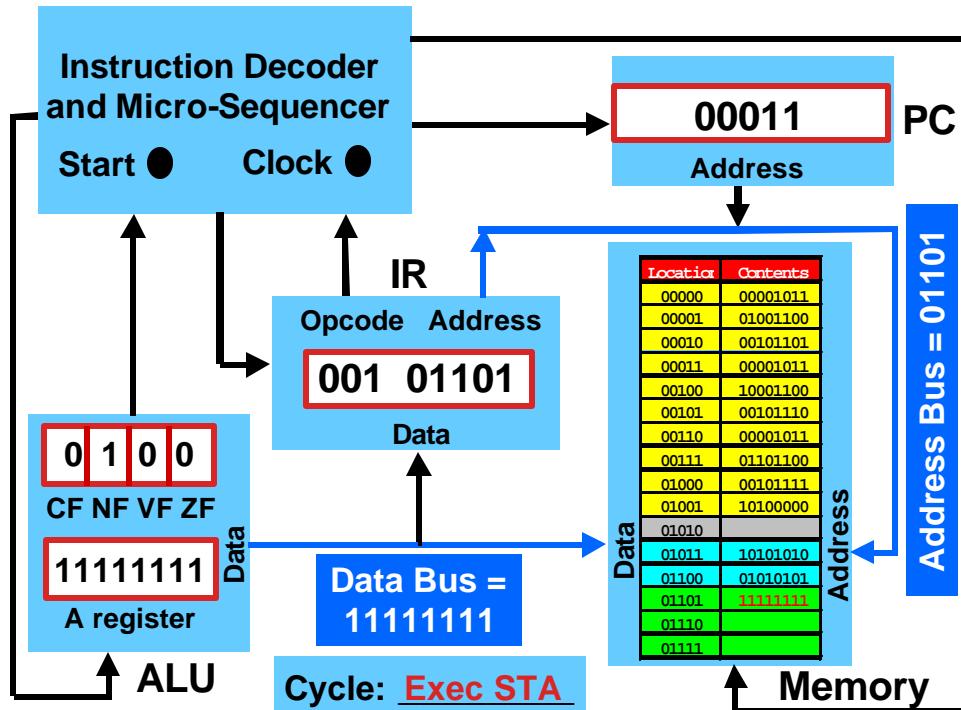


Figure 2-17 Instruction trace worksheet for third execute cycle.

During the first execute cycle, shown in Figure 2-13, the “LDA 01011” instruction in the IR is executed. When this cycle is complete, the “A” register contains the contents of memory location  $01011_2$ , i.e., the value  $10101010_2$ . Note also that the NF is set to “1” and ZF is cleared to “0”. The “execute LDA” cycle does not, however, affect the contents of any memory location, nor does it change the contents of IR or PC (condition code bits CF and VF are also unaffected).

We are now ready for the second fetch cycle (“fetch ADD”), shown in Figure 2-14. Here, the instruction at memory location  $00001_2$  is fetched and placed into the IR, and as that occurs, the value in the PC is incremented by one. The results of executing the ADD instruction are shown in Figure 2-15. Here, the contents of memory location  $01100_2$  (i.e., the value  $01010101_2$ ) are added to the value previously loaded into the “A” register. A result of  $11111111_2$  is obtained, along with condition code bits CF = “0”, NF = “1”, ZF = “0”, and VF = “0”.

This brings us to the third fetch cycle (“fetch STA”) of our tracing example, shown in Figure 2-16. Here, the instruction at memory location  $00010_2$  is fetched and placed into the IR, and as that occurs, the value in the PC is incremented by one. The results of executing the STA instruction are shown in Figure 2-17. Here, the contents of the “A” register are stored at the memory location indicated in the instruction’s address field:  $01101_2$ . When the “execute STA” cycle is complete, then, memory location  $01101_2$  contains the value  $11111111_2$ . Note, however, that the “A” register as well as the condition code bits are unchanged.

Several observations are in order. First, all of our simple computer’s fetch cycles are identical (i.e., they are independent of the instruction opcode). In fact, this *has* to be the case, since our machine basically knows nothing about the instruction being fetched until it is placed in the IR. Second, it may appear “strange” that our simple computer is incrementing the value in the PC on the same cycle that it is being used as a pointer to memory. Another way to say this is that the increment of PC is *overlapped* with the fetch of the instruction. The reason this can happen will become apparent when we start implementing each functional block in the next section. For now, though, suffice it to say that because each register will be implemented using edge-triggered flip-flops, the same clock edge that causes the IR to load the instruction being fetched also causes the PC to increment. The IR, though, will be loaded with the value on the data bus *prior to* the clock edge, while the value output by the PC (driving the address

*overlapped*

bus) will change *after* the clock edge – thus facilitating the desired overlap. This is an important point that we will revisit several times before the end of this chapter.

One final suggestion before we move to the “bottom-up” phase of our simple computer design process. Practice the “instruction tracing” process outlined in this section on other code segments to become more familiar with “what happens when” as each instruction is fetched and executed. As we say in the education industry, this is a “good test question” (GTQ)!

*good test  
question*

## 2.7 Bottom-Up Implementation of Simple Computer

Armed with a thorough understanding of how our simple computer works, we are now ready to start building it from the bottom-up. In practice, the preferred approach is to implement and test each block as it is designed. Then, when we put the various functional blocks together, we have a much better chance of the entire system working “the first time”.

### 2.7.1 Memory

The block we will start with is memory. Although most of the time we would simply choose a “memory chip” of appropriate size and speed, a knowledge of “what’s under the hood” is essential to understanding how the various functional blocks of our simple computer work together.

First, some terminology. Normally, we think of memory as an entity that, from the computer’s perspective, can be “read” or “written”. In “read” mode, the memory unit simply outputs, on its data bus lines, the contents of the location indicated on its address bus inputs. In “write” mode, the memory unit stores the bit pattern present on its data bus lines at the location indicated on its address bus inputs. The correct acronym to describe such a “read/write memory” is RWM. Despite valiant efforts, the name RWM never caught on. Instead, it is more popular to refer to these devices as “random access memories” or RAMs – so-named because any (random) location can be accessed in the same amount of time (not because something random is read after a given value is written).

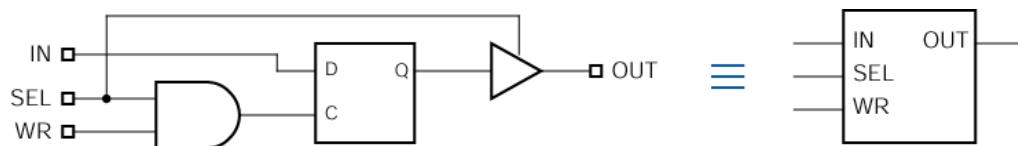
The specific type of RAM we wish to concentrate on here is *static* RAM, or SRAM. This is in contrast to *dynamic* RAM (DRAM), which

*static RAM (SRAM)  
dynamic ram (DRAM)*

requires constant refreshing to retain information. (In DRAM, data is stored as a charge on a capacitor – since the charge dissipates over time, it must be periodically refreshed.) SRAM consists of a collection of D latches that will retain data (without the need for refreshing) as long as power is applied. Once power is turned off, however, all information previously stored in the SRAM is lost (this is referred to as a *volatile* memory).

In addition to address and data bus connections (where, for our simple computer, the address bus is 5bits wide and the data bus is 8-bits wide), an SRAM needs three control signals. First, an SRAM needs an overall enable, typically called a “chip select” (CS) or “chip enable” (CE). This enable signal is needed to differentiate among multiple SRAMs or, as we will see later in this chapter, between memory and input/output devices. Second, an SRAM needs an output enable (OE) signal which, provided the SRAM is selected, turns on a series of three-state buffers that drive the data from the addressed location out onto the data bus. Finally, an SRAM needs a write enable (WE) signal which, if the SRAM is selected, opens the row of latches associated with the addressed location and allows it to take on the value presented to the SRAM on the data bus.

The basic building block of an SRAM is a memory cell, such as the one depicted in Figure 2-18, consisting of a D-latch and a three-state buffer. When the *select* (SEL) signal is asserted, the three-state buffer is enabled, placing the data stored in the latch on the cell’s OUT line. When both SEL and WR are asserted, the latch opens and accepts the data present on the IN line (by virtue of asserting the latch enable or “C” input of the D-latch). When WR is negated, the latch closes and retains the new value.



**Figure 2-18** SRAM cell (adapted from Wakerly).

A complete SRAM can be constructed by combining an array of memory cells with a (large) decoder plus some additional logic. The internal structure of an eight location, 4bit wide (or, “8x4”) SRAM is shown in Figure 2-19. Note that the number of address lines needed is  $\log_2(\text{number\_of\_locations})$ ; here,  $\log_2(8) = 3$ . Stated another way, the number of locations in an SRAM is  $2^n$ , where  $n$  is the number of

*volatile  
memory*

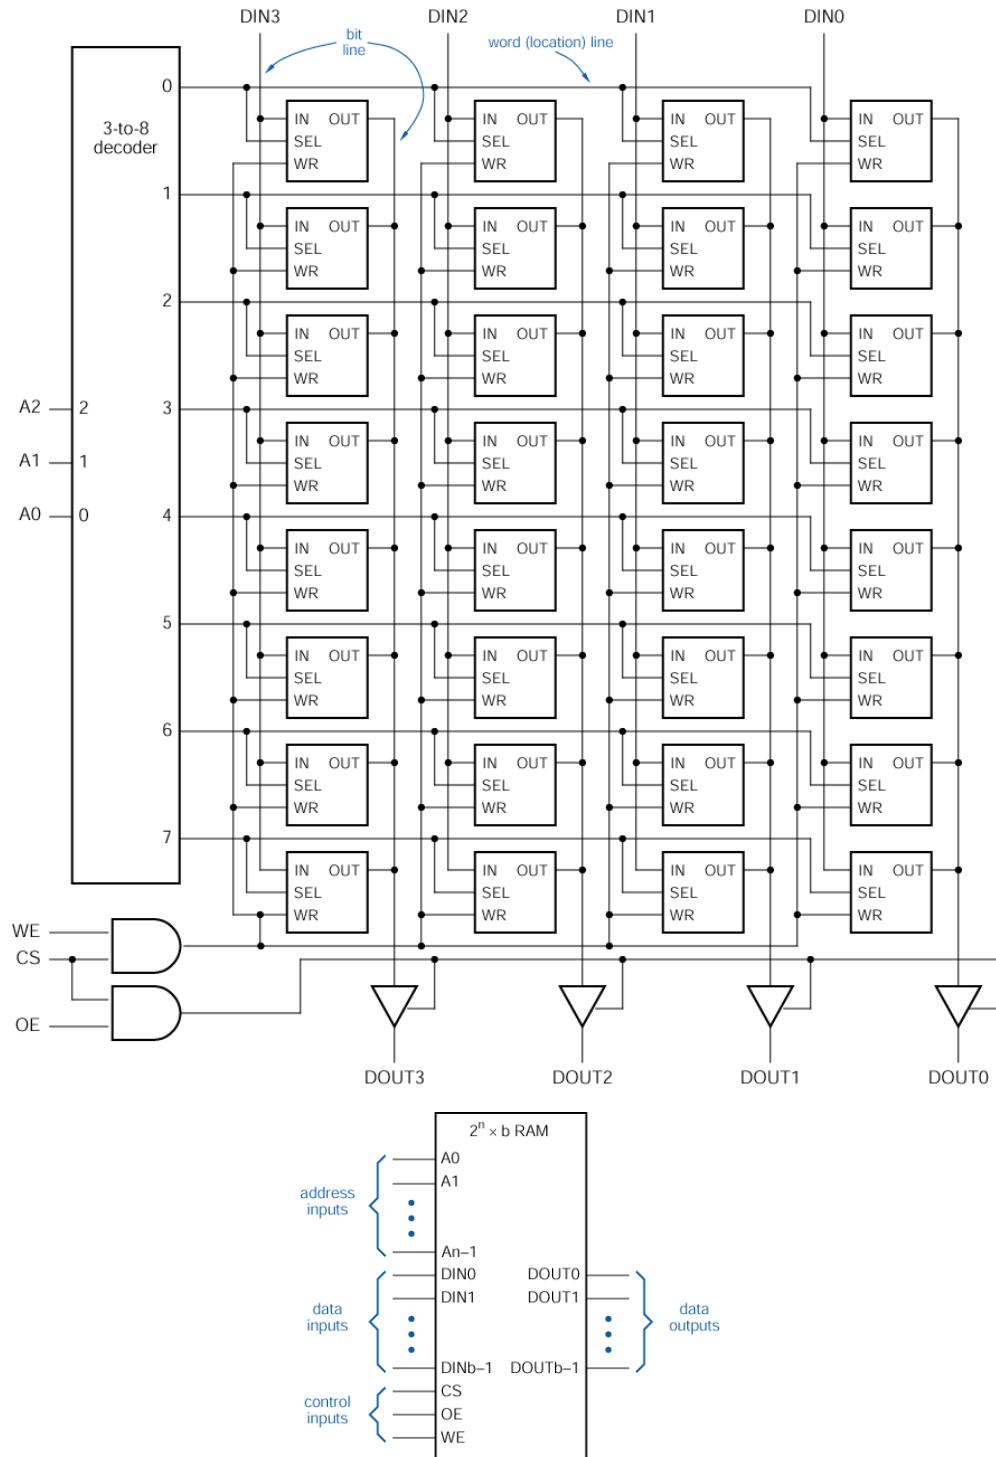
*chip select  
(CS)*

*output enable  
(OE)*

*write enable  
(WE)*

address lines. A “location” in the SRAM corresponds to a row of memory cells; to select a particular row, an  $n$ -to- $2^n$  binary decoder is needed.

*memory location*



**Figure 2-19** SRAM internal structure and symbol (adapted from Wakerly).

### GigaBiga Dittos

The prefixes K (kilo-), M (mega-), G (giga-), and T (tera-), when referring to memory sizes, mean  $2^{10} = 1024$  ("about one thousand"),  $2^{20} = 1,048,576$  ("about one million"),  $2^{30} = 1,073,741,824$  ("about one billion"), and  $2^{40} = 1,099,511,627,776$  ("about one trillion"), respectively. This brings up a very important question: Does this means the feared "Y2K bug" is yet to occur (in year 2048)? An even more important question, though, might be: Instead of calling a billion bytes a "gigabyte", wouldn't a better name be "*bigabyte*" (as in *Biga* (short for "Bigger") *Bytes of Digital Wisdom*, the subtitle for this text?)

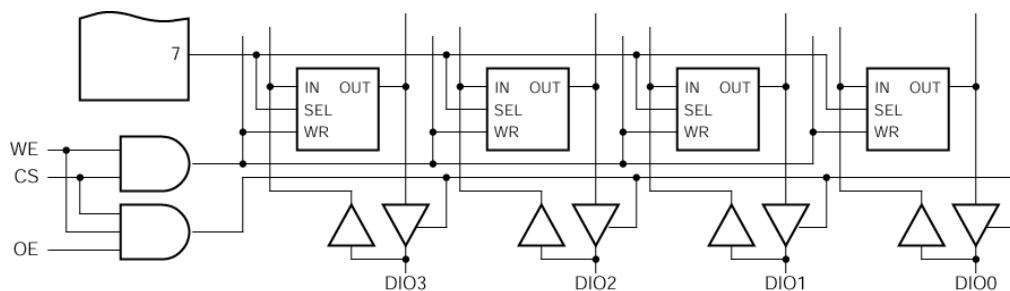
*kilo-, mega-,  
giga-, tera-*

*bigabyte*

In addition to a decoder, some logic is needed to "qualify" the actions associated with the OE and WE signals based on the assertion of CS (the overall chip enable). When WE is asserted in conjunction with CS, the data present on the DIN pins (DIN3 – DIN0) is written at the location specified on the address lines (note that the operation completes upon negation of the WE signal). When OE is asserted in conjunction with CS, the data output by a given row is routed to the three-state buffers that drive the external data lines.

Since the read and write operations are mutually exclusive, however, there is usually no need for separate data input and output lines. Instead, the data input and output lines are tied together and connected to the rest of the system using a *bi-directional* data bus. Such a configuration is shown in Figure 2-20. Note that an additional buffer is used to receive the incoming data during a write operation, to reduce the load seen by the entity driving the bus.

*bi-directional  
data bus*



**Figure 2-20** SRAM bi-directional data bus (adapted from Wakerly).

Before moving on, a few notes concerning memory timing are in order. Because an SRAM read operation is a purely combinational function, the *order* in which the address and control signals (CS and OE) are asserted is of no consequence. As we will see in Chapter 5, though, each of these signals represents a *critical timing path* with respect to receiving valid data from memory on a read cycle:  $t_{AA}$  is the address access (propagation delay) time,  $t_{CS}$  is the chip select access time, and  $t_{OE}$  is the output enable access time. When interfacing an SRAM to a computer, all of these “read” paths need to be analyzed.

*critical  
timing path*

$t_{AA}$   
 $t_{CS}$   
 $t_{OE}$

Since a “D” latch is used to store each bit of data in an SRAM, the *timing relationship* between the information on the address and data buses as well as the requisite control signals (CS and WE) is *more stringent* than for a read cycle. In particular, the address information needs to be stable, and the chip select (CS) needs to be asserted, for some time ( $t_{CW}$ ) before WE is asserted (opening the set of latches associated with the selected location). Also, the information supplied to the SRAM on the data bus must be stable  $t_{SETUP}$  *prior to* the negation of the WE signal, and  $t_{HOLD}$  *following* the negation of the WE signal. (These setup and hold timing parameters will be given specific names in Chapter 5.) The *consequence* of violating the data setup or hold timing specifications of an SRAM, or of not asserting the WE control signal for a sufficient period of time, is the *possibility of metastable behavior*. All of these “write”-related timing parameters need to be analyzed when interfacing an SRAM to a computer.

$t_{CW}$   
 $t_{SETUP}$   
 $t_{HOLD}$

*metastable  
behavior*

Returning to our simple computer, we note that by simply doubling the “width” of the SRAM depicted in Figure 2-19 (from 4-bits to 8-bits) and quadrupling the “length” (from 8 locations to 32 locations), as well as adding the bi-directional data bus interface shown in Figure 2-20, we will have the exact structure of SRAM needed. The only difference is the “unique” names we will use for our simple computer’s memory control signals: “MSL” for the memory select signal, “MOE” for the memory output enable, and “MWE” for the memory write enable.

*MSL*  
*MOE*  
*MWE*

## 2.7.2 Program Counter

The next functional block we wish to address is the program counter (PC). Basically, this is nothing more than a (5-bit) binary “up” counter with an asynchronous reset and three-state outputs. The asynchronous reset (ARS) will be connected to the START pushbutton, so that the first instruction fetched is from location  $00000_2$ . There are two other control signals needed: one that enables the PC to increment by one when a low-to-high (“positive edge”) of the system

*ARS*

CLOCK signal occurs, which we will call PCC; and one that turns on the three-state buffers that “gate” the value in the PC onto the address bus, which we will call POA. Note that if PCC is negated while a positive CLOCK edge occurs, the program counter should simply retain its current state.

PCC  
POA

To document the design of each functional block, we will present an ABEL (“Advanced Boolean Expression Language”) source file. Those unfamiliar with the ABEL language and source file format should review the material presented on this subject in Chapter 1. The ABEL source file for the program counter module is shown in Table 2-3.

ABEL

**Table 2-3** Program counter module.

```
MODULE pc

TITLE      'Program Counter Module'

DECLARATIONS

CLOCK pin;

PC0..PC4 pin istype 'reg_D,buffer';

PCC pin; " PC count enable
POA pin; " PC output on address bus tri-state enable
ARS pin; " asynchronous reset (connected to START)

EQUATIONS

"      retain state   count up by 1
PC0.d = !PCC&PC0.q # PCC&!PC0.q;
PC1.d = !PCC&PC1.q # PCC&(PC1.q $ PC0.q);
PC2.d = !PCC&PC2.q # PCC&(PC2.q $ (PC1.q&PC0.q));
PC3.d = !PCC&PC3.q # PCC&(PC3.q $ (PC2.q&PC1.q&PC0.q));
PC4.d = !PCC&PC4.q # PCC&(PC4.q $ (PC3.q&PC2.q&PC1.q&PC0.q));

[PC0..PC4].oe = POA;
[PC0..PC4].ar = ARS;
[PC0..PC4].clk = CLOCK;

END
```

Examining the source file, we see that when PCC is negated, the next state is simply the current state. When PCC is asserted, the equations for a synchronous 5-bit binary “up” counter determine the next state. Assertion of POA causes the three-state buffers associated with each

register bit to be enabled, and assertion of ARS causes each flip-flop comprising the PC to be asynchronously reset.

### 2.7.3 Instruction Register

The instruction register (IR) has a very simple mission: temporarily hold (“stage”) the instruction fetched from memory so that it can be “peeled apart” and executed. As such, it is simply a series of D flip-flops with two control signals. The first control signal, which we will call **IRL**, enables the instruction register to be loaded with the instruction read from memory; the load should occur on the positive edge of the system CLOCK. The second control signal, which we will call **IRA**, turns on the three-state buffers of the lower 5-bits of the IR, to “gate” the address field of the instruction onto the address bus.

**IRL**

**IRA**

**Table 2-4** Instruction register module.

```
MODULE ir

TITLE      'Instruction Register Module'

DECLARATIONS

CLOCK pin;

" IR4..IR0 connected to address bus
" IR7..IR5 supply opcode to IDMS

IR0..IR7 pin istype 'reg_D,buffer';
DB0..DB7 pin; " data bus

IRL pin; " IR load enable
IRA pin; " IR output on address bus enable

EQUATIONS

"           retain state          load
[IR0..IR7].d = !IRL&[IR0..IR7].q # IRL&[DB0..DB7];

[IR0..IR7].clk = CLOCK;
[IR0..IR4].oe = IRA;
[IR5..IR7].oe = [1,1,1];

END
```

Several items in the IR module source file, shown in Table 2-4, deserve explanation. First, when IRL is negated, note that the IR simply retains its current state. Second, note that, unlike the PC, there is no need to asynchronously reset the IR when the START pushbutton is pressed, since its (random) initial value is of no consequence. Finally, note that IRA only controls the three-state outputs associated with the lower 5-bits of the IR, and that the three-state buffers of the upper 3-bits (i.e., the opcode bits) are always enabled. The reason the three-state buffers associated with the upper 3-bits are always enabled is that they are connected *directly* to the IDMS module (i.e., they do not drive a bus). Recall that the IDMS uses the opcode bits to determine which system control signals are asserted on the next cycle, when the instruction is executed.

#### 2.7.4 Arithmetic Logic Unit

As mentioned earlier, the arithmetic logic unit (ALU) is so-named because it performs the arithmetic (add, subtract, etc.) and logical (“Boolean”) operations defined by the instruction set. A “real” ALU performs a wide range of arithmetic and logical functions on operands stored in either registers or in memory. Fortunately, our ALU is relatively simple: it performs four different functions on a single register (which we have called the accumulator, or “A” register) and sets four condition code bits (or flags) based on the result obtained. As such, only four control signals are needed: an overall enable, which we will call ALE; two “function select” lines, which we will call ALX and ALY; and a three-state output enable for “gating” the value in the “A” register onto the data bus, which we will call AOE. The data bus interface must be *bi-directional*, in order to input data supplied by memory on LDA, ADD, SUB, and AND operations; and to output data to memory for STA operations. The condition code bits (CF, NF, VF, ZF) are output directly to the IDMS (we will see how these flags can be used to implement conditional transfer of control instructions later).

*arithmetic and  
logical operations*

*ALE  
ALX  
ALY*

*condition  
code bits*

*nodes*

The ABEL source file for the simple computer ALU is shown in Tables 2-5, 2-6, and 2-7. Referring first to the declaration section (Tables 2-5 and 2-6), we note that signals used for “internal” purposes are declared as *nodes*. These include the carry bits and the combinational ALU outputs. In the declarations that continue in Table 2-6, the least significant bit carry-in (CIN) is defined as ALY. Noting that ALY is “0” for ADD and “1” for SUB, we realize this is exactly what is needed to add one to the diminished radix complement of the subtrahend (to obtain the radix complement) when performing a SUB operation.

**Table 2-5** Declarations section of ALU module.

```

MODULE alu

TITLE 'ALU Module'

" 8-bit, 4-function ALU with bi-directional data bus
"
" ADD:  (Q7..Q0) <- (Q7..Q0) + DB7..DB0
" SUB:  (Q7..Q0) <- (Q7..Q0) - DB7..DB0
" LDA:  (Q7..Q0) <- DB7..DB0
" AND:  (Q7..Q0) <- (Q7..Q0) & DB7..DB0
" OUT: Value in Q7..Q0 output on data bus DB7..DB0
"
" AOE  ALE  ALX  ALY  Function    CF   ZF   NF   VF
" ===  ===  ===  ===  ======  ==  ==  ==  ==
" 0    1    0    0    ADD        X    X    X    X
" 0    1    0    1    SUB        X    X    X    X
" 0    1    1    0    LDA        .    X    X    .
" 0    1    1    1    AND       .    X    X    .
" 1    0    d    d    OUT       .    .    .    .
" 0    0    d    d    <none>   .    .    .    .
"
" X -> flag affected  . -> flag not affected
"
" Note: If ALE = 0, the state of all register bits should be retained

DECLARATIONS

CLOCK pin;

" ALU control lines (enable & function select)
ALE pin; " overall ALU enable
AOE pin; " data bus tri-state output enable
ALX pin; " function select
ALY pin;

" Carry equations (declare as internal nodes)
CY0..CY7 node istype 'com';

" Combinational ALU outputs (D flip-flop inputs)
" Used for flag generation (declare as internal nodes)
ALU0..ALU7 node istype 'com';

" Bi-directional 8-bit data bus (also, accumulator register bits)
DB0..DB7 pin istype 'reg_d,buffer';

" Condition code register bits
CF pin istype 'reg_d,buffer'; " carry flag
VF pin istype 'reg_d,buffer'; " overflow flag
NF pin istype 'reg_d,buffer'; " negative flag
ZF pin istype 'reg_d,buffer'; " zero flag

```

**Table 2-6** Continuation of ALU source file declarations section.

```

" Declaration of intermediate equations

" Least significant bit carry-in (0 for ADD, 1 for SUB => ALY)
CIN = ALY;

" Intermediate equations for adder/subtractor SUM (S0..S7),
" selected when ALX = 0

S0 = DB0.q $ (DB0.pin $ ALY) $ CIN;
S1 = DB1.q $ (DB1.pin $ ALY) $ CY0;
S2 = DB2.q $ (DB2.pin $ ALY) $ CY1;
S3 = DB3.q $ (DB3.pin $ ALY) $ CY2;
S4 = DB4.q $ (DB4.pin $ ALY) $ CY3;
S5 = DB5.q $ (DB5.pin $ ALY) $ CY4;
S6 = DB6.q $ (DB6.pin $ ALY) $ CY5;
S7 = DB7.q $ (DB7.pin $ ALY) $ CY6;

" Intermediate equations for LOAD and AND,
" selected when ALX = 1

L0 = !ALY&DB0.pin # ALY&DB0.q&DB0.pin;
L1 = !ALY&DB1.pin # ALY&DB1.q&DB1.pin;
L2 = !ALY&DB2.pin # ALY&DB2.q&DB2.pin;
L3 = !ALY&DB3.pin # ALY&DB3.q&DB3.pin;
L4 = !ALY&DB4.pin # ALY&DB4.q&DB4.pin;
L5 = !ALY&DB5.pin # ALY&DB5.q&DB5.pin;
L6 = !ALY&DB6.pin # ALY&DB6.q&DB6.pin;
L7 = !ALY&DB7.pin # ALY&DB7.q&DB7.pin;

```

Intermediate equations for the full adder outputs (used for the ADD and SUB) functions as well as the “logical” functions (here, LDA and AND) are shown in Table 2-6. Note that the sole purpose of these intermediate equations is to simplify the task of writing the ALU equations. One can think of these as simply “definitions” (since they are part of the declaration section) of ‘symbols’ that will be used in “higher level” equations.

*intermediate  
equations*

The “real” equations start in Table 2-7. First are the carry equations that implement a simple ripple adder/subtractor. Next are the combinational equations that generate the ALU outputs based on the intermediate equations defined in Table 2-6. The data bus equations appear next; note that if ALE is negated, the “A” register retains its current state.

**Table 2-7** Equations section of ALU source file.

```

EQUATIONS

" Ripple carry equations (CY7 is COUT)
CY0 = DB0.q&(ALY$DB0.pin) # DB0.q&CIN # (ALY$DB0.pin)&CIN;
CY1 = DB1.q&(ALY$DB1.pin) # DB1.q&CY0 # (ALY$DB1.pin)&CY0;
CY2 = DB2.q&(ALY$DB2.pin) # DB2.q&CY1 # (ALY$DB2.pin)&CY1;
CY3 = DB3.q&(ALY$DB3.pin) # DB3.q&CY2 # (ALY$DB3.pin)&CY2;
CY4 = DB4.q&(ALY$DB4.pin) # DB4.q&CY3 # (ALY$DB4.pin)&CY3;
CY5 = DB5.q&(ALY$DB5.pin) # DB5.q&CY4 # (ALY$DB5.pin)&CY4;
CY6 = DB6.q&(ALY$DB6.pin) # DB6.q&CY5 # (ALY$DB6.pin)&CY5;
CY7 = DB7.q&(ALY$DB7.pin) # DB7.q&CY6 # (ALY$DB7.pin)&CY6;

" Combinational ALU equations
ALU0 = !ALX&S0 # ALX&L0;
ALU1 = !ALX&S1 # ALX&L1;
ALU2 = !ALX&S2 # ALX&L2;
ALU3 = !ALX&S3 # ALX&L3;
ALU4 = !ALX&S4 # ALX&L4;
ALU5 = !ALX&S5 # ALX&L5;
ALU6 = !ALX&S6 # ALX&L6;
ALU7 = !ALX&S7 # ALX&L7;

" Register bit and data bus control equations
[DB0..DB7].d = !ALE&[DB0..DB7].q # ALE&[ALU0..ALU7];

[DB0..DB7].clk = CLOCK;

[DB0..DB7].oe = AOE;

" Flag register state equations
CF.d = !ALE&CF.q # ALE&(!ALX&(CY7 $ ALY) # ALX&CF.q);

CF.clk = CLOCK;

ZF.d = !ALE&ZF.q # ALE&(!ALU7&!ALU6&!ALU5&!ALU4&!ALU3&!ALU2&!ALU1&!ALU0);

ZF.clk = CLOCK;

NF.d = !ALE&NF.q # ALE&ALU7;

NF.clk = CLOCK;

VF.d = !ALE&VF.q # ALE&(!ALX&(CY7 $ CY6) # ALX&VF.q);

VF.clk = CLOCK;

END

```

Last, but not least, are the equations that govern the four condition code bits. All of these flags retain their current state if ALE is negated. The carry flag (CF) and overflow flag (VF) are only affected by the ADD and SUB instructions. For ADD, the CF bit is set to the carry out of the most significant position (here, CY7); for SUB, the CF bit is interpreted as a *borrow* and is therefore set to the *complement* of the carry out of the sign position. The VF bit is simply the XOR of the *carry in* to the sign bit (CY6) with the *carry out* of the sign bit (CY7).

The negative flag (NF) and zero flag (ZF) are affected by all four functions implemented by our ALU. The NF bit is simply the sign bit (ALU7) of the result generated by the ALU, while the ZF bit is set to “1” if all the ALU result bits are zero.

Before moving on to the final block of our simple computer design, there is an important practical point worth noting. All of the functional blocks designed thus far – the memory, PC, IR, and ALU – can be independently implemented (or simulated) and tested (as well as debugged) before they are all “assembled together” into a completed computer. Independent testing and debugging of each functional block, in fact, is an important aspect of the “top-down, bottom-up” strategy we have espoused in this chapter.

*independent testing and debugging*

## 2.7.5 Instruction Decoder and Micro-sequencer

As described previously, there are two basic steps involved with “processing” each instruction, the combination of which is referred to as a micro-sequence. During a fetch cycle, the instruction pointed to by the PC is read from memory and loaded into the IR; the PC is incremented by one as the instruction is loaded. During the ensuing execute cycle, the instruction staged in the IR is “peeled” apart into an *opcode* field and an *operand address* field; the opcode field indicates the operation to be performed using data obtained from (or destined for) the memory location specified by the address field. The functional block that orchestrates the sequencing of these activities is called the *instruction decoder and micro-sequencer* (IDMS).

Since, in this initial version of our simple computer, there are only two different kinds of cycles (*fetch* and *execute*), a single flip-flop can be used as a *state counter* (SQ). In reality, this state counter is simply a single-bit binary counter (i.e., it simply *toggles* between “0” and “1”). Note that the state counter must be placed in the “fetch” state when START is pressed; therefore, it makes sense to assign the “reset” state

*state counter (SQ)*  
*toggles*

of the SQ flip-flop ( $SQ=0$ ) to the fetch cycle, and the “set” state of the SQ flip-flop ( $SQ=1$ ) to the execute cycle.

With the structure of the state counter established, the next step is to determine which control signals (of the functional blocks designed previously) need to be asserted when  $SQ=0$  (fetch) and  $SQ=1$  (execute). To accomplish this, we will need to refer back to each of the previous sub-sections (on the design of the individual functional blocks) as well as the instruction tracing worksheets completed previously.

Referring again to Figure 2-12, we note that the following signals need to be asserted to complete a fetch cycle. First, to “gate” the value in the PC onto the address bus, the signal POA needs to be asserted by the IDMS. To read the instruction, the memory needs to be selected (MSL asserted) and its data bus output enabled (MOE asserted). To load the instruction read from memory into the IR, the signal IRL needs to be asserted. Finally, to increment the PC as the instruction is loaded, the signal PCC needs to be asserted. A total of five system control signals, therefore, needed to be asserted by the IDMS during a fetch cycle (when  $SQ=0$ ): POA, MSL, MOE, IRL, and PCC.

The control signals that need to be asserted during an “ALU function” execute cycle (i.e., LDA, ADD, SUB, AND operation) can be inferred from Figure 2-13. First, to “gate” the operand address staged in the IR onto the address bus, the signal IRA needs to be asserted by the IDMS. To read the operand, the memory needs to be selected (MSL asserted) and its data bus output enabled (MOE asserted). To perform the operation specified by the instruction opcode (supplied to the IDMS from the upper 3-bits of the IR), ALE needs to be asserted along with the prescribed combination of ALX and ALY (based on the ALU design documented in Table 2-5).

The “store A” (STA) instruction execute cycle is similar, but notably different, than an “ALU function” execute cycle. Here, the address supplied to memory (from the IR, upon assertion of IRA) specifies the destination for the data in the “A” register. To complete the write to memory, it needs to be selected (MSL asserted) and write enabled (MWE asserted). To “gate” the data in the “A” register onto the data bus, AOE needs to be asserted. A total of four control signals need to be asserted, then, to execute a “store A” (STA) instruction: IRA, MSL, MWE, and AOE.

A succinct summary of all the system control signal assertions is provided in Table 2-8. Note that, for the sake of clarity, signal assertions are denoted using “H” (signals that are either negated or “don’t care” are left blank). By way of contrast, the control signal negations that are effected by execution of the HLT (halt) instruction are denoted using “L”.

**Table 2-8** System control table.

Decoded State	Instruction Mnemonic	MSL	MOE	MWE	PCC	POA	IRL	IRA	AOE	ALE	ALX	ALY
S0	<del>%4</del>	H	H	H	H	H						
S1	LDA	H	H					H		H	H	
S1	STA	H		H				H	H			
S1	ADD	H	H					H		H		
S1	SUB	H	H					H		H		H
S1	AND	H	H					H		H	H	H
S1	HLT	L		L	L	L			L			

The ABEL source file for the simple computer’s IDMS module is shown in Tables 2-9 and 2-10. Referring first to the declarations listed in Table 2-9, we find decoded opcode definitions (using the instruction mnemonics as pseudonyms for the corresponding opcode bit patterns) and decoded machine state definitions (S0 for fetch, S1 for execute). The purpose of defining an intermediate equation for each opcode combination is simply to make the job of writing the system control equations (that appear in Table 2-10) easier. Perhaps if we were more “clever”, we might have used the name “fetch” (instead of S0) and “execute” (instead of S1) to help make the subsequent equations a bit more clear (albeit more cumbersome to write).

Continuing with the IDMS equations in Table 2-10, we discover three basic components: the state counter, the run/stop flip-flop, and the system control equations. Looking first at the state counter, we note that if the machine RUN enable is high (i.e., the machine is “running”), the state counter flip-flop merely “toggles” each time a positive CLOCK edge occurs. If RUN is negated, SQ is reset to “0” (i.e., the “fetch” state). Pressing the START pushbutton also resets SQ to the “fetch” state.

*run/stop  
flip-flop*

**Table 2-9** Declarations section of IDMS module.

```

MODULE idms

TITLE      'Instruction Decoder and Microsequencer'

DECLARATIONS

CLOCK pin;

START pin; " asynchronous START pushbutton

OP0..OP2 pin; " opcode bits (input from IR5..IR7)

" State counter
SQ node istype 'reg_D,buffer';

" RUN/HLT state
RUN node istype 'reg_D,buffer';

" Memory control signals
MSL,MOE,MWE pin istype 'com';

" PC control signals
PCC,POA,ARS pin istype 'com';

" IR control signals
IRL,IRA pin istype 'com';

" ALU control signals (not using flags yet)
ALE,ALX,ALY,AOE pin istype 'com';

" Decoded opcode definitions
LDA = !OP2&!OP1&!OP0;   " LDA opcode = 000
STA = !OP2&!OP1& OP0;   " STA opcode = 001
ADD = !OP2& OP1&!OP0;   " ADD opcode = 010
SUB = !OP2& OP1& OP0;   " SUB opcode = 011
AND =  OP2&!OP1&!OP0;   " AND opcode = 100
HLT =  OP2&!OP1& OP0;   " HLT opcode = 101

" Decoded state definitions
S0 = !SQ.q; " fetch
S1 =  SQ.q; " execute

```

**Table 2-10** Equations section of IDMS module.

```

EQUATIONS

  " State counter
SQ.d = RUN.q&!SQ.q; " if RUN negated, resets SQ
SQ.clk = CLOCK;
SQ.ar = START;        " start in fetch state

  " Run/stop (equivalent of SR latch)
RUN.ap = START;      " start with RUN set to 1
RUN.clk = CLOCK;
RUN.d = RUN.q;
RUN.ar = S1&HLT;    " RUN is cleared when HLT executed

  " System control equations

MSL = RUN.q&(S0 # S1&(LDA # STA # ADD # SUB # AND));
MOE = S0 # S1&(LDA # ADD # SUB # AND);
MWE = S1&STA;
ARS = START;
PCC = RUN.q&S0;
POA = S0;
IRL = RUN.q&S0;
IRA = S1&(LDA # STA # ADD # SUB # AND);
AOE = S1&STA;
ALE = RUN.q&S1&(LDA # ADD # SUB # AND);
ALX = S1&(LDA # AND);
ALY = S1&(SUB # AND);

END

```

The run/stop flip-flop is defined next in Table 2-10. Here we note that pressing the START pushbutton asynchronously sets the RUN flip-flop, thereby enabling our simple computer to start executing instructions. Once set, the RUN signal remains asserted until asynchronously reset through execution of an HLT instruction.

We see how the RUN signal is used to enable/disable machine activity in the system control equations that follow. Note that if RUN is high, the system control signals are asserted according to the table in Table 2-8, as described previously. For example, MSL is asserted if a *fetch cycle* is being performed (S0 high); or, an *execute cycle* is being performed (S1 high) of an LDA instruction, an STA instruction, an ADD instruction, a SUB instruction, or an AND instruction. If RUN is low,

however, all of the pertinent system control signals are negated. Note that it is only necessary to negate the system control signals responsible for causing the various functional blocks to *change state* (i.e., it is *not necessary* to negate function select signals such as ALX and ALY, nor is it necessary to negate three-state output enables).

This completes the “bottom-up” phase of the design process for the initial version of our simple computer. All of the ABEL code described in this section could be implemented using a single, modest-size PLD. The addition of a conventional memory chip would yield a working computer. Before augmenting the instruction set with some useful extensions, though, let’s take a closer look at system timing.

## 2.8 System Timing Analysis

When we designed the program counter in Section 2.7.2, there was an appearance of “cheating” – specifically, of using the current value in the PC to access an instruction in memory while, at apparently the same time, telling the PC to increment. This is an issue that deserves further scrutiny.

To gain a better understanding of the timing relationship among different activities within our computer, we need to understand two basic hardware-imposed constraints. The first is that only one device (functional block) can drive a bus on a given bus cycle, i.e., “bus fighting” must be avoided. The second is that data can only “pass through” one edge-triggered flip-flop per cycle. Thus, it is not possible to load a value into a register and expect to “use it” (have the value available on the register’s outputs) on the same cycle.

*bus fighting*

Given these constraints, we are now prepared to examine in detail the sequence of activities that occur during a fetch cycle. A “qualitative” timing diagram is provided in Figure 2-21 for this purpose (by *qualitative* we mean that we’re not interested in the exact number of nanoseconds between one signal assertion and another, just the fact that there is a *delay*). Depicted in this diagram is the sequencing that occurs as the machine finishes an execute cycle, performs a fetch of the next instruction, and subsequently proceeds to execute the instruction just fetched. Our focus here is on the events that constitute a fetch cycle.

*qualitative  
timing diagram*

The first thing to note is that, since the functional blocks of the machine were designed using positive-edge-triggered flip-flops, the *clock edges* “drive” the machine from state-to-state. Thus, a “fetch cycle” is the

*clock edges*

time between the clock edge that drives the machine from the previous execute cycle to the current fetch cycle, and the subsequent clock edge that transitions the machine from the fetch cycle to an execute cycle. Shortly after the first clock edge in Figure 2-21, then, the control signals MSL, MOE, POA, IRL, and PCC are asserted (the delay relative to the clock edge in generating these signals is due to the propagation delay of the state counter plus the delay associated with the system control equations – see Table 2-10).

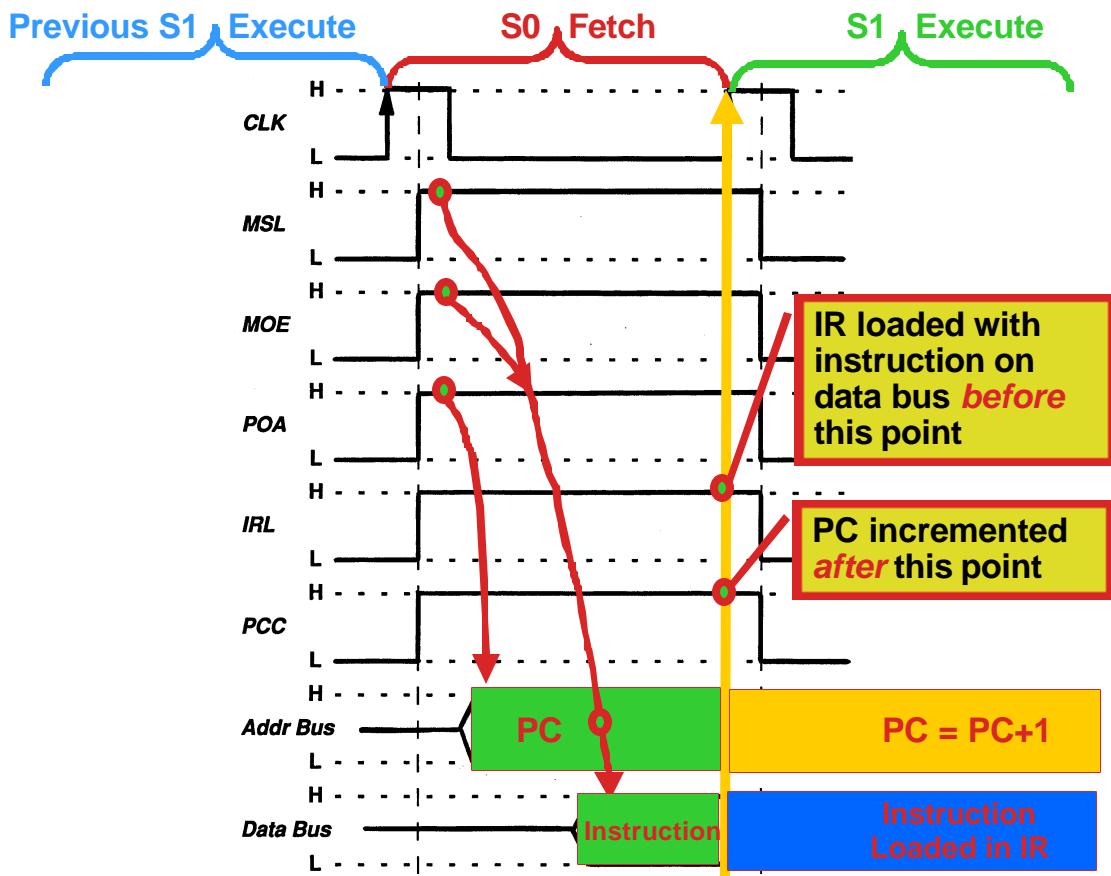


Figure 2-21 Fetch cycle event timing relationship.

The assertion of POA causes the three-state buffers of the PC to turn on and drive its value onto the address bus. The value on the address bus, in conjunction with the MSL and MOE signal assertions, causes the memory to drive the addressed instruction onto the data bus (note that, in most practical systems, this constitutes a substantial part of the cycle time). Provided the instruction is on the data bus at least  $t_{SETUP}$  (of the D flip-flop) prior to the next clock edge, it is successfully loaded into the IR (because the IRL signal is asserted) when that edge occurs.

While this may seem to be “enough” activity already, we realize that a related “housekeeping” activity can be accomplished on this cycle as well: incrementing the value in the PC, so it points to the next instruction (in preparation for the next fetch). Again, based on the use of edge-triggered flip-flops in our design, we note that the value on the data bus *just prior* to the clock edge that loads the IR determines the next state of the IR. It follows, then, that we can use that *same clock edge* to drive the PC to its next state – this is why PCC is also asserted during a fetch cycle. Note that the PC state change will occur *after* the clock edge, i.e., after the instruction has been safely loaded into the IR. This allows us to effectively *overlap* the load of the IR with the increment of the PC on the same cycle. We will make use of this same principle when we add some extensions to our machine later in this chapter.

*overlap*

One might ask at this point, “Could we have delayed the increment of the PC until the execute cycle?” In the initial version of our simple computer, it would clearly be possible: here, the “new value” in the PC would be available shortly after the commencement of the fetch cycle, thus enabling the correct instruction to be loaded into the IR (the only consequence might be a small amount of additional propagation delay for the “new” value to become stable). When we add subroutine linkage instructions to our computer, however, we will find it useful to have the “new” value of the PC available during the first execute cycle (to serve as the “return address” for a “subroutine call” instruction). In anticipation of this extension, we will include the increment of the PC as an integral part of the fetch cycle.

## 2.9 Simple Computer Extensions

When we originally designed our instruction set, we purposefully left two opcode bit patterns “uncommitted”. The reason we did this was to provide room for expansion. We will, then, add a “pair” of instructions at a time to our “base” instruction set. The “pairs” we will add include input/output (IN/OUT) instructions, transfer of control instructions (JMP/JZF), stack manipulation instructions (PSH/POP), and subroutine linkage instructions (JSR/RTS).

### 2.9.1 Input/Output Instructions

When we first drew the “big picture” of our simple computer (see Figure 2-4), we included a switch “input port” and an LED “output port”. As evident from the initial version of our instruction set, we included no

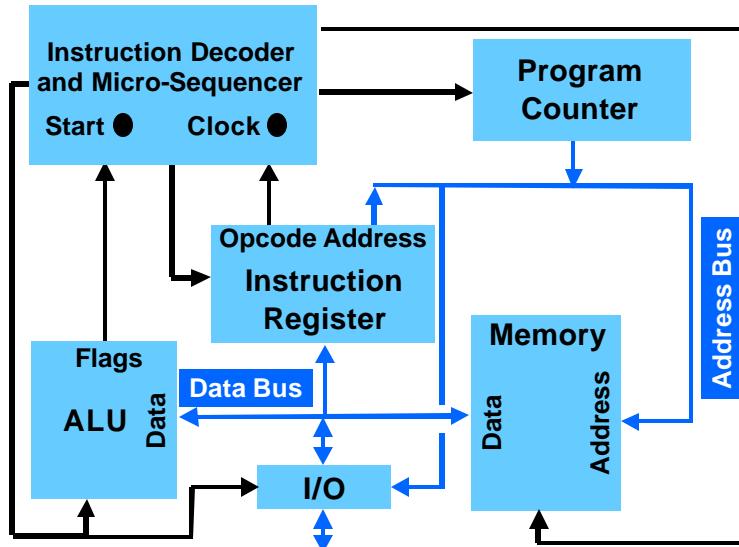
*input port*  
*output port*

provision for using these. It makes sense, then, to add instructions for providing our machine with the “modern convenience” of data input and output (“I/O”).

First, we need to establish the *destination* that will be used for data input (or *read*) from the “outside world”, as well as the *source* for data that will be output (or *written*). Given that our machine has but one register that participates in data transactions – namely, the “A” register – it is the most likely candidate to serve as the destination/source of data that is input/output, respectively. Thus, our new “IN” instruction will function in a manner similar to an LDA instruction, except the source of data will be the “outside world” and the address field will be used as a pointer to an “input device” (instead of to memory). Similarly, our new “OUT” instruction will function in a manner similar to an STA instruction, except the destination of data will be the “outside world” and the address field will be used as a pointer to an “output device”. A name commonly used for this input/output strategy is *accumulator-mapped I/O*.

Second, we need to establish how data will be communicated to/from the ubiquitous “outside world”. Basically, a “gateway” is needed between the system data bus and the external input and output devices, along with some new system control signals that enable a “read” (IOR) or a “write” (IOW) via this gateway. Also, a means of decoding the I/O addresses (typically called *port* or *device numbers*) into individual “device selects” (or *enables*) is needed. A diagram illustrating the placement of the “I/O block” is provided in Figure 2-22; an ABEL source file for a specific instance of this module is given in Table 2-11.

**IOR**  
**IOW**  
**port numbers**  
**device numbers**  
**I/O block**



**Figure 2-22** Block diagram of simple computer with I/O.

**Table 2-11** Basic I/O module.

```

MODULE io

TITLE      'Input/Output Port 00000'

DECLARATIONS

DB0..DB7 pin istype 'com';    " data bus
AD0..AD4 pin;                 " address bus
IN0..IN7 pin;                  " input port
OUT0..OUT7 pin istype 'com';   " output port

IOR pin; " Input port read
IOW pin; " Output port write

" Port select equation for port address 00000

PS = !AD4&!AD3&!AD2&!AD1&!AD0;

EQUATIONS

[DB0..DB7] = [IN0..IN7];
[DB0..DB7].oe = IOR&PS;
[OUT0..OUT7] = [DB0..DB7];
[OUT0..OUT7].oe = IOW&PS;

END

```

Referring to the ABEL file, we see that it contains a specific port address decoding equation, here for port address  $00000_2$ . When the pattern on the address bus matches this value, an I/O transaction via this port address is enabled. If an IN instruction is being executed, assertion of the IOR signal (by the IDMS) causes the value on the “IN pins” (IN0...IN7) to be gated onto the system data bus, allowing it to be loaded into the “A” register. If an OUT instruction is being executed, assertion of the IOW signal causes the value on the data bus (supplied by the “A” register) to be gated to the “OUT pins” (OUT0...OUT7).

There is a limitation, however, inherent in the I/O port design shown in Table 2-11: the value output (when an OUT instruction is executed) is only “active” for a very short time (specifically, the amount of time the IOW signal is asserted by the IDMS). For devices such as light

emitting diodes (LEDs), the brief assertion of IOW will not provide a satisfactory display. A better solution is to *latch* the value sent to the output port, and retain it until execution of a subsequent OUT instruction changes the value. An I/O module that provides a latched output port is provided in Table 2-12. Here, assertion of IOW in conjunction with the proper port address opens a transparent latch, which then assumes the new value sent on the data bus. The latch closes (retains its value) when IOW is negated.

*latched  
output port*

**Table 2-12** Latched I/O port.

```

MODULE iol

TITLE      'Input/Output Port 00000 - With Output Latch'

DECLARATIONS

DB0..DB7 pin istype 'com';      " data bus
AD0..AD4 pin;                  " address bus
IN0..INT7 pin;                 " input port
OUT0..OUT7 pin istype 'com';   " output port

IOR pin; " Input port read
IOW pin; " Output port write

" Port select equation for port address 00000

PS = !AD4&!AD3&!AD2&!AD1&!AD0;

EQUATIONS

[DB0..DB7] = [IN0..INT7];

[DB0..DB7].oe = IOR&PS;

" Transparent latch for output port

[OUT0..OUT7] = !(IOW&PS)&[OUT0..OUT7] # IOW&PS&[DB0..DB7];

END

```

The augmented system control table for our simple computer plus I/O is given in Table 2-13. Note that there are two “new” equations (for IOR and IOW), along with four equations that need to be updated (for IRA, AOE, ALE, and ALX). The updated system control equations are given in Table 2-14.

**Table 2-13** System control table modified for I/O.

Decoded State	Instruction Mnemonic	MSL	MOE	MWE	PCC	POA	IRL	IRA	AOE	ALE	ALX	ALY	IOR	IOW
S0	\$4	H	H	H	H	H								
S1	LDA	H	H					H		H	H			
S1	STA	H		H				H	H					
S1	ADD	H	H					H		H				
S1	SUB	H	H					H		H	H	H		
S1	AND	H	H					H		H	H	H		
S1	HLT	L		L	L			L						
S1	IN							H		H	H		H	
S1	OUT							H	H					H

**Table 2-14** System control equations modified for I/O.

```

" System control equations (IDMS)

MSL = RUN.q&(S0 # S1&(LDA # STA # ADD # SUB # AND));
MOE = S0 # S1&(LDA # ADD # SUB # AND);
MWE = S1&STA;
ARS = START;
PCC = RUN.q&S0;
POA = S0;
IRL = RUN.q&S0;
IRA = S1&(LDA # STA # ADD # SUB # AND # IN # OUT);
AOE = S1&(STA # OUT);
ALE = RUN.q&S1&(LDA # ADD # SUB # AND # IN);
ALX = S1&(LDA # AND # IN);
ALY = S1&(SUB # AND);

IOR = S1&IN;
IOW = S1&OUT;

END

```

## 2.9.2 Transfer-of-Control Instructions

Any program worth the silicon it runs on typically does more than execute “straight line” code. Instead, execution transfers to different parts of the program based on various conditions encountered. Generically, we refer to the instructions that allow program execution to “jump around” as *transfer-of-control* instructions.

There are two basic types of transfer-of-control instructions. If the address field of the instruction contains the (absolute) address in memory at which execution should continue, it is most often referred to as a “jump” instruction. If the address field instead represents the (signed) “distance” the next instruction is from the transfer-of-control instruction, it is referred to as a “branch”. (There is not universal agreement on this nomenclature, however – see sidebar.) Jumps (or branches) that “always happen” are called *unconditional*; those that happen only if a certain combination of condition codes exists are called *conditional*.

*straight line code*

*transfer-of-control  
instructions*

*jump instruction*

*branch instruction*

*unconditional  
conditional*

### A Branch by Any Other Name

Regrettably, there is no “universal agreement” among manufacturers of microcontrollers concerning the names used for the basic transfer-of-control instruction types. Since this is primarily a text dealing with Motorola products, we will use the names they commonly use: “jump” for absolute transfer, and “branch” for relative transfer. Be advised, though, that another “major manufacturer” (Intel) uses *just the opposite* designation: “branch” for absolute transfer, and “jump” for relative transfer. Although the author cut his “digital teeth” on Intel processors, he prefers the Motorola adopted names.

The addition of transfer-of-control instructions to our simple computer will require modifications to the PC (as well as to the IDMS). Specifically, we will need to provide a mechanism for loading a new value into the PC to implement “jump-style” instructions, or for adding a signed offset to the value in the PC to implement “branch-style” instructions. Here we will focus on the modifications necessary to implement jump-style instructions. An ABEL source file for the modified PC is provided in Table 2-15. Note that it is the same as the “original” PC (see Table 2-3), except that a “load from address bus” function (and associated control signal, PLA) has been added. Recall that the “new value” with which the PC is to be loaded is staged in the IR, and can therefore be conveniently “transported” to the PC via the address bus.

*PLA*

**Table 2-15** PC modifications to support transfer-of-control instructions.

```

MODULE pc

TITLE      'Program Counter'

DECLARATIONS

CLOCK pin;

PC0..PC4 pin istype 'reg_D,buffer';

PCC pin; " PC count enable
PLA pin; " PC load from address bus enable
POA pin; " PC output on address bus tri-state enable
ARS pin; " asynchronous reset (connected to START)

" Note: Assume PCC and PLA are mutually exclusive

EQUATIONS

"      retain state      load
PC0.d = !PCC&!PLA&PC0.q # PLA&PC0.pin
"      count up by 1
# PCC&!PC0.q;
PC1.d = !PCC&!PLA&PC1.q # PLA&PC1.pin
# PCC&(PC1.q $ PC0.q);
PC2.d = !PCC&!PLA&PC2.q # PLA&PC2.pin
# PCC&(PC2.q $ (PC1.q&PC0.q));
PC3.d = !PCC&!PLA&PC3.q # PLA&PC3.pin
# PCC&(PC3.q $ (PC2.q&PC1.q&PC0.q));
PC4.d = !PCC&!PLA&PC4.q # PLA&PC4.pin
# PCC&(PC4.q $ (PC3.q&PC2.q&PC1.q&PC0.q));

[PC0..PC4].oe = POA;
[PC0..PC4].ar = ARS;
[PC0..PC4].clk = CLOCK;

END

```

The system control table, modified to include an “unconditional jump” instruction (JMP) along with a “jump if zero flag set” (JZF) instruction, is shown in Table 2-16. As its name implies, the JZF instruction causes a transfer-of-control to the address following the opcode if the zero flag (ZF) is set, i.e., the result of the most recent ALU operation

**JMP**  
**JZF**

has generated a result of zero in the “A” register. (As it turns out, this is a fairly “popular” condition to check in practical applications.) If the condition specified by a “conditional jump” instruction (like JZF) is *not* met, however, *nothing happens* (often called a *no operation*, or “NOP”) – execution merely continues with the instruction that follows. In order to effect the load of the jump address, the IDMS needs to know the state of the various condition code bits generated by the ALU. The equations for IRA and PLA, then, will be a function of ZF for the new instructions added to the machine in Table 2-17.

*no operation*  
NOP

**Table 2-16** System control table modified for transfer-of-control instructions.

Decoded State	Instruction Mnemonic	MSL	MOE	MWE	PCC	POA	IRL	IRA	AOE	ALE	ALX	ALY	PLA
S0	%4	H	H		H	H	H						
S1	LDA	H	H					H	H	H	H		
S1	STA	H		H				H	H				
S1	ADD	H	H					H		H			
S1	SUB	H	H					H	H		H		
S1	AND	H	H					H		H	H	H	
S1	HLT	L		L	L			L					
S1	JMP							H					H
S1	JZF						ZF						ZF

**Table 2-17** IDMS modifications to support transfer-of-control.

```

" System control equations (IDMS)

MSL = RUN.q&(S0 # S1&(LDA # STA # ADD # SUB # AND));
MOE = S0 # S1&(LDA # ADD # SUB # AND);
MWE = S1&STA;
ARS = START;
PCC = RUN.q&S0;
POA = S0;
IRL = RUN.q&S0;
IRA = S1&(LDA # STA # ADD # SUB # AND # JMP # JZF&ZF);
AOE = S1&I1;
ALE = RUN.q&S1&(LDA # ADD # SUB # AND);
ALX = S1&(LDA # AND);
ALY = S1&(SUB # AND);

PLA = S1&(JMP # JZF&ZF);

END

```

One could imagine, at this point, a number of other conditions that would be useful for determining whether or not a jump or branch should be “taken”. In addition to a separate “jump on condition” instruction dedicated to each flag (CF, NF, VF, ZF), there are various *Boolean combinations* of these flags that are of interest as well (e.g., testing for “greater than” or “less than or equal to”). All of these variations will be explored when we tackle the instruction set of a “real” microcontroller in the next chapter.

*Boolean  
combinations  
of flags*

### 2.9.3 Multiple Execute Cycle Instructions

To this point, all of the instructions we originally defined or added to our simple computer required a single fetch cycle followed by a single execute cycle. As the functions performed by an individual instruction become more complex, however, additional execute cycles become necessary. On the surface, this would appear to be a relatively straightforward extension, accomplished by simply adding extra bits to the state counter in the IDMS, along with a binary decoder to decode the various states. Adding one additional bit to our original state counter would provide us with four possible states: a fetch state (S0), followed by three execute states (S1, S2, S3).

*S1  
S2  
S3*

The “complication” that arises is that, despite this addition, we want our original “single execute state” instructions to still execute in a single state. Further, we want any new instructions that require two execute states to consume only two execute states, and new instructions that require all three execute states to consume exactly three execute states. More succinctly, we want our state counter to be able to accommodate *variable-length execution cycles* (here, from 1 to 3).

*variable-length  
execution cycles*

One way this can be accomplished is by adding a *synchronous reset* capability to our (now 2-bit) state counter. For this purpose, we will add a new signal (RST) to our system control table that, when asserted, causes the state counter to *reset to zero* when the next clock edge occurs. In the system control table, this signal will be asserted on the *final execute cycle* of each instruction. For single execute cycle instructions (such as LDA, STA, ADD, AND, SUB), the RST signal will be asserted during S1 (the first execute cycle), ensuring that the next cycle will be a “fetch”. For instructions requiring two execute cycles, the RST signal will be asserted during S2 (the second execute cycle). Finally, for three-execute-cycle instructions, the RST signal will be asserted during S3 (note that, if RST is *not* asserted at this point, the

*synchronous reset*

state counter will “wrap around” to zero automatically, thus ensuring that the next cycle is a “fetch” regardless).

**Table 2-18** IDMS modifications for multi-execute-cycle instructions (declarations section).

```

MODULE idmsr

TITLE 'Instruction Decoder and Microsequencer with Multi-Execution States'

DECLARATIONS

CLOCK pin;
START pin;      " asynchronous START pushbutton
OP0..OP2 pin;  " opcode bits (input from IR5..IR7)

" State counter
SQA node istype 'reg_D,buffer'; " low bit of state counter
SQB node istype 'reg_D,buffer'; " high bit of state counter

" Synchronous state counter reset
RST node istype 'com';

" RUN/HLT state
RUN node istype 'reg_D,buffer';

" Memory control signals
MSL,MOE,MWE pin istype 'com';

" PC control signals
PCC,POA,ARS pin istype 'com';

" IR control signals
IRL,IRA pin istype 'com';

" ALU control signals
ALE,ALX,ALY,AOE pin istype 'com';

" Decoded opcode definitions
LDA = !OP2&!OP1&!OP0;    " opcode 000
STA = !OP2&!OP1& OP0;    " opcode 001
ADD = !OP2& OP1&!OP0;    " opcode 010
SUB = !OP2& OP1& OP0;    " opcode 011
AND =  OP2&!OP1&!OP0;    " opcode 100
HLT =   OP2&!OP1& OP0;    " opcode 101

" Decoded state definitions
S0 = !SQB&!SQA; " fetch state
S1 = !SQB& SQA; " first execute state
S2 =  SQB&!SQA; " second execute state
S3 =  SQB& SQA; " third execute state

```

**Table 2-19** IDMS modifications for multi-execute-cycle instructions  
(equations section).

```

EQUATIONS

" State counter
" if RUN negated or RST asserted,
" state counter is reset
SQA.d = !RST & RUN.q & !SQA.q;
SQB.d = !RST & RUN.q & (SQB.q $ SQA.q);

SQA.clk = CLOCK;
SQB.clk = CLOCK;
SQA.ar = START;           " start in fetch state
SQB.ar = START;

" Run/stop (equivalent of SR latch)
RUN.ap = START;          " start with RUN set to 1
RUN.clk = CLOCK;
RUN.d = RUN.q;
RUN.ar = S1&HLT;         " RUN is cleared when HLT executed

" System control equations
MSL = RUN.q&(S0 # S1&(LDA # STA # ADD # SUB # AND));
MOE = S0 # S1&(LDA # ADD # SUB # AND);
MWE = S1&STA;
ARS = START;
PCC = RUN.q&S0;
POA = S0;
IRL = RUN.q&S0;
IRA = S1&(LDA # STA # ADD # SUB # AND);
AOE = S1&STA;
ALE = RUN.q&S1&(LDA # ADD # SUB # AND);
ALX = S1&(LDA # AND);
ALY = S1&(SUB # AND);
RST = S1&(LDA # STA # ADD # SUB # AND);

END

```

The state counter modifications necessary to accommodate multiple execute cycles are shown in Tables 2-18 and 2-19. Following conventional notation, bit “A” of the modified state counter is the least significant bit, and bit “B” is the most significant bit. Note that if RUN is negated, or RST is asserted, the state counter is reset to “00”. Pressing the START pushbutton also resets the state counter to zero.

In the sections that follow, we will see examples of instructions that require two or three execute states. The system control tables for these “new” instruction sets will therefore include the RST signal.

## 2.9.4 Stack Manipulation Instructions

An important “modern convenience” that most “real” computers enjoy is a stack mechanism. Stacks – also referred to as *last-in, first-out* (LIFO) data structures – facilitate a number of capabilities, including expression evaluation, subroutine linkage, and parameter passing. While there are many variations on stack implementation, the most common strategy is to place the stack contents in the uppermost portion of (read/write) memory, and add a new register to the machine that serves as a pointer to the top item on the stack. Not surprisingly, this register is called the *stack pointer* (SP). An augmented system block diagram illustrating the placement of the SP register in our simple computer is given in Figure 2-23.

*last-in, first-out  
LIFO*

*expression evaluation  
subroutine linkage  
parameter passing*

*stack pointer  
SP*

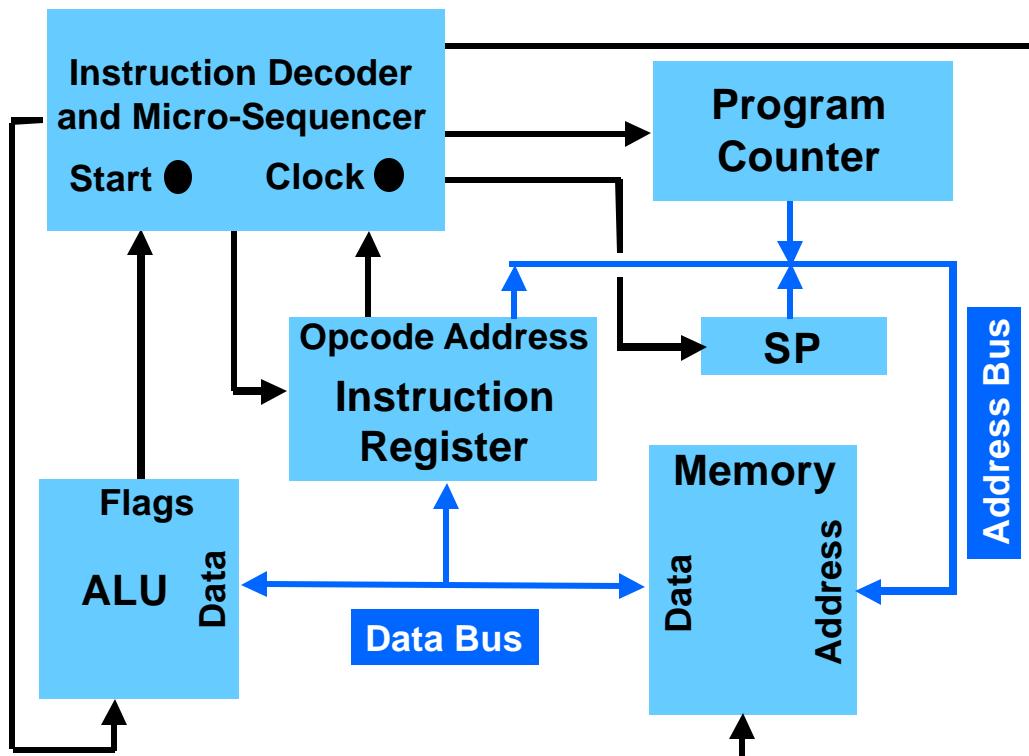


Figure 2-23 Block diagram of simple computer with stack.

Since program “growth” (or *execution direction*) is toward *increasing* addresses (starting in “low” memory), it makes sense that *stack growth* should be toward *decreasing* addresses (starting in “high” memory). The stack grows as items are “pushed” onto it, which means the SP register must decrement as it grows; conversely, as items are “popped” off the stack and its size diminishes, the SP register must increment.

*execution  
direction  
stack growth*

At this point, we realize there are two possible conventions that can be used as a “stack pointer paradigm” – we can choose to have the SP register point to the *top stack item*, or we can choose to have it point to the *next available location*. The most commonly used convention (and the one we will adopt here) is to have the SP register point to the top stack item. Based on this choice, we realize that the initial value of the SP register needs to be *one greater* than the address in which the first stack item is placed. Because the SP register points to the top stack item, it must be decremented in order to allocate space for a new item during a “push” operation. If the stack starts in the *uppermost location* of memory (for our simple computer, location  $11111_2$ ), the SP register should be initialized to  $00000_2$  (i.e., one greater than  $11111_2$ , modulo  $2^5$ ). Stack growth and retraction based on this “conventional convention” is illustrated in Figure 2-24. Note that items popped off the stack are merely *de-allocated* from the stack area, *not erased*.

*stack convention  
top stack item  
next available location*

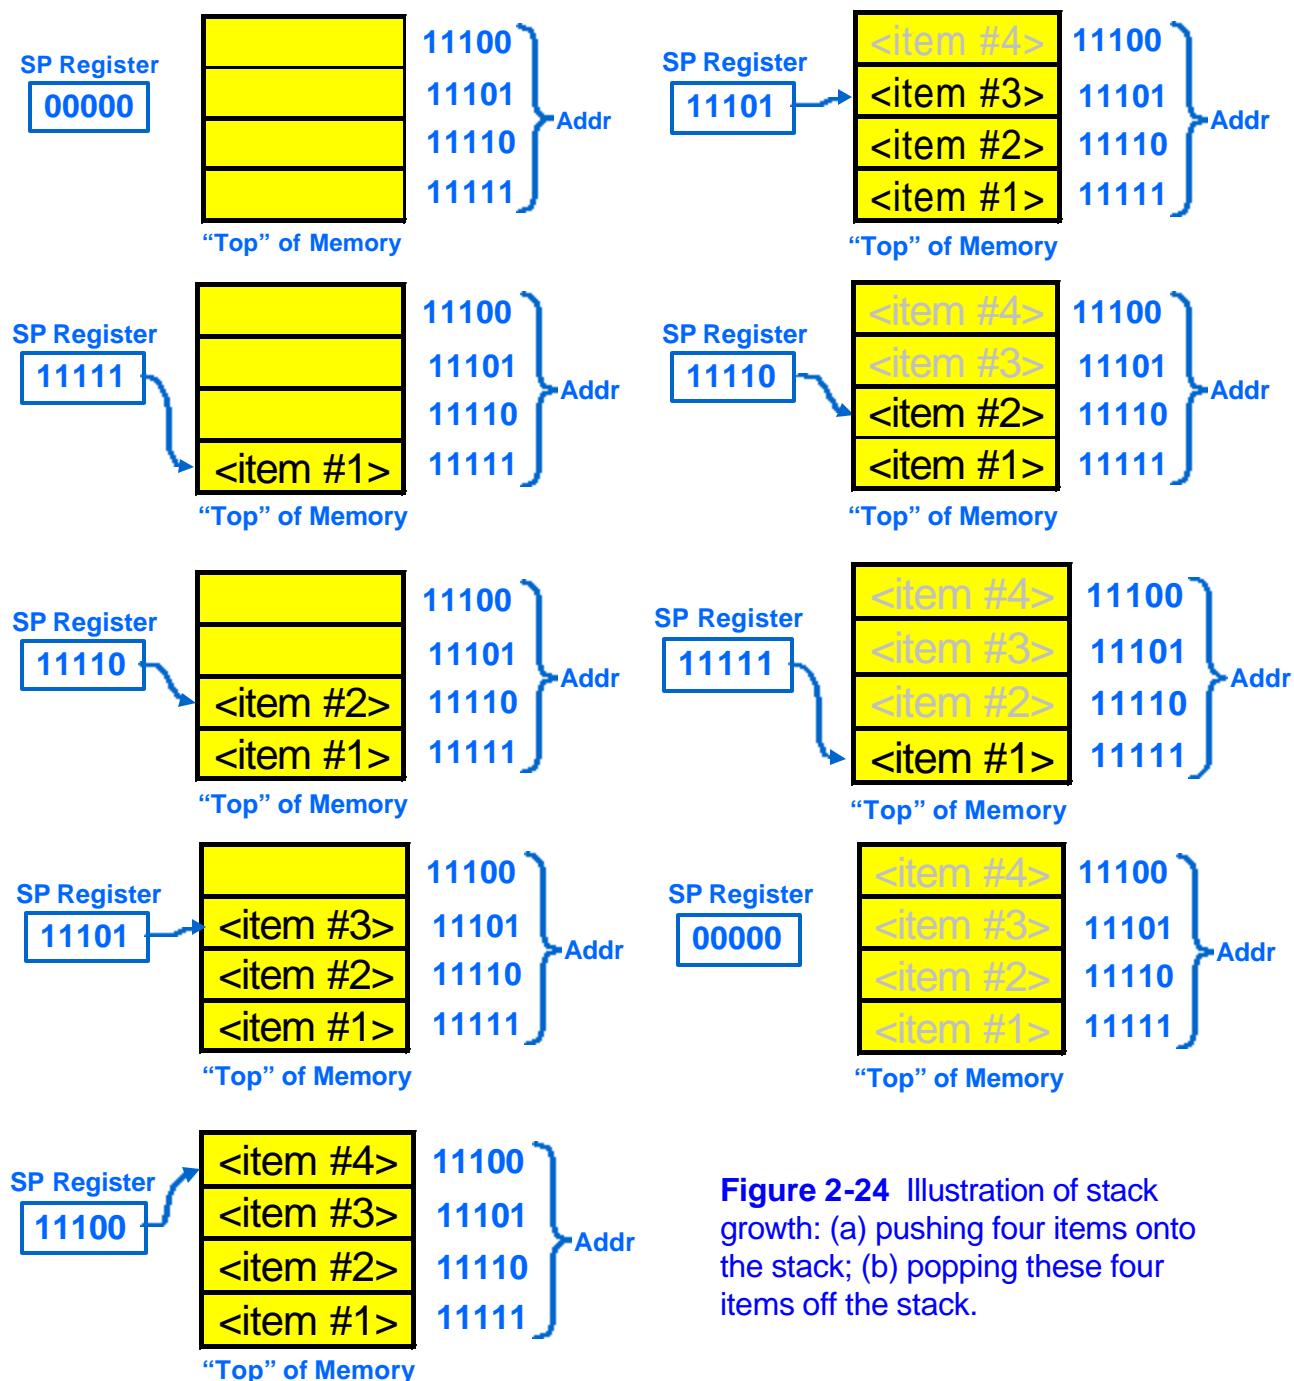
Based on an understanding of how the stack mechanism works, we can now consider the design of the SP register module, documented in Table 2-20. The first thing we note is that the SP register is simply an “up/down” binary counter, with three-state output buffers and an asynchronous reset. The IDMS, then, needs to supply the SP register with four control signals: an asynchronous reset (ARS), an increment enable (SPI), a decrement enable (SPD), and a three-state buffer enable (SPA) that gates the value in the SP register onto the address bus.

*ARS  
SPI  
SPD  
SPA*

We now have all the “ingredients” available to create two new *stack manipulation* instructions: push the contents of the “A” register onto the stack (PSH), and pop the top stack item into the “A” register (POP). One possible application for such a pair of instructions is expression evaluation. Here, intermediate results of a calculation can be placed on the stack and retrieved when needed. For example, to evaluate the expression  $(W+X) - (Y-Z)$ , we could first calculate the quantity  $(Y-Z)$  and push it onto the stack, next calculate the quantity  $(W+X)$ , and finally pop the stack and subtract that value from our “running total”. Formal methods exist for transforming an arbitrarily complex, parenthesized expression into *postfix* form.

*stack manipulation  
instructions  
PSH  
POP*

*postfix*



**Figure 2-24** Illustration of stack growth: (a) pushing four items onto the stack; (b) popping these four items off the stack.

**Table 2-20** Stack pointer module.

```

MODULE sp

TITLE      'Stack Pointer'

DECLARATIONS

CLOCK pin;

SP0..SP4 pin istype 'reg_D,buffer';

SPI pin; " SP increment enable
SPD pin; " SP decrement enable
SPA pin; " SP output on address bus tri-state enable
ARS pin; " asynchronous reset (connected to START)

" Note: Assume SPI and SPD are mutually exclusive

EQUATIONS

"      retain state      increment/decrement
SP0.d = !SPI&!SPD&SP0.q # SPI&!SP0.q
# SPD&!SP0.q;
SP1.d = !SPI&!SPD&SP1.q # SPI&(SP1.q$SP0.q)
# SPD&(SP1.q$!SP0.q);
SP2.d = !SPI&!SPD&SP2.q # SPI&(SP2.q$(SP1.q&SP0.q))
# SPD&(SP1.q$(!SP1.q&!SP0.q));
SP3.d = !SPI&!SPD&SP3.q # SPI&(SP3.q$(SP2.q&SP1.q&SP0.q))
# SPD&(SP3.q$(!SP2.q&!SP1.q&!SP0.q));
SP4.d = !SPI&!SPD&SP4.q # SPI&(SP4.q$(SP3.q&SP2.q&SP1.q&SP0.q))
# SPD&(SP4.q$(!SP3.q&!SP2.q&!SP1.q&!SP0.q));

[SP0..SP4].oe = SPA;
[SP0..SP4].ar = ARS;
[SP0..SP4].clk = CLOCK;

END

```

Implementation of the PSH instruction requires two execute states. Here, the SP register must first be *decremented* in order to allocate space for the new item (given the convention we have adopted that SP points to the *top stack item*). After the SP has been decremented, it can be used as a pointer to indicate where in memory the contents of “A” should be stored.

For POP, however, the SP register is already pointing to the “right place”, enabling the “A” register to be loaded with the contents of that location on the first execute cycle. The “bookkeeping” step of deallocating the item just popped off the stack (accomplished by incrementing the SP register) needs to follow, which at first glance appears to require a second execute cycle. Here, though, the same clock edge that is used to load the “A” register (with the value pointed to by the SP register) can be used to increment the SP register, since its value will not change until after the load has safely completed. The POP instruction, then, can be implemented using a single execute cycle. (Note the similarity between the *overlap* employed here and the overlap of the PC increment used previously in the fetch cycle.)

*de-allocation*

*overlap*

A modified system control table illustrating the addition of PSH and POP to our simple computer’s instruction set is given in Table 2-21. Here, only one of the instructions listed (PSH) requires a second execute state (S2); the remaining instructions complete in a single execute cycle. Note, therefore, that RST is not asserted until the S2 state of the PSH instruction, while for the other instructions RST is asserted during the S1 state. A modified ABEL source file for the IDMS that corresponds to this version of our instruction set is given in Table 2-22.

**Table 2-21** System control table modifications for stack manipulation instructions.

Decoded State	Instruction Mnemonic	MSL	MOE	MWE	PCC	POA	IRL	IRA	AOE	ALE	ALX	ALY	SPI	SPD	SPA	RST
S0	%4	H	H		H	H	H									
S1	LDA	H	H					H		H	H					H
S1	STA	H		H				H	H							H
S1	ADD	H	H					H		H						H
S1	SUB	H	H					H		H		H				H
S1	AND	H	H					H		H	H	H				H
S1	HLT	L		L	L			L								
S1	PSH													H		
S1	POP	H	H						H	H		H	H	H	H	H
S2	PSH	H		H				H						H	H	

**Table 2-22** IDMS modifications for stack manipulation instructions.

```

" System control equations

MSL = RUN.q&(S0 # S1&(LDA # STA # ADD # SUB # AND
    # POP) # S2&PSH);
MOE = S0 # S1&(LDA # ADD # SUB # AND # POP);
MWE = S1&STA # S2&PSH;
ARS = START;
PCC = RUN.q&S0;
POA = S0;
IRL = RUN.q&S0;
IRA = S1&(LDA # STA # ADD # SUB # AND);
AOE = S1&STA # S2&PSH;
ALE = RUN.q&S1&(LDA # ADD # SUB # AND # POP);
ALX = S1&(LDA # AND # POP);
ALY = S1&(SUB # AND);

SPI = S1&POP;
SPD = S1&PSH;
SPA = S1&POP # S2&PSH;

RST = S1&(LDA # STA # ADD # SUB # AND # POP) # S2&PSH;

END

```

Before adding our final set of simple computer extensions, some additional comments on PSH/POP are in order. Virtually every computer that has a stack mechanism implements some variation of the basic push/pop instruction pair, typically for each “important” register in the machine’s architecture. Other variations – which would be particularly useful for performing expression evaluation on our simple computer – include “pop and add” (i.e., pop the stack and add that item to the contents of the “A” register), “pop and subtract”, etc. In fact, instructions like “pop and add” are simple variations of the “basic POP” instruction, and can be implemented with only minor modifications to the ABEL source files given.

*pop and add  
pop and subtract*

## 2.9.5 Subroutine Linkage Instructions

Another important “modern convenience” that most computers enjoy is a subroutine linkage mechanism, which is the final extension to our simple computer we will explore in this chapter. A very effective way to provide this capability is to utilize a stack. While there are other ways that subroutine linkage can be implemented in practice, use of a stack is attractive because it: (a) allows *arbitrary nesting* of subroutine calls; (b) provides a mechanism for passing parameters to subroutines; (c)

*arbitrary  
nesting*

allows recursion (the ability of a subroutine to call itself); and (d) allows reentrancy (the ability of a code module to be shared among quasi-simultaneously executing tasks).

*recursion  
reentrancy*

The two subroutine-linkage instructions we will add to our “base” instruction set are “jump to subroutine” (JSR) and “return from subroutine (RTS). Generically, we can simply refer to these as (subroutine) “call” and “return” instructions. As can be seen from the “subroutine in action” illustration (Figure 2-25), one of the key things the “call” instruction must do is establish a “return path” to the calling program (hence the name “linkage”). Placing the calling program’s *return address* on the stack affords nesting of subroutine calls (i.e., one subroutine calls another, which then calls another, etc.).

*return address*

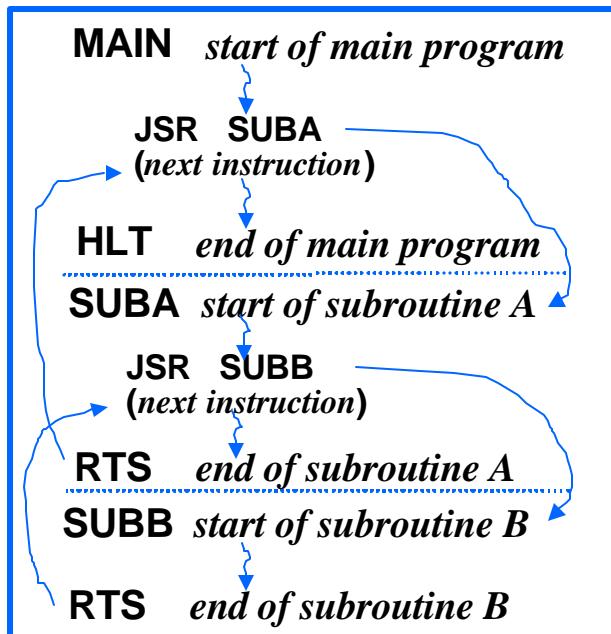
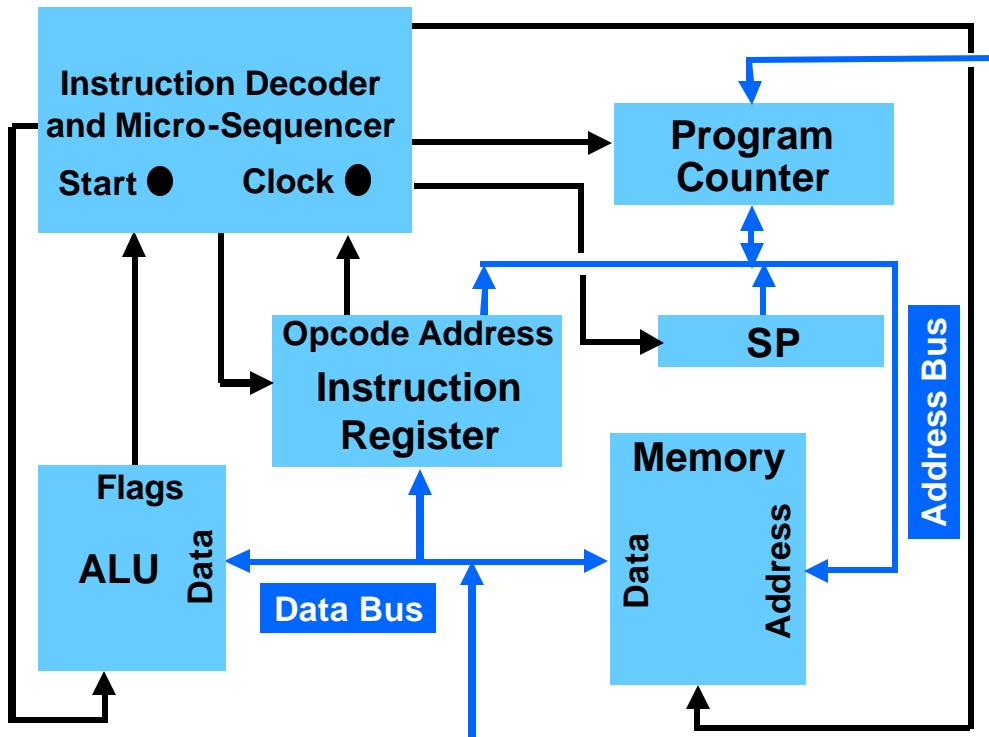


Figure 2-25 Subroutine linkage in action.

Note that the return address is simply the address of the instruction that *follows* the JSR. Recalling that the PC is automatically incremented as part of the fetch cycle, we realize that the desired return address has already been calculated. The value in the PC simply needs to be pushed onto the stack when a JSR instruction is executed. Conversely, when a return from subroutine (RTS) instruction is executed, the top stack item needs to be popped off the stack and placed into the PC.

These observations indicate that, in order to add JSR and RTS instructions to our machine, the PC register needs to be modified. Specifically, a bi-directional interface to the system data bus needs to be added so that the value in the PC can be pushed/popped. Two new control signals need to be added to the PC for this purpose: PLD, for loading the PC with the value on the data bus (popped off the stack when an RTS instruction is executed); and POD, for gating the value in the PC onto the data bus (so that it can be pushed onto the stack when a JSR instruction is executed). A block diagram depicting the modified system is given in Figure 2-26. An ABEL file for the modified PC is given in Table 2-23.



**Figure 2-26** Block diagram of simple computer with subroutine linkage mechanism.

Upon examining the block diagram of the modified system, one might initially be “disturbed” by the fact that the *width* (i.e., number of bits) of the PC register does not match that of data bus and/or memory – here, the PC register is only 5-bits wide, while the memory is 8-bits wide. In practice, though, this is of no consequence – we will simply use the lower 5-bits of the addressed memory location to store the value of the PC when it is pushed onto the stack. In most “real” computers, there is usually a better “match” between the PC and memory width (e.g., 32-bit address space and 32-bit wide memory).

**Table 2-23** Modified PC for subroutine linkage.

```

MODULE pcr

TITLE      'Program Counter with Data Bus Interface'

DECLARATIONS

CLOCK pin;

PC0..PC4 node istype 'reg_D,buffer'; " PC register bits
AB0..AB4 pin; " address bus (5-bits wide)
DB0..DB7 pin; " data bus (8-bits wide)

PCC pin; " PC count enable
PLA pin; " PC load from address bus enable
PLD pin; " PC load from data bus enable
POA pin; " PC output on address bus tri-state enable
POD pin; " PC output on data bus tri-state enable
ARS pin; " asynchronous reset (connected to START)

" Note: Assume PCC, PLA, and PLD are mutually exclusive

EQUATIONS

"      retain state      load from AB      load from DB
PC0.d = !PCC&!PLA&!PLD&PC0.q # PLA&AB0.pin # PLD&DB0.pin
"      increment
# PCC&!PC0.q;
PC1.d = !PCC&!PLA&!PLD&PC1.q # PLA&AB1.pin # PLD&DB1.pin
# PCC&(PC1.q$PC0.q);
PC2.d = !PCC&!PLA&!PLD&PC2.q # PLA&AB2.pin # PLD&DB2.pin
# PCC&(PC2.q$(PC1.q&PC0.q));
PC3.d = !PCC&!PLA&!PLD&PC3.q # PLA&AB3.pin # PLD&DB3.pin
# PCC&(PC3.q$(PC2.q&PC1.q&PC0.q));
PC4.d = !PCC&!PLA&!PLD&PC4.q # PLA&AB4.pin # PLD&DB4.pin
# PCC&(PC4.q$(PC3.q&PC2.q&PC1.q&PC0.q));

[AB0..AB4] = [PC0..PC4].q;
[DB0..DB4] = [PC0..PC4].q;

" Output logic zero on upper 3-bits of data bus
[DB5..DB7] = 0;

[AB0..AB4].oe = POA;
[DB0..DB7].oe = POD;

[PC0..PC4].ar = ARS;
[PC0..PC4].clk = CLOCK;

END

```

We are now ready to outline the steps needed to execute the JSR and RTS instructions. First, we realize there are two fundamental steps associated with performing a JSR: (a) push the return address (the value in the PC register) onto the stack, and (b) jump to the location indicated by the instruction's address field. Step (a) is accomplished in a manner similar to the PSH instruction described in Section 2.9.4: during the first execute cycle, the stack pointer is decremented; during the second execute cycle, the new item (here, the PC) is written to the location pointed to by the SP register. Step (b) is accomplished the same way as the unconditional "jump" instruction (JMP) described in Section 2.9.3: the location at which execution of the subroutine is to commence is simply transferred from the IR to the PC via the address bus. Adding it all up, we find that a total of three execute states are needed to perform a JSR instruction.

By way of contrast, execution of an RTS instruction requires only a single fundamental step: pop the return address off the stack and place it into the PC register. This is really not much different than the "basic pop" instruction (POP) described in Section 2.9.4, except here the destination is the PC rather than the "A" register. Also, because RTS is merely a "pop PC" operation, it can be performed in a single execute cycle, just like the "pop A" (POP) instruction.

**Table 2-24** System control table modifications for subroutine linkage instructions.

Dec. State	Instr. Mnem.	MSL	MOE	MWE	PCC	POA	IRL	IRA	AOE	ALE	ALX	ALY	PLA	POD	PLD	SPI	SPD	SPA	RST
S0	#4	H	H	H	H	H	H												
S1	LDA	H	H					H		H	H								H
S1	STA	H		H				H	H										H
S1	ADD	H	H					H		H									H
S1	SUB	H	H					H		H		H							H
S1	AND	H	H					H		H	H	H							H
S1	HLT	L		L	L			L											
S1	JSR																	H	
S1	RTS	H	H												H	H	H	H	H
S2	JSR	H		H										H				H	
S3	JSR							H				H							H

**Table 2-25** IDMS modifications for subroutine linkage instructions.

```

" System control equations

MSL = RUN.q&(S0 # S1&(LDA # STA # ADD # SUB # AND
    # RTS) # S2&JSR);
MOE = S0 # S1&(LDA # ADD # SUB # AND # RTS);
MWE = S1&STA # S2&JSR;
ARS = START;
PCC = RUN.q&S0;
POA = S0;

PLA = S3&JSR;
POD = S2&JSR;
PLD = S1&RTS;

IRL = RUN.q&S0;
IRA = S1&(LDA # STA # ADD # SUB # AND);
AOE = S1&STA # S2&JSR;
ALE = RUN.q&S1&(LDA # ADD # SUB # AND # RTS);
ALX = S1&(LDA # AND # RTS);
ALY = S1&(SUB # AND);

SPI = S1&RTS;
SPD = S1&JSR;
SPA = S1&RTS # S2&JSR;

RST = S1&(LDA # STA # ADD # SUB # AND # RTS) # S3&JSR;

END

```

The system control table, modified to include the new JSR and RTS instructions, is shown in Table 2-24. An ABEL file for the modified IDMS is given in Table 2-25. Note that, since the JSR consumes all three execute cycles available, it technically “doesn’t matter” whether or not the RST signal is asserted during S3 (since the 2-bit state counter will automatically “wrap around” to S0 when the next clock edge occurs). It’s probably a good idea, though, to show RTS as being asserted on S3, just in case future extensions to the instruction set require a state counter with additional bits.

## 2.9.6 Other Possibilities

Having established the “basic modern conveniences” needed to implement a very simple computer, our imaginations could “go wild” thinking up new instructions and architectural extensions. We could accommodate additional instructions (opcodes) by simply increasing the number of opcode bits (an 8-bit opcode would give us 256

possibilities). And we could incorporate a more reasonably-sized memory by simply increasing the number of address bits. We could add new registers, such as an additional accumulator or an index register, as well as new addressing modes. An index register could be used as a pointer to memory, and facilitate implementation of a variety of new addressing modes. The homework problems included at the end of this chapter will allow us to explore some useful extensions.

## 2.10 Summary and References

In this chapter we have introduced the design and implementation of a simple computer and progressively embellished it with a number of extensions. In addition to reviewing a “top-down, bottom-up” strategy for designing digital systems, we have also provided a “bridge” between the basic digital logic design topics reviewed in Chapter 1 and the microcontroller-oriented topics that commence in Chapter 3.

There are a number of texts that delve into the myriad of topics associated with computer architecture and design, written at a variety of levels. One of the best (and most widely used) introductory texts is Patterson and Hennessey’s *Computer Architecture: The Hardware-Software Interface* (Morgan Kaufmann). Their earlier text, *Computer Architecture: A Quantitative Approach* (Morgan Kaufmann), is an authoritative “advanced” text on the subject, used in numerous graduate programs.

Other highly regarded texts on computer architecture include Mano’s *Computer Engineering Hardware Design* (Prentice-Hall), Stalling’s *Computer Organization and Architecture* (Macmillan), Haye’s *Computer Architecture and Organization*, and Hamacher’s *Computer Organization*.

One of the best sources for unbiased reviews of the “latest and greatest” microprocessors is *Microprocessor Report* – a subscriber-supported periodical published by Cahners Electronics Group. Another excellent source of information on recent developments in microprocessor architecture is *IEEE Micro*, a publication of the IEEE Computer Society.

For information on embedded microcontrollers and applications, *Circuit Cellar Inc.* magazine is the source of choice. Web sites of the major manufacturers (Intel, Motorola, Texas Instruments, Hitatchi, etc.) continue to be the best sources for detailed information concerning specific microprocessors and microcontrollers.

1. Modify the section of the IDMS source file, below, to provide up to 7 execute cycles (in addition to a single fetch cycle). The original ABEL file is given in Tables 2-18 and 2-19.

```
MODULE idmsr

TITLE 'IDMS with 7 Execution States'

DECLARATIONS

" State counter
SQA node istype 'reg_D,buffer'; " low bit of state counter
SQB node istype 'reg_D,buffer';
SQC node istype 'reg_D,buffer'; " high bit of state counter

" Synchronous state counter reset
RST node istype 'com';

" RUN/HLT state
RUN node istype 'reg_D,buffer';

" Decoded state definitions

S0 =
S1 =
S2 =
S3 =
S4 =
S5 =
S6 =
S7 =

EQUATIONS

" State counter
" If RUN negated or RST asserted, state counter is reset

SQA.d =
SQB.d =
SQC.d =
```

2. The possibility of an alternate stack convention (using the SP register as a pointer to the next available location) was described in Section 2.9.4. Show how the system control table for the PSH and POP instructions would change if this alternate convention were used. Use the minimum number of execute states possible for each instruction.

Dec. State	Instr. Mnem.	MSL	MOE	MWE	PCC	POA	IRL	IRA	AOE	ALE	ALX	ALY	SPI	SPD	SPA	RST
S0	%4	H	H		H	H	H									
S1	LDA	H	H					H		H	H					
S1	STA	H		H				H	H							
S1	ADD	H	H					H		H						
S1	SUB	H	H					H		H		H				
S1	AND	H	H					H		H	H	H				
S1	HLT	L		L	L				L							
S1	PSH															
S1	POP															
S2	PSH															
S2	POP															

3. Given that a practical program has a balanced set of PSH and POP instructions (i.e., each PSH is “balanced” by a POP), are there any advantages or disadvantages inherent in the alternate stack convention used in Problem 2-2?
- 
- 
-

4. The possibility of an alternate stack convention (using the SP register as a pointer to the next available location) was described in Section 2.9.4. Show how the system control table for the JSR and RTS instructions would change if this alternate convention were used. Use the minimum number of execute states possible for each instruction.

Dec. State	Instr. Mnem.	MSL	MOE	MWE	PCC	POA	IRL	IRA	AOE	ALE	ALX	ALY	PLA	POD	PLD	SPI	SPD	SPA	RST
S0	<del>JSR</del>	H	H	H	H	H													
S1	LDA	H	H				H		H	H									
S1	STA	H		H				H	H										
S1	ADD	H	H					H		H									
S1	SUB	H	H					H		H		H							
S1	AND	H	H					H		H	H	H							
S1	HLT	L		L	L		L												
S1	JSR																		
S1	RTS																		
S2	JSR																		
S2	RTS																		
S3	JSR																		
S3	RTS																		

5. Given that a practical program has a balanced set of JSR and RTS instructions (i.e., each JSR is “balanced” by a RTS), are there any advantages or disadvantages inherent in the alternate stack convention used in Problem 2-4?

---



---



---

6. The 8-bit ALU designed in Section 2.7.4 employs a simple ripple-carry topology. Modify the ABEL source file for the adder/subtractor based on the use of two 4-bit carry look-ahead adder blocks employing a “group ripple”. The original ABEL file is listed in Tables 2-5, 2-6, and 2-7.

```
" Declaration of intermediate equations
" Generate functions

GA[0..3] = X[0..3]&Y[0..3];
GB[0..3] = X[4..7]&Y[4..7];

" Propagate functions
PA[0..3] = X[0..3]$Y[0..3];
PB[0..3] = X[4..7]$Y[4..7];

" Least significant bit carry-in (0 for ADD, 1 for SUB => ALY)
CIN = ALY;

EQUATIONS

S0 = PA0$CIN;
S1 = PA1$CA0;
S2 = PA2$CA1;
S3 = PA3$CA2;
S4 = PB0$CA3;
S5 = PB1$CB0;
S6 = PB2$CB1;
S7 = PB3$CB2;

" CLA equations (two 4-bit blocks, cascaded together)

CA0 =
CA1 =
CA2 =
CA3 =

CB0 =
CB1 =
CB2 =
CB3 =
```

7. Part of the ABEL file for the “final version” of the program counter (PC) register used in the simple computer is shown below (*reduced to 4 bits*). Add the equations necessary to complete this file, *given the declarations provided*. Recall that it is interfaced to both the Address Bus and the Data Bus, and uses the following control signals:

PCC – program counter increment enable  
 PLA – program counter load from Address Bus enable  
 POA – program counter tri-state output enable for Address Bus  
 PLD – program counter load from Data Bus enable  
 POD – program counter tri-state output enable for Data Bus  
 ARS – program counter asynchronous reset

```
MODULE pc4bit
TITLE '4-bit Version of Program Counter'

DECLARATIONS
PC0..PC3 node istype 'reg'; "PC bits - declared as internal nodes
AB0..AB3 pin istype 'com'; "Address Bus pins
DB0..DB3 pin istype 'com'; "Data Bus pins
PCC,PLA,POA,PLD,POD,ARS,CLOCK pin; "Control signals

EQUATIONS
```

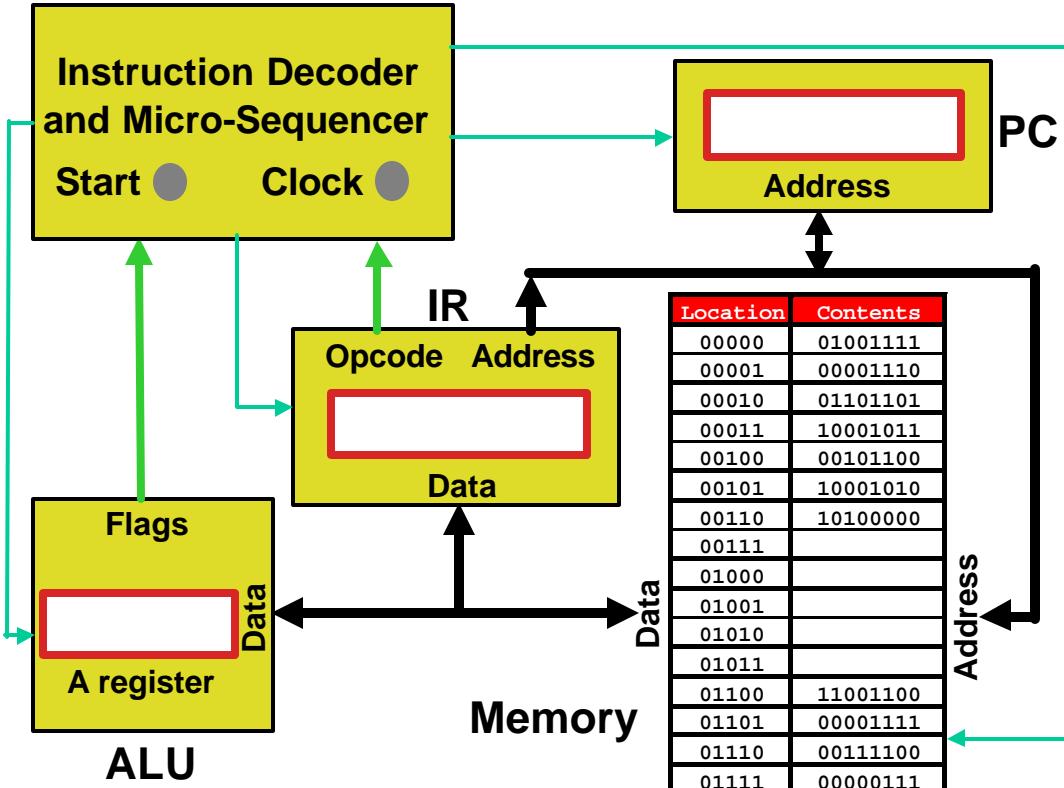
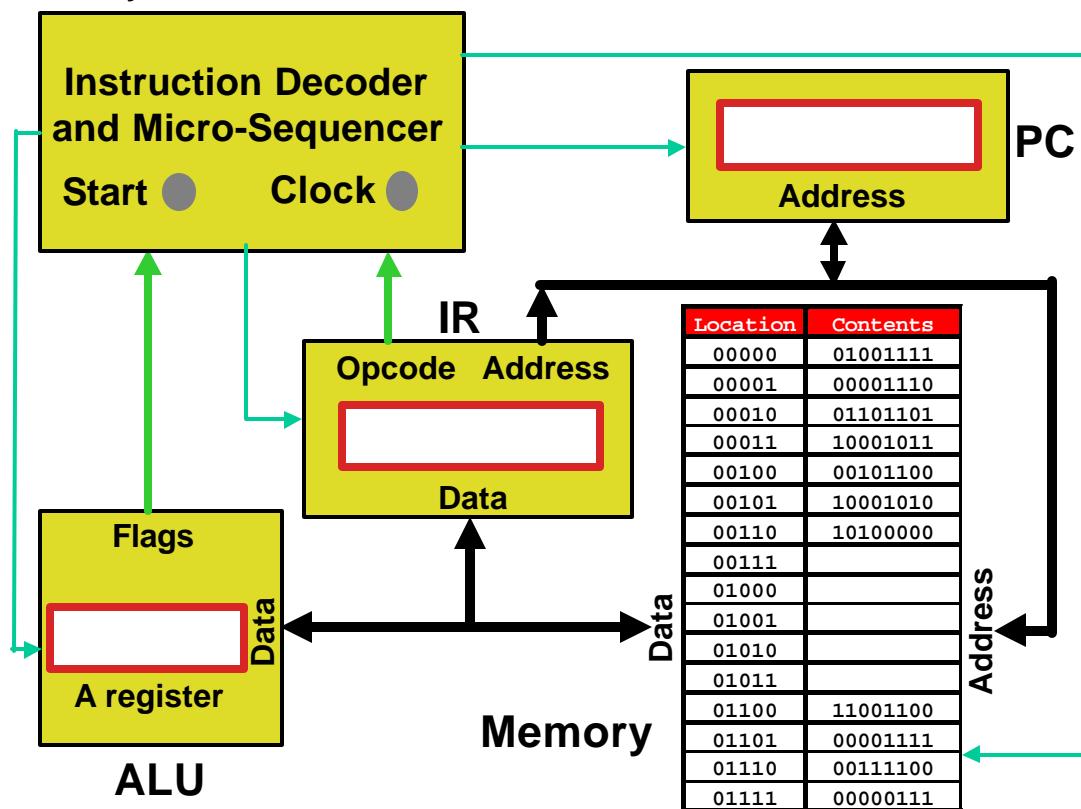
8. Assume the "simple computer" instruction set is *changed* to the following:

OPCODE	MNEMONIC	FUNCTION
000	ADD <i>addr</i>	Add contents of <i>addr</i> to contents of A register
001	SUB <i>addr</i>	Subtract contents of <i>addr</i> from contents of A register
010	LDA <i>addr</i>	Load A register with contents of location <i>addr</i>
011	AND <i>addr</i>	AND contents of <i>addr</i> with contents of A register
100	STA <i>addr</i>	Store contents of A register at location <i>addr</i>
101	HLT	(Halt) – Stop, discontinue execution

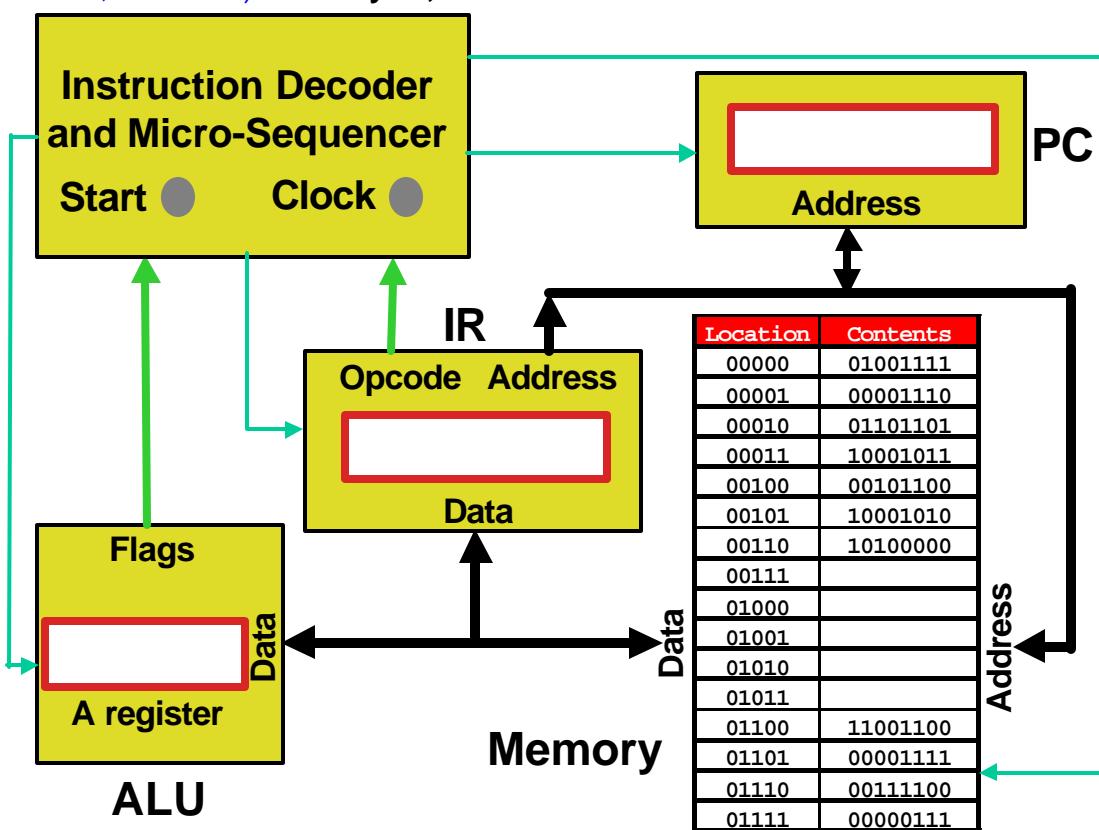
Complete the instruction trace worksheets that follow for the fetch and execute cycles of the program stored in memory (up to, but not including, the HLT instruction). Note that you will have to *disassemble* the program stored in memory to determine what it is doing.

(Problem 8, continued) Fetch Cycle, Instruction at 00000:

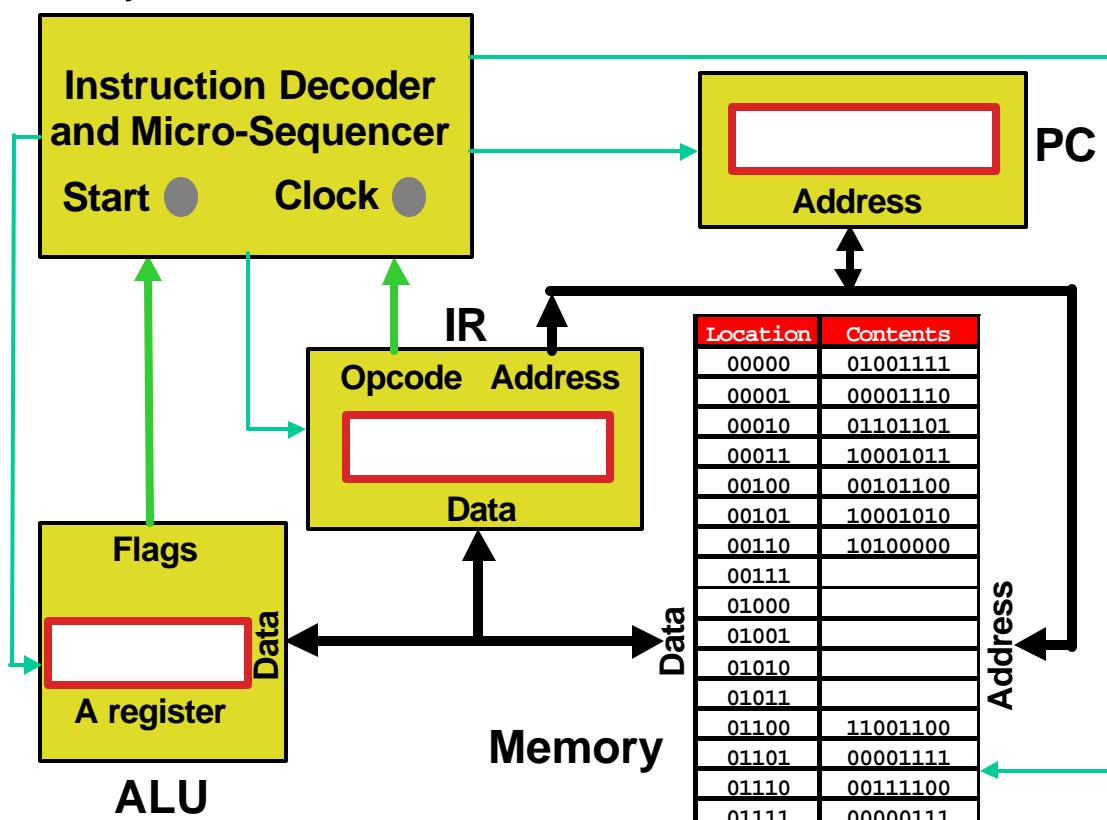
Execute Cycle, Instruction at 00000:



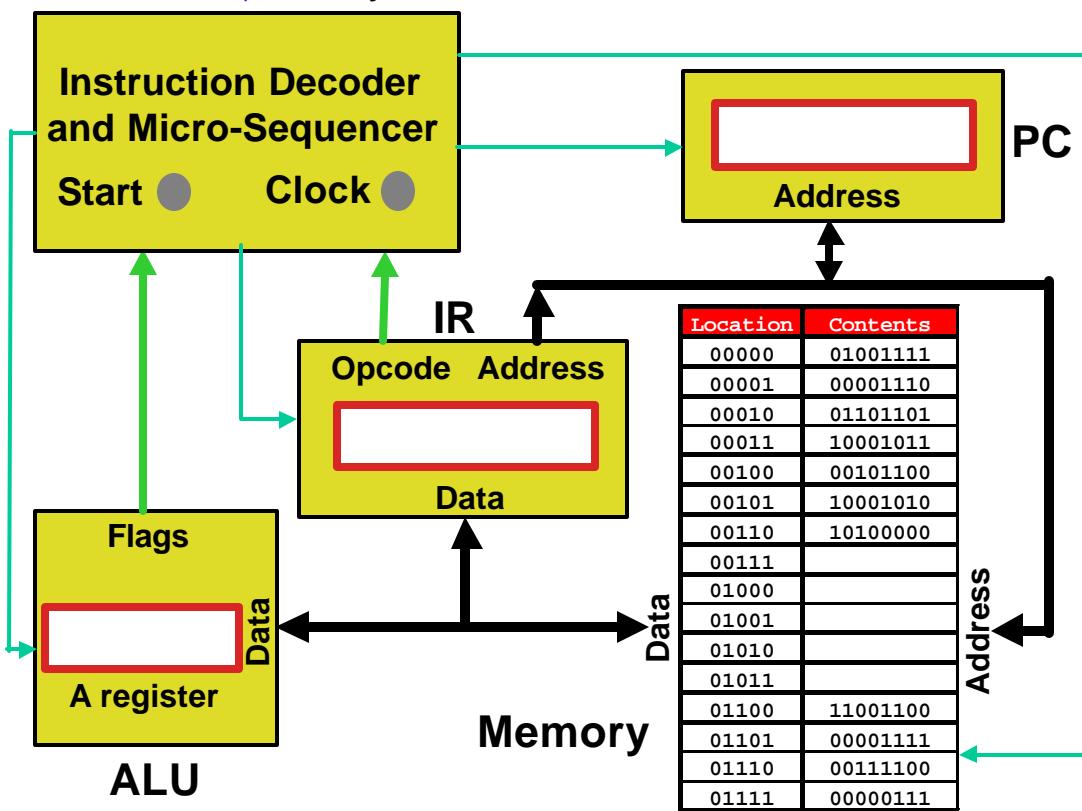
(Problem 8, continued) Fetch Cycle, Instruction at 00001:



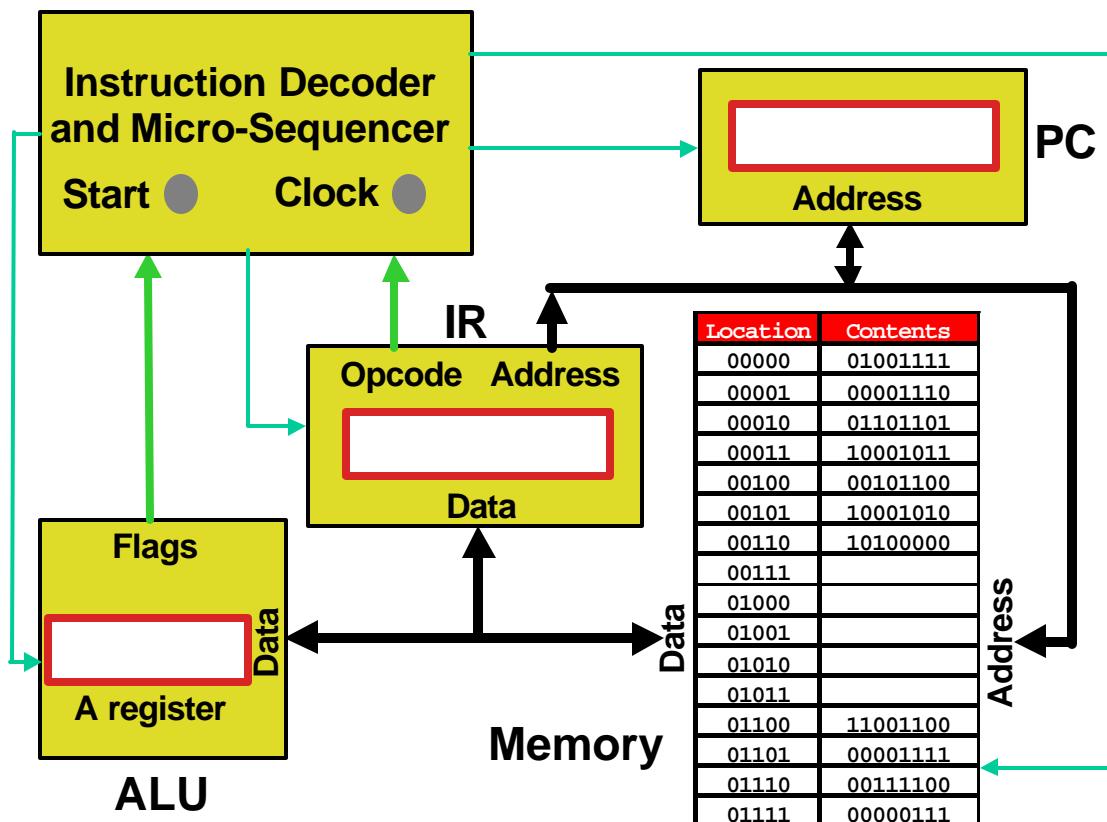
Execute Cycle, Instruction at 00001:



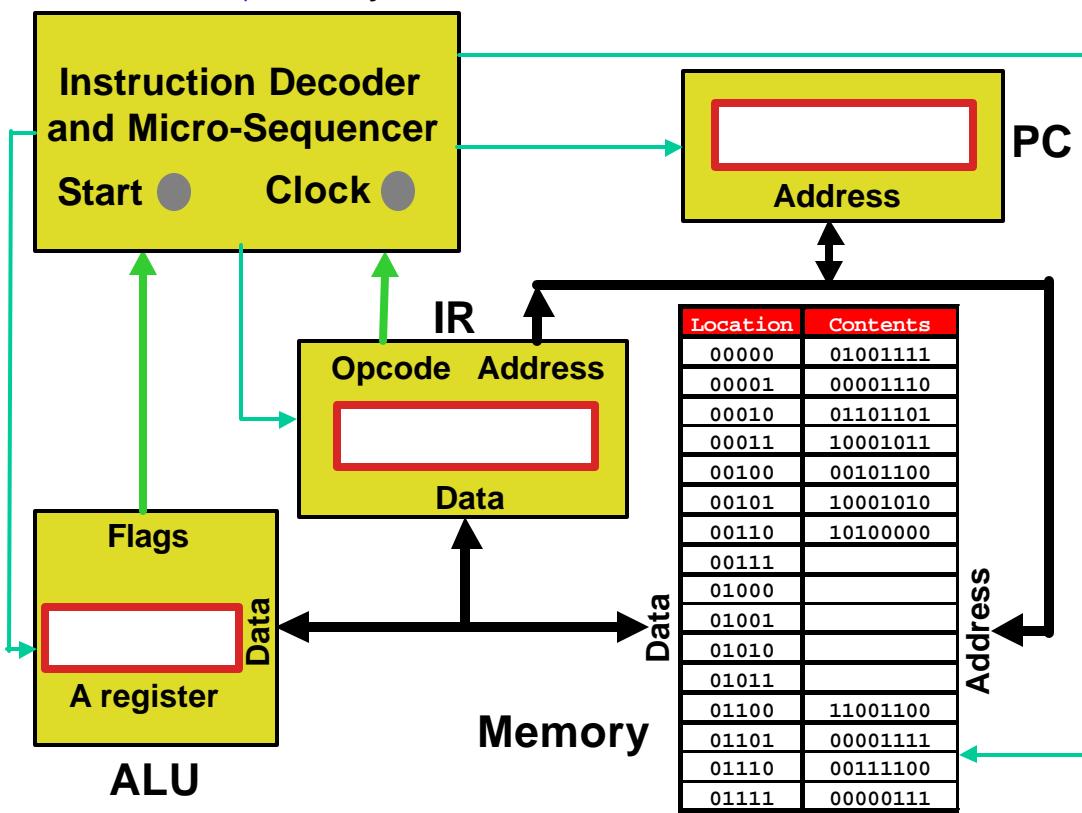
(Problem 8, continued) Fetch Cycle, Instruction at 00010:



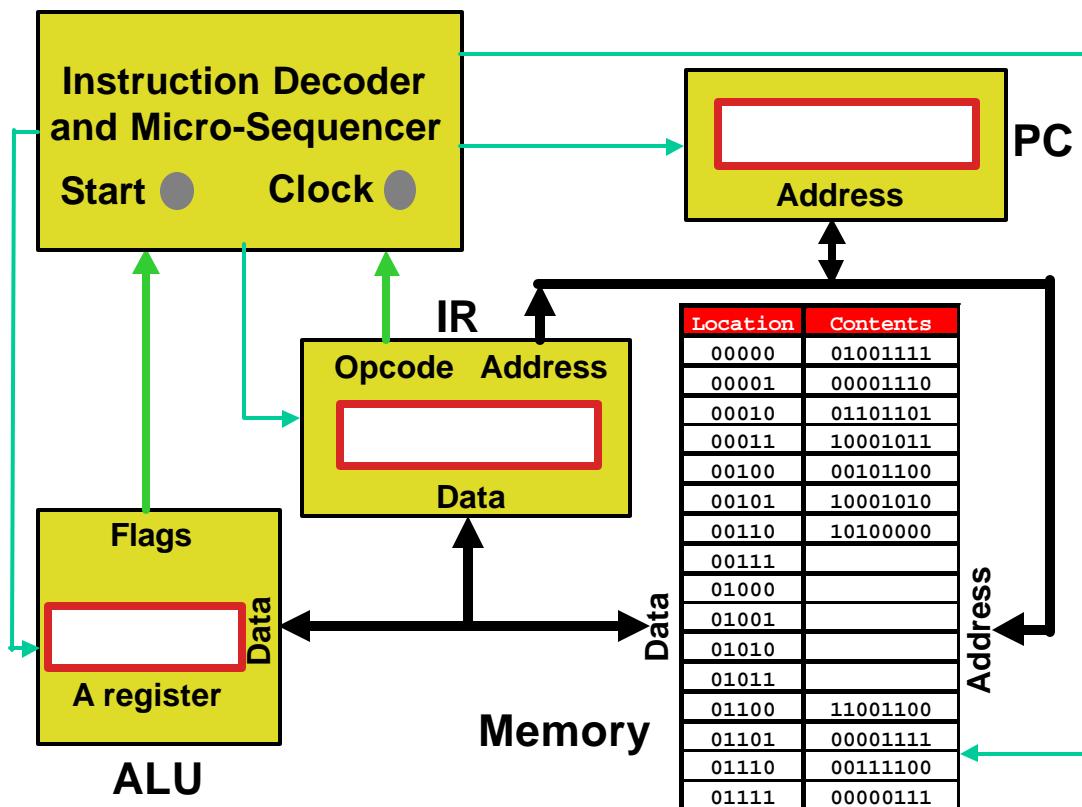
Execute Cycle, Instruction at 00010:



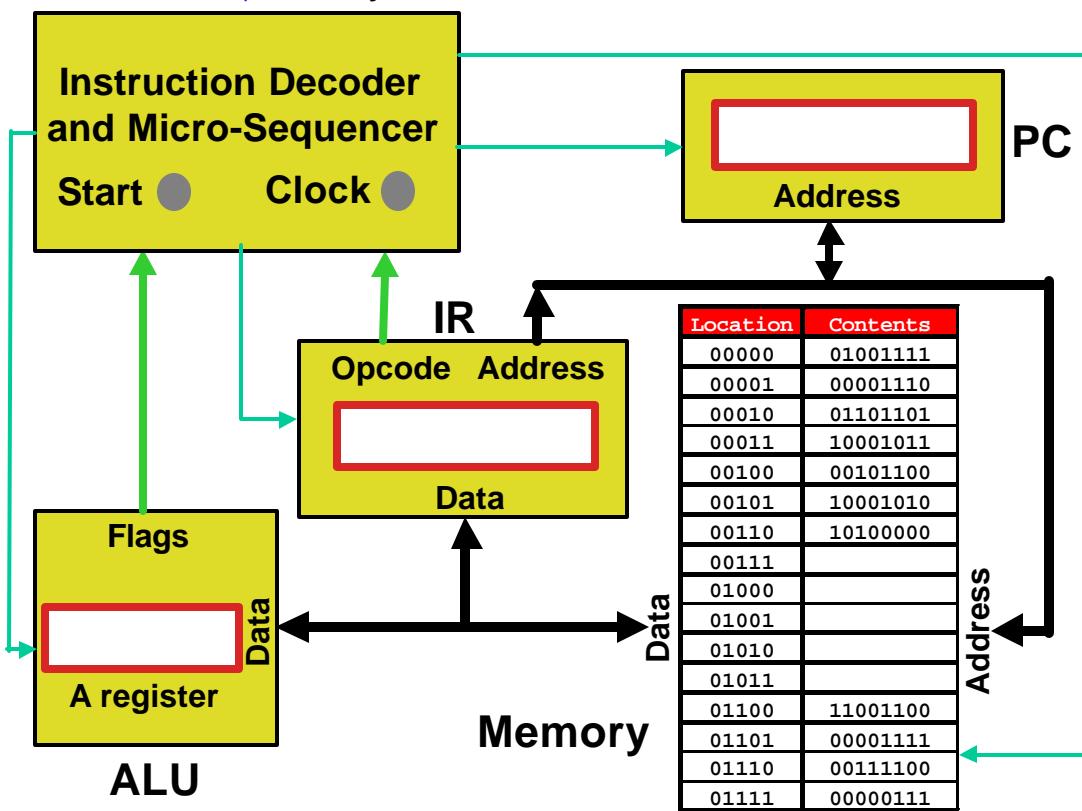
(Problem 8, continued) Fetch Cycle, Instruction at 00011:



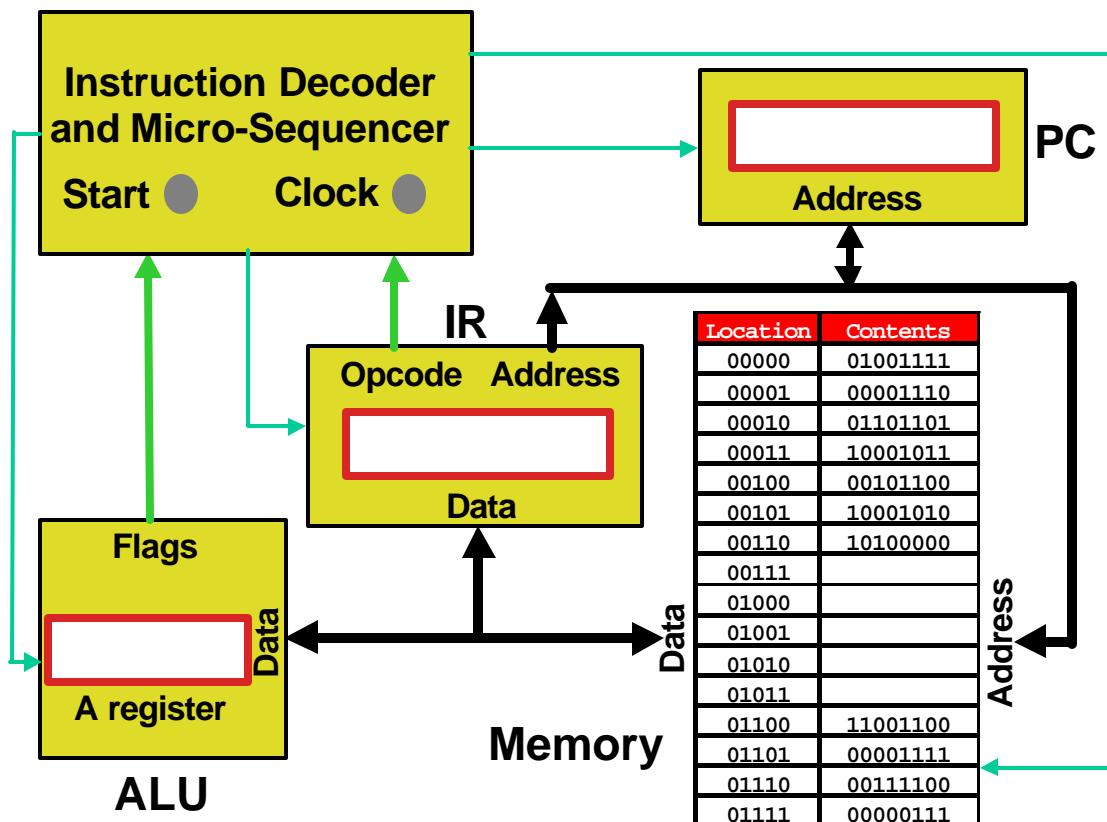
Execute Cycle, Instruction at 00011:



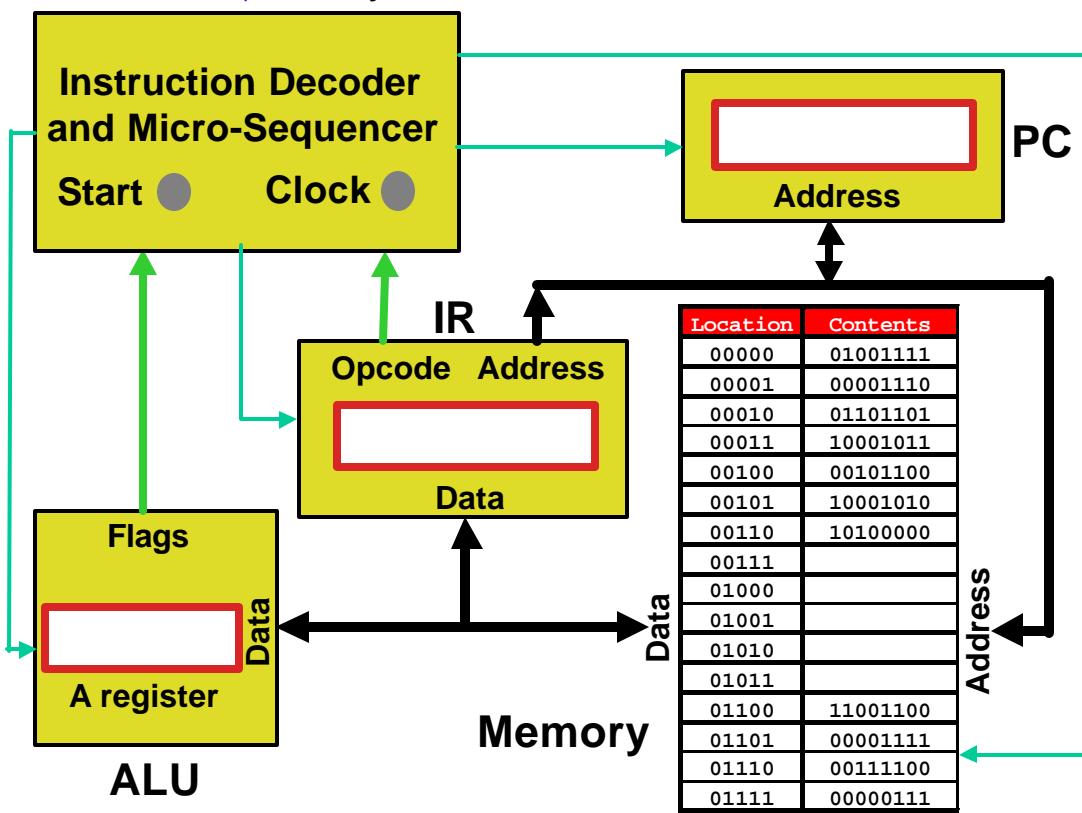
(Problem 8, continued) Fetch Cycle, Instruction at 00100:



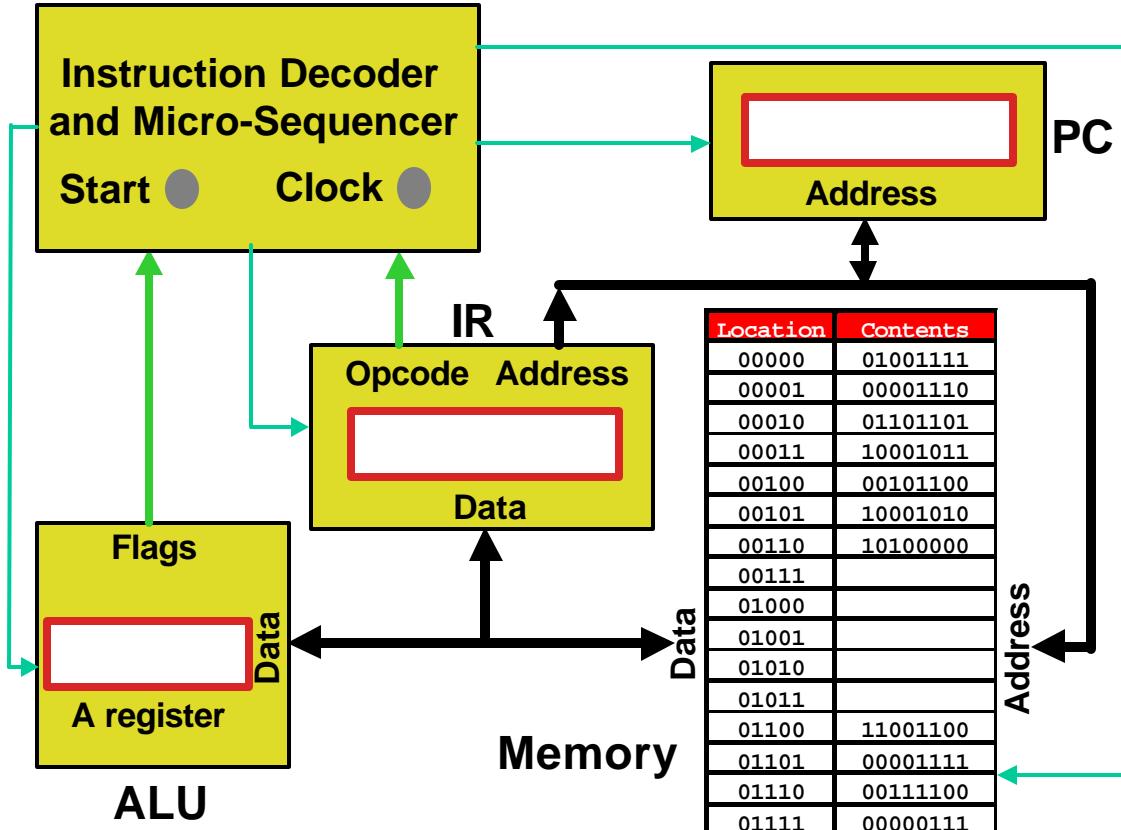
Execute Cycle, Instruction at 00100:



(Problem 8, continued) Fetch Cycle, Instruction at 00101:



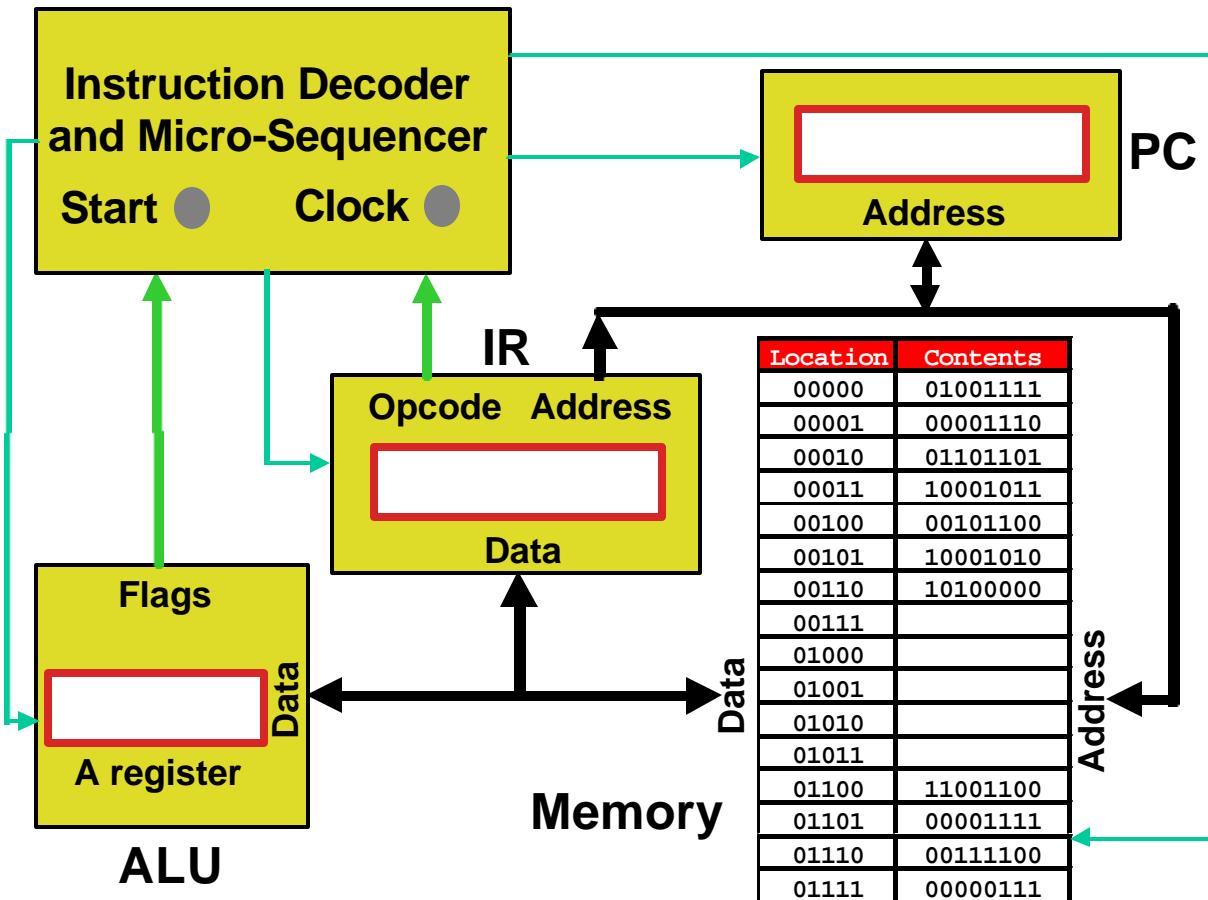
Execute Cycle, Instruction at 00101:



9. Assume the simple computer instruction set has been changed to the following:

Opcode	Mnemonic	Function Performed
0 0 0	ADD addr	Add contents of <i>addr</i> to contents of A
0 0 1	SUB addr	Subtract contents of <i>addr</i> from contents of A
0 1 0	LDA addr	Load A with contents of location <i>addr</i>
0 1 1	AND addr	AND contents of <i>addr</i> with contents of A
1 0 0	STA addr	Store contents of A at location <i>addr</i>
1 0 1	HLT	Halt – Stop, discontinue execution

On the instruction trace worksheet, below, show the *final result* of executing the program stored in memory *up to and including* the HLT instruction.



## CHAPTER 3

### INTRODUCTION TO MICROCONTROLLER ARCHITECTURE AND PROGRAMMING MODEL

A good “working analogy” useful in the study of computer instruction sets can be gleaned from a master carpenter, such as Norm Abram of *This Old House* and *New Yankee Workshop* fame. Norm would never start a construction project without first mastering the “tools in the toolbox” – an apt description of a machine’s instruction set and programming model. He would not only figure out how each tool works, but also *practice* using it before starting a project that required use of that tool. Further, Norm would not use any woodworking tool without careful adherence to safety rules, e.g., wearing safety glasses and keeping protective blade guards in place. We need to develop a similar posture as we write programs, protecting ourselves from software errors that might cause “bits to fly all over the place” – either figuratively or literally (as we will discuss in Chapter 10 when we consider ethical ramifications of product malfunctions induced by software errors).

*Norm Abram*

*tools in the toolbox*

Norm would also tell us that before, say, using a compound mitre saw or a biscuit joiner, we should *practice* (and become *good at*) making “straight cuts” with a simple table saw. Stated another way, we should master an instruction set and basic program structures before we “move up” to programming in a high-level language. Programming, like carpentry, is a *profession skill* – a skill that cannot be learned by merely reading about it or watching someone else do it. The lab experiments and homework exercises that accompany this chapter will provide an opportunity for developing these skills.

*professional skill*



<http://www.pbs.org>

**Figure 3-1** The author’s “hero” – master carpenter Norm Abram.

### 3.1 Differing World Views

A *personal computer* is perhaps the first thing that comes to mind when the word “microprocessor” is mentioned. Thanks to commercial advertising on national television and the ubiquity of PCs, virtually everyone knows what “Intel Inside™” means. If there’s one thing the much-ballyhooed “Y2K Crisis” accomplished, though, it was to make the general populace aware that embedded microprocessors are literally everywhere. The fundamental differences between microprocessors used in personal computers and those used for embedded applications are not universally appreciated, however. In fact, two basic “world views” regarding the role of microprocessors are applicable. What might be called the “general-purpose view” is that a microprocessor is an integral part of a machine that runs “shrink-wrapped” software (or, on which application programs can be written and run, most often using a high-level language or development tool). The “embedded view”, by way of contrast, is that microprocessors (or microcontrollers) are a basic building block of modern digital system design – in particular, of “intelligent” products.

Calling a computer “general-purpose” implies *user programmability*. It also implies support for an operating environment that fosters such use. Virtually all general-purpose application programs run under a *time-sharing* operating system (e.g., variants of Unix or Windows™), where the “processor’s attention” is multiplexed among multiple tasks (which is why these systems are sometimes referred to as *multi-tasking* or *multi-programming*). The amount of time it takes an application to respond to user input (*response time* or *latency*) is generally not considered “critical” in nature. Stated another way, Windows™ “doesn’t care” if the mouse pointer becomes “sluggish” in its response while the processor focuses on a more “important” activity, such as Word™’s insistence on “correcting” the author’s colorful (and sometimes questionable) use of the English language.

Embedded applications, on the other hand, are by definition *non-user-programmable*; as such, they are often referred to as “turn-key” systems (i.e., turn the key “on” and they run). Many (but not all) embedded applications are *real time* in nature – meaning they must respond within certain time constraints to external events (this is sometimes referred to as *mission critical* timing). For example, when an automobile’s antilock brake mechanism is activated, the microcontroller in charge must immediately begin to pulse the brake cylinders at a periodic rate and continue to do so until the vehicle stops. This task cannot be “rolled out” while the driver surfs the wireless web for the best buy on snowshoes.

There are several reasons why the distinction between general-purpose and embedded applications of microprocessors is important. First,

*personal  
computer*

*general-purpose  
world view*

*embedded  
world view*

*user programmability*

*time-sharing OS*

*multi-tasking  
multi-programming  
response time  
latency*

*non-user-programmable*

*turn-key system  
real time  
mission critical*

different architectural and/or organizational characteristics of microprocessors can make them more (or less) suited for the target application. One of the most challenging tasks in embedded system design is matching the requirements of a target application with the computational and peripheral interface capabilities of a candidate microcontroller. Unlike the “general-purpose” world, *more* (processing power, clock speed, I/O pins, integrated peripherals, etc.) is *not necessarily better* – rather, it is the *closeness of the “match”* between processor capability and application requirements that is key. Jaded by the impact of Moore’s Law on personal computing, this reality is hard for “beginning students” to comprehend and appreciate.

Second, to come to the conclusion that, say, a 1.5 GHz Pentium IV is a “better” processor than an 8 MHz 68HC12 – without specifying the intended *application domain* – is nonsensical. Simply stated, one would never use a 68HC12 as the “brains” of a personal computer and never use a Pentium III to control a microwave oven. Surprising as it may sound, some of the 4-bit microcontrollers currently available are “plenty powerful” for many consumer products that come to mind, such as appliance controllers, garage door openers, ceiling fan controllers, answering machines, feature phones, TV and radio tuners, etc. There are some applications, however, where the distinction is a bit less clear. For example, a point-of-sale terminal could be built around either a microcontroller like the 68HC12 or a (low-end) Pentium microprocessor (or one of its “x86” predecessors targeted for embedded applications). *The “goodness” or “badness” of a particular processor can only be evaluated in the context of a target application.*

*more is not  
necessarily better*

*Moore’s Law*

*application domain*

### A Third World View?

A relatively new “world view” that is emerging (some would say being thrust upon us) is that the personal computer is the “basic building block” of modern embedded system design. Not a conventional desktop personal computer, but a “stripped down” version running an operating system geared toward embedded applications, like Windows CE™ or variants of Linux. For the point-of-sale terminal cited in the text, one could argue that certain forms of them look “a lot like a PC” – they have a video display, a keyboard, and perhaps a bar code scanner (instead of a mouse). So, the argument goes, why not just use the “guts” of a PC as the basic building block for this device and write the application code using PC-like tools that run under a PC-like operating system? Great idea for this particular application. But what if a simpler, higher volume unit is needed of the “may I take your order” genre, where a keypad, LCD (liquid crystal display), and cash drawer release solenoid are the only forms of I/O? Here it is much harder to justify dedicating an entire PC to each terminal. As we say in the industry, some “food for thought”...

What, then, are the characteristics that distinguish processors targeted for general-purpose applications versus those targeted for embedded applications? One reason we wish to address this question is to provide rationale for choosing the “most appropriate” processor to “cut our digital teeth” on. Another reason for addressing this question is to provide a context for understanding why processors targeted for different applications are necessarily different. The discussion which follows is intended to provide a basis for this understanding. It is not, however, intended as a detailed presentation on the characteristics of general-purpose systems – complete treatment of this subject alone would fill an entire textbook!

### 3.2 Characteristics That Distinguish Microprocessors

Processors that are primarily intended for embedded applications generally possess the following characteristics. Most notably, perhaps, is they are often “smaller” (in terms of *bit width* and *address space*) than their general-purpose counterparts. Since interrupts are a “way of life” in event-driven systems, a *flexible interrupt structure* is a key characteristic of control-oriented microprocessors. And since interrupts occur frequently in event-driven systems, the *context switching overhead* must necessarily be low – generally implying the need for relatively small register sets. Because embedded systems typically involve a wide variety of interfaces, processors targeted for such applications typically provide a mixture of both digital and analog I/O on-chip. A small amount of on-chip program memory (ROM) and “scratchpad” RAM are usually sufficient, since many embedded applications are relatively “simple” in nature. Finally, due to the “real time” nature of many embedded applications, the amenability of assembly-level “patching” of time-critical code segments is important.

General-purpose applications, run under a time-sharing operating system, generally require processors with completely different characteristics and built-in features than those used for embedded applications. Due to the multi-tasking, multi-programming nature of general-purpose systems, support for *virtual memory* is typically built into the processor and its instruction set. Simply put, virtual memory provides an address space for each program or process that is not constrained by the *physical* (or actual) memory installed in the system. For example, even though a personal computer may only have 128 megabytes (MB) installed in it, a given program can have as much as a *terabyte* ( $2^{40}$  MB) of address space available to it. Coupled with protection mechanisms, virtual memory is implemented using a *hierarchy* of memory subsystems, of varying size and speed. Closest to the processor – usually on-chip – is a high-speed *cache memory* (which itself may consist of more than one level). The next level typically consists of comparatively slower dynamic RAM chips. The highest (and slowest) level is implemented with a *mass storage device*,

*bit width*  
*address space*

*flexible interrupt structure*

*context switching overhead*

*virtual memory*

*physical memory*

*memory hierarchy*

*cache memory*

*mass storage device*

such as a hard disk drive. The “illusion” of a *virtually limitless* private address space is accomplished by loading – on an “as needed” (or *demand*) basis – portions of the application and its data set that are needed at a particular instant. This *demand paging* process is managed by the time-sharing operating system: when a block of code (or data) needed is *not present* in memory, the task is “rolled out” while the code/data is retrieved from the “next higher” level(s) of the memory hierarchy. While this *page fault* is being serviced, the next task in the operating system’s process queue is started.

*demand paging*

*page fault*

Another major difference between processors targeted for general-purpose applications and embedded applications is I/O. For general-purpose systems, the *main* form of I/O is either memory-to-memory, memory-to-disk, or memory-to-network. Further, the CPU rarely “directly” participates in these I/O operations; instead, they are “delegated” to a special-purpose auxiliary processor called a *direct memory access* (DMA) controller. To perform a block transfer, the main processor simply tells the DMA controller the starting addresses of the source and destination blocks along with the size (byte count) of the transfer. For example, when the operating system wishes to update the graphics display, the DMA controller is told to copy the contents of the display buffer (in memory) to the graphics controller. The main processor can continue to execute out of its on-chip cache memory while the DMA controller uses the external address and data buses to complete the data transfer.

*direct memory access  
(DMA) controller*

Because high-level language compilation can be more effectively optimized if a number of “general-purpose” registers are available in the programming model, processors targeted for general-purpose applications often sport large register sets (where “large” is *at least eight*, and in most cases 16 or 32). The larger the register set, however, the greater the context switching overhead – thus impacting system latency. For a time-sharing operating system, though, the context switching overhead is of little consequence, since a task switch typically occurs every 5 milliseconds (i.e., at a 200 Hz rate). Since context switches are relatively infrequent (and the processing is typically not “mission critical” in nature), the increased overhead of saving and restoring large register sets is inconsequential.

*large general-purpose  
register set*

Also, because compilers are much better than humans at optimizing code targeted for large-register-set processors, assembly language patching of general-purpose application code is a practice that has largely been abandoned. Any remaining skeptics need look no further than optimized MIPS code to verify this claim – trying to “patch” this kind of code usually does more “harm” than good!

*assembly language  
patching*

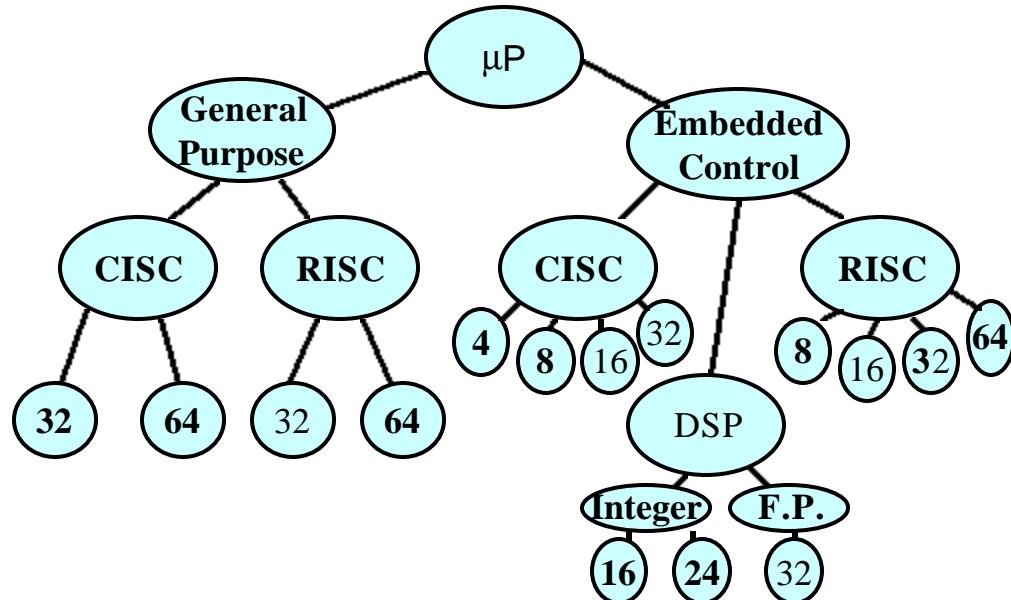
One last, but very important, distinction between processors targeted for general-purpose versus embedded applications is the “world view” of interrupts. In event-driven embedded systems, interrupts are a way of life; in general-purpose applications, they are viewed as more of an “irritation”, often (but not always) associated with something “bad” happening – e.g., “this program has performed an illegal operation and is being shut down”.

### 3.3 Taxonomy of Microprocessors

The taxonomy of processors depicted in Figure 3-2 helps put the variety of microprocessors and microcontrollers currently available into perspective. Within the major categories of “General Purpose” and “Embedded Control”, microprocessors can be further subdivided based on instruction set architecture and ALU bit-width. The “classic” classifications based on instruction set architecture are: complex instruction set computer (CISC) and reduced instruction set computer (RISC). To help understand this distinction, a brief “history lesson” is in order.

*CISC – complex instruction set computer*

*RISC – reduced instruction set computer*



**Figure 3-2** Taxonomy of Microprocessors.

The burgeoning complexity of microprocessors in the early 1980's gave rise to the “less is best” RISC mentality. The underlying principle was that a “less complex” microprocessor chip could run faster – so much so that it could run a program several times faster than a comparable CISC microprocessor, despite its lack of “powerful” instructions and addressing modes. Instead of implementing complex, multi-cycle instructions in hardware, the burden for this functionality was shifted to software. An important key requisite to code optimization was restricting memory references to “load” and “store” instructions (hence the name *load-store*

*load-store architecture*

architecture) – all other instructions (add, subtract, AND, OR, etc.) were restricted to operands contained in (and destined for) registers. The chip real estate vacated by removing large microcode ROMs common in CISCs was devoted to hardware resources that would help an optimizing compiler, such as large register sets and “register windowing” techniques to facilitate subroutine linkage. While less “compact” than a comparable CISC program, the simplicity afforded by fixed-field decoding and simple addressing modes made single-cycle execution of RISC instructions a possibility.

*microcode ROM*

*register windowing*

To be a “true RISC” back then required adherence to some rather Draconian architectural tenets: no more than 40 fixed-length, fixed-field instructions; no more than 4 addressing modes; and strictly load-store. Most so-called “RISC” machines today, however, can only be identified as such based on the last characteristic. Other than being load-store architectures, current RISC machines sport hundreds of instructions, numerous addressing modes, variable-length instructions, and non-fixed fields. Apparently concerned by this deviance from the tenets set in place by the “founding fathers” of RISC, the designers of the IBM Power architecture suggested that the acronym be changed to stand for “reduced instruction set cycles”.

*reduced instruction  
set cycles*

### High Water Mark of Complexity

*Microprocessors have become increasingly complex since their inception in the early 1970s. Perhaps a “high water mark” of complexity was the ill-fated Intel iAPX 432, that company’s attempt in 1981 to introduce the world’s first “32-bit mainframe” microprocessor. Not only did the iAPX 432 sport a sophisticated virtual memory management scheme, but it also had bit-variable length instruction opcode and operand fields. When Intel finally produced a working chip set two years later, their competitors – which included Motorola, National, and Zilog – had all produced viable 16-bit microprocessors with an inkling of virtual memory support. The problem for Intel was that the smaller competing processors were several times faster than the iAPX 432. The fate of this ambitious device was unceremoniously doomed.*

While RISCs were gradually becoming more CISC-like during the late 1980’s and early 1990’s, the world’s “most popular” CISC architecture (Intel x86) was adopting “RISC-like principles” in its design. Advances in micro-architecture and process technology have since subsumed the RISC-CISC performance debate. In essence, most contemporary microprocessors (including many microcontrollers) are in reality “CRISC” machines – *complex machines with reduced instruction set cycles*.

*CRISC*

As one might guess, much has been written about RISC versus CISC tradeoffs – a number of “classic” articles on this subject are listed at the end of this chapter. The brief account provided here is intended only to provide a context for understanding the taxonomy of microprocessors depicted in Figure 3-2. Referring once again to this figure, we note that for general-purpose applications, 32- and 64-bit machines are the basic variants currently available (the earliest devices in this category were 16-bit machines, but these are no longer considered viable for most of today’s time-sharing operating systems).

In the embedded control domain, however, there is much greater variety, including a new category: digital signal processor (DSP) devices. The primary characteristic that distinguishes a DSP from a “generic” microcontroller is the amount of hardware resources devoted to performing the “multiply-and-accumulate” (MAC) operation – a staple of most signal processing algorithms – as quickly as possible. Here there are two basic categories: integer (also called fixed point), of which there are 16- and 24-bit variants; and floating point, most of which are 32-bit devices.

*digital signal processor (DSP)*

*multiply-and-accumulate (MAC)*

### 24-bit Wonder

*In the digital world where “powers of two” rule, a 24-bit processor may seem a bit strange. What numeric-oriented applications might best be served by 24-bits of resolution? If 16-bits is insufficient for such an application, why not move up to 32-bits of resolution as the next logical choice? It turns out that the application – and it’s a big one – for which 24-bits “rule” is digital audio. So-called “CD quality” audio requires 16-bits of resolution, providing a theoretical dynamic range of 96 dB. To maintain this dynamic range in the face of various “audio processing” algorithms (filtering, equalization, reverberation, etc.), “extra bits” are required to represent intermediate results – especially in a fixed point processor. The 24-bits of resolution available in popular audio-oriented digital signal processors provide the number of bits necessary for CD-quality sound.*

CISC-style devices targeted for embedded applications range from 4- to 32-bits wide. Until recently, 4-bit devices of this genre were the highest volume parts – of *all* microprocessors and microcontrollers on the market. (Note, however, that *highest volume* does not imply *highest profit* – competition and small margins yield relatively small profits compared with, say, the “latest and greatest” microprocessors targeted for general-purpose systems, which typically enjoy a much higher “markup”.) Larger 8- and 16-bit CISC microcontrollers are the current overall volume giants, with 32-bit devices gaining ground. Many of the 16- and 32-bit CISC microprocessors targeted for embedded applications are actually “re-

*highest volume  
highest profit*

purposed" previous-generation devices formerly targeted for general-purpose systems (e.g., the Intel 386EC and 486EC as well as the Motorola 68000EC and 68020EC devices). Together, this "bubble" of 4- to 32-bit CISC devices on the taxonomy diagram represents a mammoth sales volume of components.

*re-purposed*

RISC-style devices targeted for embedded applications range from 8- to 64-bits wide. One of the newer players on the block, Microchip Corporation, has become famous for its "PIC" line of 8-bit microcontrollers. This popular, wide-ranging series of devices is the closest thing to "true RISC" currently available: they have small instruction sets, few addressing modes, small on-chip memories, and simple on-chip peripherals. Further, some of the PIC microcontrollers are housed in packages with as few as 8 pins. At the other end of the spectrum, a 64-bit MIPS RISC-style processor is very popular as well – anyone who has never heard of Nintendo 64™ either lives in Palm Beach County, or doesn't have small children! As was the case for "retired" 32-bit CISC processors, their RISC-style counterparts have also been "re-purposed" for embedded applications.

*PIC microcontrollers*

*Nintendo 64*

*Palm Beach County*

### Low Water Mark of Complexity

*Provided they "make it past" the editor, this chapter contains a number of references to Palm Beach County (Florida), which readers may recall was made famous for its use of the stupendously complex and utterly confusing "butterfly ballot" in the Election of 2000. One thing, however, that Palm Beach County and the rest of Florida deserve "partial credit" for is making the punch card ballot an artifact of the past...at least we hope!*

## 3.4 Choosing an Education-Appropriate Microprocessor

At this juncture, we are equipped to choose the computing device that will serve as the focus of our educational venture. Perhaps the only thing clear, though, is that there are a *lot* of choices – each with its own tradeoffs. And it is here where many educators choose to take different paths. Bewildered by all the tradeoffs, some simply choose to simulate a "synthetic" instruction set. This approach, however, lacks the "hands on" feel of using a "real" device that "does something". Siding with familiarity, a significant number select the Intel "x86" architecture as the vehicle of choice. A wide array of texts along with some laboratory tools have been developed for this purpose. This approach, however, can unwittingly "rob" students of the perspective that there are other, much less powerful devices available that are not only less expensive, but also much better suited for a wide range of embedded applications.

*synthetic instruction set*

*Intel x86 architecture*

Many other educators, though – motivated by the need to equip students for senior design projects in the digital systems area – choose microcontrollers as the “introductory vehicle”. This approach not only has the advantage of introducing (and reinforcing) basic concepts of computer architecture and machine instruction sets, but also of applying the hardware concepts learned in prerequisite courses to interfacing microcontrollers with external devices. Further, the same microcontroller covered in such an introductory course can be incorporated into senior design projects – where students have an opportunity to further apply what they have learned about programming and interfacing to the design of a complete system. In short, focusing on microcontrollers gives students a good opportunity to learn about and apply a “basic building block” of modern digital system design – thus the rationale for the approach embraced in this text.

We have a “slight” problem, though: microcontrollers are not designed strictly with “education” in mind (and, even if one were, it would be impossible to reach universal agreement on its instruction set, programming model, and on-chip peripherals). Rather, most have been designed under the influence of “marketing types” whose mission in life is to maximum the company’s bottom line, accomplished by making a given microcontroller as “universally applicable” as possible. The unfortunate consequence, from an educational standpoint, is an ever-increasing escalation of features and operating modes one must wade through to learn “the basics” – details that tend to confuse and confound the learning process.

*marketing types*

*universal applicability*

Accepting this dilemma (and recalling our basic mission, which is to *introduce* students not only to microcontrollers, but also to computer architecture and programming models), what considerations should be made in choosing a specific device – in particular, one that is “education appropriate” (and *friendly*)? Some key characteristics that come to mind include the following:

- straight-forward, easy-to-learn instruction set
- relatively “powerful” (i.e., CISC-like) instruction set, since we are learning to program at the “assembly level”
- enough addressing modes to make it interesting, but not so many that they become overwhelming or confusing
- variety and size of on-chip memories
- relatively few “operating modes”
- not too many bits “wide” (8- or 16-bits ideal) – we want to be able to perform reasonably powerful mathematic operations (multiply and divide), but usually don’t need (or want) the precision (and overhead) afforded by floating point
- a reasonable complement of bit manipulation instructions to facilitate control-oriented applications

- amenable to high level language compilation
- a representative set of on-chip peripherals commonly used in control-oriented applications
- appropriate, in terms of complexity (ease of use) and capability, for senior design projects
- fairly widespread application (design-ins)
- quality of documentation and support available
- commercial availability of an evaluation board and other hardware/software development tools (assemblers, debuggers, compilers)
- in-circuit debugging support
- family history/heritage
- low cost

The “bad news” is that no single commercial microcontroller possesses all the characteristics listed above. The “good news” is that a number of devices currently available satisfy many of these “education appropriate” characteristics. Among the author’s “personal favorites” are Motorola, Hitachi, and PIC devices. Forced to choose, the Motorola 68HC12 emerges as a leading candidate, with the MC68HC912B32 as the particular variant of interest.

*personal favorites*

*68HC12*

*MC68HC912B32*

### The Elusive Pedagogical Microprocessor

*Unfortunately (for educators), microprocessors and microcontrollers are created with markets in mind, not students or professors. The consequence of being market-driven (and, in most instances, “designed by committee”) is that a number of features and operating modes creep into the design of a product line – and tend to proliferate – as the availability of chip real estate increases. That plus the desire to maintain “legacy compatibility” makes it virtually impossible to find a “clean, simple, yet reasonably powerful” microcontroller ideal for education. The “hands on” appeal of using a “real” device, however, still outweighs the resignation to simply simulate a synthetic device – at least at this point in “digital history.” Hopefully, the author will have retired before the “simplest” microcontroller available is far too complex to cover in a single course!*

Why the 68HC12? It has a powerful, yet reasonably straight-forward instruction set; has a good complement of addressing modes; has multiple on-chip memories of different types (SRAM, byte-erasable EEROM, and Flash EEROM); is 16-bits wide, providing a good balance between “powerful math” and interfacing complexity; has a good set of bit manipulation instructions; has third-party “C” compilers available for it; has a great set of on-chip peripherals that are fairly easy to use; has proven itself in senior design projects the author has supervised; is gaining widespread application as the “upgrade” for its predecessor, the popular

68HC11; has good, complete documentation; has an inexpensive evaluation board available for the particular variant of interest; has in-circuit debugging capability; has a rich family heritage dating to the “humble beginnings” of microprocessors; and isn’t prohibitively expensive.

**68HC11**

The sound of whirring power tools is emanating from Norm’s *New Yankee Workshop*, so let’s start learning how to use them!

### Truth in Advertising

*The primary focus of this text is to help students learn how to design microcontroller-based systems. To accomplish this goal, it is most expedient to use a “real” microcontroller as a “working example.” And it also makes sense, along this same vein, to focus on a single representative device (here, the MC68HC912B32) rather than attempt to explain the differences (variations) among different microcontroller family members. Further, there is no pretense of providing a complete technical reference or usage guide on this particular microcontroller – these documents are readily available from the manufacturer’s web site (<http://mot-sps.com>).*

## 3.5 Tools of the Trade

The homework and lab exercises included with this text are based on use of the M68EVB912B32 Evaluation Board, shown in Figure 3-3. The EVB is packaged with printed copies of all pertinent documentation, which are also included as PDF files on the CD-ROM that accompanies this text. A disk that contains IASM12, an integrated editor and assembler program, is provided as well. This program runs under DOS on any conventional personal computer. The 68HC912B32 microcontroller on the EVB comes pre-loaded with a “debug monitor” program, called D-Bug12. This rather extensive debugging utility includes an in-line assembler, which will prove useful as we experiment with different instructions. All that needs to be added to get “up and running” are a personal computer capable of supporting DOS, a standard 9-pin serial port extension cable, and a regulated 5 VDC power supply.

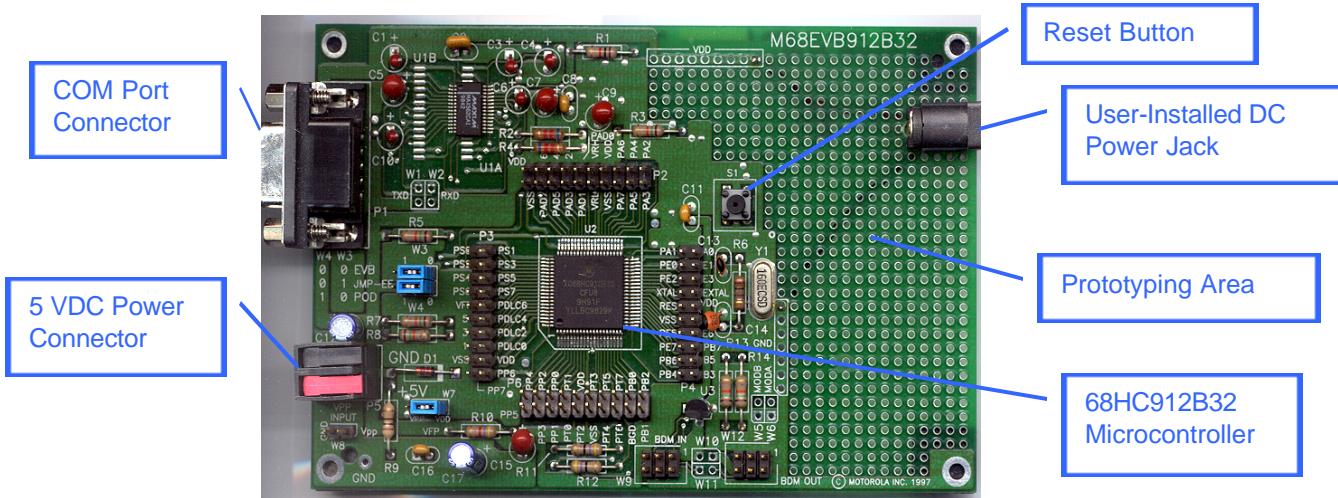
**M68EVB912B32  
Evaluation Board  
(EVB)**

**IASM12**

**D-Bug12**

Another “nice feature” of the M68EVB912B32 is a *prototyping area* that can be used to implement custom interfacing circuitry. We will make use of this provision in Chapter 8 to complete an illustrative design project. On the EVB illustrated in Figure 3-3, a standard power jack has been installed in the prototyping area to provide a convenient means of connecting a commercially available 5 VDC “wall wart” power supply.

*prototyping area*



**Figure 3-3** Motorola M68EVB912B32 Evaluation Board with power supply jack installed in prototyping area.

Before we delve into the details of the 68HC12 architecture and programming model, a few suggestions on how to make use of these “tools of the trade” are in order. There are three primary tools we will be using throughout our initial discussion of the 68HC12 instruction set: (1) the integrated editor, assembler, and communication utility; (2) the EVB, connected to the PC via a COM port; and (3) the D-Bug12 monitor program, that runs on the EVB when it is powered up.

First, some “helpful hints” on installing IASM12. After copying the contents of the diskette supplied with M68EVB912B32 to an appropriate directory on the PC’s hard drive, run the program `iasminst.exe`. For most of the options it prompts the user for, the default is fine – with some notable exceptions. Most users will want a “listing file” automatically generated, an “object file” automatically generated, “cycle counts” shown in the listing file, “macros expanded” in the listing file, and “include files expanded” in the listing file. Simply re-run the `iasminst.exe` program to verify or change any of these settings.

*installing  
IASM12*

Once installed, typing `iasm12` in a DOS window starts the program, which initially comes up in “editor” mode. To “talk” to the board, a communication (“COMM”) window must be opened; this is accomplished by pressing function key `F7`. Pressing `F8` several times will expand this window. As its name implies, the COMM window allows us to communicate directly with the EVB and the monitor program (D-Bug12) it is running. Upon powering up (or resetting) the EVB, the display shown in Figure 3-4 should be obtained. Note that “>” is the “monitor prompt”. Pressing function key `F10` closes the COMM window, returning IASM12 to its “editor” mode. A good on-line “help” capability, replete with information

*COMM window*

*monitor prompt*

*on-line help*

on how to use IASM12 as well as details about the 68HC12 instruction set (including examples), can be accessed by pressing the F1 function key.

```

F1-Help F2-Save F3-Load F4-Assemble F5-Exit F7-Comm F9-DOS shell F10-Menu
+----- COMM WINDOW -----+
|
| D-Bug12 v2.0.2
| Copyright 1996 - 1997 Motorola Semiconductor
| For Commands type "Help"
|
| >
|
|
+----- F1-Help F6-Download F7-Edit F8,F9-Resize F10-Close window -----

```

**Figure 3-4 IASM12 Communication Window to EVB.**

Once we have established communication with the EVB, we can execute any of the D-Bug12 monitor commands, described in Chapter 3 of the *M68EVB912B32 Evaluation Board User's Manual* (packaged with the EVB and included as a PDF on the CD-ROM that accompanies this text). This would be a good time to look over the various commands D-Bug12 is capable of executing, as well as the EVB setup and configuration information provided in Chapters 1 and 2 of this manual.

*Evaluation Board  
User's Manual*

Fortunately, we will only need to use a few of these commands to master the basics of the 68HC12 instruction set. In particular, we will find the *assembler/disassembler* command (*asm*) and the *trace* command (*t*) useful in understanding the functions performed by various instructions. To initialize the contents of various registers and memory locations, we will use the *register modify* (*rm*) and *memory modify* (*mm*) commands. Once we start creating assembly source files, we will use the *load* (*l*) and *go* (*g*) commands to download and execute them on the EVB.

*asm*  
*t*  
*rm*  
*mm*  
*l*  
*g*

An *assembly source file* is a text file containing a series of 68HC12 assembly instructions, along with comments that describe the program's operation; a ".asm" extension is used to distinguish the "source" version of the program file from the derivatives generated as a result of the "assembly process". Any text editor can be used to create an assembly source file: either the one integrated into IASM12 (which is somewhat cumbersome to use), or any of the standard Windows™ editors like Notepad. (Former UNIX hacks, such as the author, might prefer to use the DOS versions of *vi* or *emacs* instead.) Once an assembly source file has been created, it can be loaded into the IASM12 editor (by pressing key F3) and assembled (by pressing key F4). Provided the assembly was successful, the *object file* created (also called an "S-record" file, hence the

*assembly source file*

*UNIX hacks*

*S-record*  
*object file*

“.s19” extension) can be downloaded to the EVB for execution. As a byproduct of the assembly process, an *assembled source listing file* (“.lst”) is also created. The listing file shows the address at which each instruction is located in memory, along with the object code generated – information that will prove invaluable when debugging a program.

The first “barrier” students typically encounter is keeping track of which tool does what (and which one they are currently “talking to”) – since D-Bug12 commands to the EVB are entered through the PC’s keyboard, and the EVB’s response is displayed on the PC’s monitor. This challenge generally manifests itself the first time students attempt to create an assembly source file, assemble it, view the assembled source listing, download the object file generated to the EVB, and attempt to execute it. To help us navigate through this barrier, we will “walk” our way through a simple example based on the “simple computer” instructions we learned about in Chapter 2. We will then be prepared to test any of the 68HC12 instructions covered in the sections of this chapter that follow.

Assume we have created the assembly source file depicted in Figure 3-5, named `test.asm`, using the text editor of our choice. All that this program does is load the “A” register (accumulator) with the contents of location  $900_{16}$  in memory, add the contents of location  $901_{16}$  to it, and stores the result back in memory location  $900_{16}$ . The code that does all this “originates” at location  $800_{16}$  in memory – which is conveyed to the assembler program using the ORG pseudo-op (a *pseudo-op* is an *assembler directive* that provides information to the assembler program, but does not produce any executable code for the microcontroller). The label `MAIN` marks the beginning of the “main program” (and therefore assigned the value  $800_{16}$  by the assembler); it is used as a *symbolic reference* by the JMP instruction to transfer control back to the beginning of the instruction sequence once it completes – the astute digijock(ette) will recognize this as an “infinite loop”. The END pseudo-op simply tells the assembler program it has reached the end of the source file. Note that *comments* are delineated by a semicolon, and that “white space” may be added at will. Also note that the assembly instructions themselves are *case insensitive*, and that the *instruction fields* are separated by *tabs* (although spaces will work just as well).

Once this assembly source file has been created, start up IASM12 by typing `iasm12` in response to a DOS prompt. Press function key F3 and enter the assembly source file name (`test.asm`) followed by the ENTER key; the contents of the file should now be displayed on the screen. Next, press function key F4 to assemble the source file; the result, indicating a successful assembly, is shown in Figure 3-6. Two new files have just been created as a result of the assembly process: `test.lst` (the

*assembled source listing file*

*pseudo-op assembler directive*

*symbolic reference*

*comments*

*case insensitive*

*instruction fields*

*assembly process*

assembled source listing) and `test.s19` (the object file in S-record format).

```

        ORG      800h ; originate program at
                  ; location 800h
MAIN    LDAA    900h ; (A) = (900h)
        ADDA    901h ; (A) = (A) + (901h)
        STAA    900h ; (900h) = (A)
        JMP     MAIN   ; repeat operation

        END      ; end of assembly
                  ; source file

```

**Figure 3-5 Asssmbly source file for `test.asm`**

Let's take a moment to look at each of these files to understand what they contain. Press function key F3 and replace the ".asm" extension with ".lst" and press the ENTER key; the assembled source listing file should now be displayed on the screen, as shown in Figure 3-7. The column on the far left indicates the *address in memory* at which each instruction is destined to be stored: LDAA at location  $800_{16}$ , ADDA at  $803_{16}$ , STAA at  $806_{16}$ , and JMP at  $809_{16}$ . The number in brackets, in the next column over, indicates the *number of cycles* it takes each instruction to execute (recall that this was one of the "options" we deliberately enabled when we installed IASM12). The next column of hexadecimal numbers represent the machine code generated by the assembler program for each assembly instruction. For example, the assembly instruction LDAA 900h represents the machine code consisting of opcode byte B6<sub>16</sub> followed by the two-byte address 0900<sub>16</sub>. The bytes B6<sub>16</sub>, 09<sub>16</sub>, and 00<sub>16</sub> are stored at locations 800<sub>16</sub>, 801<sub>16</sub>, and 802<sub>16</sub>, respectively; thus, the next instruction (ADDA) starts at location 803<sub>16</sub>. The next column is the source file line number, which can be used as an aid in finding and correcting source file errors. The remaining columns are just an "echo" of the source file contents.

Appended to the end of this file is a *symbol table*, which is simply a list of each label or symbol the assembler encountered and the value that was assigned to it. Note that, as the source file is being assembled, there may be a *forward reference* to a symbol defined later in the source file; therefore, assembly requires a *two-pass* process. On the first pass, all the symbols are placed in the symbol table as they are referenced and assigned values as they are encountered; any forward references are left unresolved. On the second pass, the forward references are resolved ("filled in") based on the values determined at the completion of the first pass; if a symbol is missing or unresolved, an assembly error will occur.

*address in memory*

*number of cycles*

*symbol table*

*forward reference*

*two-pass assembly*

```

        ORG      800h ; originate program at
                  ; location 800h
MAIN    LDAA    900h ; (A) = (900h)
        ADDA    901h ; (A) = (A) + (901h)
        STAA    900h ; (900h) = (A)
        JMP     MAIN  ; repeat operation

END      +-----+ ASSEMBLE +-----+
          |
          | Assembling : (editor)
          |
          | Labels : 1
          | Lines : Total Current
          |           11      10
          |
          | Pass 2 : assembling
          | Success : Hit any key
+-----+

```

Figure 3-6 Confirmation of assembly success.

```

0800          1       ORG      800h ; originate program at
                  ; location 800h
0800 [03] B60900 3   MAIN    LDAA    900h ; (A) = (900h)
0803 [03] BB0901 4   ADDA    901h ; (A) = (A) + (901h)
0806 [03] 7A0900 5   STAA    900h ; (900h) = (A)
0809 [03] 060800 6   JMP     MAIN  ; repeat operation
                  7
080C          8       END      ; end of assembly
                  9             ; source file
                  10
                  11

Symbol Table

MAIN          0800

```

Figure 3-7 Assembled source listing file.

Let's "force" an assembly error to occur so it's not a surprise when it happens in real life. Press function key F3 and replace the ".lst" with ".asm", then press ENTER; the original source file should now be on the screen. Just for the experience of doing something useful with the IASM12 editor, use the cursor keys to move to (and subsequently change) the label MAIN to MAIN2; the source file should now look like Figure 3-8. Next, press F4 to assemble the file; note the error that occurs (the 'first parameter' – i.e., the symbol MAIN – of the JMP instruction is "unknown").

*forced error**IASM12 editor*

After pressing the ESC key, change the label `MAIN2` back to `MAIN` and reassemble the code; assembly should now be successful.

```

        ORG    800h ; originate program at
                  ; location 800h
MAIN2   LDAA   900h ; (A) = (900h)
        ADDA   901h ; (A) = (A) + (901h)
        STAA   900h ; (900h) = (A)
        JMP    MAIN  ; repeat operation

        END      ; end of assembly
                  ; source file

```

**Figure 3-8** A “forced error” in an assembly source file:  
the label `MAIN` is not defined.

Before we load and execute the S-record object file, let’s look at it. Press F3 and replace the “.asm” extension with “.s19”, then press ENTER; the screen shown in Figure 3-9 should appear. The information contained in this file is used by a *loader program*, which is part of D-Bug12 that runs on the EVB, to place the machine code in the 68HC12’s memory. It stands to reason, then, that this file must necessarily contain both address information as well as opcode and operand data. Note that the first line starts with the characters “S1”, while the second starts with the characters “S9” – hence the name “S” (for *starts with*) “19”. The “1” and “9” represent two different kinds of *records* that can be contained in a Motorola “S19” file: a “regular” one (S1) and an “ending” one (S9). The next pair of digits indicates the byte count of the line, in hexadecimal: for the S1 record (the first line), it is 0F<sub>16</sub> (or 15<sub>10</sub>), meaning that 15 bytes of information are contained in this record. The next four digits represent the two-byte starting address at which this record will be loaded into the microcontroller’s memory: 0800<sub>16</sub>. The next 24 digits represent the 12 bytes of machine code the assembler generated for this program: B60900 corresponds to the LDAA 900h instruction, BB0901 corresponds to ADDA 901h, A00900 corresponds to STAA 900h, and 060800 corresponds to JMP 800h (recall that the symbol `MAIN` was assigned the value 800<sub>16</sub>).

*loader program*

*S19*

```

S10F0800B60900BB09017A0900060800D3
S9030000FC

```

**Figure 3-9** The S-record file `test.s19`, generated by the assembler for the source file `test.asm`.

The value represented by the final pair of digits, D3, is called a *checksum*; it can be used by the loader program to check the integrity of the record as

*checksum*

it is received. The checksum is then calculated by summing, modulo  $256_{10}$ , all of the bytes in the record except the start code (S1), and then taking a *bit-wise* (or *ones'*) complement of the value. For the S1 record here, then, the checksum is found by summing  $0F_{16} + 08_{16} + \dots + 08_{16} + 00_{16} = 2C_{16}$ ; taking the bit-wise complement of  $2C_{16}$  ( $00101100_2$ ) yields  $D3_{16}$  ( $11010011_2$ ). As the D-Bug12 loader program “digests” each S-record, it sums the bytes received modulo  $256_{10}$ . When the checksum is received, it is added to the sum of the bytes received; since, on a good day, these two values should be ones’ complements of each other, their sum should yield  $FF_{16}$ . This test is performed by the loader program to check the integrity of each record as it is received.

The second S-record (that starts with S9) simply indicates the “end of file”. There are three bytes of information in an S9 record: the byte count (which, not surprisingly, is  $03_{16}$ ) followed by a two-byte address field. Here the address field is  $0000_{16}$ , but could be any value since S-record loader programs typically ignore this field. The checksum byte is calculated the same way as described above for S1-type records. We’ll have more “fun” with S-records in Chapter 4 when we write our *own* loader program!

Now that we know what an S-record is and understand the information it contains, we’re ready to actually load one into the 68HC12 microcontroller’s memory. To *download* an S-record file (on the PC) into the microcontroller’s memory (on the EVB), two things must happen: (1) D-Bug12 needs to perform a “load” command, and (2) the IASM12 program running on the PC needs to output the contents of the S-record file via the COM port connected to the EVB. Step (1) is accomplished by opening a communication window (by pressing function key F7) and, in response to the monitor prompt, typing `load`. Step (2) is accomplished by pressing function key F6 and typing the name of the S-record file to be loaded (here, `test.s19`) followed by ENTER. The contents of the S-record file will be echoed to the IASM12 COMM window as it is sent to the EVB. Pressing ENTER after the download has completed should yield a monitor prompt (>); if the message “BAD COMMAND” appears instead, something went wrong while the S-record file was being loaded. Should an error occur, check the S-record file and repeat the download process outlined above.

A quick way to check to see if an S-record file has been loaded correctly is to *disassemble* the code just loaded in the microcontroller’s memory. This can be accomplished using the D-Bug12 `asm` (assemble/disassemble) command. Since our code was loaded starting at location  $800_{16}$  in memory, type `asm 800` in response to the monitor prompt; after pressing the ENTER key four times in succession (once for each of the four instructions contained in this program), the screen shown in Figure 3-10

*download*

*bad command*

*disassemble*

should appear. Here, note that the prompt `>` has moved to the right, providing the opportunity to enter (and assemble *in-line*) a new instruction in place of the one indicated. To exit the `asm` command, type a period (`.`) – the prompt should then move back to its “normal” position.

*in-line assembly*

```

F1-Help F2-Save F3-Load F4-Assemble F5-Exit F7-Comm F9-DOS shell F10-Menu
+----- COMM WINDOW -----
|>asm 800
|0800 B60900      LDAA  $0900      >
|0803 BB0901      ADDA  $0901      >
|0806 7A0900      STAA  $0900      >
|0809 060800      JMP   $0800      >.
|>
|>
|>
+----- F1-Help F6-Download F7-Edit F8,F9-Resize F10-Close window -----

```

**Figure 3-10** Use of the D-Bug12 `asm` command.

An important limitation to note is that the `asm` command has no knowledge of the symbols used by the assembler program; thus, labels and symbols do not appear in the disassembled code. Another important limitation to keep in mind is that, if the “wrong” starting address is used (i.e., one that does not correspond to an instruction boundary), incomprehensible results will be obtained. This can be illustrated by disassembling the code, say, from location  $801_{16}$  (instead of  $800_{16}$ ) – try this to see what happens.

In the exercises and lab experiments provided for this chapter, we will primarily be investigating the function of individual instructions – or, at most, two or three instructions in succession. One way we can *empirically test* the effects of the 68HC12 instructions is to use the D-Bug12 `asm` command – here, entering the instructions we wish to test in response to the `asm` command prompt. The other way we can test instructions or instruction sequences is to place them in an assembly source file, assemble that file, and download the object file created. Most students seem to prefer the latter approach.

*empirically test*

Regardless of how the machine code has been entered into the microcontroller’s memory, we are now ready to initialize the contents of registers and memory locations in order to *trace* the execution of our program. Using the D-Bug12 register modify (`rm`) command will allow us to initialize any of the 68HC12’s registers; the only one important here is the program counter. In response to the monitor prompt, type `rm` followed by ENTER; the current value of the PC will be shown, which can be changed by typing a new value (here, 800). When ENTER is pressed, the program counter will take on value entered and subsequently prompt the

*trace*

user to update the next register in sequence (here, the stack pointer). If no change is desired, simply press ENTER. Note that the list “recycles” after the seven registers possible to change are displayed; this provides an opportunity to verify that any registers changed indeed took on the desired value. To exit the `rm` command, simply type a period followed by ENTER. The register modify sequence described above is shown in Figure 3-11.

```

F1-Help F2-Save F3-Load F4-Assemble F5-Exit F7-Comm F9-DOS shell F10-Menu
+----- COMM WINDOW -----
|
|>rm
|
|PC=0000 800
|SP=0A00
|IX=0000
|IY=0000
|A=00
|B=00
|CCR=90
|PC=0800 .
|>
+----- F1-Help F6-Download F7-Edit F8,F9-Resize F10-Close window -----

```

**Figure 3-11** Register modify sequence using D-Bug12 `rm` command.

Our illustrative program also uses some memory locations, namely  $900_{16}$  and  $901_{16}$ . Location  $900_{16}$  is used to store the “running sum” of the value calculated by this program, and location  $901_{16}$  contains the amount to add to the running sum each time it completes a “loop”. We can initialize these locations to “suitable values” using the D-Bug12 memory modify (`mm`) command. In response to the monitor prompt, type `mm 900` followed by ENTER; the current contents of memory location  $900_{16}$  should be displayed. To clear this value to zero, type `00` followed by ENTER. The `mm` command will then display the contents of the next consecutive location,  $901_{16}$ . For the purpose of testing our program, we would like this value to be one. To do this, type `01` followed by ENTER. For the moment, these are the only two locations we “care about”, so we can now exit the memory modify command by typing a period (.) followed by ENTER. The memory modify sequence described above is illustrated in Figure 3-12. Note that, depending on what has previously been loaded into or run on the EVB, the original contents of memory will vary.

We are now ready to “single step” through the execution of our program, one instruction at a time, using the trace (`t`) command. In response to the monitor prompt, press `t` followed by ENTER; the result of executing the instruction pointed to by the program counter (here, at location  $800_{16}$ ) is displayed, followed by a disassembly of the instruction *which follows* (at location  $803_{16}$ ). Referring to Figure 3-13, we note that execution of the `LDA A 900h` instruction loaded the “A” register with the contents of

*single step*

memory location  $900_{16}$  (which, using the `mm` command, we initialized to  $00_{16}$ ). Because the `LDA A 900h` instruction occupies three bytes in memory, the program counter is “bumped” to  $803_{16}$  as a result of executing this instruction.

```
F1-Help F2-Save F3-Load F4-Assemble F5-Exit F7-Comm F9-DOS shell F10-Menu
+----- COMM WINDOW -----
|>
|>
|>
|>mm 900
|0900 B7 00
|0901 56 01
|0902 20 .
|>mm 900
|0900 00
|0901 01 .
|>
|>
+---- F1-Help F6-Download F7-Edit F8,F9-Resize F10-Close window -----+
```

**Figure 3-12** Memory modify sequence using D-Bug12 `mm` command.

Pressing `t` followed by ENTER again causes the next instruction in sequence, `ADDA 901h`, to be executed. Referring to Figure 3-14, we note that this instruction adds the contents of memory location  $901_{16}$  (which, using the `mm` command, we initialized to  $01_{16}$ ) to the “A” register. Since the `ADDA 901h` instruction occupies three bytes in memory, the program counter is “bumped” to  $806_{16}$  as a result of executing this instruction.

```
F1-Help F2-Save F3-Load F4-Assemble F5-Exit F7-Comm F9-DOS shell F10-Menu
+----- COMM WINDOW -----
|>
|>
|>
|>
|>
|>t
      PC      SP      X      Y      D = A:B      CCR = SXHI NZVC
|0803  0A00  0000  0000    00:00      1011 0100
|0803  BB0901          ADDA   $0901
|>
+---- F1-Help F6-Download F7-Edit F8,F9-Resize F10-Close window -----+
```

**Figure 3-13** Result of first instruction trace using D-Bug12 `t` command.

```

F1-Help F2-Save F3-Load F4-Assemble F5-Exit F7-Comm F9-DOS shell F10-Menu
+----- COMM WINDOW -----
|>
|>t
|
| PC      SP      X      Y      D = A:B      CCR = SXHI NZVC
| 0803    0A00    0000   0000    00:00          1011 0100
| 0803    BB0901           ADDA   $0901
|>t
|
| PC      SP      X      Y      D = A:B      CCR = SXHI NZVC
| 0806    0A00    0000   0000    01:00          1001 0000
| 0806    7A0900           STAA   $0900
|>
+----- F1-Help F6-Download F7-Edit F8,F9-Resize F10-Close window -----

```

**Figure 3-14** Result of second instruction trace.

```

F1-Help F2-Save F3-Load F4-Assemble F5-Exit F7-Comm F9-DOS shell F10-Menu
+----- COMM WINDOW -----
|>t
|
| 0803    BB0901           ADDA   $0901
|>t
|
| PC      SP      X      Y      D = A:B      CCR = SXHI NZVC
| 0806    0A00    0000   0000    01:00          1001 0000
| 0806    7A0900           STAA   $0900
|>t
|
| PC      SP      X      Y      D = A:B      CCR = SXHI NZVC
| 0809    0A00    0000   0000    01:00          1001 0000
| 0809    060800           JMP    $0800
|>
+----- F1-Help F6-Download F7-Edit F8,F9-Resize F10-Close window -----

```

**Figure 3-15** Result of third instruction trace.

```

F1-Help F2-Save F3-Load F4-Assemble F5-Exit F7-Comm F9-DOS shell F10-Menu
+----- COMM WINDOW -----
|>t
|
| 0806    7A0900           STAA   $0900
|>t
|
| PC      SP      X      Y      D = A:B      CCR = SXHI NZVC
| 0809    0A00    0000   0000    01:00          1001 0000
| 0809    060800           JMP    $0800
|>t
|
| PC      SP      X      Y      D = A:B      CCR = SXHI NZVC
| 0800    0A00    0000   0000    01:00          1001 0000
| 0800    B60900           LDAA   $0900
|>
+----- F1-Help F6-Download F7-Edit F8,F9-Resize F10-Close window -----

```

**Figure 3-16** Result of fourth instruction trace.

Pressing `t` followed by `ENTER` again causes the next instruction in sequence, `STAA 900h`, to be executed. Referring to Figure 3-15, we note that this instruction stores the “updated” value in the “A” register at our “running sum” location,  $900_{16}$ . Since the `STAA 900h` instruction occupies three bytes in memory, the program counter is “bumped” to  $809_{16}$  as a result of executing this instruction.

Pressing `t` followed by `ENTER` again causes the next instruction in sequence, `JMP 800h`, to be executed. Referring to Figure 3-16, we note that execution of this instruction moves us back to the “top” of the “loop”, i.e., location  $800_{16}$ . Three more `t-ENTER` combinations will complete a second iteration of the “loop”, updating the running sum to  $02_{16}$ .

If we have large sequence of instructions that we would like to trace, pressing the `t-ENTER` combination multiple times can quickly become annoying. Fortunately, the D-Bug12 trace command can be told the number of instructions to execute in sequence. Say, for example, we wish to determine the result of executing the loop in this program five times. Since there are four instructions in the loop, we would need to execute a total of  $20_{10}$  instructions to determine the final result. This can be accomplished by simply typing `t 20` followed by `ENTER`, which causes the trace command to automatically repeat 20 times. The maximum “trace count” that can be specified this way is  $255_{10}$ .

*trace count*

To *continuously* execute our program, we could simply use the D-Bug12 “go” (`g`) command by typing `g 800` after downloading the S-record file. Try this to see what happens. Why is there “no further response” (or, why does the monitor program “appear to hang”) at this point? Because, like the infamous Election of 2000, there is no prescribed, “lawful” way for the program to terminate – it is simply an “infinite loop”! The only way to *stop* it is to press the (tiny) *reset button* on the EVB – note that doing so causes the monitor program to restart.

*continuously execute*

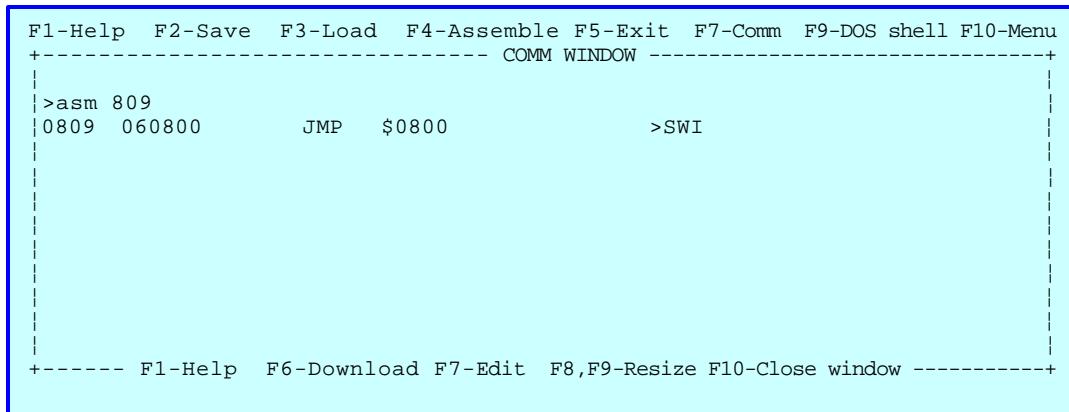
This gives us an opportunity to clear up some common misconceptions concerning what, exactly, pressing the *reset button* does (its location is shown in Figure 3-3). To explore this, use the `rm` command to view the register values after pressing the EVB reset button; note that they have all been initialized to known values. Next, use the `mm` command to check the contents of memory locations  $900_{16}$  and  $901_{16}$ ; here we find that the contents of  $900_{16}$  is some “random value” (since the loop executed literally millions of iterations between the time we started it and the time we stopped it), but the contents of location  $901_{16}$  is still  $01_{16}$ . The conclusion? Pressing the *reset button* (sometimes called performing a “hard reset”) places the processor’s registers in a known state, *but leaves memory unaffected*.

*reset button*

What if we would like to execute a series of instructions and then just “stop” so we can use various monitor commands to determine what happened? This can be accomplished by terminating the code sequence we wish to test with a *software interrupt (SWI)* instruction. Here, we can replace the `JMP 800h` instruction at the end of our program with an `SWI` instruction. To do this, we could either: (a) modify our assembly source file, re-assemble it, and download the object file; or, (b) use the D-Bug12 `asm` command to replace the `JMP` instruction with an `SWI` instruction.

*software interrupt  
SWI*

Approach (b) is probably more expedient here. Recalling that the `JMP` instruction resides at location  $809_{16}$ , we can replace it by typing `asm 809` and, in response to the prompt, type `SWI`; this is illustrated in Figure 3-17. After pressing `ENTER`, the newly inserted `SWI` instruction appears at location  $809_{16}$ ; typing a period (.) followed by `ENTER` terminates the in-line assembly process.



```

F1-Help F2-Save F3-Load F4-Assemble F5-Exit F7-Comm F9-DOS shell F10-Menu
+----- COMM WINDOW -----+
|
>asm 809
0809 060800      JMP    $0800          >SWI
|
+----- F1-Help F6-Download F7-Edit F8,F9-Resize F10-Close window -----

```

**Figure 3-17** Insertion of `SWI` instruction using `asm` command.

Once we have inserted the `SWI` instruction in place of the `JMP` (and used the `mm` command to initialize location  $900_{16}$ ), we can execute the entire program by typing `g 800`. When the `SWI` instruction is executed, the contents of machine’s registers are displayed and control is returned to D-Bug12, allowing the user to execute any monitor command.

When debugging a larger program, though, what we often wish to do is execute our code up to a certain “problematic point” and trace from there. This can be accomplished either by setting a *breakpoint* (using the D-Bug12 `br` command), or by using the “go till” (`gt`) command (which sets a temporary breakpoint). After tracing through the “questionable code”, normal execution can be resumed by simply typing `g` in response to the monitor prompt.

*breakpoint*

*questionable code*

We are now equipped with the “tools of the trade” that will help us test and execute assembly language instructions as well as code segments. With this as background, we are now prepared to learn the details of 68HC12 instruction set in the sections of this chapter that follow. From there, we will go on in Chapter 4 to learn program structures and assembly language programming techniques.

### 3.6 Motorola 68HC12 Architecture and Programming Model

In its basic form, the programming model of the Motorola 68HC12 is a fairly straight forward extension of the simple computer we designed in Chapter 2. Like our simple computer, the 68HC12 has an 8-bit accumulator register (A); a program counter register (PC), here extended to 16-bits; and a stack pointer register (SP), also extended to 16-bits. The two computers also share the same basic condition code bits: a carry/borrow flag (C), a negative flag (N), an overflow flag (V), and a zero flag (Z). These flags function in the exact same manner as those on our simple computer.

Unlike our simple computer, the 68HC12 has a second accumulator register (cleverly called “B”), which can be concatenated with the “A” register to form a *double-byte* (or “D”) accumulator. Thus, one can view the 68HC12’s accumulator as either a single 16-bit entity (referred to as “D”), or as two 8-bit “halves”, where the A register is the high byte and the B register is the low byte. There is also a “new” condition code bit, called the “half carry” flag (H), which is simply the carry out of the “lower half” (i.e., low-order 4-bits) following an ADD operation (the only time it is valid). In addition to the “arithmetic status” bits (H, N, Z, V, C), the so-called Condition Code Register (CCR) also contains three “machine control” bits: I and X are interrupt mask bits, and S is the stop disable bit. An illustration showing the position of each flag in the CCR is provided in Figure 3-18. The 68HC12 also has two 16-bit index registers (called “X” and “Y”) that primarily serve as pointers to operands. These “pointer registers” provide a number of additional ways of generating an effective address. A diagrammatic view of the 68HC12 programming model is provided in Figure 3-19.

*double-byte  
accumulator D*

*machine control bits*

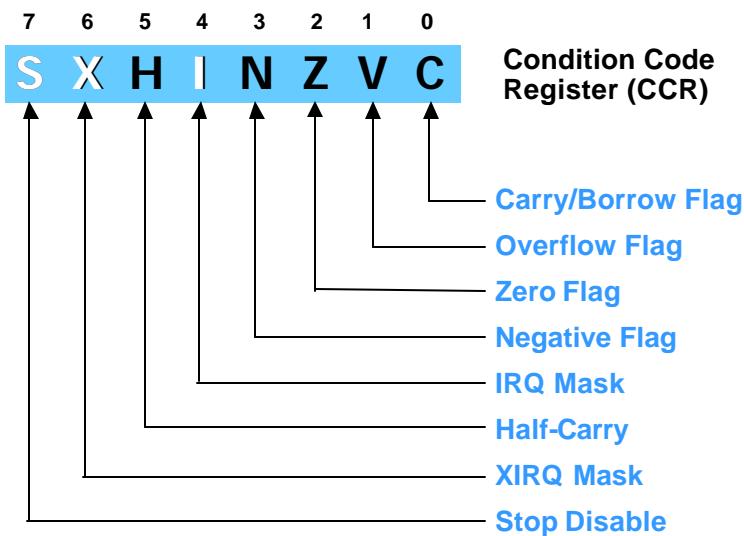
*index registers  
pointers*

*postbyte*

Another salient difference between our simple computer and the 68HC12 is that instructions can vary in length, from a single byte (8-bits) to as many as six bytes (48-bits). Opcodes are either one or two bytes, which can be followed by a “postbyte” that provides additional information about the addressing mode used. Data types supported by the 68HC12 include bit, byte, word (16-bit), double word (32-bit), packed BCD, and unsigned fractions.

### Holy War of Words

In the “early days” of microprocessors, a hot topic of contention (at one point called a “holy war”) was the ordering, in memory, of multiple-byte quantities, such as 16-bit (“word” length) addresses and data items. Intel, first to market with a commercially viable 8-bit microprocessor (the 8080), chose to place the lowest order byte of an address or operand in the lowest address at which that field was stored in memory. Using this ordering, called low-order-byte-first (or “little endian”) format, an instruction such as “JMP \$1234” would be stored in memory as XX \$34 \$12 (where “XX” is the opcode for JMP), with XX stored at location addr, \$34 at addr+1, and \$12 at addr+2. Motorola – most likely just to be “different” than Intel – chose the opposite byte ordering for their first commercial microprocessor, the 6800, that hit the market six months after the debut of the 8080. Using a high-order-byte-first (or “big endian”) format, a Motorola-style JMP \$1234 instruction would be stored in memory as XX \$12 \$34. Many claims were made (and considerable ink was spilled) concerning why one byte-ordering scheme was “better” than the other. Other manufacturers since them have “split” on the byte-ordering scheme they have chosen to use for their devices – some even have a control register bit that allows the programmer (or compiler) to select either of the two byte-ordering schemes for data items. The original claims concerning which scheme was “better” are now largely moot – especially in larger bit-width microprocessors, which generally fetch an entire instruction (or more) at once.



**Figure 3-18** Motorola 68HC12 Condition Code Register.

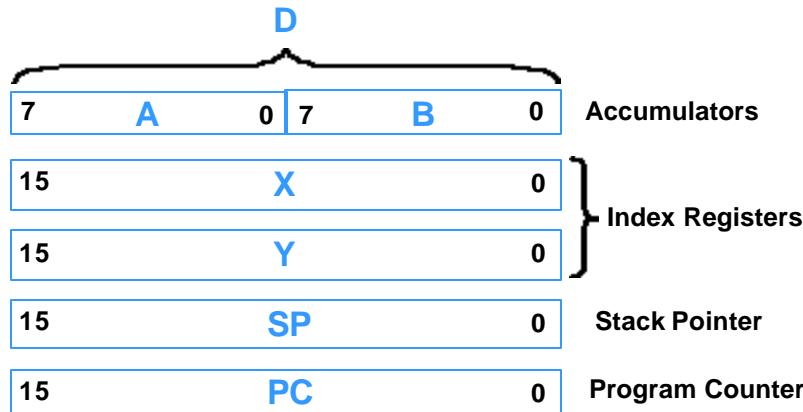


Figure 3-19 Motorola 68HC12 Programming Model.

Besides sporting a large variety of instructions, the 68HC12 also provides a number of ways of generating the effective address of the operands used by each instruction. Unlike our simple computer that had but a single (“absolute”) addressing mode, the 68HC12 can have as many as ten addressing mode variations that can be applied to each instruction. While there are on the order of 200 different instructions implemented by the 68HC12, the total number of variations possible when all the addressing modes are considered is well over 1000.

Another aspect of the 68HC12’s programming model that we need to understand before we begin to write code is its *memory map*. The 68HC912B32, the specific 68HC12 variant we will focus on here, has three different types of on-chip memory: SRAM, byte-erasable EEPROM (electronically erasable programmable read-only memory), and flash EEPROM. The relative locations of these memory modules are illustrated in Figure 3-20. A typical embedded application would most likely be placed in the 32 KB flash EEPROM which, by default, occupies the upper half of the processor’s address space (locations  $8000_{16}$  –  $FFFF_{16}$ ). On the M68EVB912B32 Evaluation Board, this area of memory is preloaded with the D-Bug12 (“debug monitor”) operating system. On the EVB, then, execution begins at location  $8000_{16}$  out of reset. (In Chapter 7, we will discuss how to create our own “turn key” embedded systems by loading our application code into the flash EEPROM.)

*memory map*

*flash EEPROM*  
*8000 - FFFF*

By default, the byte-erasable EEPROM occupies locations  $0D00_{16}$  –  $0FFF_{16}$  in the processor’s address space (which translates into a total of  $\frac{3}{4}$  KB). As its name implies, a unique feature of this non-volatile block of memory is that individual locations (bytes) can be erased and rewritten, without the need for an additional (higher) power supply voltage. (The flash EEPROM, described previously, can only be “bulk” erased, and requires a separate (higher) supply voltage to erase and reprogram.) Applications that require data that is “read mostly”, such as calibration

*byte-erasable EEPROM*  
*0D00 - 0FFF*

parameters, can make particularly effective use of this memory block. (In Chapter 7, we will see how we can dynamically change interrupt vectors by re-mapping them into the 68HC12's byte-erasable memory.)

The 68HC912B32's SRAM is primarily intended for storage of temporary variables as well as the system stack. This 1 KB block, that by default occupies locations  $0800_{16}$  –  $0BFF_{16}$ , is the area of memory in which we will place our “practice” code (as we progress, we will also begin to use the byte-erasable EEPROM, and ultimately the flash EEPROM). On the M68EVB912B32 Evaluation Board, the D-Bug12 monitor uses the upper half of SRAM ( $0A00_{16}$  –  $0BFF_{16}$ ) for temporary variables, leaving a seemingly paltry  $\frac{1}{2}$  KB ( $0800_{16}$  –  $09FF_{16}$ ) for our “fun and enjoyment”. To maximize the effectiveness of this area, the SP register is initialized by the D-Bug12 monitor to  $0A00_{16}$ . (Note that the same stack convention utilized by our simple computer is employed by the 68HC12, i.e., the SP register points to the *top stack item*, and as such, the SP register needs to be initialized to one greater than the location of the “bottom stack item”).

The questions of how to add additional (external) memory devices to a 68HC912B32 as well as how to re-map the internal memory resources will be addressed in Chapter 5.

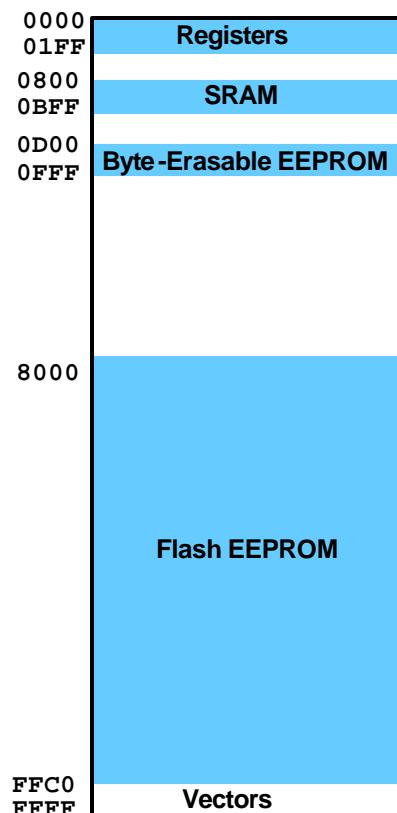


Figure 3-20 Motorola MC68HC912B32 Memory Map.

**SRAM**  
**0800 – 0BFF**

**stack convention**

## 3.7 Addressing Modes

At this point, we have what amounts to a “chicken and egg” problem: to understand all the variations of instruction formats possible, we need a firm grasp of the 68HC12 addressing modes; a good understanding of the addressing modes, however, can only be attained in the context of the 68HC12’s instruction set. To solve this dilemma, we will introduce two very basic data transfer group instructions as a “vehicle” for presenting the addressing modes. Once the addressing modes are firmly established, we will move forward with the 68HC12 instruction set details.

*chicken and egg problem*

The two instructions we will introduce first are basic “load” and “store” accumulator instructions, similar in form and function to those of our simple computer. The 68HC12 equivalent of our simple computer’s LDA instruction is **LDAA**, for **load accumulator A**; the equivalent of STA is **STAA**, for **store accumulator A**. The “absolute” addressing mode version of each of these instructions requires 3 bytes (or 24-bits): an 8-bit opcode field followed by a 16-bit operand address field. This can simply be thought of as an “expanded” version of the 3-bit opcode and 5-bit operand address used by our simple computer.

*load and store accumulator*

Recall from Chapter 2 that an *addressing mode* is used by a computer to determine the *effective address* at which an operand is stored in memory. For our purposes, the effective address can be thought of as the actual (or *absolute*) location in memory at which the data is stored. Most processors worth their silicon provide, at minimum, six basic addressing modes:

*addressing mode effective address*

1. **Absolute** (or *extended/direct*), so called because the operand field of the instruction indicates the absolute (or actual) location in memory at which the operand is stored. (This is the addressing mode implemented on our simple computer of Chapter 2.) *extended/direct*
2. **Register** (or *inherent*), so called because the operands (if any) are contained in registers – stated another way, the “name” of the operand register is included (or “inherent”) in the instruction mnemonic. *inherent/register*
3. **Immediate**, so called because the operand data immediately follows the opcode, i.e., the data is contained in the instruction itself rather than some other area of memory. *immediate*
4. **Relative**, so called because the desired location (of either data or a branch target) is *relative* to the current value in the PC register – here the operand field is viewed as a *signed offset* that, when added to the current value in the PC, yields the effective address. *relative*
5. **Indexed**, so called because the desired location is found using an index register. With indexed addressing mode comes a whole series of variants that utilize different offsets (e.g., constants or registers) to determine the effective address. *indexed*

6. *Indirect*, so called because the initial effective address calculation yields the address of a (two-byte) pointer in memory, which is then read and used to determine the actual address in memory where the desired data is stored.

Armed with two basic instructions (LDAA and STAA) along with an outline of the fundamental addressing modes supported, we can now delve into the details of the 68HC12 addressing mode variations. A word of caution plus a suggestion, though, is in order before we start. Technical documentation that describes addressing modes is often cryptic and couched in hard-to-follow notation. Further, the sheer number of addressing mode variants possible can cause one to quickly become overwhelmed. To help make our study of 68HC12 addressing modes as “painless” and effective as possible, we will develop a “simplified” notation scheme and provide several examples of each variant. As one might guess, the way to learn addressing modes and the corresponding instruction variants is to write “real code” that uses them – a task we will attend to in Chapter 4. Breaking the task of learning addressing modes into palatable parts, however, will help make the task tractable. The notation we will use in the context of describing the 68HC12 addressing modes and instruction set is provided in Table 3-1.

*simplified notation scheme*

### 3.7.1 Non-Indexed Modes

For the LDAA and STAA instructions, two basic “non-indexed” modes of addressing are relevant: “absolute” and immediate. Motorola uses two different names for what can generically be called “absolute” addressing mode, depending on the area of memory space addressed. *Extended* refers to use of a (full) 16-bit address, while *direct* refers to use of an 8-bit address (to access the machine’s register block residing in the first 256-byte block of the address space, locations  $0000_{16}$  –  $00FF_{16}$ ). The Motorola adopted names for these modes are not universally used, however.

*extended direct*

The name *immediate* is almost universally used for an addressing mode in which the operand data “immediately follows” the opcode field. In Motorola assembly code, a pound sign (#) is used to specify immediate addressing mode. A common mistake is to accidentally “forget” the pound sign, causing the assembler program to use direct or extended addressing mode instead of the desired immediate mode.

*immediate*

#### Examples:

LDAA \$FF	$; (A) \leftarrow (00FFh)$	direct mode	{ 2 bytes, 3 cycles }
LDAA \$100	$; (A) \leftarrow (0100h)$	extended mode	{ 3 bytes, 3 cycles }
LDAA #\$FF	$; (A) \leftarrow FFh$	immediate mode	{ 2 bytes, 1 cycle }
LDAA #1	$; (A) \leftarrow 1$	immediate mode	{ 2 bytes, 1 cycle }

**Table 3-1** Notation used to describe instructions and addressing modes.

Notation	How Used	Examples
<i>prefix of \$ or suffix of h or H</i>	denotes a <i>hexadecimal</i> (base 16) number	\$1234 = 1234h = 1234H = 1234 <sub>16</sub>
<i>prefix of ! or suffix of t or T</i>	denotes a <i>decimal</i> (base 10) number	!1234 = 1234t = 1234T = 1234 <sub>10</sub>
<i>prefix of % or suffix of b or B</i>	denotes a <i>binary</i> (base 2) number	%10101010 = 10101010b = 10101010B = 10101010 <sub>2</sub>
( )	denotes the <i>contents of</i> a register or memory location	(A) (0800h)
;	denotes the beginning of a <i>comment</i>	LDAA 0800h ; (A) = (0800h)
:	indicates the <i>concatenation</i> of two quantities	16-bit result in (A):(B) ≡ (D) 32-bit result in (D):(X)
<i>addr</i>	shorthand for the <i>effective address</i> in memory at which an operand is stored	LDAA <i>addr</i> ; (A) = (addr)
<i>rb</i>	shorthand for a <i>byte-length register</i> , e.g., A or B	STArb 0800h ; (0800h) = ( <i>rb</i> )
<i>rw, rwh, rwl</i>	shorthand for a <i>word-length register</i> , e.g., X, Y, D, SP, where <i>rwh</i> denotes the <i>high byte</i> of that register and <i>rwl</i> the <i>low byte</i>	LD <i>rw</i> 0800h ; ( <i>rw</i> ) = (0800h):(0801h) ; -or- ; ( <i>rwh</i> ) = (0800h) ; ( <i>rwl</i> ) = (0801h)
#	indicates use of <i>immediate addressing mode</i> when used before a constant that appears in an instructions operand field	LDAA #80h ; (A) = 80h LDAA #\$12 ; (A) = 12h LDAA #\$A5 ; (A) = A5h LDAA #10101010b ; (A) = AAh
,	indicates use of <i>indexed addressing mode</i> when placed between two entities in the operand field	LDAA 2,X ; (A) = ((X) + 2)  STAA D,Y ; ((D)+(Y)) = (A)
[ ]	indicates use of <i>indirect addressing mode</i> when used to bracket the operand field	STAA [2,X] ; (((X)+2):((X)+3)) = (A) LDAA [D,Y] ; (A) = (((D)+(Y)):((D)+(Y)+1))
← →	denotes an <i>assignment</i> or “copy” (the arrow points toward the destination)	(A) ← (B) means load the A register with the contents of the B register (the contents of B remains the same)
↔	denotes the <i>exchange</i> (or “swap”) of contents	(D) ↔ (X) means exchange the contents of the D and X registers
~	shorthand for number of instruction execution cycles	assuming an 8 MHz bus clock, each cycle is 125 ns (nanoseconds)
⌚	indicates a (bit-wise) complement	mask⌚ means the bit-wise complement of mask

### 3.7.2 Indexed Modes

The indexed addressing modes supported by the 68HC12 are numerous and diverse. X, Y, and SP are most commonly used as “index” registers, while A, B, and D are commonly used as “accumulator” offsets. While at first seemingly overwhelming (Motorola defines ten official variants), the list can be condensed to a few basic categories:

1. Indexed with (signed) constant offset, of which there are three variants: a 5-bit offset, a 9-bit offset, and a 16-bit offset.
2. Indexed with (unsigned) accumulator offset, of which there are three variants: A, B, and D.
3. Indexed with auto pre/post increment/decrement, of which there are four permutations (and eight possible values, ranging from 1 to 8, by which the indexed register can be incremented or decremented).
4. Indexed indirect, of which there are two variants: constant (16-bit offset) and accumulator (D) offset.

Encouraged by the realization that four categories are much easier to remember than ten, we can now consider the details of each.

#### *Indexed with Constant Offset*

The variants of this mode are all specified the same way: the signed offset and index register of choice (X, Y, SP, PC) are placed in the operand field of the instruction, separated by a comma. The assembler program examines the offset specified and generates one of three different instruction formats. If the offset is in the range of  $-16_{10}$  to  $+15_{10}$ , the assembler will place the 5-bit offset within the post byte that follows the opcode. A different format is used if the offset is in the range of  $-256_{10}$  to  $+255_{10}$ : here, the most significant bit (only) of the offset is placed in the post byte while the lower eight bits of the offset are placed in a single-byte extension that follows the post byte.

*signed offset*

*5-bit offset, no extension byte*

*8-bit offset, one extension byte*

If a 16-bit offset is specified, the assembler places it in a two-byte extension that follows the post byte. Normally we would construe this as an offset that ranges from  $-32,768_{10}$  to  $+32,767_{10}$ . An alternate interpretation, however, is also perfectly valid here: as an (unsigned) offset that ranges from 0 to  $65,535_{10}$ . The reason this interpretation is valid is that the offset is added to an index register modulo  $2^{16}$  (i.e., it “wraps around”). Thus, adding  $-1$  (represented as  $FFFF_{16}$ ) yields the same result as adding  $65,535_{10}$  (also represented as  $FFFF_{16}$ ), due to the “modulo nature” of the addition. This “dual” interpretation of 16-bit offsets will prove useful when we examine table lookup in Chapter 4.

*16-bit offset, two extension bytes*

*wrap around*

*dual interpretation of offsets*

Instructions using indexed with constant offset addressing mode therefore range in size from two or three bytes (one or two opcode bytes followed by a post byte) up to four or five bytes (opcode byte(s), post byte, plus one or two extension bytes). As one might guess, there are differences in the number of execution cycles associated with each variant.

### Examples:

LDAA 0,X	$; (A) \leftarrow ((X)+0)$	5-bit offset { 2 bytes, 3 cycles}
LDAA 2,X	$; (A) \leftarrow ((X)+2)$	5-bit offset { 2 bytes, 3 cycles}
LDAA 255t,Y	$; (A) \leftarrow ((Y)+255)$	9-bit offset { 3 bytes, 3 cycles}
LDAA 1000t,X	$; (A) \leftarrow ((X)+1000)$	16-bit offset { 4 bytes, 4 cycles}
STAA -1,Y	$; ((Y)-1) \leftarrow (A)$	5-bit offset { 2 bytes, 2 cycles}
STAA 1,SP	$; ((SP)+1) \leftarrow (A)$	5-bit offset { 2 bytes, 2 cycles}
STAA 100t,PC	$; ((PC)+100) \leftarrow (A)$	9-bit offset { 3 bytes, 3 cycles}

Note that the first example illustrates the assembly format used to specify “zero offset” indexed addressing (i.e., indexed addressing with no offset).

*zero offset*

The next-to-last example illustrates how the contents of the stack can be modified “in place” without pushing/popping items or disturbing the SP register – a “trick” we will find quite useful in passing parameters to/from subroutines. The final example illustrates use of the PC as an index register, which allows the creation of “position independent” code (i.e., code that is not statically bound to a given set of memory locations).

*position independent code*

### Indexed with Accumulator Offset

The variants of this mode are specified the same way: the accumulator offset (A, B, or D) and index (X, Y, SP, PC) registers of choice are placed in the operand field of the instruction, separated by a comma. The only “tricky” part associated with this addressing mode is the interpretation of the offset as an *unsigned* quantity, in contrast with the (signed) constant offset mode described previously (except for the 16-bit case, where the offset in the “D” register can be interpreted as either signed or unsigned, as described previously). The first question that comes to mind is: *Why* did the designers of the 68HC12 choose to have the (8-bit) accumulator offset interpreted as unsigned (or, stated another way, as “zero-extended” to 16-bits before being added to the index register)? It turns out that the most common application of accumulator offset indexed addressing is accessing elements in an array. Since a “negative index” is often not very meaningful in this context, interpreting the accumulator offset as *unsigned* makes sense. Further, a rather unpleasant “side effect” would occur if the offset were interpreted as being signed: incrementing a byte-length index past  $7F_{16}$  (to  $80_{16}$  and beyond) would cause a discontinuity in the accessing of array elements (recall that, interpreted as signed, the 8-bit quantity  $7F_{16}$  represents  $+127_{10}$ , while  $80_{16}$  represents  $-128_{10}$ ). Not only would this cause difficulty in reserving storage for an array (since some

*accumulator offset*

*zero extended*

*negative index*

elements could potentially be stored at locations “behind” the starting address label), but also might cause difficulty in debugging code.

### Examples:

LDAA A,X	$; (A) \leftarrow ((A)+(X))$	{ 2 bytes, 3 cycles}
LDAA B,X	$; (A) \leftarrow ((B)+(X))$	{ 2 bytes, 3 cycles}
LDAA D,Y	$; (A) \leftarrow ((D)+(Y))$	{ 2 bytes, 3 cycles}
STAA B,SP	$; ((B)+(SP)) \leftarrow (A)$	{ 2 bytes, 2 cycles}
STAA A,PC	$; ((A)+(PC)) \leftarrow (A)$	{ 2 bytes, 2 cycles}

Note, from the first example above, that the accumulator offset register and the destination of the load may be the same. Here, the “old” value of the accumulator offset is used in the effective address calculation before it takes on its new value by virtue of being the destination of the load. Since common practice is to use an accumulator offset as an array index, the second example – using distinct accumulator offset and destination registers – is often utilized.

A particularly insidious problem can occur in the third example. Recalling that “D” is merely a pseudonym for “A:B” (i.e., “D” is just shorthand for “A concatenated with B”), note that the high byte of the offset (the A register) is modified as a “byproduct” of the load operation. This is fine as long as we don’t expect to use “D” as a 16-bit accumulator offset in a subsequent instruction (and still expect it to be the same value!).

*insidious problem*

### Indexed with Auto Pre/Post Increment/Decrement

When using an index register as a pointer to elements in an array or characters in a string, a common operation is to “bump” that pointer either forward or backward in order to access the next (or previous) element. With this in mind, the designers of the 68HC12 endowed it with a powerful set of “automatic” indexed increment/decrement modes. These modes are called automatic (auto) because they occur as a “side-effect” of the instruction being executed. An auto increment (or decrement) of an index register is called a *pre-increment* (*decrement*) if the index register is modified *prior* to its use as the effective address for the operand being accessed. Conversely, an auto increment (or decrement) is called a *post-increment* (*decrement*) *after* its use as the effective address for the operand being accessed. Four permutations are therefore possible: auto pre-increment, auto pre-decrement, auto post-increment, and auto post-decrement.

*automatic  
increment/decrement*

What makes this mode particularly powerful, though, is that the amount of increment/decrement can range from 1 to 8. Thus, arrays consisting of byte, word (16-bit), or long (32-bit) data elements can be handled with equal ease.

Fortunately, the assembly language format for the “indexed auto” mode is fairly intuitive. An integer, ranging from 1 to 8, specifies the amount of increment/decrement to be performed, followed by a comma and the desired index register with a prefix or suffix of “+” or “-”. If a *pre*-increment or *pre*-decrement of the index register is to be performed, a “+” or “-” sign is placed *before* the index register name, respectively (e.g., “ $+X$ ” or “ $-X$ ”). Conversely, if a *post*-increment or *post*-decrement of the index register is to be performed, a “+” or “-” is placed *after* the index register name, respectively (e.g., “ $X+$ ” or “ $X-$ ”). Note that, due to potentially devastating (and meaningless) side effects, the PC *cannot* be used as an index register in this mode; only X, Y, and SP may be used.

### Examples:

LDAA 1,X+	$; (A) \leftarrow ((X)), (X) \leftarrow (X)+1$	<2 bytes, 3 cycles>
STAA 1,-X	$; (X) \leftarrow (X)-1, ((X)) \leftarrow (A)$	<2 bytes, 2 cycles>
LDAA 2,+Y	$; (Y) \leftarrow (Y)+2, (A) \leftarrow ((Y))$	<2 bytes, 3 cycles>
STAA 2,Y-	$; ((Y)) \leftarrow (A), (Y) \leftarrow (Y)-2$	<2 bytes, 2 cycles>
LDAA 1,SP+	$; (A) \leftarrow ((SP)), (SP) \leftarrow SP+1$	<2 bytes, 3 cycles>
STAA 1,-SP	$; (SP) \leftarrow (SP)-1, ((SP)) \leftarrow (A)$	<2 bytes, 2 cycles>

The first example illustrates the classic approach to “bumping” through an array or string consisting of single-byte data elements or ASCII characters. Taken together, the first two examples illustrate how an index register can be used as an “auxiliary” stack pointer (for a stack in which the pointer addresses the top stack item, and growth is toward decreasing addresses): “LDAA 1,X+” is equivalent to “popping A” off an auxiliary stack, while “STAA 1,-X” is equivalent to “pushing A” onto an auxiliary stack. If SP is used as the index register (as shown in the last two examples), “LDAA 1,SP+” and “STAA 1,-SP” are equivalent to “popping A” off the system stack and “pushing A” onto the system stack, respectively.

*auxiliary stack pointer*

### Asking About ASCII

A topic virtually impossible to avoid in a beginning course on microprocessors or microcontrollers is ASCII (pronounced “as-key”) code. This acronym stands for American Standard Code for Information Interchange, a 7-bit coding scheme for alphanumeric characters transmitted from keyboards or to display devices. It was originally used in conjunction with mechanical teletype machines (readers who know what an “ASR33” is are “really old”). Included in the coding scheme are a number of “control” characters, the most famous of which include: CTRL-A (\$00), the ASCII null character; CTRL-D (\$04), the end-of-transmission character; line feed (\$0A); carriage return (\$0D); CTRL-H (\$08), the backspace character; and everyone’s favorite, CTRL-G (\$07), the “bell” character.

### Indexed Indirect

At first glance, indirection appears to be (at best) completely nonsensical, and (at worst) hopelessly confusing. What purpose is served by an addressing mode that first requires a memory access to obtain a (16-bit) pointer, followed by a subsequent access using that pointer to obtain the desired operand? After all, use of a similar kind of “indirection” in football is primarily intended to confuse the opposition, not “help” it!

Fortunately, there are good uses for indirection that transcend football. A key use is the implementation of what might generically be referred to as a “jump table”, i.e., a table of pointers to different subroutines (also called a “vector table”, since it points to “where to go”). The basic idea is to access the address of the desired subroutine from a table of pointers as a function of an index variable, and then “go to” that routine. Such a transfer of control is also referred to as an *indirect jump*.

*jump/vector table*

*indirect jump*

The 68HC12 supports two variations of indirection, which Motorola includes under the category of “indexed” (since they are merely “indirect” versions of two “conventional” indexed modes described previously). *Indexed-indirect with constant offset* is simply the indirect version of indexed addressing with 16-bit constant offset, and *indexed-indirect with accumulator offset* is the indirect version of indexed addressing with 16-bit accumulator (D) offset. In both cases, brackets around the operand field signify to the assembler program that the indirect version of these indexed modes is specified. Note that the pointer accessed from memory occupies two successive bytes, with the high byte of that pointer stored in the first location and the low byte stored in the next consecutive location. These two bytes are concatenated together to form the 16-bit pointer that serves as the effective address of the operand.

*indexed-indirect with constant offset*

*indexed-indirect with accumulator offset*

### Examples:

```

LDAA [2,X]      ; (A)←(((X)+2):((X)+3))      { 4 bytes, 6 cycles}
LDAA [100t,X]   ; (A)←(((X)+100):((X)+101))  { 4 bytes, 6 cycles}
LDAA [1000t,X]  ; (A)←(((X)+1000):((X)+1001)) { 4 bytes, 6 cycles}
STAA [0,X]       ; (((X)+0):((X)+1))←(A)      { 4 bytes, 5 cycles}
LDAA [D,Y]       ; (A)←(((D)+(Y)))           { 2 bytes, 6 cycles}
STAA [D,Y]       ; (((D)+(Y)))←(A)            { 2 bytes, 5 cycles}

```

Note from the examples above that all the constant offset modes are four bytes in length (opcode byte, post byte, and two extension bytes for the 16-bit offset), while the accumulator offset version occupies only two bytes (opcode byte plus post byte). As was the case for the non-indirect versions of these addressing modes, valid index registers include X, Y, SP, and PC.

### 3.7.3 Addressing Mode Summary

We are now equipped with the background to understand all the addressing mode variants possible for each 68HC12 instruction. We have also begun to see the impact of the addressing mode utilized on both the length of the instruction in memory (byte count) as well as the total number of cycles needed for execution (cycle count). A summary of all the 68HC12 addressing modes that generally apply to data manipulation instructions is provided in Table 3-2.

In an effort to help our trek through the 68HC12 instruction set be a bit less overwhelming and somewhat more intuitive, we will use some “icons” to denote the addressing mode possibilities for each instruction type. These icons will provide a “visual” way to remember the addressing mode variations, in place of the somewhat obtuse “official abbreviations” published by Motorola (here, highlighted in blue). We will use the “ring dot” symbol (◎) as an icon for *inherent* (INH) addressing, based on the “self-contained” nature of this mode (a “better” name for this mode, in some instances, is *register* addressing). For *immediate* (IMM) mode, we will use a *pound sign* (#) as the icon, since it is the symbol used in assembly language source statements to specify that mode. *Direct* (DIR) and *extended* (EXT) modes are lumped together because, from a functional point of view, they work the same way: they allow the instruction to “directly dial” the address of the operand in memory. What better icon, then, to represent *direct* (“local”) or *extended* (“long distance”) addressing modes than a telephone (☎).

While there is quite a bit of variety in the indexed modes, they are all based on use of an index register as a pointer; given this commonality, we will use an “index finger” (☞) icon to represent it. In general, if a given 68HC12 instruction supports indexed addressing, all of the variants (constant offset with one extension byte, constant offset with two extension bytes, accumulator offset, auto pre/post increment/decrement, etc.) are supported – with very few exceptions. Motorola distinguishes among the indexed modes based on the number of extension bytes (beyond the postbyte) used: IDX is shorthand for modes with no extension bytes, IDX1 for modes with one extension byte, and IDX2 for modes with two extension bytes.

Finally, as a natural extension to use of an “index finger” as the icon for indexed addressing, we will place brackets around it ([☞]) to represent the indexed-indirect modes. Motorola distinguishes between the two possibilities here based on the number of extension bytes: the “indirect form” of the two-extension-byte indexed mode is abbreviated [IDX2]; while the indirect form of the accumulator offset indexed mode (where the “D” register is the only possibility) is abbreviated [D,IDX].

*intuitive icons*

*inherent/register*  
INH ◎

*immediate*  
IMM #

*direct/extended*  
DIR/EXT ☎

*indexed* ☞

*no extension bytes* IDX  
*one extension byte* IDX1  
*two extension bytes* IDX2

*indexed-indirect* [☞]

*two extension bytes* [IDX2]

*accumulator offset* [D,IDX]

**Table 3-2 Addressing Mode Summary for Data Manipulation Instructions.**

Icon	Abbrev.	Name	Description	Examples
◎	INH	Inherent/Register	Operand(s) is (are) contained in registers; “inherent” means name of register part of instruction mnemonic	DAA
#	IMM	Immediate	Operand data “immediately follows” opcode; pound sign (#) denotes use of immediate data	LDAA #\$FF LDAA #1
☞	DIR/EXT	Direct/Extended	Effective address of operand (“absolute” location in memory) follows opcode; called “direct” if the address can be contained in a single byte, or “extended” if two bytes are required	LDAA \$FF ;direct STAA 900h ;extended
☞	IDX IDX1 IDX2	Indexed with Constant Offset	Effective address is determined by adding a (signed) constant offset (5-bit, 8-bit, or 16-bit) to an index register (which may be X, Y, SP, or PC)	LDAA 0,X STAA 1,Y LDAA 5,SP STAA 2,PC
	IDX	Indexed with Accumulator Offset	Effective address is determined by adding an (unsigned) accumulator (A, B, or D) to an index register (X, Y, SP, or PC)	LDAA B,X STAA B,Y LDAA D,X
	IDX	Indexed with Auto Pre-/Post-Increment or Decrement	Effective address is determined by an index register (X, Y, or SP) that can be modified prior to its use (pre-inc/dec) or following its use (post-inc/dec); the amount of pre/post modification possible ranges from 1 to 8	STAA 1,-X ;pre-dec LDAA 1,X+ ;post-inc STAA 8,+X ;pre-inc LDAA 8,X- ;post-dec
☞	[IDX2]	Indexed-Indirect with Constant Offset	Indexed with constant offset addressing mode is used to access a 16-bit pointer in memory, which is then used as the effective address of the operand; brackets denote use of indirection	LDAA [4,X] STAA [2,Y]
	[D,IDX]	Indexed-Indirect with Accumulator Offset	Indexed with accumulator (D) offset mode is used to access a 16-bit pointer in memory, which is then used as the effective address of the operand; brackets denote use of indirection	LDAA [D,Y] STAA [D,X]

## 3.8 Motorola 68HC12 Instruction Set Overview

Continuing with the “Norm” analogy introduced at the beginning of this chapter, the best way to view a machine’s instruction set is as a collection of “tools in a toolbox.” Just as there are basic “tool types” available to a carpenter (e.g., saws, hammers, screwdrivers, wrenches, routers, biscuit joiners, etc.), so too are there basic “instruction types” available to a programmer. The basic instruction types supported by most computers include: data transfer, arithmetic, logical, transfer-of-control, machine control, and “special” (i.e., atypical instructions for specialized applications such as graphics or signal processing). And just as there is a wide variety of different “saw group” tools (table saws, band saws, hack saws, etc.) available to a carpenter, there is a wide variety of “arithmetic group” instructions (add, subtract, multiply, divide, etc.) available to a programmer.

*tool types*

*instruction types*

Our approach, then, will be to break the 68HC12’s instruction set into the six major groups listed above. Because we are already familiar with the addressing mode variants possible for data manipulation instructions, we will describe the syntax of each instruction independent of the addressing mode variants (the abbreviation *addr* will be used to denote the effective address). The addressing mode possibilities for each instruction will be indicated using the icons (Ⓐ, #, ☎, ↳, [ ]) described in the previous section. To help make the discussion a bit more tractable, we will focus our attention on the variants of a given instruction that are most commonly used – as always, the “rest of the story” (instruction cycle counts and “weird” but legal variants) can be obtained from the official Motorola documentation (see <http://mot-sps.com> for complete details).

One disclaimer before we embark on the classifications. Admittedly, some of the classifications represent a “judgment call” – for example, the “sign extend” instruction can be construed as either a “data transfer” instruction or an “arithmetic” instruction. Remember, though, that our objective is to develop a *framework* that will help us remember the instructions based on function. Returning to the “Norm” analogy for a moment, if our objective is to drive a nail, both a hammer and a socket wrench will “work” – the fact that we have classified the latter as a “wrench group” tool has no bearing on this utility.

*judgment call*

*framework*

### 3.8.1 Data Transfer Group Instructions

As its name implies, the function that links members of this group is *transfer of data* – which includes *load*, *store*, *move*, *exchange*, and *stack manipulation* operations. In general, this group of instructions has a limited effect on the machine’s condition codes (“CC” or “flags”). Move (also called “transfer”) and exchange instructions have *no effect* on the

*transfer of data*

condition code bits, while load and store instructions affect only the negative (N), zero (Z), and overflow (V) flags. Note that the carry/borrow (C) flag is purposefully *not* affected by load and store instructions, since a common application of the “C” condition code bit is to propagate a carry (or borrow) forward in an extended precision arithmetic routine.

Load ([LD](#)) and store ([ST](#)) instructions are listed in Table 3-3. Note that all applicable variants of the addressing modes are supported, with the exception of immediate mode for stores (which would be meaningless). Also note that store instructions affect the condition code bits just like the load instructions, even though this would appear to be “unnecessary” and perhaps even counterintuitive (recall that the simple computer we designed in Chapter 2 did *not* affect the flags when a store was executed). In fact, the first time the author noted that the 68HC12 affects flags as a “side effect” of store instructions, he thought it was a mistake (and didn’t believe it until he tried it out on a “live” microcontroller)!

[load register  
LD](#)

[store register  
ST](#)

**Table 3-3 Data Transfer Group: Load and Store Registers.**

Description	Mnemonic	Operation	CC	Examples	Mode	~
Load Register	LD <i>rb</i> <i>addr</i> <i>rb</i> = A, B  <i>addr</i> = #   [ <td data-kind="parent" data-rs="5">(rb) ← (addr)</td> <td data-kind="parent" data-rs="5">N ← <math>\ddagger</math> Z ← <math>\ddagger</math> V ← 0</td> <td>LDAA #1</td> <td>#</td> <td>1</td>	(rb) ← (addr)	N ← $\ddagger$ Z ← $\ddagger$ V ← 0	LDAA #1	#	1
		LDAA \$FF		3		
		LDAB 900h		3		
		LDAA 1,X		3		
		LDAA B,Y		3		
	LD <i>rw</i> <i>addr</i> <i>rw</i> = D, X, Y, S  <i>addr</i> = #   [ <td data-kind="parent" data-rs="5">(rw) ← (addr)</td> <td data-kind="parent" data-rs="5">N ← <math>\ddagger</math> Z ← <math>\ddagger</math> V ← 0</td> <td>LDAB 2,Y+</td> <td></td> <td>3</td>	(rw) ← (addr)	N ← $\ddagger$ Z ← $\ddagger$ V ← 0	LDAB 2,Y+		3
		LDAA [0,Y]	[ <td>6</td>	6		
		LDAA [D,X]	[ <td>6</td>	6		
		LDD #1	#	2		
		LDS #\$A00	#	2		
Store Register	ST <i>rb</i> <i>addr</i> <i>rb</i> = A, B  <i>addr</i> =   [ <td data-kind="parent" data-rs="5">(addr) ← (rb)</td> <td data-kind="parent" data-rs="5">N ← <math>\ddagger</math> Z ← <math>\ddagger</math> V ← 0</td> <td>LDX 900h</td> <td></td> <td>3</td>	(addr) ← (rb)	N ← $\ddagger$ Z ← $\ddagger$ V ← 0	LDX 900h		3
		LDY A,X		3		
		LDX [D,Y]	[ <td>6</td>	6		
		STAA \$FF		2		
		STAB 900h		3		
	ST <i>rw</i> <i>addr</i> <i>rw</i> = D, X, Y, S  <i>addr</i> =   [ <td data-kind="parent" data-rs="5">(addr) ← (rw)</td> <td data-kind="parent" data-rs="5">N ← <math>\ddagger</math> Z ← <math>\ddagger</math> V ← 0</td> <td>STAA 1,X</td> <td></td> <td>2</td>	(addr) ← (rw)	N ← $\ddagger$ Z ← $\ddagger$ V ← 0	STAA 1,X		2
		STAA B,Y		2		
		STAB 2,Y+		2		
		STAA [0,Y]	[ <td>5</td>	5		
		STAA [D,X]	[ <td>5</td>	5		
	STD 900h STX 2,Y STY A,X STX [2,Y] STS [D,Y]			STD 900h		3
				STX 2,Y		2
				STY A,X		2
				STX [2,Y]	[ <td>5</td>	5
				STS [D,Y]	[ <td>5</td>	5

Load effective address ([LEA](#)), one of the 68HC12's most non-intuitive (and confusing) instructions, is documented in Table 3-4. This instruction loads the named index register (X, Y, or SP) with the *effective address* generated by the indexed mode specified in the operand field (note the absence of parenthesis around "addr" in the description). The reason most thinking adults have trouble with this is that generally an effective address, once generated, is used to access an operand from memory; here, though, the *effective address itself* is loaded into the named index register. Why (and where) would one use such a capability?

[load effective address](#)

[LEA](#)

A “less intimidating” way to understand what the LEA instruction does is to think of it as a powerful way to modify the contents of an index register – through the addition of a signed constant (up to 16-bits in length), an (unsigned) accumulator, or even an auto-increment/decrement mode. Any indexed addressing mode can be used to specify the modification desired, and any index register (X, Y, SP, PC) can serve as the “source” of the modification. Note, however, that certain variants have no “socially redeeming value”. For example, if the source and destination index registers are the same, auto post-increment/decrement does not affect that register’s contents (e.g., LEAX 1,X+ and LEAY 2,Y+ have no effect on the contents of X or Y, respectively). This is because the effective address generated is based on the current value of the index register specified, *not* the “post-modified” version.

Returning to the question posed above, the LEA instruction is typically used to add/subtract an arbitrary constant to/from an index register or, stated another way, to increment/decrement an index register by an arbitrary amount. It is also used to initialize an index register relative to another (e.g., Y initialized to one greater than X). While somewhat arcane, the LEA instruction will prove quite useful in many applications.

**Table 3-4 Data Transfer Group: Load Effective Address.**

Description	Mnemonic	Operation	CC	Examples	Mode	~
Load Effective Address	LEArw <i>addr</i> <i>rw</i> = X, Y, S  <i>addr</i> =	( <i>rw</i> ) ← <i>addr</i>	–	LEAX 2, Y LEAY B, X LEAX D, SP LEAS 1, X+ LEAY 2, -X LEAS 200t, SP LEAX 1000t, SP	2	

The exchange ([EXG](#)) instruction variants are listed in Table 3-5. Most of the time, this instruction is used to “swap” the contents of two like-sized registers. “Mismatched” swaps are “legal”, though, and included for the sake of completeness (the author has yet to find a good use for this “feature”, however). In a mismatched swap, the byte-register ([rb](#)) is

[exchange/swap](#)

[EXG](#)

swapped with the low byte of the word-register (*rw*), and the high byte of the word-register (*wh*) is cleared to zero. Note that all variations of EXG execute in a single cycle, occupy two bytes (an opcode byte followed by a post byte that indicates the registers involved), and do not affect any of the condition code bits. While Motorola officially calls the addressing mode used by this instruction *inherent*, the author believes this to be a misnomer. Since the registers involved are indicated by a *post byte* rather than “inherently” specified by the instruction opcode, a more accurate name for the addressing mode used here would be “register”.

*post byte**register addressing***Table 3-5** Data Transfer Group: Exchange Instructions.

Description	Mnemonic	Operation	CC	Examples	Mode	$\sim$
Exchange Register Contents	EXG <i>rb1,rb2</i> <i>rb</i> = A, B, CCR	$(rb1) \leftrightarrow (rb2)$	—	EXG A, B	◎	1
	EXG A, CCR	◎	1			
	EXG D, X	◎	1			
	EXG X, Y	◎	1			
	EXG A, X	◎	1			
	EXG B, Y	◎	1			
	EXG CCR, D	◎	1			
	EXG X, A	◎	1			
	EXG Y, B	◎	1			
	EXG D, CCR	◎	1			

What Motorola calls “transfer” (*TFR*) instructions – which the rest of the civilized world calls “move” instructions, but might more appropriately be called “copy” instructions – are listed in Table 3-6. The main difficulty here is keeping track of which register is the *source* of the transfer and which is the *destination*. Long ago (where “long” is about 30 years), someone at Motorola decided that the first register name in the operand field should be the source of the transfer and the second the destination. (This, of course, was done with the primary intention of being “different than Intel”, that had adopted a “destination followed by source” format for their “MOV” instructions.) Thus, “TFR A,B” means transfer (or *copy*) the contents of register A to register B. As is the case with the EXG instruction, transfers of mismatched size are also legal for TFR: “byte-to-word” transfers are zero-extended (“padded with zeroes”), and “word-to-byte” transfers are merely truncated. Also like the EXG instruction, all variants of TFR execute in a single cycle, occupy two bytes (an opcode byte followed by a post byte), and do not affect any condition code bits. Again, even though Motorola officially calls the addressing mode used by the TFR instruction *inherent*, a better name would be “register”.

*move/copy registers**TFR*

**Table 3-6** Data Transfer Group: Transfer (Move) Register Instructions.

Description	Mnemonic	Operation	CC	Examples	Mode	$\sim$
Transfer (Move) Register	TFR <i>rb1,rb2</i> <i>rb</i> = A, B, CCR	$(rb1) \rightarrow (rb2)$	—	TFR A, B	◎	1
	TFR <i>rw1,rw2</i> <i>rw</i> = D, X, Y, S	$(rw1) \rightarrow (rw2)$	—	TFR A, CCR	◎	1
	TFR <i>rw,rb</i> <i>rw</i> = D, X, Y, S <i>rb</i> = A, B, CCR	$(rw) \rightarrow (rb)$	—	TFR X, D	◎	1
				TFR D, Y	◎	1
				TFR X, A	◎	1
				TFR Y, B	◎	1
				TFR X, CCR	◎	1
	TFR <i>rb,rw</i> <i>rb</i> = A, B, CCR <i>rw</i> = D, X, Y, S	\$00:( <i>rb</i> ) $\rightarrow$ ( <i>rw</i> )	—	TFR A, X	◎	1
				TFR B, Y	◎	1
				TFR CCR, D	◎	1

The so-called “sign extend” (**SEX**) instruction, described in Table 3-7, can be thought of as a specialized version of a “mismatched” (byte-to-word) TFR. Instead of padding the upper byte of the destination word-register with zeroes, the “sign extend” instruction pads it with the sign (most significant bit) of the source byte-register (as such, a better mnemonic for this operation might have been “TFRS”). The SEX instruction can therefore be used to sign extend an 8-bit offset before adding it to a 16-bit index register. Note that despite being a “legal” variant, sign extending the condition code register (CCR) makes absolutely no sense.

*sign extend  
SEX*

**Table 3-7** Data Transfer Group: Sign Extend Instruction.

Description	Mnemonic	Operation	CC	Examples	Mode	$\sim$
Sign Extend Byte Register	SEX <i>rb,rw</i> <i>rb</i> = A, B, CCR <i>rw</i> = D, X, Y, S	$(rb) \rightarrow (rw)$ <i>rwh</i> padded with sign of <i>rb</i>	—	SEX B, Y	◎	1

The next set of data transfer group instructions, “move memory” (**MOV**), is listed in Table 3-8. These “new” instructions (not included in Motorola 68xx predecessor instruction sets) provide a convenient way to transfer a byte or word of data from one memory location to another, replacing the “LD-ST” sequence previously required with a single instruction. We will find them particularly useful for initializing the peripheral device registers (located in the first 256-byte block in the processor’s address space). Like the TFR assembly mnemonic, the source operand address is listed first, followed by the destination address. Source operands can be specified using immediate, extended, or any “short form” indexed mode (i.e., indexed modes that do not utilize extension bytes); destination operands are limited to extended and “short form” indexed modes. A total of six source-destination addressing mode permutations are therefore possible; an example of each is given in Table 3-8. MOV instructions can occupy

*move memory  
MOV*

as many as six bytes, and take as long as six cycles to execute; they can also be “tricky” to interpret, given there can be as many as four items (separated by commas) in the operand field. Like the EXG and TFR instructions, MOV instructions do not affect any of the condition code bits.

**Table 3-8 Data Transfer Group: Move Memory Instructions.**

Description	Mnemonic	Operation	CC	Examples	Mode	~
Move Memory	MOVB <i>addr1,addr2</i> <i>addr1</i> = # <i>addr2</i> =	(addr1) → (addr2)	-	MOVB #\$FF,\$900	# →	4
				MOVB #2,0,X	# →	4
				MOVB \$900,\$901	→	6
				MOVB \$900,1,X	→	5
				MOVB 1,X-, \$900	→	5
				MOVB 1,X+, 2,Y+	→	5
	MOVW <i>addr1,addr2</i> <i>addr1</i> = # <i>addr2</i> =	(addr1) → (addr2) (addr1+1) → (addr2+1)	-	MOVW #\$FFFF,\$900	# →	5
				MOVW #1,0,X	# →	4
				MOVW \$900,\$902	→	6
				MOVW \$900,2,X	→	5
				MOVW 2,X-, \$900	→	5
				MOVW 2,X+, 4,Y+	→	5

Note: Only indexed modes () that employ *no extension bytes* (beyond the post byte) can be used with the move memory instructions; this implies that only short constant offsets (-15 to +16) are valid.

The final set of data transfer instructions, listed in Table 3-9, perform stack-related data transfers. In our simple computer of Chapter 2, we called these operations “push” and “pop” – the names for these operations used by virtually every other manufacturer of microprocessors...except Motorola. Again, just to be “different than Intel”, Motorola chose the mnemonics “push” (PSH) and “pull” (PUL), respectively, for stack-related data transfers.

*push*  
**PSH**

*pull (pop)*  
**PUL**

### Push Pulling

Notable by their absence are instructions that allow the PC or SP to be pushed onto or pulled off the stack. While pushing either of these registers onto the stack is of no consequence, pulling either of them off the stack would most likely cause “anomalous behavior” (i.e., cause “bits to fly all over the place”). For example, if the PC could be pulled from the stack, execution would continue at the location specified by the top stack item – this only makes sense if a “return address” has been placed on the stack by a calling program (recall the simple computer’s RTS instruction); otherwise, a program could quickly arrive at an “unknown location”. A somewhat more insidious problem might occur if the SP could be pulled from the stack. Here, the location of the entire stack would change, effectively canceling all bets as to the stack’s current contents! In summary, there are good reasons why the PC and SP are not included in the list of registers that can be pushed or pulled.

There are two basic variants of PSH and PUL: one for byte-registers (A, B, CCR) and another for word-registers (D, X, Y). Note that neither SP nor PC can be pushed or pulled. Note also that the same stack convention used by our simple computer is used here: the *stack pointer* (SP) points to the location in which the *top stack item* is stored (for *word-length* items, SP points to the *high byte* of the top stack item). Generally, PSH and PUL do not affect any of the condition code bits – with the obvious exception of PULC, which affects *all* the condition code bits.

**Table 3-9 Data Transfer Group: Stack Manipulation Instructions.**

Description	Mnemonic	Operation	CC	Examples	Mode	~
Push register onto stack	PSH <sub>rb</sub> <i>rb</i> = A, B, C	(SP) $\leftarrow$ (SP) – 1 ((SP)) $\leftarrow$ ( <i>rb</i> )	–	PSHA	◎	2
				PSHB	◎	2
				PSHC	◎	2
	PSH <sub>rw</sub> <i>rw</i> = D, X, Y	(SP) $\leftarrow$ (SP) – 1 ((SP)) $\leftarrow$ ( <i>rw</i> ) (SP) $\leftarrow$ (SP) – 1 ((SP)) $\leftarrow$ ( <i>rw</i> )	–	PSHD	◎	2
				PSHX	◎	2
				PSHY	◎	2
Pull (pop) register from stack	PUL <sub>rb</sub> <i>rb</i> = A, B, C	( <i>rb</i> ) $\leftarrow$ ((SP)) (SP) $\leftarrow$ (SP) + 1	*	PULA	◎	3
				PULB	◎	3
				PULC	◎	3
	PUL <sub>rw</sub> <i>rw</i> = D, X, Y	( <i>rw</i> ) $\leftarrow$ ((SP)) (SP) $\leftarrow$ (SP) + 1 ( <i>rw</i> ) $\leftarrow$ ((SP)) (SP) $\leftarrow$ (SP) + 1	–	PULD	◎	3
				PULX	◎	3
				PULY	◎	3

\* PULC affects *all* the condition code bits, with the exception of X, which cannot be set by a software instruction once it is cleared.

### 3.8.2 Arithmetic Group Instructions

Instructions that perform an arithmetic operation (add, subtract, multiply, divide) are broadly classified here as belonging to the arithmetic group. As one might guess, *most* of these instructions affect *all* of the condition code bits (with a few notable exceptions).

Table 3-10 lists the variations of add (ADD) and subtract (SUB) of which the 68HC12 is capable. The “with carry” versions (ADC and SBC) are provided for implementing extended (or “infinite”) precision add or subtract routines; in Chapter 4, we will learn how to write such routines. For the ADC instruction, the “C” bit of the condition code register is interpreted as a *carry propagated forward*, and is therefore *added* to the result. For the SBC instruction, the “C” bit is interpreted as a *borrow propagated forward*, and is therefore *subtracted* from the result. The “astute digijock(ette)” will

<i>add</i>	ADD
<i>add with carry</i>	ADC
<i>subtract</i>	SUB
<i>subtract with carry</i>	SBC

realize this is equivalent to adding the complement of the “C” bit to the result, i.e., the way we did it in hardware.

In addition to the “memory plus register” add instructions described above, there are two register-to-register add instructions. As documented in Table 3-11, the first of these adds the contents of the two byte accumulators (A and B) and places the result in the A register (**ABA**), while the second adds the (zero-extended) contents of the B register to the X or Y register (**ABX** or **ABY**). The ABX and ABY instructions are artifacts of the “original” 6800 instruction set (*circa* 1975). These instructions have been supplanted by the “LEA” instruction (described previously); 68HC12 assembler programs convert ABX and ABY mnemonics into “LEAX B,X” and “LEAY B,Y” instructions, respectively.

*add B to A  
ABA*

*add B to X  
ABX*

*add B to Y  
ABY*

**Table 3-10** Arithmetic Group: Add/Subtract Instructions.

Description	Mnemonic	Operation	CC	Examples	Mode	$\sim$
Add contents of memory location to register	ADDrb <i>addr</i> <i>rb</i> = A, B	$(rb) \leftarrow (rb) + (\text{addr})$	N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$ H $\leftarrow \downarrow$	ADDA #1 ADDB \$900 ADDA 1,X ADDB A,X ADDA [2,Y]	# \$ 1,X A,X [2,Y]	1 3 3 3 6
	ADCrb <i>addr</i> <i>rb</i> = A, B	$(rb) \leftarrow (rb) + (\text{addr}) + (\text{C})$	N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$ H $\leftarrow \downarrow$	ADCA #1 ADCB \$900 ADCA 1,X ADCB A,X ADCA [2,Y]	# \$ 1,X A,X [2,Y]	1 3 3 3 6
	ADDD <i>addr</i>	$(D) \leftarrow (D) + (\text{addr}):(addr+1)$	N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$	ADDD #1 ADDD \$900 ADDD 1,X ADDD [2,Y]	# \$ 1,X [2,Y]	2 3 3 6
	<i>addr</i> = #   []					
	SUBrb <i>addr</i> <i>rb</i> = A, B	$(rb) \leftarrow (rb) - (\text{addr})$	N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$	SUBA #1 SUBB \$900 SUBA 1,X SUBB A,X SUBA [2,Y]	# \$ 1,X A,X [2,Y]	1 3 3 3 6
	SBCrb <i>addr</i> <i>rb</i> = A, B	$(rb) \leftarrow (rb) - (\text{addr}) - (\text{C})$	N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$	SBCA #1 SBCB \$900 SBCA 1,X SBCB A,X SBCA [2,Y]	# \$ 1,X A,X [2,Y]	1 3 3 3 6
	<i>addr</i> = #   []					
	SUBD <i>addr</i>	$(D) \leftarrow (D) - (\text{addr}):(addr+1)$	N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$	SUBD #1 SUBD \$900 SUBD 1,X SUBD [2,Y]	# \$ 1,X [2,Y]	2 3 3 6
	<i>addr</i> = #   []					

**Table 3-11** Arithmetic Group: Register-to-Register Adds.

Description	Mnemonic	Operation	CC	Examples	Mode	$\sim$
Add registers	ABA	$(A) \leftarrow (A) + (B)$	$N \leftarrow \downarrow$ $Z \leftarrow \downarrow$ $V \leftarrow \downarrow$ $C \leftarrow \downarrow$ $H \leftarrow \downarrow$	ABA	◎	2
	AB <sub>rw</sub> rw = X, Y	$(rw) \leftarrow \$00:(B) + (rw)$	—	ABX ABY	◎	2

Note that ADD, ADC, and ABA are the *only* 68HC12 instructions that (meaningfully) affect the so-called “half carry” (H) condition code bit. That’s because the only instruction that uses the “H” bit is the “decimal adjust A (after add)” (DAA) instruction, described in Table 3-12 (note that an appropriate five-letter mnemonic would be “DAAAA”). The purpose of this instruction is to “correct” the result of an add operation performed on two (packed) binary-coded decimal (BCD) operands, to produce a BCD result (plus a BCD carry, for extended precision applications). “Packed BCD” means that two (4-bit) BCD digits are placed in a single (8-bit) byte.

*decimal adjust A  
DAA*

**Table 3-12** Arithmetic Group: Decimal Adjust “A” Register.

Description	Mnemonic	Operation	CC	Examples	Mode	$\sim$
Decimal Adjust A	DAA	decimal adjust the result of ADD, ADC, or ABA	$N \leftarrow \downarrow$ $Z \leftarrow \downarrow$ $V \leftarrow ?$ $C \leftarrow \downarrow$	DAA	◎	3

When a pair of packed BCD operands is added together, the “H” condition code bit represents the carry out of the “one’s position”, while the “C” condition code bit represents the carry out of the “ten’s position”. Note that this often-misunderstood instruction *does not* “convert” binary operands to BCD format; instead, it simply applies a “correction” to the result obtained from directly adding packed BCD operands (similar in function to the BCD adder circuit reviewed in Chapter 1). The action performed by DAA is illustrated in Figure 3-21. Note that DAA does not produce a meaningful result following a subtract operation, and that the 68HC12 does *not* have an instruction dedicated to performing decimal adjust after subtraction.

*correction function*

Closely associated with add/subtract are instructions that can be used to complement the contents of a register or memory location. The 68HC12 provides two possibilities: a “ones’ complement” (COM) instruction and a “two’s complement” (NEG) instruction, documented in Table 3-13. Both of these instructions support all applicable addressing modes.

*complement  
COM*

*negate  
NEG*

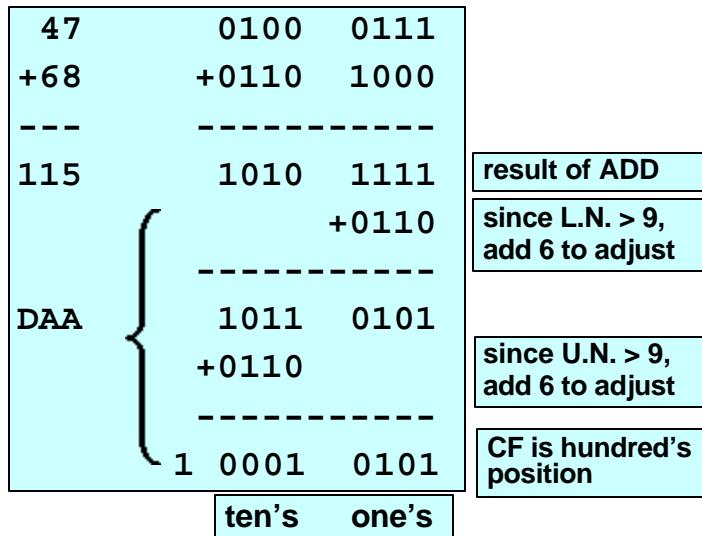


Figure 3-21 Illustration of DAA.

Table 3-13 Arithmetic Group: Complement.

Description	Mnemonic	Operation	CC	Examples	Mode	$\sim$
Ones' complement	COMrb rb = A, B	$(rb) \leftarrow \$FF - (rb)$	N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow 0$ C $\leftarrow 1$	COMA	$\odot$	1
	COM addr addr = $\overline{R} \overline{X} [R]$	$(addr) \leftarrow \$FF - (addr)$	N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow 0$ C $\leftarrow 1$	COM \$900 COM 1,X COM B,X COM [D,Y]	$\overline{R}$ $\overline{X}$ [R]	4 3 3 6
Two's complement	NEGrb rb = A, B	$(rb) \leftarrow \$00 - (rb)$	N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$	NEGB	$\odot$	1
	NEG addr addr = $\overline{R} \overline{X} [R]$	$(addr) \leftarrow \$00 - (addr)$	N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$	NEG \$900 NEG 1,X NEG B,X NEG [D,Y]	$\overline{R}$ $\overline{X}$ [R]	4 3 3 6

The manner in which these two instructions affect the condition code bits deserves some explanation. For the COM instruction, the N and Z flags are set according to the new contents of the affected register or memory location. The overflow (V) flag is cleared and, strictly for “legacy compatibility” reasons, the carry/borrow (C) flag is set (there is no compelling reason, however, for the COM instructions to affect the V and C bits this way). For the NEG instruction, the two’s complement negation of the operand is formed by subtracting it from \$00; the condition code bits are simply set or cleared based on the results of this subtraction.

Also closely related to the subtract instructions is the “compare and test” subgroup, listed in Table 3-14. The “compare” (CMP or CP) instructions work the same as the subtract instructions except the difference calculated is not stored; instead, only the condition codes (N, Z, V, C) are affected. As such, compare instructions are intended for use prior to conditional transfer of control instructions (covered in Section 3.8.4). It is important to note that the condition code bits are set or cleared based on a subtract operation and, in particular, that the C bit (“carry/borrow flag”) is interpreted as a *borrow*. We will discuss the ramifications of this when we cover the “transfer of control” group of instructions.

*compare  
CMP*

A somewhat more “specialized” version of compare is the “test” (TST) instruction, which sets or clears the condition code bits based on subtracting zero from a byte-register or memory location. One might argue that this less general variant of compare really isn’t necessary, given that “TSTA” and “TSTB” are functionally equivalent to “CMPA #0” and “CMPB #0”, respectively. Both TST<sub>rb</sub> and CMP<sub>rb</sub> execute in a single cycle, although the TST<sub>rb</sub> instructions occupy a single byte while the immediate mode version of CMP<sub>rb</sub> occupies two. The “test memory” variant, however, is a bit more useful, since the “compare memory” equivalent would require loading an accumulator with zero. An interesting thing to note about this subgroup is that, since zero is subtracted from the operand, the overflow (V) and carry (C) flags are always cleared (since overflow cannot occur, and there can never be a borrow). The only meaningful condition code bits following a “test” instruction are N and Z.

*test for zero  
TST*

**Table 3-14** Arithmetic Group: Compare/Test.

Description	Mnemonic	Operation	CC	Examples	Mode	~
Compare Accumulators	CBA	set CCR based on (A) – (B)	N $\leftarrow \uparrow$ Z $\leftarrow \uparrow$ V $\leftarrow \uparrow$ C $\leftarrow \uparrow$	CBA	⊕	2
Compare Register with Memory	CMP <sub>rb</sub> addr <i>rb</i> = A, B	set CCR based on ( <i>rb</i> ) – (addr)	N $\leftarrow \uparrow$ Z $\leftarrow \uparrow$ V $\leftarrow \uparrow$ C $\leftarrow \uparrow$	CMPA #2	#	1
	addr = #  []		CMPB \$900		3	
	CP <sub>rw</sub> addr <i>rw</i> = D, X, Y, S		CMPA 2, X		3	
	addr = #  []		CMPB [2, Y]	[]	6	
	TST <sub>rb</sub> <i>rb</i> = A, B	set CCR based on ( <i>rb</i> ) – \$00	N $\leftarrow \uparrow$ Z $\leftarrow \uparrow$ V $\leftarrow 0$ C $\leftarrow 0$	TSTA	⊕	1
	TST addr addr =  []		TSTB	⊕	1	
	TST \$900					
	TST 1, X					
	TST [2, Y]					
	[]					6

The next set of arithmetic group instructions, documented in Table 3-15, provide the capability to increment (**INC**) or decrement (**DEC**) the contents of a register or memory location. The byte-increment/decrement subset affect the N, Z, and V condition code bits; the carry/borrow (C) flag is *not affected* “on purpose” to facilitate use of INC/DEC instructions as loop counters (or “pointer bumpers”) in extended precision arithmetic routines. The word-register (X, Y, SP) increment/decrement subset affects (at most) the Z flag (INS and DES do not affect any flags). Recall that the LEA instruction provides a considerably more powerful and flexible means of incrementing or decrementing a word-register.

*increment*  
**INC**

*decrement*  
**DEC**

Multiply and divide operations comprise the next set of arithmetic group instructions, listed in Tables 3-16 through 3-18. Here there are a number of permutations, depending on the size of the operands (8-, 16-, or 32-bits) and whether or not the operands are signed. Special variants include a fractional divide plus a “multiply-and-accumulate”.

**Table 3-15** Arithmetic Group: Increment/Decrement.

Description	Mnemonic	Operation	CC	Examples	Mode	$\sim$
Increment	<b>INCr</b> <i>r</i> = A, B	$(r) \leftarrow (r) + 1$	N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$	INCA	⊕	1
	<b>INrw</b> <i>rw</i> = X, Y, S	$(rw) \leftarrow (rw) + 1$	Z $\leftarrow \downarrow$	INX INY	⊕	1
			—	INS	⊕	1
	<b>INC addr</b> <i>addr</i> =	$(addr) \leftarrow (addr) + 1$	N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$	INC \$900 INC 1,X INC B,X INC [D,Y]		4 3 3 6
Decrement	<b>DECr</b> <i>r</i> = A, B	$(r) \leftarrow (r) - 1$	N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$	DECB	⊕	1
	<b>DERw</b> <i>rw</i> = X, Y, S	$(rw) \leftarrow (rw) - 1$	Z $\leftarrow \downarrow$	DCX DCY	⊕	1
			—	DCS	⊕	1
	<b>DEC addr</b> <i>addr</i> =	$(addr) \leftarrow (addr) - 1$	N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$	DEC \$900 DEC 1,X DEC B,X DEC [D,Y]		4 3 3 6

Looking first at the multiply instructions in Table 3-16, the basic multiply (**MUL**) instruction – that had its humble beginnings back in the late 1970s with the venerable Motorola 6809 – performs an 8-bit by 8-bit unsigned integer multiply. The A and B registers are used as the source operands, which are overwritten with the result (high byte in A, low byte in B). Only the carry flag (C) is affected by this instruction, which (if desired) can be used to “round” the upper byte (contained in the A register). This rounding capability, which can be implemented by following the MUL instruction

*8x8-bit multiply*  
**MUL**

with an “ADCA #0” instruction (which simply adds the carry bit to the value in the A register), is useful in cases where the operands are construed as (unsigned) binary fractions. We might wish to truncate or round the result if it is destined for an 8-bit digital-to-analog converter.

### Star \* Wars

*In the late 1970's, when both Motorola and Intel were introducing their “second-and-a-half” generation 8-bit microprocessors (the 6809 and 8085, respectively), Motorola attempted to “trump” the 8085 (which beat the 6809 to market) by adding a feature its fiercest competitor (and market dominator) did not have: a multiply instruction. It's not clear how much the much-vaunted MUL instruction affected the 6809's market share, but it was certainly a novel feature for a microprocessor of that era.*

**Table 3-16** Arithmetic Group: Multiply.

Description	Mnemonic	Operation	CC	Examples	Mode	$\sim$
8x8 unsigned integer multiply	MUL	$(D) \leftarrow (A) \times (B)$	$C \leftarrow \$$	MUL	$\odot$	3
16x16 unsigned integer multiply	EMUL	$(Y):(D) \leftarrow (D) \times (Y)$	$N \leftarrow \$$ $Z \leftarrow \$$ $C \leftarrow \$$	EMUL	$\odot$	3
16x16 signed integer multiply	EMULS	$(Y):(D) \leftarrow (D) \times (Y)$	$N \leftarrow \$$ $Z \leftarrow \$$ $C \leftarrow \$$	EMULS	$\odot$	3

**Table 3-17** Arithmetic Group: Multiply and Accumulate.

Description	Mnemonic	Operation	CC	Examples	$\sim$
16x16 integer multiply and accumulate	EMACS <i>addr</i> <i>addr</i> = special	$(addr):(addr+1):(addr+2):(addr+3) \leftarrow (addr):(addr+1):(addr+2):(addr+3) + ((X) \times ((Y)))$	$N \leftarrow \$$ $V \leftarrow \$$ $Z \leftarrow \$$ $C \leftarrow \$$	EMACS \$900	13

Recall from Chapter 1 that, for a binary fraction, the radix point is to the “far left”, making the most significant bit of weight  $2^{-1}$  ( $1/2 = 0.5_{10}$ ), the next most significant bit of weight  $2^2$  ( $1/4 = 0.25_{10}$ ), and so on. Multiplying the bit pattern 10000000b (1/2) by 01000000b (1/4) yields the 16-bit result 00100000 00000000b in (A):(B), or 1/8 ( $0.125_{10}$ ). Here, the result could be truncated to the 8-bit value in the A register with no loss of precision; the C condition code bit is therefore cleared by the MUL instruction to nullify the effect of an ensuing “ADCA #0” instruction.

Consider, however, the case of multiplying the bit pattern 11111111b ( $255/256 = 0.99609375_{10}$ , or “the largest possible 8-bit unsigned fraction”)

*largest possible  
unsigned fraction*

by 01000000b (1/4), which yields 00111111 11000000b in (A):(B), or 255/1024 (0.2490234375<sub>10</sub>). Here, truncating the result to the 8-bit value in the A register produces the result 63/256 (0.24609375<sub>10</sub>), while rounding the result (as described above) produces the result 01000000b in the A register, or (0.25<sub>10</sub>). To enable rounding, the MUL instruction sets the C bit so that an ensuing “ADCA #0” instruction can increment the value in the A register by one. The “astute digijock(ette)” will recognize that rounding should be performed when the most significant bit of the B register (the lower byte of the result) is one, which is exactly how the C condition code bit is affected by the MUL instruction.

*truncating*

Before leaving the MUL instruction, it is important to note that the operands in A and B can also be construed as simply unsigned integers. For example, multiplying 32<sub>10</sub> (00010000b) by 64<sub>10</sub> (00100000b) yields 00000010 00000000b in (A):(B), or 2048<sub>10</sub>. For multiplication of integers, the C condition code bit holds “no social significance”.

Continuing with the “extended” (16-bit × 16-bit) multiply instructions in Table 3-16, we find that they basically work the same as the “original” MUL instruction, but with some notable differences. Here, the D and Y registers are used to contain the two 16-bit operands, while the 32-bit result is placed in (Y):(D). Like the MUL instruction, EMUL and EMULS use the C condition code bit to facilitate rounding of binary fractions: here, C is set to the most significant bit of the result in the D register (i.e., the low-word of the result). Unlike MUL, though, both extended multiply instructions affect the N and Z condition code bits. The only difference between EMUL and EMULS is that the latter instruction assumes the operands are *signed* (two’s complement) integers or fractions.

*extended 16x16-bit multiply**EMUL (unsigned)  
EMULS (signed)*

The 68HC12’s “multiply and accumulate” (EMACS) instruction, described in Table 3-17, is rarely found in “generic” microcontrollers. Rather, it is an instruction that is typically found only in so-called digital signal processor (DSP) chips. The “MAC” (multiply and accumulate) operation is a staple of common signal processing applications such as digital filters and Fast Fourier Transforms (FFTs). In the 68HC12 implementation of EMACS, two 16-bit signed operands (pointed to by the X and Y registers) are multiplied together; the 32-bit intermediate result obtained is then added to a 32-bit “running sum” stored in memory.

*multiply and accumulate**EMACS*

The main difference between the 68HC12’s EMACS instruction and an equivalent that might be found on a 16-bit integer DSP chip is *speed*: on the 68HC12, execution of the EMACS instruction consumes 13 cycles; while on a DSP chip, the equivalent operation is typically executed in a *single* cycle. The primary impediment to speed on the 68HC12 is lack of a sufficient number of registers – not only to contain the 32-bit accumulated result, but also to provide pointers for the operand arrays. Short of adding

additional registers, the only solution was to use four consecutive memory locations as the 32-bit “accumulator”. Given that the (starting) address of this 32-bit accumulator is specified using extended addressing mode and that the X and Y registers are used as pointers to the two operand arrays, there is no “conventional” addressing mode name that is applicable (hence the designation *special*).

The various possibilities for performing a “divide” operation on the 68HC12 are documented in Table 3-18. One important thing to note, in contrasting this set of instructions to the “multiply” sub-group, is that integers and fractions are handled differently. With that in mind, let’s examine the integer divide (**DIV** and **IDIVS**) instructions first. Here, the D register is used to contain a 16-bit dividend (unsigned for IDIV, signed for IDIVS) and the X register is used to contain a 16-bit (unsigned or signed) divisor. The resulting 16-bit quotient is placed in the X register, while the 16-bit remainder is placed in the D register. If a “divide-by-zero” is attempted, the C condition code bit is set and the quotient is set to \$FFFF (the remainder is indeterminate). For both IDIV and IDIVS, the Z condition code bit is set when a quotient of zero is generated. IDIV and IDIVS differ, however, in how they affect the N and V bits. The N bit is not affected by the unsigned divide (IDIV), but is affected as expected (set to the sign of the quotient) by the signed divide (IDIVS). The V bit is simply cleared by the IDIV instruction, but is set by IDIVS if two’s complement overflow occurs. An example of where two’s complement overflow occurs is attempting to divide the “largest negative 16-bit signed integer” ( $-32,768_{10} = \$8000$ ) by minus one (\$FFFF). Theoretically, the result  $+32,768_{10}$  should be produced, but since the “largest positive 16-bit signed integer” is  $+32,767_{10}$  (\$7FFF), overflow occurs.

*integer 16x16-bit divide*

**IDIV (unsigned)**  
**IDIVS (signed)**

The “extended” divides (**EDIV** and **EDIVS**) are so-called because the dividend is extended to 32-bits; the divisor, quotient, and remainder, however, are limited to 16-bits. The Y register concatenated with the D register is used to contain the 32-bit dividend, while the X register is used to contain the 16-bit divisor. The 16-bit quotient is placed in the Y register, and the 16-bit remainder is placed in the D register. EDIVS (the “signed” version) affects the condition code bits (N, Z, V, C) the same way IDIVS does, but EDIV (the “unsigned” version) differs from IDIV – primarily due to the disparity between the length of the dividend and quotient. Instead, EDIV affects the condition code bits the same way EDIVS does, except for the overflow (V) bit. Since the quotient is limited to 16-bits, an unsigned result exceeding  $65,535_{10}$  (\$FFFF) can be generated (e.g., dividing anything with a non-zero “upper-word” by one).

*extended 32x16-bit integer divide*

**EDIV (unsigned)**  
**EDIVS (signed)**

**Table 3-18** Arithmetic Group: Divide.

Description	Mnemonic	Operation	CC	Examples	Mode	$\sim$
16÷16 unsigned integer divide	IDIV	$(X) \leftarrow (D) \div (X)$ $(D) \leftarrow \text{remainder}$	$V \leftarrow 0$ $Z \leftarrow \uparrow$ $C \leftarrow \uparrow$	IDIV	$\odot$	12
16÷16 signed integer divide	IDIVS	$(X) \leftarrow (D) \div (X)$ $(D) \leftarrow \text{remainder}$	$N \leftarrow \uparrow$ $V \leftarrow \uparrow$ $Z \leftarrow \uparrow$ $C \leftarrow \uparrow$	IDIVS	$\odot$	12
32÷16 unsigned integer divide	EDIV	$(Y) \leftarrow (Y):(D) \div (X)$ $(D) \leftarrow \text{remainder}$	$N \leftarrow \uparrow$ $V \leftarrow \uparrow$ $Z \leftarrow \uparrow$ $C \leftarrow \uparrow$	EDIV	$\odot$	11
32÷16 signed integer divide	EDIVS	$(Y) \leftarrow (Y):(D) \div (X)$ $(D) \leftarrow \text{remainder}$	$N \leftarrow \uparrow$ $V \leftarrow \uparrow$ $Z \leftarrow \uparrow$ $C \leftarrow \uparrow$	EDIVS	$\odot$	12
32÷16 unsigned fraction divide	FDIV	$(X) \leftarrow (D) \div (X)$ $(D) \leftarrow \text{remainder}$	$V \leftarrow \uparrow$ $Z \leftarrow \uparrow$ $C \leftarrow \uparrow$	FDIV	$\odot$	12

The final member of the “divide” sub-group, fractional divide ([FDIV](#)), is also perhaps the most misunderstood. The key is to remember that the two 16-bit operands are construed as unsigned binary fractions (i.e., with the radix point to the “far left”): the dividend is contained in the D register, and the divisor is contained in the X register. After execution, the quotient is placed in the X register and the remainder is placed in D. The remainder can be resolved into the next-most-significant 16 fractional result bits through execution of another FDIV instruction.

*fractional divide  
FDIV*

As an illustrative example, if the dividend is 1/8 (\$2000) and the divisor is 1/2 (\$8000), the result will be 1/4 (\$4000). The Z condition code bit, as expected, is set if the quotient is zero; and, like the other 68HC12 divides, the C bit is set if a “divide-by-zero” is attempted. If the divisor is less than or equal to the dividend, the V bit is set and the quotient is set to \$FFFF (the remainder is indeterminate). “Reversing” the example cited above – i.e., using a dividend of 1/2 (\$8000) and divisor of 1/8 (\$2000) – will produce a result of “overflow”.

One last note about the “divide” sub-group: they are all “cycle hogs”, consuming 11-12 clock ticks to execute. This is in contrast to the 3 cycles consumed by each of the various multiply instructions.

The “min/max” instructions ([MIN/MAX](#), [EMIN/EMAX](#)), listed in Table 3-19, constitute the final subset of arithmetic group instructions. These instructions compare two unsigned operands – one of which is an accumulator (“A” for the 8-bit version, “D” for the 16-bit version) and the other of which resides in memory – and places the larger/smaller of the two in the named accumulator or in memory. These instructions only use

*8-bit unsigned  
min/max  
MIN  
MAX*

the indexed and indexed-indirect addressing modes. There are eight permutations, based on: the size of the operands (8- or 16-bits), whether the destination is memory or the accumulator, and whether a “min” or “max” is performed. The condition codes (N, Z, V, C) are affected based on subtracting the value in memory from the named accumulator.

**16-bit unsigned  
min/max  
EMIN  
EMAX**

**Table 3-19 Arithmetic Group: Minimum/Maximum.**

Description	Mnemonic	Operation	CC	Examples	Mode	~
Unsigned 8-bit Minimum	MINA <i>addr</i> <i>addr</i> = $\ddagger [ \ddagger ]$	$(A) \leftarrow \min \{(A), (\text{addr})\}$	$N \leftarrow \ddagger$ $Z \leftarrow \ddagger$ $V \leftarrow \ddagger$ $C \leftarrow \ddagger$	MINA 0 , X	$\ddagger$	4
				MINA 2 , X+	$\ddagger$	4
				MINA 1000t , Y	$\ddagger$	5
				MINA [ D , X ]	$[ \ddagger ]$	7
				MINA [ 2 , Y ]	$[ \ddagger ]$	7
	MINM <i>addr</i> <i>addr</i> = $\ddagger [ \ddagger ]$	$(\text{addr}) \leftarrow \min \{(A), (\text{addr})\}$	$N \leftarrow \ddagger$ $Z \leftarrow \ddagger$ $V \leftarrow \ddagger$ $C \leftarrow \ddagger$	MINM 0 , X	$\ddagger$	4
				MINM 2 , X+	$\ddagger$	4
				MINM 1000t , Y	$\ddagger$	5
				MINM [ D , X ]	$[ \ddagger ]$	7
				MINM [ 2 , Y ]	$[ \ddagger ]$	7
Unsigned 8-bit Maximum	MAXA <i>addr</i> <i>addr</i> = $\ddagger [ \ddagger ]$	$(A) \leftarrow \max \{(A), (\text{addr})\}$	$N \leftarrow \ddagger$ $Z \leftarrow \ddagger$ $V \leftarrow \ddagger$ $C \leftarrow \ddagger$	MAXA 0 , X	$\ddagger$	4
				MAXA 2 , X+	$\ddagger$	4
				MAXA 1000t , Y	$\ddagger$	5
				MAXA [ D , X ]	$[ \ddagger ]$	7
				MAXA [ 2 , Y ]	$[ \ddagger ]$	7
	MAXM <i>addr</i> <i>addr</i> = $\ddagger [ \ddagger ]$	$(\text{addr}) \leftarrow \max \{(A), (\text{addr})\}$	$N \leftarrow \ddagger$ $Z \leftarrow \ddagger$ $V \leftarrow \ddagger$ $C \leftarrow \ddagger$	MAXM 0 , X	$\ddagger$	4
				MAXM 2 , X+	$\ddagger$	4
				MAXM 1000t , Y	$\ddagger$	5
				MAXM [ D , X ]	$[ \ddagger ]$	7
				MAXM [ 2 , Y ]	$[ \ddagger ]$	7
Unsigned 16-bit Minimum	EMIND <i>addr</i> <i>addr</i> = $\ddagger [ \ddagger ]$	$(D) \leftarrow \min \{(D), (\text{addr}):(addr+1)\}$	$N \leftarrow \ddagger$ $Z \leftarrow \ddagger$ $V \leftarrow \ddagger$ $C \leftarrow \ddagger$	EMIND 0 , X	$\ddagger$	4
				EMIND 2 , X+	$\ddagger$	4
				EMIND 1000t , Y	$\ddagger$	5
				EMIND [ D , X ]	$[ \ddagger ]$	7
				EMIND [ 2 , Y ]	$[ \ddagger ]$	7
	EMINM <i>addr</i> <i>addr</i> = $\ddagger [ \ddagger ]$	$(addr):(addr+1) \leftarrow \min \{(D), (\text{addr}):(addr+1)\}$	$N \leftarrow \ddagger$ $Z \leftarrow \ddagger$ $V \leftarrow \ddagger$ $C \leftarrow \ddagger$	EMINM 0 , X	$\ddagger$	4
				EMINM 2 , X+	$\ddagger$	4
				EMINM 1000t , Y	$\ddagger$	5
				EMINM [ D , X ]	$[ \ddagger ]$	7
				EMINM [ 2 , Y ]	$[ \ddagger ]$	7
Unsigned 16-bit Maximum	EMAXD <i>addr</i> <i>addr</i> = $\ddagger [ \ddagger ]$	$(D) \leftarrow \max \{(D), (\text{addr}):(addr+1)\}$	$N \leftarrow \ddagger$ $Z \leftarrow \ddagger$ $V \leftarrow \ddagger$ $C \leftarrow \ddagger$	EMAXD 0 , X	$\ddagger$	4
				EMAXD 2 , X+	$\ddagger$	4
				EMAXD 1000t , Y	$\ddagger$	5
				EMAXD [ D , X ]	$[ \ddagger ]$	7
				EMAXD [ 2 , Y ]	$[ \ddagger ]$	7
	EMAXM <i>addr</i> <i>addr</i> = $\ddagger [ \ddagger ]$	$(addr):(addr+1) \leftarrow \max \{(D), (\text{addr}):(addr+1)\}$	$N \leftarrow \ddagger$ $Z \leftarrow \ddagger$ $V \leftarrow \ddagger$ $C \leftarrow \ddagger$	EMAXM 0 , X	$\ddagger$	4
				EMAXM 2 , X+	$\ddagger$	4
				EMAXM 1000t , Y	$\ddagger$	5
				EMAXM [ D , X ]	$[ \ddagger ]$	7
				EMAXM [ 2 , Y ]	$[ \ddagger ]$	7

In summary, the arithmetic group includes add/subtract, decimal adjust, complement/negate, compare/test, increment/decrement, multiply/divide, and min/max instructions. While there is no “direct” support for floating point numbers, software libraries are available for this purpose.

### 3.8.3 Logical Group Instructions

Instructions that perform logical manipulation and testing of data – including AND, OR, XOR, shifts, and rotates – are members of this group. We will find this group of instructions particularly useful for interrogating or manipulating individual bits (or sets of bits) contained in peripheral device registers. There is a variety of “arithmetic applications” of these instructions as well.

**Table 3-20** Logical Group: Boolean Operations.

Description	Mnemonic	Operation	CC	Examples	Mode	$\sim$
AND	AND <i>rb</i> <i>addr</i> <i>rb</i> = A, B  <i>addr</i> = #   []	$(rb) \leftarrow (rb) \cap (addr)$	N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow 0$	ANDA #1 ANDA \$FF ANDB 900h ANDA 1, X ANDA B, Y ANDB 2, Y+ ANDA [0, Y] ANDA [D, X]	#  3  3  3  3  3 [] 6 [] 6	1 3 3 3 3 3 6 6
ANDCC	ANDCC <i>addr</i> <i>addr</i> = #	$(CC) \leftarrow (CC) \cap data$	all*	ANDCC #\$FE	#	1
OR	OR <i>rb</i> <i>addr</i> <i>rb</i> = A, B  <i>addr</i> = #   []	$(rb) \leftarrow (rb) \cup (addr)$	N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow 0$	ORA #1 ORA \$FF ORB 900h ORA 1, X ORA B, Y ORB 2, Y+ ORA [0, Y] ORA [D, X]	#  3  3  3  3  3 [] 6 [] 6	1 3 3 3 3 3 6 6
ORCC	ORCC <i>addr</i> <i>addr</i> = #	$(CC) \leftarrow (CC) \cup data$	all*	ORCC #1	#	1
XOR	EOR <i>rb</i> <i>addr</i> <i>rb</i> = A, B  <i>addr</i> = #   []	$(rb) \leftarrow (rb) \oplus (addr)$	N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow 0$	EORA #1 EORA \$FF EORB 900h EORA 1, X EORA B, Y EORB 2, Y+ EORA [0, Y] EORA [D, X]	#  3  3  3  3  3 [] 6 [] 6	1 3 3 3 3 3 6 6

\* Any condition code bit can potentially be cleared by an ANDCC instruction or set by an ORCC instruction, with the exception of the “X” bit (non-maskable interrupt mask bit), which cannot be set by a software instruction – more on this in Chapter 5.

Perhaps the first subgroup of logical instructions that comes to mind is Boolean. The 68HC12 implements the most useful and basic of these, listed in Table 3-20: [AND](#), [OR](#), and [XOR](#) ([EOR](#)). These instructions perform a bit-wise Boolean operation on the named byte-register and operand in memory; the result is stored in the named register. The overflow (V) flag is cleared while the negative (N) and zero (Z) flags are affected based on the result obtained.

[Boolean operations](#)  
[AND](#)  
[OR](#)  
[EOR](#)

There are two special variants contained in this subgroup: ANDCC and ORCC. These instructions provide a generic way to clear or set any of the condition code bits (well, almost any – the “X” bit, the non-maskable interrupt mask, can be *cleared* but *cannot be set* by a software instruction – the “machine control” portion of the CCR will be discussed in Chapter 5). Note that the only addressing mode available is *immediate*.

[ANDCC](#)  
[ORCC](#)

ANDCC and ORCC can be used in place of the “vintage” (legacy) set/clear instructions dedicated to specific condition code register bits. These instructions, listed in Table 3-21, provide a “direct” means for setting or clearing the carry flag (C), the overflow flag (V), or the system interrupt mask bit (I).

[vintage CCR set/clear instructions](#)

**Table 3-21** Logical Group: Condition Code Bit Set/Clear.

Description	Mnemonic	Operation	CC	Examples	Mode	$\sim$
Clear C bit of CCR	CLC	$(C) \leftarrow 0$	$(C) \leftarrow 0$	CLC	⊕	1
Set C bit of CCR	SEC	$(C) \leftarrow 1$	$(C) \leftarrow 1$	SEC	⊕	1
Clear V bit of CCR	CLV	$(V) \leftarrow 0$	$(V) \leftarrow 0$	CLV	⊕	1
Set V bit of CCR	SEV	$(V) \leftarrow 1$	$(V) \leftarrow 1$	SEV	⊕	1
Clear I bit of CCR	CLI	$(I) \leftarrow 0$	$(I) \leftarrow 0$	CLI	⊕	1
Set I bit of CCR	SEI	$(I) \leftarrow 1$	$(I) \leftarrow 1$	SEI	⊕	1

The “complement and clear” sub-group, documented in Table 3-22, provides a means for clearing and setting byte-registers or memory locations (CLRA followed by COMA will set (A) to \$FF). The astute digijock(ette) will realize that the COM instruction was also included as a member of the arithmetic group. Like Florida in the 2000 election, this one was “too close to call”. (Conversely, a case could be made for calling the “CLR” instruction an arithmetic instruction – a “hand recount” might be necessary to sort this one out, or maybe just a high-priced lawyer.)

[clear](#)  
[CLR](#)

[complement](#)  
[COM](#)

**Table 3-22** Logical Group: Byte Clear and Complement.

Description	Mnemonic	Operation	CC	Examples	Mode	$\sim$
Clear	CLR <i>rb</i> <i>rb</i> = A, B	$(rb) \leftarrow \$00$	N $\leftarrow 0$ Z $\leftarrow 1$ V $\leftarrow 0$ C $\leftarrow 0$	CLRA	⊕	1
	CLR <i>addr</i> <i>addr</i> =   [D, Y]	$(addr) \leftarrow \$00$	N $\leftarrow 0$ Z $\leftarrow 1$ V $\leftarrow 0$ C $\leftarrow 0$	CLR \$900		3
			CLR 1, X		2	
			CLR B, X		2	
			CLR [D, Y]	[D, Y]	5	
Complement	COM <i>rb</i> <i>rb</i> = A, B	$(rb) \leftarrow \$FF - (rb)$	N $\leftarrow \uparrow$ Z $\leftarrow \uparrow$ V $\leftarrow 0$ C $\leftarrow 1$	COMA	⊕	1
	COM <i>addr</i> <i>addr</i> =   [D, Y]	$(addr) \leftarrow \$FF - (addr)$	N $\leftarrow \uparrow$ Z $\leftarrow \uparrow$ V $\leftarrow 0$ C $\leftarrow 1$	COM \$900		4
				COM 1, X		3
				COM B, X		3
				COM [D, Y]	[D, Y]	6

Even more useful than the byte clears and sets are the *bit* clear and set instructions ([BCLR](#) and [BSET](#)), listed in Table 3-23. These instructions provide a convenient, powerful means for setting or clearing individual bits or groups of bits within a byte. The bit positions to be set or cleared are indicated by a *mask pattern* (that follows the address field): bits of the mask pattern that are “1” indicate the bits to be cleared or set by BSET and BCLR, respectively. For example, execution of a “BCLR *addr,\$01*” instruction *clears* the bit position corresponding to the mask pattern 00000001b, i.e., the least significant position (bit position 0). Execution of a “BSET *addr,\$F0*” instruction *sets* the bit positions corresponding to the mask pattern 11110000b, i.e., the most significant four bits (bit positions 7 through 4).

*bit clear*  
[BCLR](#)*bit set*  
[BSET](#)*mask pattern***Table 3-23** Logical Group: Bit Clear and Set.

Description	Mnemonic	Operation	CC	Examples	Mode	$\sim$
Bit clear	BCLR <i>addr,mask</i> <i>addr</i> =	$(addr) \leftarrow (addr) \cap \text{mask}_8'$	N $\leftarrow \uparrow$ Z $\leftarrow \uparrow$ V $\leftarrow 0$	BCLR \$50, \$FE		4
				BCLR \$900, \$FE		4
				BCLR 1, X, \$01		4
				BCLR 2, X+, \$F0		4
				BCLR 1000t, Y, \$02		6
Bit set	BSET <i>addr,mask</i> <i>addr</i> =	$(addr) \leftarrow (addr) \cup \text{mask}_8$	N $\leftarrow \uparrow$ Z $\leftarrow \uparrow$ V $\leftarrow 0$	BSET \$50, \$FE		4
				BSET \$900, \$FE		4
				BSET 1, X, \$01		4
				BSET 2, X+, \$F0		4
				BSET 1000t, Y, \$02		6

Another important tool for bit-oriented operations is the “bit test” (**BIT**) instruction, documented in Table 3-24. This instruction is analogous to the TST instruction, except here the bit test is performed by ANDing the named byte-register with the contents of a memory location and setting the condition code bits accordingly. Like the TST instruction, the result of the AND operation is not stored; only the condition codes are affected.

**bit test**  
**BIT**

**Table 3-24** Logical Group: Bit Test.

Description	Mnemonic	Operation	CC	Examples	Mode	$\sim$
Bit test	BIT <i>rb</i> <i>addr</i> <i>rb</i> = A, B <i>addr</i> = #   []	set CCR based on ( <i>rb</i> ) $\cap$ ( <i>addr</i> )	N $\leftarrow$ $\updownarrow$ Z $\leftarrow$ $\updownarrow$ V $\leftarrow$ 0	BITA #1 BITA \$FF BITB 900h BITA 1,X BITA B,Y BITB 2,Y+ BITA [0,Y] BITA [D,X]	# 1 # 3 # 3 # 3 # 3 # 3 [ ] 6 [ ] 6	

The final subgroup of logical instructions is the “shift and rotate” group. The first question that comes to mind is: “What’s the difference between a shift and a rotate?” Shifts are generally regarded as *arithmetic* operations: a (sign-preserving) multiply-by-two (shift left) or divide-by-two (shift right). Rotates generally involve a “wrap-around” effect, i.e., the bit “rotated out” at one end gets “rotated in” at the other end. Therefore, if an N-bit register is rotated N times right or N times left, it will return to its “original state”. This is in contrast with their “shifty” cousins, which are classically “end-off” shifts – i.e., bits shifted out wind up in the proverbial “bit bucket”. An N-bit register shifted left arithmetically N (or more) times will be filled with zeroes, while that same register shifted right arithmetically N (or more) times will be filled with the sign of the original operand (i.e., *all* zeroes if the original value was *positive*, or *all* ones if the original value was *negative*).

*sign-preserving arithmetic shift*

*end-off shift*

*bit bucket*

Starting with the rotates, the first thing to note is that these instructions operate on a 9-bit value consisting of the C condition code bit concatenated with the named register or memory location. (Including “C” in the instruction mnemonics – what Intel did for similar instructions in their microprocessors – would have perhaps made this fact a bit easier to remember!) The “proper names” for these instructions, documented in Table 3-25, are therefore “rotate left through carry” (**ROL**) and “rotate right through carry” (**ROR**). Note that since the C-bit is construed as an integral part of the value being rotated, it is usually important that this flag be placed in a *known initial state* prior to a rotate; otherwise, “strange bits” may appear in the rotated result.

*9-bit rotate through C*

*ROL (left)*  
*ROR (right)*

**Table 3-25** Logical Group: Shift and Rotate.

Description	Mnemonic	Operation	CC	Examples	Mode	$\sim$
Rotate left through carry	ROLrb rb = A, B		N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$	ROLA		1
	ROL addr addr =   []		N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$	ROL \$900 ROL 1,X ROL B,X ROL [D,Y]	   	4 3 3 6
Rotate right through carry	RORrb rb = A, B		N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$	RORA		1
	ROR addr addr =   []		N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$	ROR \$900 ROR 1,X ROR B,X ROR [D,Y]	   	4 3 3 6
Arithmetic shift left*	ASLrb rb = A, B		N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$	ASLA		1
	ASLrw rw = D		N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$	ASLD		1
	ASL addr addr =   []		N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$	ASL \$900 ASL 1,X ASL B,X ASL [D,Y]	   	4 3 3 6
Arithmetic shift right	ASRrb rb = A, B		N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$	ASRA		1
	ASR addr addr =   []		N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$	ASR \$900 ASR 1,X ASR B,X ASR [D,Y]	   	4 3 3 6
Logical shift left*	LSLrb rb = A, B		N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$	LSLA		1
	LSLrw rw = D		N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$	LSLD		1
	LSL addr addr =   []		N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$	LSL \$900 LSL 1,X LSL B,X LSL [D,Y]	   	4 3 3 6
Logical shift right	LSRrb rb = A, B		N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$	LSRA		1
	LSRrw rw = D		N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$	LSRD		1
	LSR addr addr =   []		N $\leftarrow \downarrow$ Z $\leftarrow \downarrow$ V $\leftarrow \downarrow$ C $\leftarrow \downarrow$	LSR \$900 LSR 1,X LSR B,X LSR [D,Y]	   	4 3 3 6

\*ASL and LSL instruction mnemonics generate identical machine code.

For a *rotate left through carry* (**ROL**), the entire contents of the targeted register or memory location is translated left one position; the “vacated” low-order bit is loaded with the value that was in the C bit just prior to the ROL, and the C-bit is loaded with the value that rotated out of the high order bit. If a series of *nine* ROL instructions is executed, the original state of the targeted register or memory location as well as the C bit will be restored.

**9-bit rotate left through C**

**ROL**

A rotate right through carry (**ROR**) works the same as ROL, except the contents of the targeted register or memory location is translated right one position. Here, the vacated high-order bit is loaded with the value that was in the C bit just prior to the ROR, and the C bit is loaded with the value that rotated out of the low order bit. As was the case for ROL, a series of nine ROR instructions yields the original state. Note that while ROL and ROR affect all of the flags (N, Z, V, C), the only one of “social significance” is the C bit.

**9-bit rotate right through C**

**ROR**

At this point, one might properly ask: “Why was this strange ‘9-bit rotate through the carry bit’ implemented instead of a more intuitive 8-bit rotate within the targeted register or memory location?” It turns out that a classic (and useful) application of the “rotate through carry” mechanism is to “pick off bits” and subsequently make decisions (through execution of conditional transfer-of-control instructions) based on the state of individual bits as they are encountered.

Continuing with the shifts (also listed in Table 3-25), we find that an *arithmetic shift left* (**ASL**) translates the entire contents of the targeted register or memory location one position left. Here, the “vacated” low order bit is filled with a zero, and the bit that shifts out of the most significant position is preserved in the C flag (for the purpose of determining whether or not “overflow” occurred). The original contents – whether originally positive or negative – is thus multiplied by two, within the precision afforded by the targeted register or memory location. For example, if the original contents of the A register is \$01 ( $1_{10}$ ), the result will be \$02 ( $2_{10}$ ) after one ASLA instruction is executed, \$04 ( $4_{10}$ ) after a second ASLA instruction is executed, up to a (positive) maximum of \$40 ( $64_{10}$ ) after six consecutive ASLA instructions are executed. Here, note that execution of one additional ASLA instruction would produce the value \$80, or  $-128_{10}$ , thus changing the sign and causing “overflow” to occur. Conversely, if the original contents of the A register is \$FF ( $-1_{10}$ ), the result will be \$FE ( $-2_{10}$ ) after one ASLA instruction is executed, \$FC ( $-4_{10}$ ) after two ASLA instructions are executed, up to a maximum (in magnitude) of \$80 ( $-128_{10}$ ) after seven consecutive ASLA instructions are executed. Note that the overflow (V) flag is set if there is a “disagreement” between the sign bit (reflected by the N flag) and the carry (C) flag, i.e.,  $V = N \oplus C$ ,

**arithmetic shift ASL (left)**

which would occur here if one more ASLA instruction were executed (producing a result of \$00 with the C bit set).

An *arithmetic shift right* ([ASR](#)) translates the contents of the targeted register or memory location one position right. Here, the vacated high-order bit is filled with a *copy* of its original value, i.e., the sign bit is *replicated*. The bit that shifts out of the least significant position is preserved in the C bit, to facilitate rounding the result – which is effectively the original contents divided by two. For example, if the original contents of the A register is \$7F ( $+127_{10}$ ), the result will be \$3F ( $+63_{10}$ ) after one ASRA instruction is executed, \$1F ( $31_{10}$ ) after two ASRA instructions are executed, down to \$01 ( $+1_{10}$ ) after six ASRA instructions are executed, and \$00 after seven (or more) ASRA instructions are executed. (Note that if the result after the first ASRA, \$3F, had been rounded to \$40, the contents of the A register would not reach \$00 until a total of *eight* or more ASRA instructions had been executed.)

[arithmetic shift  
ASR \(right\)](#)

Unlike ASL, though, the overflow (V) flag has no meaning for ASR since the sign of the result cannot “flip” as a consequence of shifting “one too many” times. For example, if the original contents of the A register is \$80 ( $-128_{10}$ ), the result will be \$C0 ( $-64_{10}$ ) after one ASRA instruction is executed, \$E0 ( $-32_{10}$ ) after two ASRA instructions are executed, down to \$FE ( $-2_{10}$ ) after six ASRA instructions are executed, and \$FF ( $-1_{10}$ ) after seven (or more) ASRA instructions are executed. Note that, after the eighth ASRA instruction, the C bit is set, enabling the result to be rounded to \$00. In either case (i.e., rounded or not), execution of additional ASRA instructions will not change the contents of the A register (i.e., it will “freeze” at either \$FF or \$00).

In addition to arithmetic shifts, the 68HC12 provides “logical shifts” – defined as “end-off” shifts with *zero fill*. Thus, an *arithmetic shift left* and *logical shift left* ([LSL](#)) are identical – in fact, the ASL and LSL assembly mnemonics generate the same object code (machine instruction). Further, an *arithmetic shift right* produces the same result as a *logical shift right* ([LSR](#)) for positive operands. Only for the case of *negative operands* will an arithmetic shift right produce a different result than a logical shift right. A logical shift, then, translates the contents of the targeted register or memory location one position left or right; the vacated position is filled with a zero and the position that “shifts out” is preserved in the C bit. Therefore, if an N-bit register is logically shifted left or right N (or more) times, the resulting value will be all zeroes. An interesting (and useful) variant provided for the logical shifts (and, by association, the arithmetic shift left) is a 16-bit shift of the double-byte (D) accumulator: [LSD](#) and [LSRD](#).

[logical \(zero fill\) shift](#)

[LSL \(left\)  
LSR \(right\)](#)

[16-bit logical shift](#)  
[LSD \(left\)  
LSRD \(right\)](#)

In summary, the logical group includes Boolean, complement/clear, bit set/clear, bit test, and shift/rotate instructions. Some of the instructions included in this group – by virtue of their “bit-oriented” nature – are ostensibly arithmetic, however.

### 3.8.4 Transfer-of-Control Group Instructions

As its name implies, this group includes all the 68HC12 instructions that facilitate transfer of control from one location of a program to another. The major variants available include an unconditional jump instruction, conditional and unconditional branch instructions, compound test and branch instructions, and subroutine linkage instructions.

In Chapter 2, we defined the difference between a “jump” and a “branch” as follows. If the address field of the instruction contains the (absolute) address in memory at which execution should continue, it is usually referred to as a “jump” instruction. If the address field instead represents the (signed) “distance” the next instruction to execute is from the transfer-of-control instruction, it is referred to as a “branch”. (There is not universal agreement on this nomenclature, however – Intel typically uses the *opposite* definitions for jump and branch.) Jumps (or branches) that “always happen” are called *unconditional*; those that happen only if a certain combination of condition codes exists are called *conditional*.

Beginning with the unconditional jump (**JMP**) instruction listed in Table 3-26, we find that the 68HC12, through the variety of addressing modes supported, provides a very powerful transfer-of-control mechanism that includes use of indexed modes (for “computing” the address of the next instruction) and indirection (for “looking up” the address of the next instruction). We will make extensive use of so-called “jump tables” in the programming examples that follow in Chapter 4.

*unconditional jump  
JMP*

**Table 3-26** Transfer-of-Control Group: Unconditional Jump.

Description	Mnemonic	Operation	CC	Examples	Mode	~
Jump	JMP <i>addr</i> <i>addr</i> =	(PC) ← <i>addr</i>	–	JMP \$900 JMP 0,X JMP 100t,Y JMP 1000t,S JMP [D,Y] JMP [1000t,S]	3 3 3 4 6 6	

Branch instructions – including the unconditional branch (**BRA**) listed in Table 3-27 as well as the plethora of conditional branches that follow – all have two forms: “short”, for which the signed offset ranges from  $-128_{10}$  to  $+127_{10}$ ; and “long”, for which the signed offset ranges from  $-32,768_{10}$  to  $+32,767_{10}$ . A prefix of “L” in the assembly mnemonic is used to specify the “long version” of a particular branch. In general, the “short” branches

*short unconditional  
branch  
BRA*

*long unconditional  
branch  
LBRA*

(unconditional and conditional) are two bytes long (one opcode byte plus one offset byte); the “long” branches are all four bytes in length (two opcode bytes plus two offset bytes). Because the destination of the branch is determined “in relation” to the current location (i.e., the location pointed to by the PC), the addressing mode is called *relative* (for which we will use the icon ).

*relative addressing mode icon*  


**Table 3-27 Transfer-of-Control Group: Unconditional Branch.**

Description	Mnemonic	Operation	CC	Examples	Mode	~
(Short) Branch	BRA <i>rel8</i>	$(PC) \leftarrow (PC) + rel8^*$	—	BRA <i>label1</i>		2
Long Branch	LBRA <i>rel16</i>	$(PC) \leftarrow (PC) + rel16^*$	—	LBRA <i>label1</i>		4

\*Calculation of the two’s complement relative offset must take into account the byte-length of the branch instruction. The “short” branch (BRA) instruction occupies two bytes while the “long” branch (LBRA) instruction occupies four bytes. Because the program counter is automatically incremented as a by-product of the instruction fetch, the offset calculation must compensate for this.

A “tricky” (and perhaps confusing) aspect of calculating the signed offset for a branch instruction is compensating for the PC increment that occurs as a byproduct of the instruction fetch. Just as was the case for our simple computer in Chapter 2, the PC points to the next instruction once the current instruction has been fetched (and is about to be executed). For the “short” branches, this means that the PC has already been incremented by *two* before the offset is added; for the “long” branches, the value is *four*. To implement the equivalent of an “infinite loop” with a BRA instruction (i.e., a “branch to itself”), then, an offset of *-2* (or \$FE) must be used. For a LBRA instruction, an offset of *-4* (or \$FFFC) must be used to obtain the same result.

```

0800          1      org     800h
              2
0800 [01] 20FE 3 short   bra    short
              4
0802 [04] 1820FFFC 5 long    lbra   long
              6
0806          7      end
              8
              9

Symbol Table

LONG          0802
SHORT         0800

```

**Figure 3-22 Comparison of short and long branch offsets.**

Fortunately, the offset calculation usually does not need to be done “by hand” – assembler programs use symbols for labels and calculate the offset field of branch instructions automatically. So even though a “hard” number (like \$FE or \$FFFC) *could* be placed in the address field of a branch instruction, we will virtually never do this in practice. Instead, we will use the symbol *label* to denote the destination of the branch, as shown in Figure 3-22, based on the tacit assumption that an assembler program can calculate the relative offset much more accurately than we could ever do “by hand”. This will certainly come as good news for the poll workers in Palm Beach County!

*branch offset  
calculation*

*label*

### The Long and Short of It (Locality of Reference)

A question that is sure to come to mind when studying the 68HC12 instruction set is: “Why are there both ‘short’ and ‘long’ branches?” Back in the early 1970’s when the “grandfather” of the MC68xx series was conceived, just “short” (unconditional and conditional) branches plus a “long” (unconditional) jump were included in the instruction set. Short branches work well for a large percentage of applications due to the principle of locality of reference. According to this principle, there is a high probability that the next instruction will be fetched from a location relatively close to the current instruction. For typical application code, the percentage of time this is true is greater than 95%. But on occasions when a “short” branch isn’t quite “long enough”, there is not a “pretty” solution. A complete set of long (unconditional and conditional) branches was therefore one of the key features added when the MC6809 was introduced in the late 1970’s.

**Table 3-28 Transfer-of-Control Group: Subroutine Linkage.**

Description	Mnemonic	Operation	CC	Examples	Mode	~
Jump to Subroutine	JSR <i>addr</i> <i>addr</i> =	$(SP) \leftarrow (SP) - 2$ $((SP)) \leftarrow (PC_h)$ $((SP)+1) \leftarrow (PC_l)$ $(PC) \leftarrow addr$	-	JSR \$20 JSR \$900 JSR 0,X JSR 100t,Y JSR 1000t,S JSR [D,Y] JSR [1000t,S]	      	4 4 4 4 5 7 7
Branch to Subroutine	BSR <i>rel8</i>	$(SP) \leftarrow (SP) - 2$ $((SP)) \leftarrow (PC_h)$ $((SP)+1) \leftarrow (PC_l)$ $(PC) \leftarrow (PC) + rel8^*$	-	BSR label		4
Return from Subroutine	RTS	$(PC_h) \leftarrow ((SP))$ $(PC_l) \leftarrow ((SP)+1)$ $(SP) \leftarrow (SP) + 2$	-	RTS		4

\*Calculation of the two’s complement relative offset must take into account the byte-length of the BSR instruction, which is two bytes.

The subroutine linkage instructions provided by the 68HC12 are listed in Table 3-28. In the spirit of the unconditional jump and branch described above, subroutines can be “called” using either a jump ([JSR](#)) or a branch ([BSR](#)). Both instructions push the *return address* – effectively the current value in the PC after the JSR or BSR has been fetched – onto the stack in a similar fashion. One simply follows the “push PC” operation with a *jump* to the subroutine address (JSR), while the other performs a *branch* using an 8-bit signed offset (BSR). Like the JMP instruction, the JSR supports a replete set of addressing modes. Note, however, that there is not a “long” version of BSR (and other than “legacy compatibility”, there is not compelling reason for even having the BSR itself). The return from subroutine ([RTS](#)) instruction simply “pops” (uh, *pulls*) the return address off the stack and loads it into the PC, enabling program execution to continue at the location following the JSR or BSR that previously “called” the subroutine.

[subroutine linkage](#)

[call: JSR/BSR](#)

[return: RTS](#)

### Puddle Jumping

*Imagine a world without long branches. Greater than 95% of the time, not a problem. But when a single byte signed offset just won't reach, there's no great solution. Similar to a frog attempting to cross a stream via a collection of strategically-placed lilypads (or the author attempting to fly from his adopted hometown of Lafayette, Indiana, to virtually anywhere else in the civilized world), the only way to get from point A to point B is by “puddle jumping”. For the 6800 (and, unfortunately, also the more recent 68HC11), this is precisely the kind of technique that must be employed. This is “bad enough” when attempting to program in assembly language, but even more of a nightmare for a compiler!*

We are now ready to consider the rather overwhelming collection of conditional branch instructions implemented on the 68HC12. The first set of instructions, listed in Table 3-29, are appropriately called “simple” conditionals since each involves the testing of a single flag (C, Z, N, V). The “carry condition” ([BCC/BCS](#)) is based on the state of the C flag: “clear” (BCC) means that the branch is taken if the carry flag is zero, and “set” (BCS) means that the branch is taken if the carry flag is one. The “test for equality” ([BNE/BEQ](#)) is based on taking the difference of two operands (using a previous CMP or TST instruction) and obtaining a result of zero, thus setting the Z flag – a condition we will use quite often in the code writing exercises ahead in Chapter 4. The “plus/minus” test ([BPL/BMI](#)) is based on the state of the N flag, while the “overflow” test ([BVC/BVS](#)) is based on the state of the V flag. Referring to the cycle (~) column, note that more cycles are required to execute a branch that is “taken” compared with a branch that is “not taken”. The reason for this disparity is the need to “flush” and “refill” the processor’s instruction queue each time a transfer-of-control takes place.

[simple conditionals](#)

(clear/set)

[C: BCC/BCS](#)

[Z: BNE/BEQ](#)

[N: BPL/BMI](#)

[V: BVC/BVS](#)

[instruction queue flush and refill](#)

**Table 3-29** Transfer-of-Control Group: Simple Conditional Branches.

Description	Mnemonic	Operation*	CC	Examples	Mode	~**
Branch if carry clear $C = 0$	BCC rel8	$(PC) \leftarrow (PC) + rel8$	—	BCC label		3/1
	LBCC rel16	$(PC) \leftarrow (PC) + rel16$	—	LBCC label		4/3
Branch if carry set $C = 1$	BCS rel8	$(PC) \leftarrow (PC) + rel8$	—	BCS label		3/1
	LBCS rel16	$(PC) \leftarrow (PC) + rel16$	—	LBCS label		4/3
Branch if not equal $Z = 0$	BNE rel8	$(PC) \leftarrow (PC) + rel8$	—	BNE label		3/1
	LBNE rel16	$(PC) \leftarrow (PC) + rel16$	—	LBNE label		4/3
Branch if equal $Z = 1$	BEQ rel8	$(PC) \leftarrow (PC) + rel8$	—	BEQ label		3/1
	LBEQ rel16	$(PC) \leftarrow (PC) + rel16$	—	LBEQ label		4/3
Branch if positive $N = 0$	BPL rel8	$(PC) \leftarrow (PC) + rel8$	—	BPL label		3/1
	LBPL rel16	$(PC) \leftarrow (PC) + rel16$	—	LBPL label		4/3
Branch if negative $N = 1$	BMI rel8	$(PC) \leftarrow (PC) + rel8$	—	BMI label		3/1
	LBMI rel16	$(PC) \leftarrow (PC) + rel16$	—	LBMI label		4/3
Branch if overflow clear $V = 0$	BVC rel8	$(PC) \leftarrow (PC) + rel8$	—	BVC label		3/1
	LBVC rel16	$(PC) \leftarrow (PC) + rel16$	—	LBVC label		4/3
Branch if overflow set $V = 1$	BVS rel8	$(PC) \leftarrow (PC) + rel8$	—	BVS label		3/1
	LBVS rel16	$(PC) \leftarrow (PC) + rel16$	—	LBVS label		4/3
Branch never (No-op)	BRN rel8	—	—	BRN label		1
	LBRN rel16	—	—	LBRN label		3

\*Operation performed if branch is *taken*. If branch is *not taken*, the instruction effectively becomes a “no operation” (NOP). Calculation of the two’s complement relative offset must take into account the byte-length of the branch instruction itself (2 for short, 4 for long).

\*\*The first number indicates the number of cycles consumed if the branch is *taken*; the second number indicates the number of cycles consumed if the branch is *not taken*.

**Table 3-30** Transfer-of-Control Group: Signed Conditional Branches.

Description	Mnemonic	Operation*	CC	Examples	Mode	~**
Branch if greater than $Z + (N \oplus V) = 0$	BGT rel8	$(PC) \leftarrow (PC) + rel8$	–	BGT label		3/1
	LBGT rel16	$(PC) \leftarrow (PC) + rel16$	–	LBGT label		4/3
Branch if less than or equal to $Z + (N \oplus V) = 1$	BLE rel8	$(PC) \leftarrow (PC) + rel8$	–	BLT label		3/1
	LBLT rel16	$(PC) \leftarrow (PC) + rel16$	–	LBLT label		4/3
Branch if greater than or equal $N \oplus V = 0$	BGE rel8	$(PC) \leftarrow (PC) + rel8$	–	BGE label		3/1
	LBGE rel16	$(PC) \leftarrow (PC) + rel16$	–	LBGE label		4/3
Branch if less than $N \oplus V = 1$	BLT rel8	$(PC) \leftarrow (PC) + rel8$	–	BLT label		3/1
	LBLT rel16	$(PC) \leftarrow (PC) + rel16$	–	LBLT label		4/3

**Table 3-31** Transfer-of-Control Group: Unsigned Conditional Branches.

Description	Mnemonic	Operation*	CC	Examples	Mode	~**
Branch if higher than $C + Z = 0$	BHI rel8	$(PC) \leftarrow (PC) + rel8$	–	BHI label		3/1
	LBHI rel16	$(PC) \leftarrow (PC) + rel16$	–	LBHI label		4/3
Branch if lower than or same $C + Z = 1$	BLS rel8	$(PC) \leftarrow (PC) + rel8$	–	BLS label		3/1
	LBLS rel16	$(PC) \leftarrow (PC) + rel16$	–	LBLS label		4/3
Branch if higher than or same $C = 0$	BHS rel8	$(PC) \leftarrow (PC) + rel8$	–	BHS label		3/1
	LBHS rel16	$(PC) \leftarrow (PC) + rel16$	–	LBHS label		4/3
Branch if lower than $C = 1$	BLO rel8	$(PC) \leftarrow (PC) + rel8$	–	BLO label		3/1
	LBLO rel16	$(PC) \leftarrow (PC) + rel16$	–	LBLO label		4/3

\*Operation performed if branch is *taken*. If branch is *not taken*, the instruction effectively becomes a “no operation” (NOP). Calculation of the two’s complement relative offset must take into account the byte-length of the branch instruction itself (2 for short, 4 for long).

\*\*The first number indicates the number of cycles consumed if the branch is *taken*; the second number indicates the number of cycles consumed if the branch is *not taken*.

Compound conditionals – so-called because they typically involve more than one flag – are comprised of two subsets: one that construes the operands as *signed* (listed in Table 3-30), and the other that construes them as *unsigned* (listed in Table 3-31). Both the signed and unsigned conditional branches must be preceded by either a CMP or SUB instruction. Recall that these instructions set or clear the flags (C, Z, N, V) based on the subtraction of an operand (specified by the effective

*compound conditionals*

*signed - unsigned*  
*BGT - BHI*  
*BGE - BHS*  
*BLT - BLO*  
*BLE - BLS*

address) from the named register, i.e.,  $(\text{register}) - (\text{address})$ . For example, the sequence “CMPA #5” followed by “BGT *label*” would cause a branch to the instruction at address *label* if  $(A) \geq 5_{10}$ . Stated another way, if the calculation  $(A) - 5_{10}$  yields a result *greater than* zero, the branch to address *label* will be “taken” by the BGT instruction.

Comparing identical bit patterns, however, can cause a “greater than” (the *signed BGT* or *unsigned BHI*) conditional branch to be taken or not taken, depending on the interpretation of the bit patterns as signed or unsigned. Consider the case of  $(A) = \$01$  with a “CMPA \$FF” instruction performed. Note that \$FF, when interpreted as *signed*, is the two’s complement representation for  $-1$ ; when interpreted as *unsigned*, however, \$FF is the representation for  $255_{10}$ . Because  $(A)$  is greater than  $-1$ , a subsequent “BGT *label*” instruction would cause the branch to address *label* to be taken. But because  $(A)$  is *not* greater than  $255_{10}$ , a subsequent “BHI *label*” instruction would *not* cause a branch to address *label*.

*BGT  
BHI*

For the compound conditional branches, it’s a bit challenging to remember the variety of signed and unsigned instruction mnemonics as well as the differences in how they work. The “naming convention” adopted by Motorola is to use “greater/less than” to denote the signed conditionals, and “higher/lower than” to denote the unsigned conditionals. An interesting aspect of how the conditionals are evaluated centers around the Boolean expressions used (see Tables 3-32 and 3-33). This is a subject that the author confesses to “glossing over” for many years, when temporarily embarrassed by questions such as: “Why is  $Z + (N \oplus V) = 0$  used as the Boolean expression to determine the BGT conditional?”

*conditional branch  
naming convention*

The best way to understand where these Boolean expressions “come from” is to derive them based on the “2-bit” case (i.e., the simplest case that enumerates all the possibilities of both signed and unsigned comparisons). The derivations for the signed and unsigned cases are given in Tables 3-32 and 3-33, respectively. The 2-bit operands loaded in the named register are designated  $R_1R_0$ , and the 2-bit operands residing at the effective address in memory are designated  $M_1M_0$ . The flag settings ( $C$ ,  $Z$ ,  $N$ ,  $V$ ) are based on performing the operation  $(R) - (M)$ . Here’s a critical point: the SUB or CMP instruction that performs  $(R) - (M)$  could care less if the operands being compared are construed as signed or unsigned. In fact, note that Tables 3-32 and 3-33 are basically identical except for interpretation of the bit patterns and resulting comparisons.

**Table 3-32** Derivation of Signed Comparisons.

R <sub>1</sub>	R <sub>0</sub>	(R)	M <sub>1</sub>	M <sub>0</sub>	(M)	?	C	Z	N	V
0	0	0	0	0	0	(R) = (M)	0	1	0	0
0	0	0	0	1	+1	(R) < (M)	1	0	1	0
0	0	0	1	0	-2	(R) > (M)	1	0	1	1
0	0	0	1	1	-1	(R) > (M)	1	0	0	0
0	1	+1	0	0	0	(R) > (M)	0	0	0	0
0	1	+1	0	1	+1	(R) = (M)	0	1	0	0
0	1	+1	1	0	-2	(R) > (M)	1	0	1	1
0	1	+1	1	1	-1	(R) > (M)	1	0	1	1
1	0	-2	0	0	0	(R) < (M)	0	0	1	0
1	0	-2	0	1	+1	(R) < (M)	0	0	0	1
1	0	-2	1	0	-2	(R) = (M)	0	1	0	0
1	0	-2	1	1	-1	(R) < (M)	1	0	1	0
1	1	-1	0	0	0	(R) < (M)	0	0	1	0
1	1	-1	0	1	+1	(R) < (M)	0	0	1	0
1	1	-1	1	0	-2	(R) > (M)	0	0	0	0
1	1	-1	1	1	-1	(R) = (M)	0	1	0	0

$$Z = 1 \\ (R) = (M)$$

$$Z + (N \oplus V) = 0 \\ (R) > (M)$$

$$N \oplus V = 1 \\ (R) < (M)$$

**Table 3-33** Derivation of Unsigned Comparisons.

R <sub>1</sub>	R <sub>0</sub>	(R)	M <sub>1</sub>	M <sub>0</sub>	(M)	?	C	Z	N	V
0	0	0	0	0	0	(R) = (M)	0	1	0	0
0	0	0	0	1	+1	(R) < (M)	1	0	1	0
0	0	0	1	0	+2	(R) < (M)	1	0	1	1
0	0	0	1	1	+3	(R) < (M)	1	0	0	0
0	1	+1	0	0	0	(R) > (M)	0	0	0	0
0	1	+1	0	1	+1	(R) = (M)	0	1	0	0
0	1	+1	1	0	+2	(R) < (M)	1	0	1	1
0	1	+1	1	1	+3	(R) < (M)	1	0	1	1
1	0	+2	0	0	0	(R) > (M)	0	0	1	0
1	0	+2	0	1	+1	(R) > (M)	0	0	0	1
1	0	+2	1	0	+2	(R) = (M)	0	1	0	0
1	0	+2	1	1	+3	(R) < (M)	1	0	1	0
1	1	+3	0	0	0	(R) > (M)	0	0	1	0
1	1	+3	0	1	+1	(R) > (M)	0	0	1	0
1	1	+3	1	0	+2	(R) > (M)	0	0	0	0
1	1	+3	1	1	+3	(R) = (M)	0	1	0	0

$$Z = 1 \\ (R) = (M)$$

$$C + Z = 0 \\ (R) > (M)$$

$$C = 1 \\ (R) < (M)$$

To derive the Boolean expression for a given conditional, the function defined by the corresponding shaded area can be mapped and minimized. For example, the BGT conditional corresponds to the dark blue portion of Table 3-32. Realizing that only a subset of the 16 possible combinations of C-Z-N-V can occur in practice (and marking the ones that *can't occur* as "don't cares"), we obtain the K-map depicted in Figure 3-23. Grouping zeroes provides the minimal solution for this function, which turns out to be the expression for the "complement" of the BGT conditional, namely BLE. Here, we find that the "BLE taken condition" can be expressed by the function  $Z + N \cdot V + NV' = Z + (N \oplus V)$ , which is the same as saying the BLE "is taken" when  $Z + (N \oplus V) = 1$ . The "BGT taken condition", then, is just the complement of this, or  $(Z + (N \oplus V))'$ , which is the same as saying that the BGT "is taken" when  $Z + (N \oplus V) = 0$ . Don't feel bad if this isn't "instantly obvious" – it wasn't to the author either!

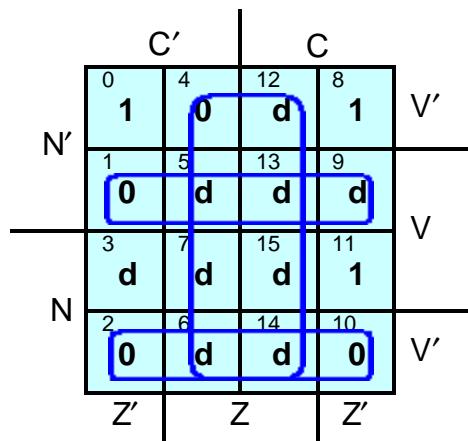


Figure 3-23 Derivation of BGT/BLE functions.

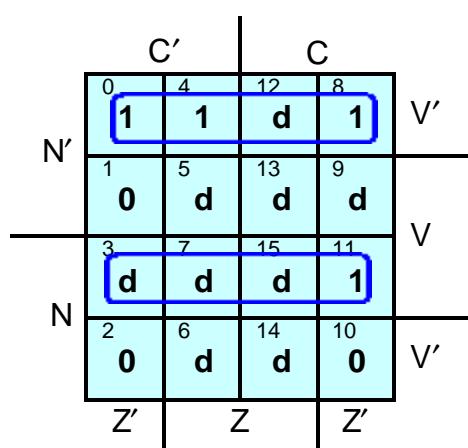


Figure 3-24 Derivation of BGE/BLT functions.

We can do a similar derivation for the **BGE/BLT** “pair” of conditionals, as shown in Figure 3-24. Grouping the ones, we find the “BGE taken condition” to be  $N \cdot Z' + NZ = (N \oplus Z)'$ , which is the same as saying the BGE “is taken” when  $N \oplus Z = 0$ . Conversely, we can say that the BLT “is taken” when the opposite condition is true, i.e.,  $N \oplus Z = 1$ .

**BGE**  
**BLT**

For the **BHI/BLS** pair – the “unsigned cousins” of the BGT/BLE pair – the K-map in Figure 3-25 (derived from Table 3-33) applies. Here, grouping zeroes leads to the minimal function, which is simply  $C + Z$ . Since this corresponds to the “complement” function (BLS), the BLS “is taken” when  $C + Z = 1$ ; conversely, the BHI “is taken” when the function  $C + Z = 0$ .

**BHI**  
**BLS**

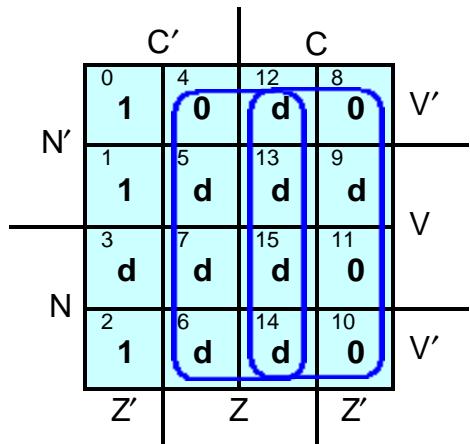


Figure 3-25 Derivation of BHI/BLS functions.

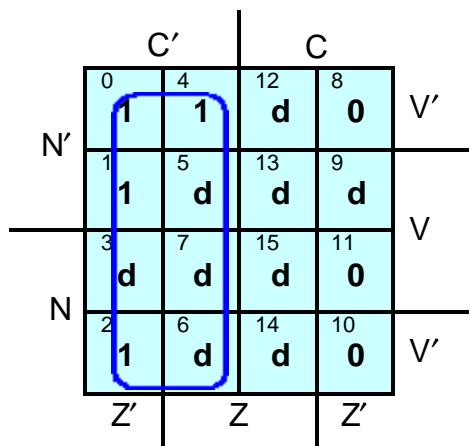


Figure 3-26 Derivation of BHS/BLO functions.

Finally, for the **BHS/BLO** pair – the “unsigned cousins” of the BGE/BLT pair – the K-map in Figure 3-26 applies. Grouping ones yields the function for BHS, which is simply C. Thus, the BHS “is taken” when C = 0, while the BLT “is taken” when C = 1. As such, since BHS and BLO are merely tests of the carry flag, they are synonyms for (and produce the same opcodes as) BCC and BCS, respectively. Fortunately, assembler programs accept the mnemonics BHS/BLO to prevent any confusion associated with trying to remember that BHS is the same as BCC, and that BLO is the same as BCS.

**BHS  
BLO**

The conditional branches covered thus far are primarily “legacy” instructions, carried over from earlier MC68xx family members. A common “feature” of these legacy conditionals is that they must be preceded by a CMP or SUB instruction. At some point, with more silicon at their disposal, microcontroller design engineers realized that the compare and branch operations could be combined into a single instruction. The 68HC12 provides three basic types of so-called “compare and branch” instructions: those that branch based on bit tests (listed in Table 3-34), those that branch based on register tests (listed in Table 3-35), and those that increment/decrement a register and subsequently branch based on a test of that register (listed in Table 3-36). Note that, since all of these instructions are essentially “self-contained”, there is no need for them to affect any of the condition code bits.

*legacy instructions*

**Table 3-34** Transfer-of-Control Group: Bit Test and Branch.

Description	Mnemonic	Operation	CC	Examples	M	-
Branch if bits clear	BRCLR <i>addr,mask8,rel8</i> <i>addr</i> =	IF ( <i>addr</i> ) $\cap$ <i>mask8</i> = 0  THEN (PC) $\leftarrow$ (PC) + <i>rel8</i>	–	BRCLR \$50, <b>01</b> ,label BRCLR \$900, <b>01</b> ,label BRCLR 0,X, <b>\$FF</b> ,label BRCLR 10t,X, <b>01</b> ,label BRCLR 100t,Y, <b>02</b> ,label BRCLR 1000t,S, <b>03</b> ,label	4 5 4 4 6 8	
Branch if bits set	BRSET <i>addr,mask8,rel8</i> <i>addr</i> =	IF ( <i>addr</i> )' $\cap$ <i>mask8</i> = 0  THEN (PC) $\leftarrow$ (PC) + <i>rel8</i>	–	BRSET \$50, <b>01</b> ,label BRSET \$900, <b>01</b> ,label BRSET 0,X, <b>\$FF</b> ,label BRSET 10t,X, <b>01</b> ,label BRSET 100t,Y, <b>02</b> ,label BRSET 1000t,S, <b>03</b> ,label	4 5 4 4 6 8	

The first subset of these instructions, **BRCLR** and **BRSET**, test individual bits (or sets of bits) of a memory location and, if the test is successful, branch to a new location based on an 8-bit signed offset. The bits participating in the test are specified by an 8-bit mask pattern, where a “1” in the mask pattern means that the corresponding bit position in the operand is tested. For the BRCLR (“branch if bits clear”) instruction, the branch is taken if all the bit positions specified by the mask pattern are zeroes. This is accomplished by ANDing the mask pattern with the

**BRCLR  
BRSET**

contents of the memory location; if the result of the bit-wise AND is all zeroes, the branch conditional is true. For the BRSET (“branch if bits set”) instruction, the branch is taken if all the bit positions specified by the mask pattern are *ones*. This is accomplished by ANDing the mask pattern with the *complement* of the memory location contents; if the result of the bit-wise “complement-and-AND” operation yields all zeroes, the branch is taken.

Direct, extended, and indexed addressing modes can be used by BRSET and BRCLR to access the desired location in memory. Instruction lengths vary from four to six bytes, with execution times as high as eight cycles. We will find these instructions extremely useful for performing conditional branches based on the state of various bits in the 68HC12’s peripheral device registers.

The next subset of what we have broadly called “compare and branch” instructions combines the equivalent of a TST instruction with either a BEQ or BNE. These instructions, listed in Table 3-35, are **TBEQ** (“test register and branch if zero”) and **TBNE** (“test register and branch if not zero”). These “compound” instructions are actually a bit more powerful than the “simple” predecessors that inspired them: not only can they use any of the machine’s registers (A, B, D, X, Y, SP), but also the relative branch offset has been extended to 9-bits (effectively doubling the range of the signed offset).

**TBEQ**  
**TBNE**

**Table 3-35 Transfer-of-Control Group: Register Test and Branch.**

Description	Mnemonic	Operation	CC	Examples	Mode	$\sim$
Test Register and Branch if Zero  <i>r</i> = A,B,D,X,Y,S	TBEQ <i>r,rel9</i>	IF ( <i>r</i> ) = 0 THEN (PC) $\leftarrow$ (PC) + <i>rel9</i>	–	TBEQ A, label1		3
				TBEQ Y, label1		3
Test Register and Branch if Not Zero  <i>r</i> = A,B,D,X,Y,S	TBNE <i>r,rel9</i>	IF ( <i>r</i> ) $\neq$ 0 THEN (PC) $\leftarrow$ (PC) + <i>rel9</i>	–	TBNE X, label1		3
				TBNE SP, label1		3

The final subset of “compare and branch” instructions allows the named register to be incremented or decremented, and causes the branch to be taken based on whether or not the register has reached zero. The four variants – **IBEQ** (“increment register and branch if zero”), **IBNE** (“increment register and branch if not zero”), **DBEQ** (“decrement register and branch if zero”), and **DBNE** (“decrement register and branch if not zero”) – are listed in Table 3-36. These instructions are quite useful in creating programs with simple, low overhead loop structures.

**IBEQ**  
**IBNE**  
**DBEQ**  
**DBNE**

It's probably safe to say that the 68HC12 has one of the most versatile sets of conditional branch instructions out there – certainly more than a typical Palm Beach County poll worker could accurately count...especially the ones quoted as saying, "What should I do when I run out of hands?"

**Table 3-36 Transfer-of-Control Group: Increment/Decrement Register, Test, and Branch.**

Description	Mnemonic	Operation	CC	Examples	Mode	~
Inc Register and Branch if Zero <i>r</i> = A,B,D,X,Y,SP	IBEQ <i>r,rel9</i>	$(r) \leftarrow (r) + 1$ IF $(r) = 0$ THEN $(PC) \leftarrow (PC) + rel9$	–	IBEQ A, label		3
				IBEQ Y, label		3
Inc Register and Branch if Not Zero <i>r</i> = A,B,D,X,Y,SP	IBNE <i>r,rel9</i>	$(r) \leftarrow (r) + 1$ IF $(r) \neq 0$ THEN $(PC) \leftarrow (PC) + rel9$	–	IBNE X, label		3
				IBNE SP, label		3
Dec Register and Branch if Zero <i>r</i> = A,B,D,X,Y,SP	DBEQ <i>r,rel9</i>	$(r) \leftarrow (r) - 1$ IF $(r) = 0$ THEN $(PC) \leftarrow (PC) + rel9$	–	DBEQ A, label		3
				DBEQ Y, label		3
Dec Register and Branch if Not Zero <i>r</i> = A,B,D,X,Y,SP	DBNE <i>r,rel9</i>	$(r) \leftarrow (r) - 1$ IF $(r) \neq 0$ THEN $(PC) \leftarrow (PC) + rel9$	–	DBNE X, label		3
				DBNE SP, label		3

### 3.8.5 Machine Control Group Instructions

This group, as it turns out, might be more palatable in Palm Beach County than the one just completed since we can literally count its members "by hand" (i.e., there are fewer than ten). The purpose and function of most of these instructions will not become clear until we formally introduce the topic of interrupts in Chapter 5. For the sake of discussion here, an *interrupt* can be viewed as an asynchronous (or "unexpected"), *hardware*-induced subroutine call. This is in contrast to what is sometimes called an *exception*, which is also "unexpected" but typically not induced by a "hardware signal". Rather, an exception is induced by a run-time anomaly encountered as the program executes. (Unfortunately, the terms "interrupt" and "exception" are sometimes used interchangeably – see sidebar.)

*interrupt*  
*exception*

Some examples may be helpful here. Pressing a key on a keypad, requesting transmission of the next character, and signaling completion of a data conversion are classic examples of asynchronous "events" that might trigger the execution of an *interrupt service routine*. Here, assertion of a hardware signal causes the processor to *alter* its fetch cycle. Instead of processing the next instruction pointed to by the PC, it looks up the address of the routine dedicated to servicing the interrupt request (from an "interrupt vector table"), saves the machine state (or "context"), and transfers control to that routine. In other words, the equivalent of a "subroutine call" takes place, along with saving the machine state, in

*interrupt service routine*  
*context*  
*interrupt vector table*

response to a hardware signal. Note that the machine state, which consists of all the program-visible registers *except SP*, must be saved on the stack so that interrupt handling occurs *transparently*, i.e., the “interrupted program” is oblivious to having been interrupted.

**Table 3-37** Machine Control Group.

Description	Mnemonic	Operation	CC	Examples	M	~
Return from Interrupt	RTI	$(CCR) \leftarrow ((SP)), (SP) \leftarrow (SP) + 1,$ $(D) \leftarrow ((SP)), (SP) \leftarrow (SP) + 2,$ $(X) \leftarrow ((SP)), (SP) \leftarrow (SP) + 2,$ $(Y) \leftarrow ((SP)), (SP) \leftarrow (SP) + 2,$ $(PC) \leftarrow ((SP)), (SP) \leftarrow (SP) + 2$	all <sup>1</sup>	RTI	◎	$8/10^2$
Unimplemented Opcode Trap	TRAP	$(SP) \leftarrow (SP) - 2, ((SP)) \leftarrow (PC),$ $(SP) \leftarrow (SP) - 2, ((SP)) \leftarrow (Y),$ $(SP) \leftarrow (SP) - 2, ((SP)) \leftarrow (X),$ $(SP) \leftarrow (SP) - 2, ((SP)) \leftarrow (D),$ $(SP) \leftarrow (SP) - 1, ((SP)) \leftarrow (CCR),$ <b>I bit of CCR ← 1,  <math>(PC) \rightarrow (\text{Trap Vector})</math></b>	—	\$18 tn <sup>3</sup>	◎	11
Software Interrupt	SWI	$(SP) \leftarrow (SP) - 2, ((SP)) \leftarrow (PC),$ $(SP) \leftarrow (SP) - 2, ((SP)) \leftarrow (Y),$ $(SP) \leftarrow (SP) - 2, ((SP)) \leftarrow (X),$ $(SP) \leftarrow (SP) - 2, ((SP)) \leftarrow (D),$ $(SP) \leftarrow (SP) - 1, ((SP)) \leftarrow (CCR),$ <b>I bit of CCR ← 1,  <math>(PC) \rightarrow (\text{SWI Vector})</math></b>	—	SWI	◎	9
Enter Background Debug Mode	BGND	Like a software interrupt, but no registers are stacked – routines in the BDM ROM control operation	—	BGND	◎	5
Wait for Interrupt	WAI	$(SP) \leftarrow (SP) - 2, ((SP)) \leftarrow (PC),$ $(SP) \leftarrow (SP) - 2, ((SP)) \leftarrow (Y),$ $(SP) \leftarrow (SP) - 2, ((SP)) \leftarrow (X),$ $(SP) \leftarrow (SP) - 2, ((SP)) \leftarrow (D),$ $(SP) \leftarrow (SP) - 1, ((SP)) \leftarrow (CCR),$ <b>Stop CPU Clocks</b>	—	WAI	◎	$8/5^4$
Stop Processing	STOP	$(SP) \leftarrow (SP) - 2, ((SP)) \leftarrow (PC),$ $(SP) \leftarrow (SP) - 2, ((SP)) \leftarrow (Y),$ $(SP) \leftarrow (SP) - 2, ((SP)) \leftarrow (X),$ $(SP) \leftarrow (SP) - 2, ((SP)) \leftarrow (D),$ $(SP) \leftarrow (SP) - 1, ((SP)) \leftarrow (CCR),$ <b>Stop All Clocks</b>	—	STOP	◎	$9/5^4$
No-operation	NOP	—	—	NOP	◎	1

<sup>1</sup> RTI affects *all* the condition code bits, with the *exception of X*, which cannot be set by a software instruction once it is cleared.

<sup>2</sup> Normal execution requires 8 cycles. If another interrupt is pending when the RTI is executed, 10 cycles are consumed.

<sup>3</sup> Unimplemented 2-byte opcodes are those where the first opcode byte is \$18 and the second opcode byte ranges from \$30 to \$39 or \$40 to \$FF.

<sup>4</sup> The cycles listed correspond to entering and exiting WAI or STOP.

At the conclusion of an interrupt service routine, a “special” version of the “return” instruction is needed – one that restores the machine state in addition to resuming the “main-line” program at the point it was interrupted. This leads us to our first Machine Control Group instruction, return from interrupt (**RTI**), listed in Table 3-37. This instruction simply restores each register from the copy saved previously on the stack. Note that restoring the PC causes the interrupted program to resume where it left off.

Interrupts provide a convenient framework for constructing “real-time” (or “event-driven”) embedded control systems. Stated another way, interrupts are a “way of life” in the design of microcontroller-based products. This is in contrast to *exceptions*, which are typically associated with “something bad” happening. Overflow, dividing by zero, or attempting to execute an invalid opcode are examples of exceptions.

On the 68HC12, attempting to execute an invalid opcode will cause a “trap” to occur. As such, a trap can be construed as an exception. Similar to an interrupt, a trap causes the processor to save its state on the stack and transfer control to a “trap handling” routine. Note that the **TRAP** mnemonic, listed in Table 3-37, is not recognized by assembler programs; rather, it simply documents the processor’s response to an unrecognized (“unimplemented”) opcode. Perhaps somewhat insidiously, TRAP can be used to advantage, allowing a system designer to define “new” instructions comprised of unused opcodes. Here the TRAP handling routine would be used to *emulate*, in software, the processing of these new instructions. An example of where this might be used is for “higher-level” functions such as floating-point arithmetic.

Sometimes it is useful to “force” an exception to occur in the normal software execution stream. This is particularly useful in debugging code, where one might wish to temporarily “interrupt” a program (by virtue of hitting a “breakpoint”) to check the state of registers and/or memory locations. On the 68HC12, this can be accomplished using the “software interrupt” (**SWI**) instruction, also listed in Table 3-37. Like a TRAP, execution of an SWI instruction saves the machine state on the stack and transfers control to an SWI handling routine. For program debugging, SWI instructions can be ‘manually’ inserted in code, or “automatically” inserted by a “debug monitor” (e.g., Motorola’s D-Bug12).

**RTI****TRAP****SWI**

### I Take Exception to Your Interrupt

*The distinction between interrupts and exceptions is sometimes blurred. While the name of the “software interrupt” (SWI) instruction aptly describes what it does (i.e., interrupts the normal flow of software execution), it’s really not an interrupt as defined earlier – rather, it is an exception.*

In addition to the “software breakpoint” capability afforded by the SWI instruction, the 68HC12 has an even more powerful debugging capability, called *background debug mode*. Here, a target microcontroller system running an application can be interrogated by a “pod” (a second 68HC12) via a single-wire serial interface. The “pod” 68HC12, operating in “BDM mode”, can start or stop the target application as well as retrieve the state of registers or memory locations while the application is running. Background debug mode is commenced through execution of the BGND instruction, listed in Table 3-37. Hardware-assisted debugging is now a common feature in many modern microprocessors and microcontrollers.

*background  
debug mode  
BDM*

BGND

The “wait” (WAI) instruction, listed next in Table 3-37, provides a means for allowing the processor to “pause” execution (effected by stopping the CPU clock) until an interrupt occurs. When a WAI instruction is executed, the machine state is saved on the stack and the CPU clock is stopped (the clock signals provided to the on-chip peripherals continue to run, however). The WAI instruction is useful in applications where the CPU, at a given point in a program, doesn’t have anything meaningful to do until an interrupt occurs.

WAI

The “stop” (STOP) instruction is similar to WAI, but a bit more “drastic”. Like WAI, execution of a STOP instruction causes the machine state to be saved on the stack. After that occurs, *all* the clocks are stopped (including those supplied to the on-chip peripherals), effectively putting the 68HC12 in “standby” mode. While in standby mode, the internal state is maintained along with the states of I/O pins; power consumption, though, is greatly reduced. Asserting RESET or an interrupt input ends standby mode. For STOP to be executed, the “stop disable” (S) bit in the condition code register must be cleared; if the S bit is set, execution of STOP simply consumes two cycles. The STOP instruction is useful in battery-powered applications where there is a benefit from putting the processor “to sleep” for extended periods of inactivity to maximize battery life.

STOP

The final machine control instruction listed in Table 3-37 does nothing! The only purpose in life for “no-operation” (NOP) is to consume an execution cycle, sometimes useful in so-called “delay loops”. Examples of no-ops by other names include “branch never” (BRN), that also consumes one cycle; and “long branch never” (LBRN), that consumes three cycles. Recall that some addressing mode variants of the LEA instruction also accomplish nothing more than consuming cycles.

NOP

BRN

LBRN

### 3.8.6 Special Group Instructions

Special, as its name implies, is used to refer to instructions that are not ordinarily included on “generic” microcontrollers. Unfortunately, this distinction is far from absolute, given the tendency of manufacturers to

continuously expand “features” based on the increasing availability of chip “real estate”.

The 68HC12 sports several subsets of instructions that might be deemed “special”. The MIN/MAX instructions and EMACS instruction, covered previously as part of the arithmetic group, could be called “special” since few “generic” microcontrollers have these capabilities. With a bit more confidence, though, we could claim that the “lookup and interpolate” (TBL) and “fuzzy logic” instructions are indeed “special” – they are not only “more rare” among mainstream microcontrollers, but also fit “less nicely” into the broad categories of instructions previously defined. Our special group, then, will consist only of these latter two subsets.

The “lookup and interpolate” (TBL) instruction is documented in Table 3-38. This instruction, or its “extended cousin” (ETBL), can be used to perform a linear interpolation on values that fall between a pair of data entries in a lookup table stored in memory. A lookup table is simply an array of values that can be used to perform data translations or conversions. The TBL instruction facilitates very compact storage of lookup tables that are piece-wise linear.

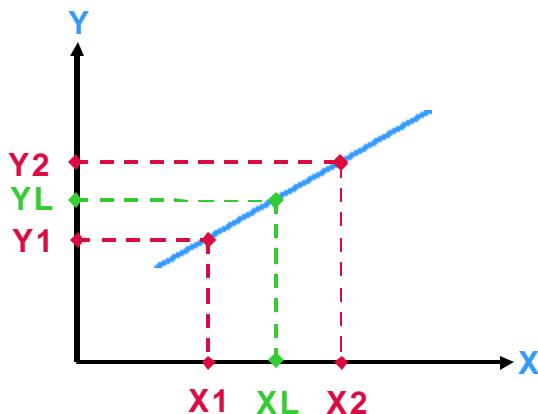
**TBL**

**ETBL**

**Table 3-38 Special Group: Table Lookup and Interpolate.**

Description	Mnemonic	Operation	CC	Examples	Mode	~
Table Lookup and Interpolate	TBL <i>addr</i> <i>addr</i> = $\oplus^*$	$(A) \leftarrow (addr) + \{ (B) X ((addr+1) - (addr)) \}$	N $\leftarrow \ddagger$ Z $\leftarrow \ddagger$ C $\leftarrow ?$	TBL 0,X TBL 2,X+ TBL 2,Y- TBL -16t,PC TBL 15t,SP	$\oplus$	8
	ETBL <i>addr</i> <i>addr</i> = $\oplus^*$	$(D) \leftarrow (addr):(addr+1) + \{ (B) X \{ (addr+2):(addr+3) - (addr):(addr+1) \} \}$	N $\leftarrow \ddagger$ Z $\leftarrow \ddagger$ C $\leftarrow ?$	ETBL 0,X ETBL 2,X+ ETBL 2,Y- ETBL -16t,PC ETBL 15t,SP	$\oplus$	10

\*Only indexed modes with “short” constant offsets (requiring no extension bytes) can be used.



**Figure 3-27** Illustration of TBL Parameters.

Successful use of TBL involves a multi-step process – perhaps another reason for calling it “special”. Referring to Figure 3-27, the desired “lookup point” ( $XL$ ) is in-between (the nearest) two table entries stored in memory:  $X1$  and  $X2$ . Given these points along the X-axis, the calculations  $XL-X1$  and  $X2-X1$  are then made. Using FDIV, a binary fraction is calculated based on dividing  $XL-X1$  by  $X2-X1$ ; the resulting unsigned fraction is then placed in the B register. The last step before executing TBL is to set an index register (X, Y, SP, or PC) to point to the first table entry,  $X1$ . Execution of TBL then produces the following result in the A register:  $(A) \leftarrow (addr) + \{ (B) X \{ (addr+1) - (addr) \} \}$ .

As an example, consider the function represented by the (base 10) X-Y data points (1,10), (2,20), (4,50), and (5,80). Assume the “Y” value corresponding to  $XL = 2.5$  is desired. Here,  $X1 = 2$  and  $X2 = 4$ . Plugging in the numbers,  $XL-X1 = 0.5$  and  $X2-X1 = 2$ ; therefore,  $(XL-X1)/(X2-X1) = 0.25$ . With an index register pointed to the “ $X1$ ” table entry (i.e., the value 20) and the binary fraction 01000000b ( $0.25_{10}$ ) in the B register, TBL performs the following calculation:  $(A) = 20 + \{ 0.25 X \{ 50 - 20 \} \} = 20 + 7 = 27$ . Note that the intermediate value resulting from the fractional multiplication is not rounded, and therefore truncated to 7, yielding an interpolated value of  $27_{10}$  in the A register as TBL’s “final answer”.

**Table 3-39 Special Group: Fuzzy Logic.**

Description	Mnemonic	Operation	CC	Examples	~
Determine Grade of Membership	MEM	$((Y)) \leftarrow \text{grade of membership}$ $(Y) \leftarrow (Y) + 1$ $(X) \leftarrow (X) + 4$	$N \leftarrow ?$ $Z \leftarrow ?$ $V \leftarrow ?$ $C \leftarrow ?$ $H \leftarrow ?$	MEM	5
Fuzzy Logic Rule Evaluation	REV	MIN – MAX rule evaluation	$N \leftarrow ?$ $Z \leftarrow ?$ $V \leftarrow 1$ $C \leftarrow ?$ $H \leftarrow ?$	REV	*
Fuzzy Logic Rule Evaluation (Weighted)	REVVW	MIN – MAX rule evaluation with optional rule weighting; C bit in CCR selects weighted (1) or unweighted (0) rule evaluation	$N \leftarrow ?$ $Z \leftarrow ?$ $V \leftarrow 1$ $C \leftarrow ?$ $H \leftarrow ?$	REVVW	*
Weighted Average	WAV	Performs weighted average calculations on values stored in memory	$N \leftarrow ?$ $Z \leftarrow 1$ $V \leftarrow ?$ $C \leftarrow ?$ $H \leftarrow ?$	WAV	*

\*Number of cycles varies based on number of elements in rule list.

There are, at this point, only four 68HC12 instructions that remain: those that support fuzzy logic. These instructions, listed in Table 3-39, are **MEM**, which evaluates trapezoidal membership functions; **REV** and **REVVW**, which perform unweighted or weighted MIN-MAX rule evaluation;

and [WAV](#), which performs weighted average defuzzification on singleton output membership functions. The actions associated with these instructions are relatively involved and complex compared with other 68HC12 instructions. To fully understand them requires a background on fuzzy logic. We will illustrate their use in a programming example in the chapter that follows.

*fuzzy logic*

[MEM](#)  
[REV](#)  
[REVV](#)  
[WAV](#)

### 3.9 Summary and References

We began this chapter with the “Norm analogy” – that machine instructions available to a computer engineer are like the “tools in the toolbox” available to a master carpenter. Our objective was to learn *what* tools we had in our “instruction set” toolbox along with some basics on *how* to use them. The lab experiments and homework problems included with this chapter will help you learn this material. There is no substitute for “hands on” practice!

The authoritative reference for the material covered in this chapter is Motorola’s *CPU12 Reference Manual*. A “soft copy” of this manual is included as a PDF on the CD-ROM that accompanies this text; a printed copy can be obtained directly from Motorola’s Literature Distribution Center (LDC). A printed copy is also bundled with the M68EVB912B32 Evaluation Board.

Students who purchase the EVB will also want to become familiar with the material covered in the first three chapters of Motorola’s *M68EVB912B32 Evaluation Board User’s Manual*. A “soft copy” of this manual is included as a PDF on the CD-ROM that accompanies this text; a printed copy can be obtained directly from Motorola’s Literature Distribution Center (LDC). A printed copy is also bundled with the M68EVB912B32 Evaluation Board. Looking through the *IASM12 User’s Guide*, included as a “.doc” file on the IASM12 diskette bundled with the EVB, will also prove helpful.

Readers interested in a more complete account of the “RISC-CISC” debate, summarized at the beginning of this chapter, may want to review several key papers written on the subject:

- Patterson, D., “Reduced Instruction Set Computers,” *Communications of the ACM*, January 1985, pp. 8-21.
- Colwell, R., et. al., “Computers, Complexity, and Controversy,” *IEEE Computer*, September 1985, pp. 8-19.
- Wallich, P., “Toward Simpler, Faster Computers,” *IEEE Spectrum*, August 1985, pp. 38-45.

A thorough (as well as entertaining) summary and analysis of the “byte-ordering” debate can be found in the article, “On Holy Wars and a Plea for Peace”, which can be found at <http://www.op.net/docs/RFCs/ien-137>.

## Problems

The CD-ROM that accompanies this text includes a printable version of the problems that follow in PDF format. Selected problems can be printed from this file and completed on the "full size" sheets produced.

- 3-1. Disassemble the 68HC12 machine code listed below and "single step" through it by hand, completing the chart below. Write the disassembled instructions under the *Disassembled Instructions* heading, clearly indicating the instructions associated with the specific memory contents. Each "step" refers to the execution of one instruction. Assume the first opcode byte is at location 0800h.

<i>Address</i>	<i>Contents</i>	<i>Disassembled Instructions</i>
----------------	-----------------	----------------------------------

0800	86
0801	E2
0802	C6
0803	42
0804	18
0805	06
0806	86
0807	43
0808	8B
0809	71
080A	18
080B	07
080C	36
080D	E0
080E	B0
080F	3F

Execution Step	(PC)	(A)	(B)	(CC)
Initial Values	0800	00	00	90
After Single Step 1				
After Single Step 2				
After Single Step 3				
After Single Step 4				
After Single Step 5				
After Single Step 6				
After Single Step 7				
After Single Step 8				
After Single Step 9				
After Single Step 10				

- 3-2. Disassemble the 68HC12 machine code listed below and "single step" through it by hand, completing the chart below. Write the disassembled instructions under the *Disassembled Instruction* heading, clearly indicating the instructions associated with the specific memory contents. Each "step" refers to the execution of one instruction. Assume the first opcode byte is at location 0900h.

<i>Address</i>	<i>Contents</i>	<i>Disassembled Instruction</i>
0900	86	
0901	53	
0902	8B	
0903	97	
0904	18	
0905	07	
0906	C6	
0907	87	
0908	37	
0909	AB	
090A	80	
090B	18	
090C	07	
090D	86	
090E	19	
090F	A0	
0910	B0	
0911	18	
0912	07	
0913	3F	

<i>Address</i>	<i>Contents</i>	<i>Disassembled Instruction</i>
0900	86	
0901	53	
0902	8B	
0903	97	
0904	18	
0905	07	
0906	C6	
0907	87	
0908	37	
0909	AB	
090A	80	
090B	18	
090C	07	
090D	86	
090E	19	
090F	A0	
0910	B0	
0911	18	
0912	07	
0913	3F	

Execution Step	(PC)	(A)	(B)	(CC)
Initial Values	0900	00	00	90
After Single Step 1				
After Single Step 2				
After Single Step 3				
After Single Step 4				
After Single Step 5				
After Single Step 6				
After Single Step 7				
After Single Step 8				
After Single Step 9				
After Single Step 10				

- 3-3. Assemble the 68HC12 instructions listed below into machine code. Place the assembled machine code (corresponding with the instructions) into memory under the *Contents* heading. Assume an **ORG 0802h** precedes the instructions listed below. Be sure to clearly indicate how the instructions and memory contents correspond.

<i>Address</i>	<i>Contents</i>	<i>Instructions</i>
0800		LDAB #\\$8A
0801		ORAB \\$0954
0802		LDAA #\$AB
0803		ABA
0804		STAA \$09DE
0805		PSHA
0806		PULB
0807		LDAA #\$02
0808		STAB A,X
0809		SWI
080A		
080B		
080C		
080D		
080E		
080F		
0810		
0811		
0812		
0813		
0814		

- 3-4. Write a specific example of 12 additional 68HC12 addressing mode variations of an LDAB instruction. Write the name of each specific addressing mode, the instruction byte count, and the instruction cycle count.

Assembly Source Form	Formal (Complete) Addressing Mode Name	Motorola Abbreviation	Byte Count	Cycle Count
LDAB \$091E	Extended	EXT	3	3

- 3-5. The following table shows the data initially stored in a 68HC12's memory, starting at location 0900h. The initial value of the registers is also given. Assume the five instructions listed in parts (a) – (e) are stored elsewhere in memory, and executed in the order listed (i.e., execution of a given instruction *may* affect the execution of a subsequent instruction). Complete the blanks for each instruction.

ADDRESS	CONTENTS	ADDRESS	CONTENTS
0900	08	0908	67
0901	01	0909	2E
0902	FD	090A	BC
0903	9D	090B	9E
0904	09	090C	43
0905	0D	090D	24
0906	7E	090E	09
0907	F3	090F	02

Initial Values: (A) = 00, (B) = 00, (CC) = 91, (X) = 0906, (Y) = 0900

(a) LDAA 2, X

(A) = \_\_\_\_ h NF = \_\_\_\_ Cycles = \_\_\_\_

Addressing Mode = \_\_\_\_\_

(b) ADCA [4, Y]

(A) = \_\_\_\_ h CF = \_\_\_\_ Cycles = \_\_\_\_

Addressing Mode = \_\_\_\_\_

(c) LDAB 3, X+

(B) = \_\_\_\_ h ZF = \_\_\_\_ Cycles = \_\_\_\_

Addressing Mode = \_\_\_\_\_

(d) STAB 3, X

(B) = \_\_\_\_ h VF = \_\_\_\_ Cycles = \_\_\_\_

Addressing Mode = \_\_\_\_\_

(e) EORA 2, X

(A) = \_\_\_\_ h NF = \_\_\_\_ Cycles = \_\_\_\_

Addressing Mode = \_\_\_\_\_

- 3-6. The following table shows the data initially stored in a 68HC12's memory, starting at location 0800h. The initial value of the registers is also given. Assume the five instructions listed in parts (a) – (e) are stored elsewhere in memory, and executed in the order listed (i.e., execution of a given instruction *may* affect the execution of a subsequent instruction). Complete the blanks for each instruction.

ADDRESS	CONTENTS	ADDRESS	CONTENTS
0800	11	0808	08
0801	22	0809	01
0802	33	080A	08
0803	44	080B	02
0804	55	080C	08
0805	66	080D	03
0806	77	080E	08
0807	88	080F	04

Initial Values: (A) = 00, (CC) = 91, (X) = 0804, (Y) = 0808

(a) ADCA -2 , X

(A) = \_\_\_\_\_ h CF = \_\_\_\_\_ Cycles = \_\_\_\_\_

Addressing Mode = \_\_\_\_\_

(b) SBCA [ 6 , Y ]

(A) = \_\_\_\_\_ h CF = \_\_\_\_\_ Cycles = \_\_\_\_\_

Addressing Mode = \_\_\_\_\_

(c) LDAA 3 , X

(A) = \_\_\_\_\_ h NF = \_\_\_\_\_ Cycles = \_\_\_\_\_

Addressing Mode = \_\_\_\_\_

(d) EORA [ 0 , Y ]

(A) = \_\_\_\_\_ h ZF = \_\_\_\_\_ Cycles = \_\_\_\_\_

Addressing Mode = \_\_\_\_\_

(e) ANDA 1 , +X

(A) = \_\_\_\_\_ h NF = \_\_\_\_\_ Cycles = \_\_\_\_\_

Addressing Mode = \_\_\_\_\_

- 3-7. The following table shows the data initially stored in a 68HC12's memory, starting at location 0800h. The initial value of the registers is also given. Assume the five instructions listed in parts (a) – (e) are stored elsewhere in memory, and executed in the order listed (i.e., execution of a given instruction *may* affect the execution of a subsequent instruction). Complete the blanks for each instruction.

ADDRESS	CONTENTS	ADDRESS	CONTENTS
0800	11	0808	08
0801	22	0809	01
0802	33	080A	08
0803	44	080B	02
0804	55	080C	08
0805	66	080D	03
0806	77	080E	08
0807	88	080F	04

Initial Values: (A) = 00, (CC) = 91, (X) = 0803, (Y) = 080E

(a) ADCA \$0805

(A) = \_\_\_\_ h CF = \_\_\_\_ Cycles = \_\_\_\_

Addressing Mode = \_\_\_\_\_

(b) SBCA #\$99

(A) = \_\_\_\_ h CF = \_\_\_\_ Cycles = \_\_\_\_

Addressing Mode = \_\_\_\_\_

(c) LDAA -2,X

(A) = \_\_\_\_ h NF = \_\_\_\_ Cycles = \_\_\_\_

Addressing Mode = \_\_\_\_\_

(d) ORAA [-2,Y]

(A) = \_\_\_\_ h ZF = \_\_\_\_ Cycles = \_\_\_\_

Addressing Mode = \_\_\_\_\_

(e) ANDA 2,X+

(A) = \_\_\_\_ h NF = \_\_\_\_ Cycles = \_\_\_\_

Addressing Mode = \_\_\_\_\_

- 3-8. For the program listing shown below, show the contents of the PC, SP, D, X, and Y registers as well as the contents of the memory locations indicated (reserved for the stack area) after the execution of each *marked* instruction. Initially, (CC) = 90. Any stack locations that are “don’t cares” should be designated “XX”. The assembly source file for this problem is available on the CD-ROM that accompanies this text.

```

0800          1      ORG      $800
0800 [02] CD09D7 2      LDY      #$09D7
0803 [02] 35    3      PSHY
0804 [02] 0702 4      BSR      SUBR    ; *** 1 ***
0806 [03] 30    5      PULX
0807 [09] 3F    6      SWI
0808 [03] EC82  7
080A [02] 36    8      SUBR    LDD     2 , SP
080B [01] 46    9      PSHA
080C [03] EBB0 10     RORA
080E [01] 55    11    ADDB    1 , SP+   ; *** 3 ***
080F [02] 6C82 12     ROLB
0811 [05] 3D    13     STD     2 , SP   ; *** 4 ***
0812          14     RTS
0812          15     END

```

Registers	Initial	After *1*	After *2*	After *3*	After *4*	After *5*
(PC)	0800					
(SP)	0A00					
(D)	0000					
(X)	0000					
(Y)	0000					
Stack	Initial	After *1*	After *2*	After *3*	After *4*	After *5*
(09FA)	00					
(09FB)	00					
(09FC)	00					
(09FD)	00					
(09FE)	00					
(09FF)	00					
(0A00)	00					

- 3-9. For the program listing shown below, show the contents of the PC, SP, D, X, and Y registers as well as the contents of the memory locations indicated (reserved for the stack area) after the execution of each *marked* instruction. Initially, (CC) = 90. Any stack locations that are “don’t cares” should be designated “XX”. The assembly source file for this problem is available on the CD-ROM that accompanies this text.

```

0800          1      ORG    $800
0800 [02] CE9876 2      LDX    #$9876
0803 [02] 34     3      PSHX
0804 [02] 0702  4      BSR    SUBR   ; *** 1 ***
0806 [03] 31     5      PULY
0807 [09] 3F     6      SWI
0808 [03] EC82  7
080A [01] 45     8      SUBR   LDD    2,SP
080B [03] A682  9      ROLA
080D [01] 55     10     LDAA   2,SP   ; *** 2 ***
080E [01] 45     11     ROLB   ; *** 3 ***
080F [02] 6C82  12     ROLA
0811 [05] 3D     13     STD    2,SP   ; *** 4 ***
0812           14     RTS
0812           15     END

```

Registers	Initial	After *1*	After *2*	After *3*	After *4*	After *5*
(PC)	0800					
(SP)	0A00					
(D)	0000					
(X)	0000					
(Y)	0000					
Stack	Initial	After *1*	After *2*	After *3*	After *4*	After *5*
(09FA)	00					
(09FB)	00					
(09FC)	00					
(09FD)	00					
(09FE)	00					
(09FF)	00					
(0A00)	00					

- 3-10. For the program listing shown below, show the contents of the PC, SP, D, X, and Y registers as well as the contents of the memory locations indicated (reserved for the stack area) after the execution of each *marked* instruction. Initially, (CC) = 90. Any stack locations that are “don’t cares” should be designated “XX”. The assembly source file for this problem is available on the CD-ROM that accompanies this text.

```

0800          1      ORG    $800
0800 [02] CEFEDC 2      LDX    #$FEDC
0803 [02] CD1234 3      LDY    #$1234
0806 [02] 34     4      PSHX
0807 [02] 35     5      PSHY      ; *** 1 ***
0808 [02] 0703  6      BSR    SUBR   ; *** 2 ***
080A [03] 31     7      PULY
080B [03] 30     8      PULX      ; *** 5 ***
080C [09] 3F     9      SWI
10
080D [03] EC82  11     SUBR   LDD    2,SP
080F [01] 59     12     LS LD   STD    2,SP      ; *** 3 ***
0810 [02] 6C82  13     STD    4,SP
0812 [03] EC84  14     LDD    4,SP
0814 [01] 59     15     ASLD
0815 [02] 6C84  16     STD    4,SP      ; *** 4 ***
0817 [05] 3D     17     RTS
0818          18     END

```

Registers	Initial	After *1*	After *2*	After *3*	After *4*	After *5*
(PC)	0800					
(SP)	0A00					
(D)	0000					
(X)	0000					
(Y)	0000					
Stack	Initial	After *1*	After *2*	After *3*	After *4*	After *5*
(09FA)	00					
(09FB)	00					
(09FC)	00					
(09FD)	00					
(09FE)	00					
(09FF)	00					
(0A00)	00					

3-11. Describe the actions caused by the following lines of code, such that the differences among them are clear.

- STAA -2,X
- STAA 2,-X
- STAA 2,X-

3-12. Describe the actions caused by the following lines of code, such that the differences among them are clear.

- LDAB 3,Y-
- LDAB -3,Y
- LDAB 3,-Y

3-13. For each of the following lines of code, write an instruction that performs the equivalent function.

- LDAB 1,SP+
- STAB 1,-SP
- ASLB

3-14. Show how, using LDAA and STAA instructions in conjunction with the 68HC12's auto increment/decrement addressing modes, the X index register can be used as a "software" stack pointer for implementing the equivalent of the "PSHA" and "PULA" instructions, here using the same convention as the SP register (which points to the *top stack item*).

3-15. Show how, using LDD and STD instructions in conjunction with the 68HC12's auto increment/decrement addressing modes, the Y index register can be used as a "software" stack pointer for implementing the equivalent of "PSHD" and "PULD", here using the convention that the software stack pointer (Y) points to the *next available location*.

3-16. Indicate the D-Bug12 monitor command that should be used to accomplish each of the following operations:

- set the serial port baud rate
- load user program S-record object file
- reset the 68HC12
- modify the 68HC12 register contents
- modify memory (SRAM) contents
- begin execution of a user program
- execute a single instruction and display register contents
- set/display user breakpoints
- clear user breakpoints
- enter assembly instruction mnemonics line-by-line
- display contents of memory
- display contents of registers
- bulk erase byte-erasable EEPROM
- execute a user subroutine
- set a temporary breakpoint and begin execution of a user program

3-17. Provide a single-sentence explanation of the four modes in which the M68EVB912B32 can begin operation: EVB mode, JUMP-EE mode, POD mode, and BOOTLOAD mode.

# Notes